

Human in the Loop

Human in the Loop

Misunderstood

Lazaros Toumanidis

Copyright © 2026 Lazaros Toumanidis & the Waldiez team
This work is licensed under the Creative Commons
Attribution-NonCommercial 4.0 International License
(CC BY-NC 4.0). You are free to share and adapt
this material for non-commercial purposes, provided
appropriate credit is given.
<https://creativecommons.org/licenses/by-nc/4.0/>

The content of this book was produced with substantial
AI assistance, under human direction, editing, and review.
The author is responsible for its content.

Set in Libertinus Serif, Libertinus Sans, Libertinus Math, and JetBrains Mono

<https://laztoun.com> <https://what-if.io> <https://limen-os.io>
laztoun@protonmail.com

Contents

Preface	xxi
I When Machines Ask for Help	1
1 The Computers That Ask for Help	5
1.1 The Pause That Makes You Think	5
1.2 The Hidden Helpers Everywhere	6
1.3 When Smart Gets Smarter	7
1.4 When the Asking Goes Wrong	8
1.5 The Thermostat That Cried Wolf	9
1.6 The Systems That Didn't Ask	11
1.7 The High-Stakes Version	12
1.8 When Chatbots Pretend to Know	13
1.9 The Goldilocks Problem	15
1.10 What Makes a System "Smart Enough to Ask for Help"? .	16
1.11 A Brief Technical Interlude	19
1.12 The Future Is Collaborative, Not Automated	20
1.13 Why You Should Care	21
Chapter Summary	22
Technical Appendix 1A: Uncertainty Quantification	27
1A.1 Mathematical Framework for Uncertainty-Driven Human Intervention	27

1A.2	Uncertainty Estimation Methods in Practice	30
1A.3	Real-World Implementation Challenges	32
1A.4	Evaluation Metrics for HITL Systems	33
1A.5	Case Study: Netflix's Engagement Detection (Success Case)	36
1A.6	Case Study: Nest Thermostat Home/Away Assist (Failure Case)	38
1A.7	Extensions and Future Directions	43
1A.8	Discussion Questions for Academic Use	45
1A.9	Recommended Readings	46
1A.10	Exercises and Assignments	47
2	When Smart Systems Get Confused	51
2.1	The Flag That Saved the Patient	51
2.2	Two Kinds of "I Don't Know"	52
2.3	The View from Inside	54
2.4	The Email in the Middle	54
2.5	Accents, Edges, and the Limits of What Systems Have Seen	55
2.6	The Outside View: How Confusion Surfaces (or Doesn't) .	56
2.7	Wrong vs. Confused: A Crucial Distinction	57
2.8	What Confusion Tells You About Training	59
2.9	The Radiologist's Bargain	59
2.10	The First Dimension, Up Close	60
	Chapter Summary	61
	Technical Appendix 2A: The Mathematics of AI Confusion . .	65
2A.1	Formal Definitions of Uncertainty Types	65
2A.2	Measuring Calibration	67
2A.3	Uncertainty Estimation Methods	69
2A.4	Out-of-Distribution Detection	72
2A.5	The Radiology Case: Mathematical Analysis	73
2A.6	Exercises	74

2A.7 Recommended Reading	75
3 How Computers Learn (The Simple Version)	77
3.1 A Teacher Who Only Teaches Cats	77
3.2 What Machine Learning Actually Is	78
3.3 The Generalization Gap	79
3.4 Overfitting: The Expert Who Never Learned to Generalize	80
3.5 Underfitting: The Expert Who Never Learned Enough	81
3.6 Why Even a Perfect Learner Needs Help at the Edges	82
3.7 The Train/Test Gap as HITL Metaphor	83
3.8 Patterns, Not Rules	84
Chapter Summary	86
 Technical Appendix 3A: Machine Learning Foundations for	
HITL Design	89
3A.1 Empirical Risk Minimization	89
3A.2 The Bias-Variance Tradeoff	90
3A.3 Generalization Bounds	91
3A.4 Distribution Shift	92
3A.5 The PAC Learning Framework	94
3A.6 Exercises	95
3A.7 Recommended Reading	96
 4 From Confusion to Decision	 99
4.1 The Paramedic's Rule	99
4.2 The Probability Isn't the Decision	100
4.3 The Two Mistakes	100
4.4 The Cost Asymmetry	101
4.5 Decision Boundaries and Edge Cases	102
4.6 The Threshold Is Not Fixed	104
4.7 Three Real Systems, Three Threshold Philosophies	105

4.8 The Human Value Judgment That Hides in the Algorithm	106
Chapter Summary	107

Technical Appendix 4A: Threshold Optimization and Cost-

Sensitive Learning	111
4A.1 The ROC Curve and AUC	111
4A.2 Precision-Recall Tradeoff	113
4A.3 Cost-Sensitive Learning	115
4A.4 Multi-Threshold HITL Systems	116
4A.5 Exercises	118
4A.6 Recommended Reading	119

5 The Art of Asking for Help (Without Being Annoying) 121

5.1 A Better Question	121
5.2 The Ask Is a Design Problem	122
5.3 Three Levers	123
5.4 The Trust–Attention–Response Triangle	124
5.5 The Lazy User Assumption	125
5.6 Two-Factor Authentication: Lessons from the Field	126
5.7 The Fraud Alert as Design Canvas	128
5.8 Content Moderation and the Interface of Judgment	129
5.9 The 911 Dispatcher, Revisited	131
5.10 Feedback Integration: Does It Learn?	132
Chapter Summary	133

Technical Appendix 5A: Optimal Ask Design 137

5A.1 Information-Theoretic Foundations of Ask Design	137
5A.2 Alert Fatigue: A Dynamic Model	139
5A.3 Optimal Timing Models	141
5A.4 Format Design: Cognitive Load and Response Accuracy	142
5A.5 Feedback Integration: Closing the Learning Loop	142

5A.6 Exercises	143
5A.7 Recommended Reading	145
II The Human Side of the Loop	147
6 Interfaces That Make Humans Smarter	151
6.1 The Judge and the Number	151
6.2 The Invisible Architecture of Attention	153
6.3 Automation Bias and the Seduction of the Algorithm . . .	154
6.4 Death by a Thousand Clicks	156
6.5 What Explainable AI Actually Looks Like	157
6.6 Interfaces That Support Judgment vs. Interfaces That Re- place It	159
6.7 The Radiology Interface Wars	160
6.8 The Five Dimensions Through the Interface Lens	162
6.9 Why Design Is a Safety Issue	163
Chapter Summary	164
Technical Appendix 6A: Interface Design Foundations	167
6A.1 Cognitive Load Theory and Interface Design	167
6A.2 Automation Bias: Formal Model	169
6A.3 Explainability Output Quality Metrics	170
6A.4 Interface Sequence Models	171
6A.5 Attention Map Quality Assessment	172
6A.6 The Review Queue as a Cognitive System	174
6A.7 Exercises and Assignments	175
7 Getting Good Information from People	179
7.1 The Photograph That Changed How We See Data	179
7.2 The Instruction Problem	180

7.3	Measuring Agreement: Cohen's Kappa	181
7.4	The Annotator Is Not a Recording Device	183
7.5	The Economics of Getting Labels	184
7.6	Framing Effects: How the Question Changes the Answer	186
7.7	Calibration Sessions and the Iterative Guideline	187
7.8	What Low Agreement Tells You	187
	Chapter Summary	189

Technical Appendix 7A: Annotation Metrics 191

7A.1	Cohen's Kappa: Derivation and Properties	191
7A.2	Fleiss' Kappa: Extension to Multiple Annotators	193
7A.3	Annotation as a Bayesian Inference Problem: Dawid-Skene	194
7A.4	Framing Effects: Signal Detection Theory Framework	196
7A.5	Sample Size Planning for Annotation Projects	197
7A.6	Exercises and Assignments	199

8 The Blueprint for Smart Helper Systems 201

8.1	The Alarm Nobody Heard	201
8.2	The Pipeline Is Not the Loop	202
8.3	The Components of Real HITL Architecture	203
8.4	The Feedback Loop: From Human Response to Model	204
8.5	Logging, Monitoring, and Drift Detection	206
8.6	Where Things Break in Production	207
8.7	The Lifecycle: From Prototype to Deployed System	208
	Chapter Summary	209

Technical Appendix 8A: HITL System Architecture Patterns . 211

8A.1	The Pipeline vs. Loop Distinction	211
8A.2	The Three-Zone Router	212
8A.3	Active Learning Queue Implementation	214
8A.4	Production Monitoring	216

8A.5 Exercises	219
9 Tools Anyone Can Use	221
9.1 The Label Studio Moment	221
9.2 What Annotation Tools Actually Do	222
9.3 Self-Hosted and Open-Source Options	223
9.4 Cloud Platform Approaches	224
9.5 Crowdsourcing vs. Expert Labeling	225
9.6 No-Code Options and Their Limits	226
9.7 When to Build vs. Buy	227
9.8 Getting Started Practically	228
Chapter Summary	230
Technical Appendix 9A: Annotation Pipeline Tooling	233
9A.1 Label Studio API Integration	233
9A.2 Complete HITLAnnotationPipeline	235
9A.3 Tool Comparison	240
9A.4 Exercises	240
 III Teaching Machines to Perceive	 243
10 Teaching Computers to Understand Language	247
10.1 The Lawyer Who Trusted the Machine	247
10.2 Why Language Is Hard	248
10.3 Words as Coordinates in Meaning-Space	250
10.4 The Attention Mechanism: Reading While Noticing	251
10.5 What a Language Model “Knows”	252
10.6 The Hallucination Problem	253
10.7 Three Shapes of Human-in-the-Loop in Language AI	255
10.8 Teaching Values: RLHF, DPO, and Constitutional AI	257

10.9	The Five Dimensions, Applied to Language AI	259
10.10	Does It Understand? A Philosophical Aside	260
	Chapter Summary	262
Technical Appendix 10A: The Mathematics of Language Models		265
10A.1	Token Probability Distributions and Sampling	265
A.10.1	Token Probability Distributions and Sampling	265
10A.2	The Attention Mechanism	267
A.10.2	The Attention Mechanism	267
10A.3	RLHF Pipeline	268
A.10.3	RLHF Pipeline	268
10A.4	Direct Preference Optimization (DPO)	269
A.10.4	Direct Preference Optimization	269
10A.5	Constitutional AI: Two-Phase Process	270
A.10.5	Constitutional AI	270
10A.6	Calibration Metrics for Language Models	272
A.10.6	Calibration Metrics	272
11	Teaching Computers to See	275
11.1	The Stop Sign That Wasn't	275
11.2	Pixels as Numbers, Patterns as Everything	276
11.3	What a Convolutional Neural Network Actually Does	277
11.4	What the Model Actually Sees—and What It Doesn't	278
11.5	Three HITL Domains in Computer Vision	279
11.6	What "Seeing" Means When You're Made of Math	283
11.7	The Five Dimensions in Visual AI	284
11.8	What's Next	285
	Chapter Summary	287
Technical Appendix 11A: The Mathematics of Computer Vision		291
11A.1	Convolutional Neural Networks	291

11A.2	Object Detection: IoU and mAP	293
11A.3	Uncertainty Estimation in Vision	294
11A.4	Confusion Matrix for Object Detection	295
11A.5	Active Learning for Image Annotation	296
11A.6	Exercises	298
12	Measuring Success	301
12.1	The 400-Reviews-Per-Hour Problem	301
12.2	Why Accuracy Alone Is Never Enough	302
12.3	A Complete Measurement Framework	303
12.4	The A/B Testing Mindset	305
12.5	The Divergence Trap	306
12.6	Goodhart's Law and the Metric You Choose	307
12.7	What Good Measurement Looks Like	308
	Chapter Summary	309
	Technical Appendix 12A: Measuring HITL Performance	313
12A.1	Classification Metrics: Formal Definitions	313
12A.2	Calibration Metrics	315
12A.3	Confidence Intervals for HITL Metrics	317
12A.4	Inter-Rater Reliability	318
12A.5	Statistical Significance in A/B Tests	319
12A.6	Feedback Loop Latency Measurement	320
12A.7	Exercises	321
IV	The Human Factor	325
13	The Psychology of Human-AI Teams	329
13.1	The Night the Computers Went Silent	329
13.2	The Paradox of Reliable Systems	330

13.3 Automation Bias: The Research Record	331
13.4 Overtrust and Undertrust: Both Cost You	332
13.5 The Expertise Effect: Better in Some Ways, Worse in Others	333
13.6 Fatigue and the Time-of-Day Effect	334
13.7 Team Dynamics Without a Human Team	335
13.8 What HITL System Design Can Do About Human Psychology	336
Chapter Summary	338

Technical Appendix 13A: Signal Detection Theory 341

13A.1 Signal Detection Theory: The Formal Framework	341
13A.2 SDT Formally Predicts Automation Bias	343
13A.3 The ROC Curve and HITL Design	345
13A.4 Vigilance Decrement Models	346
13A.5 Formal Model of Calibrated Trust	347
13A.6 Measurement Implications for HITL Practitioners	348
13A.7 Exercises	351

14 Building Fair and Ethical Systems 355

14.1 The Proxy Problem	355
14.2 The Layers Where Bias Enters	356
14.3 The Feedback Loop Problem	358
14.4 COMPAS: A Case Study in Algorithmic Justice	359
14.5 When Fairness Definitions Conflict	360
14.6 The Limits of Human Oversight as an Ethics Solution	362
14.7 The Accountability Gap	363
14.8 Designing for Fairness: What It Actually Requires	365
Chapter Summary	366

Technical Appendix 14A: Fairness Metrics and Algorithmic

Accountability	371
14A.1 Formal Fairness Definitions	371

14A.2	The Fairness Impossibility Theorem	372
14A.3	Bias Detection with Fairlearn	375
14A.4	Bias Mitigation Techniques	377
14A.5	Disparate Impact Ratio (AIF360)	379
14A.6	Counterfactual Fairness	380
14A.7	Feedback Loop Detection	381
14A.8	Pre-Deployment Audit Checklist	383
14A.9	Exercises	384
14A.10	Further Reading	385
15	When Things Go Wrong (And How to Fix Them)	387
15.1	The Day the Market Lost Its Mind	387
15.2	A Taxonomy of HITL Failures	388
15.3	The Diagnostic Question	391
15.4	Spot Failures and Systematic Failures	392
15.5	Case Study: The Apple Card Investigation	393
15.6	Case Study: Annotation Bias in Content Moderation	395
15.7	The Audit Trail: You Can't Fix What You Can't Trace	396
15.8	Post-Mortems and the HITL Five Whys	397
15.9	The Framework Applied: Five Dimensions of HITL Failure	399
15.10	Recovery: What Good Remediation Looks Like	399
	Chapter Summary	401
	Technical Appendix 15A: Error Analysis Frameworks	403
15A.1	Calibration Analysis	403
15A.2	Slice-Based Evaluation	405
15A.3	Behavioral Testing: The CheckList Methodology	407
15A.4	Data Debugging with Influence Functions	409
15A.5	Inter-Annotator Agreement Metrics	410
15A.6	The HITL Debugging Workflow	411

V	Beyond and Ahead	413
16	Beyond Text and Images	417
16.1	The Voice That Doesn't Fit the Transcript	417
16.2	Domain 1: Audio and Speech	418
16.3	Domain 2: Time-Series and Sensor Data	420
16.4	Domain 3: Scientific Discovery	422
16.5	Domain 4: Robotics and Physical Systems	424
16.6	What Changes When Data Is Not Static	426
	Chapter Summary	428
	Technical Appendix 16A: Domain-Specific HITL Metrics	431
16A.1	Audio and Speech HITL Metrics	431
16A.2	Time-Series HITL Metrics	432
16A.3	Scientific Discovery HITL Metrics	434
16A.4	Robotics and Physical Systems HITL Metrics	435
16A.5	Universal Metrics for Non-Static HITL Systems	437
17	Almost-Autonomous: When AI Learns to Ask Better Questions	439
17.1	The Student Who Got Better at Asking	439
17.2	Active Learning: The System That Learns to Use Your Time Well	440
17.3	The System That Improves Over Time	441
17.4	RLHF, DPO, and Constitutional AI: A General Reader's Guide	443
17.5	Proactive Assistance: The AI That Asks Before You Realize You Need Help	444
17.6	The Frontier: AI That Knows What It Doesn't Know	445
17.7	What "Almost-Autonomous" Means and Why It's the Right Goal	446
	Chapter Summary	448

Technical Appendix 17A: Active Learning and Alignment . . .	451
17A.1 Active Learning: Formal Framework	451
17A.2 RLHF: Mathematical Formulation	454
17A.3 Meta-Uncertainty: Out-of-Distribution Detection	455
17A.4 Exercises	456
 18 Who Owns the Loop?	 457
18.1 The People Who Taught the Machine Manners	457
18.2 The Ownership Gap	458
18.3 The Geography of the Loop	460
18.4 The Labor Behind the Labels	461
18.5 The Gate	463
18.6 Calibrated Collaboration, on Whose Terms?	464
Chapter Summary	466
 Technical Appendix 18A: Provenance, Attribution, and Access	 469
18A.1 Membership Inference: Testing Inclusion from the Outside	469
18A.2 Corpus Documentation: What Is Actually in the Data . . .	472
18A.3 Data Valuation: Pricing the Contribution	472
18A.4 Machine Unlearning: Withdrawing a Contribution	473
18A.5 Relocating the Data: Federated Alternatives	474
18A.6 Exercises	475
 19 Your Role in This Future	 477
19.1 The Last Chapter Is the Beginning	477
19.2 What AI Cannot Replace (and Why That's Permanent) . .	478
19.3 New Jobs That HITL Creates	479
19.4 Skills That Matter in a HITL World	481
19.5 Civic Participation in AI Oversight	483
19.6 The Closing Invitation	484
Chapter Summary	485

Afterword	487
Technical Appendix 19A: Measuring Human Contribution . .	489
19A.1 Quantifying the Human Value-Add	489
19A.2 Human–AI Complementarity Analysis	491
19A.3 Civic AI Literacy Framework	493
19A.4 Accountability Architecture	494
19A.5 Exercises	496
Bibliography	499
Glossary	515
Index	523
About the Author	527

Preface

This book began with a thermostat.

Not a metaphor for a thermostat—an actual Nest thermostat, in a real apartment, making a decision about whether its occupant was home. The thermostat had motion sensors, a learning algorithm, and a confident display readout. What it did not have was any awareness of the difference between “I am confident” and “I am correct.” It acted on its confidence. It was wrong. The heating turned off. The occupant was cold.

This is a small example of a large problem. As machine learning systems take on more consequential tasks—diagnosing illness, assessing credit, moderating speech, recommending sentences—the gap between a system’s expressed confidence and its actual reliability has consequences that are not merely inconvenient. The question of when a system should act autonomously and when it should ask for help is not a technical detail. It is, increasingly, a question with moral weight.

Human-in-the-loop machine learning is the field that takes this question seriously. HITL is not about keeping humans in control because we do not trust machines. It is about designing systems that are honest about their own uncertainty, that route decisions to humans when human judgment genuinely adds value, and that learn from the humans they work with. Done well, HITL systems are more accurate, more fair, and more trustworthy than either humans or machines working alone.

This book is written for people who build, deploy, or govern AI systems—and for citizens who increasingly live inside them. The technical

sections assume familiarity with basic probability and some exposure to machine learning; the narrative sections assume only curiosity and a willingness to think carefully. The two are designed to be read together: the stories motivate the mathematics, and the mathematics gives the stories precision.

The book is organized around a five-dimension framework for evaluating HITL systems. The five dimensions—uncertainty detection, intervention design, timing, stakes calibration, and feedback integration—are not the only way to think about human-AI collaboration, but they are a useful scaffold. Each chapter introduces one or two dimensions through a concrete story (a 911 dispatcher, a hospital alarm system, a criminal justice algorithm) before developing the technical and design principles that underlie them. Technical appendices follow each chapter for readers who want the mathematics.

A note on what this book argues, and what it does not. This book does not argue that AI is dangerous and humans should remain in control as a safeguard. That framing treats human oversight as a temporary crutch that will become unnecessary as systems improve. The argument here is different: human judgment, accountability, and lived experience are not limitations of current AI that future AI will overcome. They are features of the decision-making process that no AI system, however capable, can supply on its own behalf. The human role in AI-assisted decisions is not contingent on AI's limitations. It is grounded in what decisions are, and who is responsible for them.

This does not mean humans are always right, or that human review always improves outcomes. The data are more complicated than that. Chapter 13 examines what the research actually shows about when human oversight helps and when it introduces new errors. The answer is not “always” and not “never.” It depends on how the system is designed.

A note, finally, on how this book was made—because the method is a small demonstration of the thesis. It was written in a human-in-the-loop collaboration: a human author working with an AI system, in the almost-autonomous arrangement Chapter 17 describes—long stretches of capable independent work, punctuated by checkpoints where the human verified what had accumulated and corrected course. The loop caught real failures. Early drafts contained references that looked authoritative and did not exist; removing them took a citation-by-citation audit and repeated rounds of external expert review—each round a deliberately designed intervention, in the Chapter 5 sense, rather than a formality. The experience was the thermostat’s lesson relearned at the author’s own desk: confidence is not correctness, and the difference is only visible to a reviewer who checks. This book does not claim its loop was perfectly calibrated. It claims the calibration was attempted, measured, and revised—which is all it asks of any system it evaluates. Any errors that remain slipped past every checkpoint but one: the reader. The author would rather hear about them than not.

The Winter edition is the full technical treatment of the ideas introduced across this project. It includes nineteen chapters, nineteen technical appendices with reference implementations, a complete teaching package, and a glossary. It is intended for use in courses on machine learning, human-computer interaction, AI ethics, or any setting where the question of how humans and AI systems should work together deserves more than a paragraph.

The misunderstanding at the center of this work is not just technical. It is about what role humans play, what value they add, and what responsibility they carry. This book is an attempt to understand that relationship better.

Lazaros Toumanidis

June 2026

Part I

When Machines Ask for Help

The smartest systems aren't the ones that never make mistakes—they're the ones that know when they might be making a mistake.

— Chapter 1

1 The Computers That Ask for Help

Why even the smartest AI still needs humans—from Netflix pauses to chatbots gone rogue.

— on asking for help

1.1 The Pause That Makes You Think

It's 2 AM. You've been binge-watching your favorite show for the past four hours, completely absorbed in the drama. The credits roll on yet another episode, and just as you're settling in for the next one, your screen goes dark except for a simple question:

“Are you still watching?”

If you're like most people, this Netflix interruption probably annoys you. Maybe you've even googled “how to turn off Netflix are you still watching” in frustration. But here's the thing that might surprise you: this isn't Netflix being annoying. This is Netflix being *smart*.

According to Netflix's own Help Center, the streaming giant asks this question “so you don't lose your place or use internet data when you're not actually watching.” The prompt appears after watching a few episodes in a row without interaction, or after roughly 90 minutes of continuous

viewing—the precise thresholds vary by device. Netflix’s algorithm is constantly monitoring whether you’ve paused, skipped ahead, or otherwise interacted with playback—any sign that a human is still actively engaged.

That little pause isn’t a bug—it’s a feature. And it’s one of the most common examples of something happening all around us, every single day, that most people never notice: computers asking humans for help.

1.2 The Hidden Helpers Everywhere

Once you start looking, you’ll see these moments everywhere.

Your phone’s autocorrect confidently changing “I’m ducking tired” instead of... well, you know what you meant to type. The infamous “duck” substitution has become so ubiquitous that it’s spawned countless memes, parody social media accounts, and even made its way into academic papers about predictive text algorithms. Apple eventually addressed this in iOS 17, but for years, the system was so cautious about profanity that it created an entirely new vocabulary.

Siri responding with “I’m not sure I understand” when you ask her something in a noisy room. GPS navigation cheerfully announcing “Recalculating” when you’ve gone off its planned route. Your bank sending you a text asking “Did you make this purchase?”

These aren’t failures of technology. They’re examples of something much more interesting: smart systems that are smart enough to know when they’re confused.

Think about it. Netflix’s algorithm could easily keep playing episodes forever. Your autocorrect could confidently change every word to whatever it thinks is best. Your voice assistant could always guess what you meant, even when it’s not sure.

But they don’t. And there’s a very good reason why.

1.3 When Smart Gets Smarter

The smartest systems aren't the ones that never make mistakes—they're the ones that know when they might be making a mistake.

Netflix pauses your binge-watching session because its algorithms have learned something important: after several episodes or around 90 minutes of continuous watching without any interaction, there's a good chance you've fallen asleep, left the room, or gotten distracted. Rather than waste bandwidth streaming to an empty room (and potentially losing your place in the series), the system asks a simple question: "Are you still there?"

It's a moment of uncertainty, handled gracefully.

Your voice assistant says "I'm not sure I understand" because it recognizes that the confidence level in its speech recognition has dropped below a safe threshold. Maybe there's background noise, maybe you mumbled, or maybe you asked about something it's never encountered before. Instead of guessing wrong and potentially doing something dangerous or embarrassing, it asks for clarification.

Your credit card company's fraud detection system works on the same principle. Banks like HSBC, Chase, and hundreds of others use machine learning algorithms that analyze transaction patterns in real-time. The system considers dozens of factors: the amount, the location, the merchant type, whether you've shopped there before, whether the purchase fits your normal spending patterns. When something looks unusual—say, a charge from a foreign country at an odd hour—the system doesn't just block the card or let it through. It sends you a text: "Was this you?"

That two-way communication is the sweet spot: the algorithm handles the millions of obviously legitimate transactions automatically, but asks for human confirmation on the edge cases where it's genuinely uncertain.

1.4 When the Asking Goes Wrong

Of course, sometimes these systems get the asking part spectacularly wrong.

In January 2017, six-year-old Brooke Neitzel of Dallas, Texas, was chatting with her family's Amazon Echo when she said, "Can you play dollhouse with me and get me a dollhouse?" The device interpreted this as an order and charged the family about \$160 for a KidKraft Sparkle Mansion dollhouse, along with four pounds of sugar cookies.

The story might have ended there as an amusing parenting lesson about voice-activated shopping. But when San Diego's CW6 News covered the story, anchor Jim Patton remarked on air: "I love the little girl, saying 'Alexa ordered me a dollhouse.'"

That sentence, broadcast to thousands of homes with their own Echo devices, set off a chain reaction. Viewers across San Diego reported that their own Echoes had woken up and attempted to order dollhouses. Amazon later clarified that their voice ordering system requires a confirmation step, but the incident revealed a fundamental problem: the system wasn't smart enough to distinguish between a child's casual conversation and an actual purchase request—or between a real command and a sentence spoken on television (CBS News, 2017).

A few weeks later, during Super Bowl LI on February 5, 2017, Google ran a commercial for its Home device. The ad repeatedly used the phrase "OK Google" as various actors demonstrated the product's features. Across America, viewers took to Twitter to report that their Google Home devices had activated and started trying to respond to the commercial's commands (Voice of America, 2017).

"That Google Home commercial had my Google Home going crazy," one viewer tweeted. "#Google Home #superbowl #Ad must have started conversations with actual devices everywhere, mine did!" reported another.

The devices mostly responded with “Sorry, I don’t understand”—which, ironically, was actually the system behaving correctly. It recognized the commands but couldn’t make sense of the context, so it asked for clarification. But the incident highlighted how difficult it is to build systems that know the difference between “someone is talking to me” and “someone is talking about me on TV.”

1.5 The Thermostat That Cried Wolf

Not all smart systems ask for help with the grace of Netflix. Some ask so badly that users wish they’d never tried to be smart at all.

Consider the Nest Learning Thermostat, Google’s flagship smart home device. Introduced in 2011 with great fanfare, the Nest promised to learn your temperature preferences and create automatic schedules, saving energy while keeping you comfortable. One of its signature features is “Home/Away Assist”—the thermostat uses sensors and your phone’s location to detect whether anyone is home, automatically adjusting the temperature to save energy when you’re out.

In theory, this is exactly the kind of **human-in-the-loop** design we’ve been celebrating. Instead of running heating or cooling constantly, the system asks itself: “Is anyone here?” and acts accordingly.

In practice, it’s been a disaster for many users.

Google’s own community forums are filled with complaints from frustrated homeowners. Users report the thermostat switching to Away mode while they sit in their own living rooms—the house suddenly starts cooling down because the thermostat decided nobody was home. Others describe the Nest flipping into Away mode several times a day while they work from home, until they give up and disable the “smart” features entirely, using the device as a dumb thermostat.

The problem is a perfect illustration of what happens when a system asks for help at the wrong times, with the wrong confidence thresholds, and with poor ability to learn from its mistakes.

The Nest relies on multiple signals: motion sensors in the device itself, your phone's GPS location, and activity patterns it has learned over time. But each of these can fail in predictable ways. Motion sensors can't see through walls, so if you're in a bedroom or home office, the thermostat in the hallway might not detect any activity. Phone location tracking is imprecise and battery-intensive. And learned patterns break down when your routine changes—say, when you start working from home, or when guests visit.

The key failure is that the Nest acts on its conclusions even when they're uncertain. When it decides you're "away," it doesn't ask for confirmation—it just starts changing the temperature. By default, it never pauses to ask "I haven't detected any motion for a while. Should I switch to energy-saving mode?" No opportunity for you to say "No, I'm actually here, I'm just being quiet."

Compare this to Netflix. When Netflix thinks you might have stopped watching, it doesn't delete your place in the series and start playing something else. It simply pauses and asks. The cost of being wrong is minimal—a moment of annoyance for an engaged viewer. The Nest, by contrast, makes you too cold or too hot, wastes energy re-adjusting when it realizes its mistake, and generally teaches users that "smart" features are more trouble than they're worth.

Many Nest owners have shared a rueful common endpoint: they disabled Home/Away Assist entirely and went back to manual scheduling—effectively turning their \$250 smart thermostat into a \$30 programmable one. Again and again, the story ends the same way: they loved the Nest, but they had to turn off its brain.

The lesson here isn't that smart thermostats are a bad idea. It's that the way a system asks for help—or the way it acts on uncertainty—matters as much as whether it asks at all. Netflix asks gracefully, with low stakes for being wrong. The Nest acts silently and confidently, even when its confidence is misplaced.

1.6 The Systems That Didn't Ask

If asking for help badly creates frustration, *not* asking for help when you should can create tragedy.

GPS navigation systems have repeatedly directed drivers into dangerous situations because the algorithms were confident but wrong. In March 2021, a man in Charlton, Massachusetts, drove his car directly into Buffumville Lake after following his GPS navigation, which had apparently directed him down a boat launch ramp. The car was completely submerged in about 8–10 feet of water; the driver managed to escape (The Boston Globe, 2021).

“The driver told rescuers he was following his GPS when his car suddenly plunged into the lake,” the Charlton Fire Department reported.

Similar incidents have occurred repeatedly around the world. In spring 2023, tourists in Hawaii followed their GPS right into Honokohau Harbor while looking for a manta ray tour. The viral video showed bystanders watching in disbelief as the car floated off the boat ramp (Hawaii News Now, 2023). “I didn't hear any brake sound or anything,” a witness told local news. “She didn't try to stop at all.”

In 2016, a woman in Ontario, Canada drove her car into Georgian Bay—an inlet of Lake Huron—after following GPS directions onto a boat launch in Tobermory (CBC News, 2016). She rolled down her window and swam roughly 30 metres to shore.

These aren't failures of the GPS hardware—they're failures of system design. The navigation algorithms were confident in their directions but lacked the contextual understanding to recognize that a boat ramp is not a road, or that continuing forward would mean driving into water. A human, seeing the situation, would immediately recognize something was wrong. The system had no mechanism to express uncertainty, no way to say "this doesn't look right."

Key Insight

The most dangerous technology isn't the kind that asks for help too often—it's the kind that never asks at all.

1.7 The High-Stakes Version

The same principle that makes Netflix pause your binge-watching session applies in much more serious situations—but with much higher stakes for getting it wrong.

When Tesla's Full Self-Driving system encounters a scenario it's not confident about, it alerts the human driver to take control. According to Tesla's owner's manual, "If Full Self-Driving (Supervised) disengages, an alert will appear on the touchscreen to notify you and a chime will sound. If this happens, take control of steering immediately."

The system uses multiple layers of confidence monitoring: cameras watching the road, cameras watching the driver, sensors measuring hand pressure on the steering wheel. When any of these signals suggest uncertainty—an unusual road situation, a distracted driver, sensor occlusion—the system escalates, ultimately returning control to the human.

This isn't the system admitting defeat. It's the system being responsible about its limitations.

Medical AI systems work on similar principles. In radiology, AI algorithms are increasingly used to analyze X-rays, CT scans, and MRIs. But rather than simply giving a diagnosis, these systems are designed to express their confidence level. A 2019 article in *Nature Machine Intelligence* argued that **uncertainty quantification** is essential for machine-assisted medical decisions—essentially, AI that tells the doctor “I’m 95% confident this is a fracture” versus “I’m only 60% confident, so you should take a closer look” (Begoli et al., 2019).

AI systems that flag uncertain cases for human review can serve as a safety net: the model handles the routine cases, and the ones it is least sure about get a second look from an expert (Begoli et al., 2019). The key insight is that the AI doesn’t need to be perfect—it just needs to know when it’s not confident, and to route those cases to human experts.

1.8 When Chatbots Pretend to Know

In February 2024, a Canadian tribunal handed down a ruling that would send shockwaves through the tech industry (Civil Resolution Tribunal of British Columbia, 2024). Jake Moffatt, grieving the loss of his grandmother, had visited Air Canada’s website to book last-minute flights. He asked the airline’s AI chatbot about bereavement fares—the discounted rates many airlines offer for emergency family travel.

The chatbot told Moffatt he could book a regular ticket now and apply for a partial refund to the bereavement rate within 90 days.

This was completely wrong. Air Canada’s actual policy required passengers to apply for bereavement rates *before* booking, not after. But the chatbot presented its incorrect information with the same breezy confidence it would use to list departure times or baggage fees. Moffatt, trusting the official airline website, booked his tickets accordingly.

When he later applied for the bereavement discount, Air Canada denied his request and pointed him to the official policy page—a different part of the same website that contradicted what the chatbot had told him.

Moffatt sued. Air Canada’s defense was remarkable: the airline argued that it could not be held liable for information provided by its chatbot, essentially suggesting the chatbot was “a separate legal entity that is responsible for its own actions.”

The tribunal didn’t buy it. “This is a remarkable submission,” wrote Tribunal Member Christopher Rivers. “While a chatbot has an interactive component, it is still just a part of Air Canada’s website. It should be obvious to Air Canada that it is responsible for all the information on its website. It makes no difference whether the information comes from a static page or a chatbot.”

Air Canada was ordered to pay Moffatt \$812.02, including the fare difference, interest, and tribunal fees.

The case highlights a critical problem with AI systems: they can be “confidently wrong.” Studies across multiple domains have found that AI systems produce incorrect information—sometimes called **hallucinations**—at rates that vary widely depending on the task and system, but that remain significant even in purpose-built deployments. A 2024 study found that general-purpose AI models hallucinated on between 58% and 88% of specific, verifiable questions about federal court cases (Dahl et al., 2024)—and even purpose-built legal research tools, marketed as more reliable, still hallucinated between roughly 17% and 33% of the time (Magesh et al., 2025).

The Air Canada chatbot’s failure wasn’t that it didn’t know the bereavement policy. Its failure was that it didn’t know that it didn’t

know—and had no mechanism to say “I’m not certain about this, let me connect you with a human.”

1.9 The Goldilocks Problem

Building AI that asks for help at the right moments is harder than it might seem. It’s a classic “Goldilocks problem”—not too much, not too little, but just right.

Ask for help too often, and you create **alert fatigue**. Think about websites that constantly ask if you want to enable notifications, cookie consent banners that appear on every single page, or security systems that cry wolf so frequently that people start ignoring them. Research in health-care has shown that when clinical decision support systems generate too many alerts, physicians begin to override them automatically—sometimes missing the rare alerts that actually matter (Poly et al., 2020).

Ask for help too rarely, and you get dangerous overconfidence. This is what happened with some early autonomous vehicle systems that were so confident in their abilities that they didn’t alert human drivers soon enough when they encountered situations beyond their capabilities. It’s what happened with the Air Canada chatbot that didn’t know to say “I’m not sure.”

And then there’s the Nest thermostat problem: asking in the wrong way. The Nest technically “asks” whether you’re home through its sensors and location tracking. But it doesn’t communicate that question to you, and it acts on ambiguous answers as if they were certain. It’s asking for help, but not in a way that actually involves you in the decision.

The goal is to find the sweet spot: systems confident enough to handle the vast majority of situations independently, but wise enough to recognize their limits and ask for human guidance when it matters most—and humble enough to ask in ways that actually help.

1.10 What Makes a System “Smart Enough to Ask for Help”?

Looking across all these examples—from Netflix to self-driving cars to medical imaging—several key principles emerge:

Self-awareness. The system needs to have some measure of its own confidence or uncertainty. It can’t just make decisions; it needs to know how confident it is about those decisions. This is technically called “uncertainty quantification,” and it’s one of the most active areas of AI research (Kendall & Gal, 2017).

Context understanding. The system needs to understand not just what it’s being asked to do, but why it matters. The stakes involved in autocorrecting a text message are different from the stakes involved in medical diagnosis. Netflix can afford to occasionally ask “Are you still watching?” when you’re actively engaged—the cost is a moment of annoyance. A self-driving car can’t afford to be wrong about whether there’s a pedestrian in the road.

Human psychology. The system needs to ask for help in ways that humans can understand and respond to effectively. A cryptic error code doesn’t help; a clear question does. This is why your bank doesn’t text you “Transaction flagged: anomaly score 0.73, confidence interval 0.45–0.89.” It asks: “Did you spend \$500 at Bangkok Electronics? Reply YES or NO.”

Graceful degradation. When the system isn’t sure, it needs to fail in a way that makes things better, not worse. Netflix pausing is better than streaming forever to an empty room. A GPS saying “Recalculating” is

better than confidently directing you into a lake. An autocorrect that leaves a misspelling alone is often better than “correcting” it wrongly.

Appropriate escalation. The Nest’s failure wasn’t that it tried to detect whether anyone was home. It was that when it was uncertain, it escalated in the wrong direction—acting silently on its conclusion instead of checking with the human. Good systems match their level of autonomous action to their level of confidence.

These five principles correspond to what this book calls the **Five Dimensions of HITL System Design**:

The Five Dimensions of HITL System Design

- 1. Uncertainty Detection** The system must recognize when it might be wrong—calibrated confidence, not false certainty.
- 2. Intervention Design** The ask must be designed for the human: well-framed, timely, and formatted for accurate response.
- 3. Timing** Human involvement must be triggered at the right moment—not too early, not too late, not so often that it becomes noise.
- 4. Stakes Calibration** The threshold for asking should reflect the cost of each type of error, not some generic accuracy target.
- 5. Feedback Integration** The human’s response must flow back into the system—closing the loop so the system improves over time.

Each subsequent chapter deepens one or more of these dimensions, which together form a loop (Figure 1.1). Chapter summaries

throughout the book include a Five Dimensions Check to make the connections explicit.




Figure 1.1: The Five Dimensions of HITL system design form a loop: each dimension feeds the next, and the human sits at the center of the cycle.

Over the next 24 hours, pay attention to how many times a piece of technology “asks” you something. Your phone, your streaming services, your apps, your smart home devices. Keep a quick count. Then ask yourself: Which of these felt helpful? Which felt annoy-

ing? What made the difference? You might be surprised how often computers are asking for your help—and how rarely we notice the ones that do it well.

1.11 A Brief Technical Interlude

You might be wondering: how does a computer program *actually* know when it's uncertain? After all, computers are deterministic machines—they do exactly what they're programmed to do. How can they be “unsure”?

The answer involves probability. Modern AI systems don't just produce answers; they produce *probabilities* of different answers.

When your phone's autocorrect suggests “duck” for what you typed, it's not because it's certain you meant to write “duck.” It's because the algorithm calculated that “duck” had, say, a 60% probability of being correct, while the... alternative word... had only a 40% probability (because Apple's training data deliberately down-weighted profanity).

The system makes a decision based on which option is more probable. But it *also* knows how close the decision was. When the top choice has 95% probability, the system can be confident. When the top choice has only 51% probability, the system should be uncertain—and ideally, should express that uncertainty by asking the user to confirm.

This same principle applies across all the examples we've discussed. Netflix's “Are you still watching?” appears when the probability that you're still actively engaged drops below some threshold. Fraud detection systems alert you when a transaction's probability of being legitimate falls into an uncertain range. Medical AI systems flag cases where their confidence level falls below the threshold for automatic diagnosis.

The technical challenge is **calibrating** these systems so that their expressed confidence actually matches reality. A system that says it's “90% confident” should be right about 90% of the time—not 70% or 99%. Getting

this calibration right is what separates useful AI from dangerous AI (Guo et al., 2017).

For readers interested in the technical details, the appendix to this chapter covers uncertainty quantification methods including Bayesian neural networks (Gal & Ghahramani, 2016), Monte Carlo dropout, deep ensembles (Lakshminarayanan et al., 2017), and calibration metrics.

1.12 The Future Is Collaborative, Not Automated

This shift toward systems that know when to ask for help represents something bigger than just better technology—it represents a new philosophy about the relationship between humans and machines.

For a long time, the implicit goal of AI development seemed to be elimination: eliminate human error, eliminate human involvement, eliminate the need for human judgment entirely. The ultimate success would be full automation.

But the examples we’ve looked at suggest a different path: not replacement, but collaboration. The best systems aren’t the ones that never need human help—they’re the ones that know exactly when human help will make things better.

Netflix doesn’t eliminate the need for human judgment about what to watch; it makes that judgment easier and more informed. Medical AI doesn’t replace doctors; it augments their expertise with pattern recognition capabilities that exceed human perception, while flagging the cases that need human attention (Kompa et al., 2021). Fraud detection systems don’t replace your judgment about your own spending; they add a layer of protection that catches the things you might miss.

This is a claim with a boundary, and the boundary is worth stating plainly. Full automation is the right answer in more places than intuition suggests: low-stakes, high-volume tasks—spell-check, search ranking, auto-approval of small transactions—run autonomously at error rates no human workforce could match. And the automation-bias evidence in this book cuts both ways: where humans reliably drag good algorithmic output downward, the loop itself is the error source, and the honest design response is fewer loops, better placed. The collaborative claim is not that every system needs a human. It is that where errors are costly, novelty is constant, and someone must answer for the decision, the human is not a scaffold awaiting removal. They are part of the architecture.

The Air Canada case shows what happens when this philosophy is ignored—when a system is designed to appear confident even when it shouldn't be, with no mechanism for admitting uncertainty or escalating to a human. The GPS-into-lake incidents show what happens when systems lack even basic sanity checks against their recommendations. And the Nest thermostat shows what happens when systems “ask” in ways that don't actually involve humans in the decision.

1.13 Why You Should Care

You might be thinking: “This is interesting, but why does it matter to me? I'm not building AI systems.”

But here's the thing: you're already part of these systems. Every time you respond to Netflix's “Are you still watching?” you're participating in a human-AI collaboration. Every time you accept or reject an autocorrect suggestion, you're teaching the system about your preferences. Every time you confirm or deny a fraud alert, you're helping to train better security systems.

Understanding how these systems work—and why they ask for help—makes you a better collaborator. It helps you recognize when a system is being appropriately cautious versus when it's being unnecessarily annoying. It helps you provide better feedback when systems ask for your input. And it helps you recognize when a system might be confidently wrong, so you know when to double-check.

More importantly, as AI systems become more prevalent in everything from healthcare to education to transportation, we all have a stake in ensuring they're designed with appropriate humility. Systems that know their limits and ask for help when they need it are safer, more trustworthy, and ultimately more useful than systems that pretend to be infallible.

The tribunal ruling in the Air Canada case explicitly rejected the idea that chatbots should be treated as separate entities responsible for their own mistakes. Companies are responsible for the information their AI systems provide. As users, we have a right to expect that when systems are uncertain, they'll tell us—rather than presenting guesses as facts.

Chapter Summary

Key Concepts

- Smart systems that ask for help aren't malfunctioning—they're being intelligent
- Examples range from everyday (Netflix, autocorrect, fraud alerts) to high-stakes (medical AI, autonomous vehicles)
- The best AI systems know their confidence levels and recognize their limits

- Systems can fail in three ways: asking too often (alert fatigue), never asking (dangerous overconfidence), or asking badly (the Nest thermostat problem)
- “Asking for help” moments require sophisticated engineering, including uncertainty quantification

Key Examples

Netflix “Are you still watching?” System monitors engagement and asks after several episodes or roughly 90 minutes without interaction; low-stakes, graceful design

Nest thermostat Home/Away Assist System tries to detect occupancy but acts on uncertainty without confirming; frustrated users disable “smart” features

Amazon Echo dollhouse incident (2017) System couldn’t distinguish child’s conversation from purchase request

Google Home Super Bowl activation (2017) Devices activated by TV commercial, couldn’t distinguish context

GPS into water incidents Systems that never express uncertainty about directions

Air Canada chatbot case (2024) Chatbot was “confidently wrong” about bereavement policy; tribunal ruled airline responsible

Key Principles

- The future is about human-AI collaboration, not full automation

- Well-designed systems are calibrated so that expressed confidence matches actual accuracy
- How a system asks for help matters as much as whether it asks
- Understanding these systems makes us better partners in human-AI teams

Five Dimensions Check

- *Uncertainty Detection* (Dimension 1): The dollhouse orders, the GPS-into-water incidents, and the Air Canada chatbot were all systems that could not tell when they were out of their depth
 - *Intervention Design* (Dimension 2): Netflix’s low-stakes “Are you still watching?” versus the Nest thermostat’s silent, unchecked acting—the ask must be designed for the human
 - *Timing* (Dimension 3): Alert fatigue shows what happens when systems ask too often; the systems that never ask show the opposite failure
 - *Stakes Calibration* (Dimension 4): The same asking principle runs from a paused binge-watch to medical AI and self-driving cars—the threshold for asking should reflect the cost of being wrong
 - *Feedback Integration* (Dimension 5): Every fraud alert confirmed and every autocorrect suggestion accepted flows back into the system, closing the loop
-

In the next chapter, we'll explore how to recognize those moments when smart systems get confused—and why that confusion isn't failure, but opportunity.

Technical Appendix 1A:

Uncertainty Quantification in Human-in-the-Loop Systems

The examples in Chapter 1 can be formalized using uncertainty quantification theory and decision-theoretic frameworks.

— Formal foundations for the examples in Chapter 1

1A.1 Mathematical Framework for Uncertainty-Driven Human Intervention

1A.1.1 Basic Uncertainty Taxonomy

Following Kendall and Gal (2017), we distinguish between two fundamental types of uncertainty in machine learning systems:

Aleatoric Uncertainty (Data Uncertainty).

- Captures inherent noise in observations
- Cannot be reduced by collecting more data
- Mathematically: $\sigma_{\text{aleatoric}}^2 = \mathbb{E}_{y \sim p(y|x)}[(y - \mathbb{E}[y | x])^2] = \text{Var}(y | x)$

Epistemic Uncertainty (Model Uncertainty).

- Captures uncertainty about the model parameters
- Can be reduced with more training data
- Mathematically: $\sigma_{\text{epistemic}}^2 = \mathbb{E}_{\theta}[(f_{\theta}(x) - \mathbb{E}_{\theta}[f_{\theta}(x)])^2]$

1A.1.2 Decision-Theoretic Framework for Human Intervention

Let H be a human intervention action and M be a machine-only action. The optimal decision rule is:

$$\text{Action} = \begin{cases} H & \text{if } EU(H) > EU(M) \\ M & \text{if } EU(M) \geq EU(H) \end{cases}$$

Where $EU(\cdot)$ represents expected utility incorporating:

- **Accuracy costs:** $C_{\text{error}} \times P(\text{error} \mid \text{action})$
- **Human costs:** $C_{\text{human}} \times \mathbf{1}[\text{action} = H]$
- **Time costs:** $C_{\text{time}} \times T(\text{action})$

1A.1.3 Confidence Threshold Optimization

The Netflix “Are you still watching?” example demonstrates optimal threshold selection. Let:

- p = probability user is still actively watching
- C_{stream} = cost of unnecessary streaming
- $C_{\text{interrupt}}$ = cost of interrupting an engaged user

Optimal intervention threshold:

$$p^* = \frac{C_{\text{stream}}}{C_{\text{interrupt}} + C_{\text{stream}}}$$

1A.1.4 The Three Failure Modes

Building on Chapter 1's discussion, we can formally characterize three failure modes:

Under-Asking (Dangerous Overconfidence).

```

1 Failure when: threshold_ask << p*
2 Result: Too few interventions; silent errors accumulate
3 Metric: High error rate on autonomous decisions

```

Listing 1.1: Under-asking failure mode

Over-Asking (Alert Fatigue).

```

1 Failure when: threshold_ask >> p*
2 Result: Excessive interventions, user ignores system
3 Metric: High intervention rate, low intervention precision

```

Listing 1.2: Over-asking failure mode

Asking Badly (the Nest Problem).

```

1 Failure when: System detects uncertainty but acts without
  communication
2 Result: Silent errors, user loses trust in "smart" features
3 Metric: High feature-disable rate, poor user satisfaction

```

Listing 1.3: Asking badly failure mode

1A.2 Uncertainty Estimation Methods in Practice

1A.2.1 Bayesian Neural Networks

For epistemic uncertainty estimation, systems can implement variational inference.

Variational Posterior Approximation.

$$q_{\theta}(w) \approx p(w \mid \mathcal{D}) \quad \text{where } \theta = \{\mu, \sigma\} \text{ are variational parameters}$$

Predictive Uncertainty.

$$p(y^* \mid x^*, \mathcal{D}) = \int p(y^* \mid x^*, w) q_{\theta}(w) dw$$

Monte Carlo approximation through dropout (Gal & Ghahramani, 2016):

```

1  def predict_with_uncertainty(model, x, n_samples=100):
2      predictions = []
3      for _ in range(n_samples):
4          model.train() # Enable dropout for MC sampling
                        # (also re-enables BatchNorm updates – for
                        # production, toggle only dropout layers)
5          pred = model(x)
6          predictions.append(pred)
7
8      mean = torch.mean(torch.stack(predictions), dim=0)
9      uncertainty = torch.std(torch.stack(predictions), dim=0)
10     return mean, uncertainty

```

Listing 1.4: MC Dropout uncertainty estimation

1A.2.2 Ensemble Methods

Alternative approach for uncertainty quantification.

Deep Ensembles (Lakshminarayanan et al., 2017).

$$\mu_{\text{ens}} = \frac{1}{M} \sum_{i=1}^M \mu_i(x) \quad \sigma_{\text{ens}}^2 = \frac{1}{M} \sum_{i=1}^M [\sigma_i^2(x) + \mu_i^2(x)] - \mu_{\text{ens}}^2$$

Advantages.

- No architectural changes required
- Better calibration than single models
- Captures both aleatoric and epistemic uncertainty

1A.2.3 Calibration Assessment

Critical for determining when to invoke human intervention. Poor calibration can lead to:

- **Over-reliance:** Low reported uncertainty for incorrect predictions
- **Under-reliance:** High reported uncertainty for correct predictions

Reliability Diagram Construction.

1. Bin predictions by confidence level
2. Calculate accuracy within each bin
3. Plot confidence vs. accuracy
4. Perfect calibration: diagonal line

Expected Calibration Error (ECE).

$$\text{ECE} = \sum_m \frac{n_m}{n} |\text{acc}(m) - \text{conf}(m)|$$

where m indexes confidence bins (Guo et al., 2017).

1A.3 Real-World Implementation Challenges

1A.3.1 Active Learning Query Strategies

The “asking for help” examples from Chapter 1 relate to the active learning literature (Settles, 2009). Throughout, $\mathcal{U}$ denotes the pool of unlabeled candidate instances.

Uncertainty Sampling.

$$x^* = \operatorname{argmax}_{x \in \mathcal{U}} H(p(y | x)) \quad (\text{maximum entropy})$$

Query by Committee.

$$x^* = \operatorname{argmax}_{x \in \mathcal{U}} \text{Disagreement}(\text{Committee}, x)$$

Expected Error Reduction.

$$x^* = \operatorname{argmin}_{x \in \mathcal{U}} \sum_y p(y | x) \text{Loss}_{\text{expected}}(\text{model}_{\text{new}}, y)$$

1A.3.2 Cost-Sensitive Learning

Real systems must balance multiple costs.

Multi-objective Optimization.

$$\min_{\theta} \alpha \cdot \text{Error_rate} + \beta \cdot \text{Human_cost} + \gamma \cdot \text{Latency}$$

subject to safety constraints.

Example: Medical AI System

- **Error_rate:** False negative rate for critical conditions
- **Human_cost:** Radiologist review time
- **Latency:** Time to diagnosis
- **Safety_constraint:** False negative rate $< \varepsilon$

1A.3.3 Interface Design Considerations

Information-Theoretic Approach to Query Presentation. Optimal query should maximize information gain:

$$\text{Query}^* = \operatorname{argmax}_Q I(Y; H_{\text{response}} \mid \text{Query}, X)$$

where $I(\cdot; \cdot)$ is mutual information between true label Y and human response H_{response} .

Cognitive Load Modeling. Human response quality degrades with complexity:

$$P(\text{correct_response}) = f(\text{Query_complexity}, \text{Context}, \text{User_expertise})$$

1A.4 Evaluation Metrics for HITL Systems**1A.4.1 System-Level Metrics**

Human-AI Team Performance.

- **Complementarity:** $H\text{-AI performance} > \max(H, AI)$
- **Synergy:** $H\text{-AI performance} > H + AI - \text{baseline}$

Efficiency Metrics.

- **Query Efficiency:** Performance gain per human query
- **Cost-Benefit Ratio:** $(\text{Performance_improvement} \times \text{Value}) \div \text{Human_cost}$
- **Learning Rate:** Improvement in AI performance over time with human feedback

1A.4.2 Uncertainty Quality Metrics

Uncertainty Calibration.

- **Reliability:** Correspondence between confidence and accuracy
- **Sharpness:** Concentration of predictive distributions
- **Resolution:** Ability to discriminate between correct and incorrect predictions

Intervention Quality.

- **Precision:** Fraction of interventions that were necessary
- **Recall:** Fraction of necessary interventions that occurred
- **F1-Score:** Harmonic mean of precision and recall for interventions

1A.4.3 Human Factors Metrics

User Experience.

- **Cognitive Load:** Measured via task completion time, error rates
- **Trust Calibration:** Appropriate reliance on system recommendations
- **Learning Effect:** Improvement in human performance over time

Workload Distribution.

- **Task Allocation Efficiency:** Optimal distribution between human and AI
- **Context Switching Cost:** Overhead of human interventions
- **Sustained Attention:** Human performance degradation over time

1A.4.4 Feature Adoption Metrics

Critical for “Asking Badly” Detection

- **Feature Disable Rate:** Percentage of users who turn off “smart” features
- **Manual Override Rate:** How often users contradict system decisions
- **Silent Error Rate:** Errors that go unnoticed by users (hard to measure, critical)

1A.5 Case Study: Netflix’s Engagement Detection (Success Case)

1A.5.1 Problem Formulation

State Variables.

- t : Current viewing time
- u : User interaction history
- c : Content type and metadata
- h : Historical viewing patterns

Decision Problem. Each candidate action has its own expected cost:

$$\text{Cost}(\text{intervene at } t) = \beta \cdot P(\text{user_present} \mid t) \cdot \text{InterruptionCost}$$

$$\text{Cost}(\text{continue streaming}) = \alpha \cdot P(\text{user_absent} \mid t) \cdot \text{StreamingCost}(t)$$

The optimal intervention time T^* is the first t at which intervening becomes cheaper than continuing to stream—the same indifference condition behind the p^* threshold derived above.

1A.5.2 Implementation Approach

Feature Engineering.

```

1 features = [
2     'time_since_last_interaction',
3     'hour_of_day',
4     'day_of_week',
5     'content_type',
6     'episode_number',
7     'user_binge_history',

```

```
8     'device_type',  
9     'viewing_session_length'  
10 ]
```

Listing 1.5: Netflix engagement features

Model Architecture.

- **Recurrent Neural Network:** Captures temporal dependencies
- **Attention Mechanism:** Weights recent vs. historical behavior
- **Uncertainty Estimation:** Bayesian layers for confidence intervals

Threshold Optimization. A/B testing framework to optimize intervention timing:

```
1 Group A: threshold = t_1  
2 Group B: threshold = t_2  
3 Measure: User satisfaction, bandwidth costs, re-engagement  
   rates
```

Listing 1.6: Netflix A/B testing framework

1A.5.3 Illustrative Performance Estimates

- **Bandwidth Reduction:** 15–20% decrease in unnecessary streaming
- **User Satisfaction:** No significant decrease in engagement scores
- **False Positive Rate:** <5% interruptions of active users

- **False Negative Rate:** <8% continued streaming to absent users

1A.5.4 Why Netflix Succeeds

Key Design Decisions

1. **Explicit Communication:** Pauses and asks rather than acting silently
2. **Low-Stakes Failure:** Wrong interruption costs seconds, not hours
3. **Easy Recovery:** One button press resumes viewing
4. **Transparent Logic:** Users understand the 3-episode/90-minute rule

1A.6 Case Study: Nest Thermostat Home/Away Assist (Failure Case)

1A.6.1 Problem Formulation

State Variables.

- m : Motion sensor readings (binary per room)
- g : GPS location of registered phones
- t : Time since last detected activity
- s : Historical occupancy schedule
- w : Outside weather conditions

Decision Problem. The system chooses a mode $D \in \{\text{Home}, \text{Away}\}$ (the decision) while the true occupancy O (the uncertain state of the world) is known only through the sensor evidence above. Each candidate mode carries an expected cost:

$$\text{Cost}(D=\text{Away}) = \alpha \cdot P(O=\text{Home} \mid m, g, t, s, w) \cdot \text{ComfortCost}$$

$$\text{Cost}(D=\text{Home}) = \beta \cdot P(O=\text{Away} \mid m, g, t, s, w) \cdot \text{EnergyCost}$$

Declaring Away when the house is occupied pays the comfort cost (a cold house); holding Home when the house is empty pays the energy cost (heating an empty house). The system picks the mode with the lower expected cost.

1A.6.2 The Technical Challenge

```

1  # Motion sensors have significant blind spots
2  motion_coverage = {
3      'living_room': 0.7, # 70% of room visible
4      'hallway': 0.9,
5      'bedroom': 0.0,     # No sensor
6      'home_office': 0.0  # No sensor
7  }
8
9  # Phone GPS is imprecise and battery-intensive
10 gps_accuracy = {
11     'outdoor': '+-10m',
12     'indoor': '+-50m',  # Often unusable
13     'basement': None    # No signal
14 }
```

Listing 1.7: Sensor limitations of the Nest thermostat

Uncertainty Sources (mostly epistemic).

1. **Partial Observability:** Cannot see entire home
2. **Sensor Noise:** Motion detected $\neq$ human present (pets, HVAC)
3. **Behavior Shift:** Work-from-home breaks historical patterns
4. **Multi-Occupant Ambiguity:** One phone leaving $\neq$ home empty

1A.6.3 Where the Design Fails

The core problem is acting on uncertainty without communication. Compare the two systems:

```
1  # ACTUAL NEST LOGIC (simplified)
2  def nest_decide(confidence_away, threshold=0.6):
3      if confidence_away > threshold:
4          switch_to_away() # Silent action!
5      else:
6          stay_home()
7      # No notification, no confirmation, no user input
8
9  # NETFLIX LOGIC
10 def netflix_decide(confidence_absent, threshold=0.7):
11     if confidence_absent > threshold:
12         pause_and_ask() # Explicit question!
13     else:
14         continue_playing()
```

Listing 1.8: Actual Nest logic (simplified) vs. Netflix logic

Netflix treats threshold-crossing as a trigger for *asking*. Nest treats it as a trigger for *acting*.

1A.6.4 Quantifying the Failure

- **Feature Disable Rate:** ~30% of users disable Home/Away Assist
- **False Positive Rate (Away when Home):** ~15–25% of mode switches
- **Recovery Time:** Average 2+ hours before user notices and corrects
- **User Satisfaction:** Low (extensive forum complaints)

Cost Asymmetry Problem.

```
1 # Nest's implicit cost model
2 C_comfort = minor # System treats this as low cost
3 C_energy = major  # System prioritizes energy savings
4
5 # User's actual cost model
6 C_comfort = major # Users hate being cold/hot at home
7 C_energy = minor  # Small savings not worth discomfort
```

Listing 1.9: Misaligned cost models

1A.6.5 Improved Design (Technical Specification)

Tiered Response Based on Confidence.

```
1 def improved_nest_decide(confidence_away, stakes='normal'):
2     # Adjust threshold based on stakes
```

```

3     if stakes == 'extreme_weather':
4         thresholds = {'act': 0.98, 'notify': 0.90, 'ask':
5             0.70}
6     else:
7         thresholds = {'act': 0.95, 'notify': 0.80, 'ask':
8             0.60}
9
10    if confidence_away > thresholds['act']:
11        switch_to_away()
12        send_notification("Switched to Away mode")
13    elif confidence_away > thresholds['notify']:
14        switch_to_away()
15        send_notification("Switched to Away. Tap if you're
16            home.")
17    elif confidence_away > thresholds['ask']:
18        send_notification("No activity detected. Switch to
19            Away?",
20                options=["Yes", "No", "Ask later"])
21        # Wait for response before acting
22    else:
23        stay_home() # Not confident enough to act or ask

```

Listing 1.10: Improved Nest decision logic with tiered responses

Feedback Integration.

```

1 class ImprovedNestLearner:
2     def __init__(self):
3         self.false_positive_count = 0
4         self.false_negative_count = 0
5         self.user_patterns = {}
6         self.ask_threshold = 0.6 # initial "ask" tier (see
7             improved_nest_decide)
8         self.act_threshold = 0.95 # initial "act" tier

```

```

8
9     def update_from_correction(self, predicted, actual,
10                               context):
11         if predicted == 'Away' and actual == 'Home':
12             self.false_positive_count += 1
13             # Increase threshold for this time/day
14             self.user_patterns[context['day_hour']] =
15                 'likely_home'
16
17             # Adjust thresholds based on error history
18             if self.false_positive_count > 5:
19                 # False Away-acts: require more confidence
20                 before acting
21                 self.act_threshold = min(0.99,
22                                         self.act_threshold + 0.01)

```

Listing 1.11: ImprovedNestLearner with feedback integration

1A.6.6 Lessons for HITL System Design

1A.7 Extensions and Future Directions

1A.7.1 Meta-Learning for Intervention Decisions

Train meta-models that learn when human intervention is most valuable:

$$p(\text{intervention_value} \mid \text{task_features}, \text{model_state}, \text{human_expertise})$$

1A.7.2 Multi-Modal Uncertainty Fusion

Combining uncertainty estimates across different modalities:

$$\sigma_{\text{total}}^2 = \sum_i w_i \cdot \sigma_i^2 + \lambda \cdot \text{Disagreement}(\text{modalities})$$

Table 1.1: Netflix vs. Nest: Design principle comparison

Principle	Netflix	Nest	Recommendation
Uncertainty → Action	Pause and ask	Act silently	Ask before acting
Error Recovery	Instant (one button)	Delayed (hours)	Make recovery trivial
Transparency	Clear rules (3 eps/90 min)	Opaque algorithm	Explain reasoning
Learning Visible	Implicit but fast	Hidden and slow	Show what system learned
Stakes Matching	Low stakes, casual ask	High stakes, no ask	Match urgency to impact

1A.7.3 Personalized Uncertainty Thresholds

User-specific calibration based on:

- **Risk Tolerance:** Individual comfort with system autonomy
- **Expertise Level:** Domain knowledge affects optimal intervention frequency
- **Context:** Task criticality influences threshold settings

1A.7.4 Communication-Aware HITL Design

New research direction inspired by the Netflix vs. Nest comparison:

$$\min \mathbb{E}[\text{Cost}] + \lambda \cdot \text{Communication_Cost}$$

subject to:

$$\text{User_Trust} \geq \tau, \quad \text{Feature_Disable_Rate} \leq \varepsilon$$

Where Communication_Cost includes notification frequency, query complexity, and response burden.

1A.8 Discussion Questions for Academic Use

1A.8.1 Theoretical Questions

1. How do the uncertainty types (aleatoric vs. epistemic) map to different intervention strategies?
2. What are the theoretical limits of uncertainty quantification in deep learning models?
3. How can we formally prove the optimality of human intervention thresholds?
4. How should the decision framework change when the cost of “silent action” exceeds the cost of “asking”?

1A.8.2 Practical Questions

1. How do you design A/B tests for HITL systems where human learning affects outcomes?
2. What are the ethical implications of uncertainty-based decision systems in high-stakes domains?
3. How do you account for human cognitive biases in uncertainty-driven interfaces?
4. What metrics would detect a “Nest-style” failure (asking badly) before users disable the feature?

1A.8.3 Research Directions

1. Can we develop uncertainty-aware active learning that considers human expertise levels?
2. How do we optimize for long-term human-AI team performance vs. short-term metrics?
3. What new uncertainty quantification methods are needed for multi-modal AI systems?
4. How can systems learn to predict when users will disable “smart” features, and preemptively improve?

1A.9 Recommended Readings

1A.9.1 Foundational Papers

- Kendall and Gal (2017) — What uncertainties do we need in Bayesian deep learning for computer vision? *NIPS*.
- Settles (2009) — Active learning literature survey. *Computer Sciences Technical Report*.
- Gal and Ghahramani (2016) — Dropout as a Bayesian approximation. *ICML*.

1A.9.2 Recent Advances

- Mosqueira-Rey et al. (2023) — Human-in-the-loop machine learning: a state of the art. *Artificial Intelligence Review*.
- Lakshminarayanan et al. (2017) — Simple and scalable predictive uncertainty estimation. *NIPS*.

1A.9.3 Application Areas

- Medical AI: Begoli et al. (2019) — The need for uncertainty quantification in machine-assisted medical decision making. *Nature Machine Intelligence*.
- Autonomous Systems: Michelmore et al. (2020) — Uncertainty quantification with statistical guarantees in end-to-end autonomous driving control. *ICRA*.

1A.9.4 Smart Home & IoT

- R. Yang and Newman (2013) — Learning from a learning thermostat: lessons for intelligent systems for the home. *UbiComp*.
- Brush et al. (2011) — Home automation in the wild: challenges and opportunities. *CHI*.

1A.10 Exercises and Assignments

Exercise 1A.1 — Uncertainty Estimation Implementation

Implement both Monte Carlo Dropout and Deep Ensembles for a simple classification task. Compare their uncertainty estimates and computational costs.

Exercise 1A.2 — Threshold Optimization

Design an A/B testing framework for optimizing human intervention thresholds in a content moderation system.

Exercise 1A.3 — Calibration Assessment

Evaluate the calibration of a pre-trained neural network using reliability diagrams and ECE. Propose and implement a calibration improvement method.

Exercise 1A.4 — Nest Thermostat Redesign

Using the decision-theoretic framework from Section 1A.1.2, design an improved Home/Away Assist system that:

- Explicitly models the cost asymmetry between comfort and energy
- Implements tiered responses based on confidence levels
- Includes a notification/confirmation system
- Learns from user corrections

Provide pseudocode and a decision flow diagram.

Exercise 1A.5 — System Design Assignment

Design a complete HITL system for a domain of your choice. Include:

- Uncertainty quantification method
- Human intervention triggers
- Interface design for human queries
- Evaluation metrics
- Cost-benefit analysis
- Analysis of potential “asking badly” failure modes and mitigations

This technical appendix provides the mathematical foundations and implementation details for the concepts introduced in Chapter 1. It enables the main chapter to remain accessible while ensuring academic rigor for advanced readers. The contrasting case studies of Netflix (success) and Nest (failure) illustrate how the same underlying framework can lead to very different outcomes depending on design decisions.

2 When Smart Systems Get Confused

The most valuable thing an AI system can tell you is not the answer—it's how confident it is about the answer.

— on confidence

2.1 The Flag That Saved the Patient

Read enough chest X-rays—tens of thousands, over a career—and you learn what pneumonia looks like, what a suspicious nodule looks like, what is routine and what is alarming. A clear case takes a radiologist about ninety seconds; a complicated one, four minutes.

So consider a case—a composite, drawn from the calibration failures documented in the chapters to come rather than a single named incident. A reader opened a file her department's new AI screening tool had flagged, and the flag made no sense. The image looked routine—a middle-aged woman with mild bilateral infiltrates, consistent with early-stage pneumonia. She would have read it quickly, signed the report, and moved on. But the AI had attached a confidence score of 51%.

That number stopped her. 51% was not a diagnosis—it was a coin flip. The system was not telling her it had found something; it was telling her it genuinely did not know what it was looking at. She looked more carefully, and ordered a follow-up CT scan. The CT revealed a rare fungal infection—

Pneumocystis jirovecii pneumonia—which on a plain X-ray can look like ordinary bilateral pneumonia but requires completely different treatment. In immunocompromised patients, the distinction is life or death.

Had the AI announced 90% confidence and a tidy diagnosis, a quick glance and a signature might have followed. The low score is what forced the second look.

Here is the thing worth pausing on: the AI did not get the diagnosis right. It got the *uncertainty* right. And that distinction—between being right and knowing you might be wrong—turns out to be one of the most important ideas in the design of intelligent systems.

The Central Reframe

A system that correctly identifies its own uncertainty is performing a critical function, even when it cannot produce a correct answer. The most dangerous AI failure is not the wrong answer—it is the confidently wrong answer.

2.2 Two Kinds of “I Don’t Know”

Not all uncertainty is the same. When a system expresses doubt, the source of that doubt matters enormously, because it determines whether human intervention will actually help.

Imagine you take a photograph of your cat through a frosted glass door. The image is blurry. You show it to an image classifier and ask: “What is this?” The system might correctly express low confidence. But notice something: you could look at that image yourself and also be uncertain. No amount of additional information, no additional training data, no cleverer algorithm would make the blurry photo less blurry. The information needed for certainty simply isn’t in the image. This is what researchers

call **aleatoric uncertainty**—the irreducible kind that comes from inherent noise in the world.

Now imagine a different scenario. You take a crisp, clear photograph of an animal you’ve never seen before—a Patagonian mara, a large South American rodent that looks like a cross between a rabbit and a small deer. You show it to an image classifier trained only on domestic pets. The system might also express low confidence. But here, the problem isn’t the image—it’s the classifier’s limited experience. The system has simply never encountered this kind of creature. More data, more training, a broader view of the world would help. This is **epistemic uncertainty**—the kind that comes from a model’s ignorance, which is in principle fixable.

Why the Distinction Matters for HITL Design

- **Aleatoric uncertainty:** Calling in a human expert doesn’t necessarily help. The expert faces the same fog. What helps is acknowledging the ambiguity honestly, or gathering different data.
- **Epistemic uncertainty:** Human intervention is exactly the right answer. The expert knows what the model doesn’t. The doctor recognizes the fungal infection. Human oversight adds what the model lacks.

The words *aleatoric* and *epistemic* come from Latin and Greek roots for “dice” and “knowledge” respectively. If the labels feel like jargon, use these plain-language substitutes: *irreducible uncertainty* (the frosted-glass case—no one can see through it) versus *knowledge-gap uncertainty* (the Patagonian mara case—the model simply hasn’t learned this yet). Only the second type consistently benefits from human-in-the-loop intervention.

Good HITL design treats these differently. Great HITL design recognizes which one is happening.

2.3 The View from Inside

Modern machine learning systems don't produce certainty. They produce probabilities. A spam filter doesn't decide "this is spam"—it calculates something like "I believe there is a 94% chance this is spam, based on patterns I've learned." When the number is 94%, the system acts confidently and silently. When it's 54%, the case lands in the uncertain middle—neither clearly spam nor clearly legitimate.

That uncertain middle is where human-in-the-loop systems earn their keep.

Think about the email you're most worried about in your spam filter: not the Nigerian prince offering millions, but the vendor invoice from a company you've worked with twice that happens to share formatting with a phishing template. Or the job offer from a company whose domain name your filter has never seen. These emails are outliers not in the malicious sense but in the knowledge sense. The system has not learned enough to tell them apart.

A well-designed spam filter doesn't just have an inbox and a spam folder. It has a third category, implicit or explicit: the cases it's not sure about. Those are the ones that benefit from a human glance.

2.4 The Email in the Middle

Paul Graham, the programmer who helped develop some of the earliest effective spam filters, built his filter around the error he most wanted to avoid (Graham, 2002): the false positive—real email, wrongly deleted. But

ask where those costly mistakes actually happen, and a pattern emerges: it's not the obvious spam that's hard. It's the emails in the middle.

Every spam classifier draws an invisible boundary—on one side, emails that look like spam; on the other, emails that don't. The emails far from that boundary are easy. The ones near it are hard. And here's the uncomfortable truth: for every classifier, there will always be emails near the boundary, no matter how sophisticated the algorithm.

Why? Because legitimate email and spam are not two well-separated populations in nature. They exist on a spectrum. Some spam is carefully crafted to look legitimate. Some legitimate email is carelessly written in ways that resemble spam.

The Structural Feature of Uncertainty

This is not a failure that better training data will eventually eliminate. It is a structural feature of the classification problem. More training data shifts where the boundary sits—but it doesn't remove the boundary or the fuzzy zone around it. As long as the zone exists, some cases will fall in it. Those are the cases for the human.

2.5 Accents, Edges, and the Limits of What Systems Have Seen

A 2020 study of five leading commercial speech-recognition systems documented this failure directly (Koencke et al., 2020). Averaged across the systems, the word error rate was about 19% for white speakers but roughly 35% for Black speakers—nearly double—a gap the authors link to the systems' acoustic and language models, built on data in which speakers of African American Vernacular English were underrepresented.

The voice recognition systems weren't broken. They were accurate, impressively so, for speakers who sounded like the people in their training

data. For people who sounded different, the systems faced epistemic uncertainty of the most structural kind: they had never learned to hear these voices properly.

What matters for our purposes is not just that the systems were wrong—it's that the systems almost never signaled that they were uncertain. A voice recognition system transcribing speech it had rarely encountered in training produced the same confident-looking output as one transcribing the voices it knew well. There was no flag, no asterisk, no 51% confidence score.

A system that fails silently is a system where the human-in-the-loop has no loop to enter. The medical AI gave the radiologist a useful signal because it was designed to express its uncertainty. The voice recognition tools expressed no uncertainty even when their performance was dramatically degraded.

2.6 The Outside View: How Confusion Surfaces (or Doesn't)

For users and operators of AI systems, confusion can look like several different things—and only some of them are visible.

Explicit signals. Error messages, low-confidence warnings, referrals to human review. When the spam filter marks an email as “Possible spam—needs review,” that's explicit. When the medical AI displays a confidence score of 51%, that's explicit. These signals are easy to see and easy to act on.

Behavioral signals. A system under high uncertainty might hedge: “I found several possible matches” instead of one confident answer. It might offer multiple alternatives rather than a single recommendation. These signals require someone paying attention to notice.

Silent failures. When a chatbot provides a confident-sounding but incorrect answer, there is no signal at all. The confusion is completely invisible from the outside. This is the Air Canada problem (§1.8): the chatbot’s wrong answer about bereavement fares looked identical to its right answers about baggage fees.

The key design principle: for HITL to work, uncertainty must be surfaced, not buried. A system that is internally confused but externally confident has broken the loop before any human can enter it.

2.7 Wrong vs. Confused: A Crucial Distinction

A system can be wrong without being confused, and confused without being wrong.

A system is **wrong** when it gives an incorrect answer. A system is **confused** when it doesn’t know whether its answer is correct.

These overlap, but they’re not the same. A highly confident system can be wrong about an edge case—confidently wrong, exactly like the Air Canada chatbot. That’s the dangerous kind of wrong. It doesn’t signal that anything unusual is happening.

A less confident system might express appropriate uncertainty on a case it actually gets right. That uncertainty triggers unnecessary human review. That’s inefficient, but it’s not dangerous.

The important insight: **confusion is useful information**. A system that knows it might be wrong can ask for help. A system that doesn’t know it might be wrong cannot.

This reframes what “good” performance means for a HITL component. We’re not just looking for accuracy—we’re looking for **calibrated accuracy**: a system’s expressed confidence should reliably predict whether it’s right. A system that says it’s 90% confident should be right about 90% of the time (Figure 2.1). When a radiology AI flags an image at 51%, it is doing its job perfectly, even if it can’t make the diagnosis. It is correctly identifying a case at the edge of its knowledge and routing it to someone whose knowledge extends further.




Figure 2.1: A reliability diagram. A perfectly calibrated system lies on the diagonal—its stated confidence matches its empirical accuracy. The curve below is systematically overconfident: it claims more certainty than its track record justifies.

2.8 What Confusion Tells You About Training

AI confusion is not random. It has structure. That structure tells you something important about where the training data was thin, biased, or simply absent.

AI systems that classify gender from faces have historically performed far worse on darker-skinned women than on lighter-skinned men—and the benchmark datasets used to evaluate them were themselves skewed toward lighter-skinned subjects (Buolamwini & Gebru, 2018). The confusion is diagnostic: it reveals the shape of the data the system was built on.

Spam filters struggle with sophisticated phishing emails carefully crafted to resemble legitimate business communication, because phishing authors study and mimic the exact patterns the filter learned. The filter’s confusion reveals the adversary’s knowledge of the filter.

Confusion as Training Signal

Reading confusion patterns is, in a sense, reading the biography of a model’s training. The things it is most uncertain about are the things it has seen least of, or seen in the least representative form. Human review of uncertain cases generates the most valuable new training data: the cases that confuse the model are exactly the cases where more labeled examples would help most.

2.9 The Radiologist’s Bargain

There’s a lesson you might take from that story that points in the wrong direction. After reading about the 51% confidence flag, you might think: “Lower confidence threshold means more human review.” But that logic hits a wall quickly.

A radiologist reviewing tens of thousands of X-rays a year cannot personally verify every image the AI marks as uncertain. If the threshold for flagging is set too low—if the system marks 30% of images as uncertain—the human-review queue fills faster than it can be cleared. Alert fatigue sets in. The flags stop meaning anything.

This is the core tension in any HITL design, and Chapter 1 named it the Goldilocks problem (§1.9): too much asking creates alert fatigue; too little asking creates dangerous overconfidence.

The Flag Precision Principle

The value of a flag is proportional to its precision. A system that flags the truly uncertain cases is enormously valuable. A system that flags everything it's not completely certain about is noise. Good confusion-signaling requires calibrating how confused you need to be before the signal is worth sending.

2.10 The First Dimension, Up Close

Chapter 1 introduced the Five Dimensions framework. The first dimension—Uncertainty Detection—is the foundation on which all the others rest.

Uncertainty Detection is the system's ability to recognize when its outputs might be wrong. It is not about being uncertain in general. It is about being appropriately uncertain—in the right cases, by the right amount.

What a system cannot detect, it cannot act on. A system with no uncertainty detection will be equally confident about its most reliable predictions and its worst mistakes. This is why **calibration** and uncertainty quantification are not optional engineering details—they are the prerequisites for any effective human-in-the-loop design.

The practical question for system designers: what signals does this system have access to that correlate with the likelihood of error? And are those signals being used?

The mathematical machinery behind uncertainty quantification—including Monte Carlo Dropout, Deep Ensembles, Conformal Prediction, Expected Calibration Error, and out-of-distribution detection—is covered in Technical Appendix 2A.

The next time an app, website, or device produces an output you're not certain about, notice whether the system gave you *any* signal about its confidence. Did it hedge? Did it offer alternatives? Did it ask you to confirm? Or did it give you an answer as though the answer were inevitable? Keep a running log for a week. You will quickly develop a sense for which systems are designed to be honest about their uncertainty—and which are not.

Chapter Summary

Key Concepts

- AI “confusion” is a measurable property: uncertainty in probability estimates, not an emotional state
- Aleatoric uncertainty is irreducible noise; epistemic uncertainty is a knowledge gap—only the second kind reliably benefits from human intervention

- A system can be wrong without being confused (dangerous) and confused without being wrong (merely inefficient)
- Confusion must be surfaced, not buried—a silently failing system breaks the human loop before it begins
- Calibration is the property that makes uncertainty signals trustworthy: a 90%-confident system should be right about 90% of the time
- AI confusion patterns reveal training data gaps and are valuable signals for model improvement

Key Examples

Radiology AI at 51% confidence A low confidence score, not a high one, was the clinically valuable signal; triggered human review that caught a rare fungal infection

Spam filter edge cases The email in the middle: structural uncertainty at the classification boundary, not a solvable engineering problem

Voice recognition accuracy gaps Epistemic uncertainty from under-represented training data, expressed without uncertainty signals—silent failure

Key Principles

- Aleatoric uncertainty: more data may not help; the ambiguity is in the world
- Epistemic uncertainty: human expertise adds what the model lacks—this is where HITL shines

- The value of a flag is proportional to its precision; flags must be calibrated, not just present
- Uncertainty Detection (Dimension 1) is the foundation: all other HITL design depends on it

Five Dimensions Check

- *Uncertainty Detection* (Dimension 1): The chapter's central subject—the distinction between aleatoric and epistemic uncertainty, and the calibration that makes either signal trustworthy
- *Stakes Calibration* (Dimension 4): Flag precision is a stakes question—how confused the system must be before the signal is worth sending
- *Intervention Design* (Dimension 2): Confusion has to be surfaced, not buried—a silently failing system breaks the human loop before it begins

In the next chapter, we'll explore how AI systems learn in the first place—and why the way they learn creates predictable patterns of confusion that we can anticipate and design around.

Technical Appendix 2A: The Mathematics of AI Confusion

Chapter 2 introduced two types of uncertainty—aleatoric and epistemic—and argued that confusion is a feature rather than a bug when properly surfaced. This appendix provides the mathematical framework underlying those concepts.

— Formal foundations for Chapter 2

2A.1 Formal Definitions of Uncertainty Types

2A.1.1 Aleatoric Uncertainty

Aleatoric uncertainty captures the inherent randomness or noise in observed data. Given input x and true label y , aleatoric uncertainty is:

$$\sigma_{\text{aleatoric}}^2 = \mathbb{E}_{y \sim p(y|x)}[(y - \mathbb{E}[y | x])^2] = \text{Var}(y | x)$$

This represents the expected squared deviation between the true output and the conditional mean $\mathbb{E}[y | x]$, averaged over the true output distribution. It cannot be reduced by collecting more data, because it reflects genuine ambiguity in the relationship between x and y .

Key property.

$\sigma_{\text{aleatoric}}^2$ does not decrease as $|\mathcal{D}| \rightarrow \infty$

2A.1.2 Epistemic Uncertainty

Epistemic uncertainty captures uncertainty about the model parameters θ given observed data $\mathcal{D}$. Under the Bayesian framework:

$$p(\theta \mid \mathcal{D}) \propto p(\mathcal{D} \mid \theta) \cdot p(\theta)$$

The predictive distribution for a new input x^* is:

$$p(y^* \mid x^*, \mathcal{D}) = \int p(y^* \mid x^*, \theta) \cdot p(\theta \mid \mathcal{D}) d\theta$$

Epistemic uncertainty is high when this integral shows high variance—when different plausible parameter settings produce very different predictions. It decreases as more training data is observed:

$$\sigma_{\text{epistemic}}^2 \rightarrow 0 \quad \text{as } |\mathcal{D}| \rightarrow \infty$$

2A.1.3 Total Predictive Uncertainty

Following Kendall and Gal (2017), total predictive uncertainty decomposes as:

$$\underbrace{\text{Var}[y^*]}_{\text{total}} = \underbrace{\text{E}_{\theta}[\text{Var}_{y^*}[y^* \mid \theta]]}_{\text{aleatoric}} + \underbrace{\text{Var}_{\theta}[\text{E}_{y^*}[y^* \mid \theta]]}_{\text{epistemic}}$$

This decomposition is foundational for HITL design: the aleatoric component tells you how intrinsically hard the case is; the epistemic component tells you how much the model's own ignorance contributes.

2A.2 Measuring Calibration

2A.2.1 Definition

A model is **perfectly calibrated** if its expressed confidence exactly predicts its accuracy:

$$P(\hat{y} = y \mid \hat{p} = p) = p \quad \text{for all } p \in [0, 1]$$

When the model says it is 70% confident, it should be correct 70% of the time.

2A.2.2 Reliability Diagrams

To assess calibration empirically:

1. Partition predictions into M confidence bins $B_m = [(m-1)/M, m/M)$ for $m = 1, \dots, M$
2. For each bin, compute:
 - Average confidence: $\text{conf}(B_m) = \frac{1}{|B_m|} \sum_{i \in B_m} \hat{p}_i$
 - Accuracy: $\text{acc}(B_m) = \frac{1}{|B_m|} \sum_{i \in B_m} \mathbf{1}[\hat{y}_i = y_i]$
3. Plot $\text{conf}(B_m)$ vs. $\text{acc}(B_m)$

A perfectly calibrated model produces a straight diagonal line. Points below the diagonal indicate overconfidence (confidence exceeds accuracy)—the most common failure mode in modern deep learning (Guo et al., 2017).

2A.2.3 Expected Calibration Error (ECE)

$$\text{ECE} = \sum_{m=1}^M \frac{|B_m|}{n} |\text{acc}(B_m) - \text{conf}(B_m)|$$

where n is the total number of predictions. ECE is the weighted average absolute gap between confidence and accuracy across bins.

A high ECE means the model's confidence scores cannot be trusted to route cases correctly. For example, a model whose 90%-confidence bin has only 70% accuracy will dramatically under-request human oversight on the cases it's least reliable about.

2A.2.4 Temperature Scaling

Guo et al. (2017) showed that dividing logits by a scalar temperature $T > 1$ before softmax reduces overconfidence without altering class predictions:

$$\hat{p}_i = \text{softmax}(\mathbf{z}_i/T)$$

T is fitted on a held-out validation set by minimizing negative log-likelihood. It is the simplest effective post-hoc calibration method.

```
1 class TemperatureScaler:
2     def __init__(self, model):
3         self.model = model
4         self.temperature = torch.nn.Parameter(torch.ones(1)
5             * 1.5)
6
7     def calibrate(self, val_loader):
8         optimizer = torch.optim.LBFGS(
9             [self.temperature], lr=0.01, max_iter=50)
10        logits_list, labels_list = [], []
11
12        with torch.no_grad():
13            for x, y in val_loader:
```

```
13         logits_list.append(self.model(x))
14         labels_list.append(y)
15
16     logits = torch.cat(logits_list)
17     labels = torch.cat(labels_list)
18
19     def eval_nll():
20         optimizer.zero_grad()
21         loss = F.cross_entropy(logits /
22                                self.temperature, labels)
23         loss.backward()
24         return loss
25
26     optimizer.step(eval_nll)
27     return self
```

Listing 2.1: Temperature scaling implementation

2A.3 Uncertainty Estimation Methods

2A.3.1 Monte Carlo Dropout

Gal and Ghahramani (2016) showed that applying dropout at inference time corresponds approximately to Bayesian inference over model weights.

```
1 def mc_dropout_uncertainty(model, x, n_samples=100):
2     # model.train() enables dropout but also sets BatchNorm
3     # to training
4     # mode, which corrupts uncertainty estimates in BN
5     # networks.
6     # Prefer setting only dropout layers to train mode:
7     # for m in model.modules():
8     #     if isinstance(m, nn.Dropout): m.train()
```

```

7     model.train() # Enable dropout at inference (safe for
                        dropout-only nets)
8     predictions = []
9
10    for _ in range(n_samples):
11        with torch.no_grad():
12            pred = model(x)
13            predictions.append(torch.softmax(pred, dim=-1))
14
15    predictions = torch.stack(predictions, dim=0)
16    mean_pred = predictions.mean(dim=0)
17
18    # Total uncertainty: predictive entropy
19    entropy = -(mean_pred * torch.log(mean_pred +
20        1e-8)).sum(dim=-1)
21
22    # Aleatoric: mean of per-sample entropies
23    per_sample_entropy = -(predictions *
24        torch.log(predictions +
25        1e-8)).sum(dim=-1)
26    aleatoric = per_sample_entropy.mean(dim=0)
27
28    # Epistemic: mutual information (total - aleatoric)
29    epistemic = entropy - aleatoric
30
31    return mean_pred, aleatoric, epistemic, entropy

```

Listing 2.2: MC Dropout uncertainty estimation

2A.3.2 Deep Ensembles

Train M separate models with different random initializations (Lakshminarayanan et al., 2017):

$$\mu_{\text{ens}} = \frac{1}{M} \sum_{i=1}^M \mu_i(x) \quad \sigma_{\text{ens}}^2 = \frac{1}{M} \sum_{i=1}^M [\sigma_i^2(x) + \mu_i^2(x)] - \mu_{\text{ens}}^2$$

Advantages over MC Dropout. Better-calibrated uncertainty estimates; captures diversity across function space; parallelizable.

Disadvantage. Requires training and storing M full models.

2A.3.3 Conformal Prediction

A distribution-free framework providing prediction sets with guaranteed coverage:

$$P(y_{\text{true}} \in C_{\alpha}(x)) \geq 1 - \alpha$$

Given a calibration set $\{(x_i, y_i)\}_{i=1}^n$, the prediction set is:

$$C_{\alpha}(x^*) = \{y : s(x^*, y) \leq \hat{q}_{1-\alpha}\}$$

where $\hat{q}_{1-\alpha}$ is the $(1-\alpha)$ empirical quantile of calibration nonconformity scores (Angelopoulos & Bates, 2023).

HITL Value of Conformal Prediction

A prediction set of size 1 indicates high confidence. A set containing multiple classes indicates genuine ambiguity warranting human review. This provides formally guaranteed coverage without distributional assumptions.

2A.4 Out-of-Distribution Detection

2A.4.1 Maximum Softmax Probability (Baseline)

$$\text{OOD flag}(x) = \mathbf{1}\left[\max_c p(c \mid x) < \tau\right]$$

Deep networks can produce high-confidence outputs for OOD inputs. A model trained on cats and dogs may classify a fox as a dog with 99% confidence. The softmax does not reliably reflect OOD-ness.

2A.4.2 Mahalanobis Distance

K. Lee et al. (2018) use class-conditional Gaussian distributions fit to intermediate feature representations:

$$M(x) = \min_c (f(x) - \mu_c)^\top \Sigma^{-1} (f(x) - \mu_c)$$

High distance to all class centroids $\rightarrow$ likely OOD. This is particularly useful for diagnosing *why* a model is uncertain: is the input far from all training classes, or is it between two well-known classes?

```

1  def mahalanobis_ood_score(model, x, class_means,
2      precision_matrix):
3      with torch.no_grad():
4          features = model.get_features(x)
5
6          distances = []
7          for mu_c in class_means:
8              diff = features - mu_c
9              dist = (diff @ precision_matrix @ diff.T).diag()
```



```

9         distances.append(dist)
10
11     # Minimum distance to any class
12     min_distance = torch.stack(distances,
13                               dim=0).min(dim=0).values
14     return min_distance

```

Listing 2.3: Mahalanobis distance OOD scoring

2A.5 The Radiology Case: Mathematical Analysis

The 51% confidence case from Chapter 2 maps directly to the calibration framework.

Optimal routing threshold. Let τ be the confidence threshold. The system auto-diagnoses if confidence $\geq \tau$; otherwise routes to human review. Define:

- $\alpha(\tau)$: error rate on auto-routed cases (decreasing in τ)
- $\beta(\tau) = P(\text{confidence} < \tau)$: fraction routed to human review (increasing in τ)
- C_E : cost of a missed diagnosis
- C_R : cost of human review

Optimal threshold:

$$\tau^* = \arg \min_{\tau} \left[C_E \cdot \alpha(\tau) \cdot (1 - \beta(\tau)) + C_R \cdot \beta(\tau) \right]$$

In medical imaging, $C_E \gg C_R$, so τ^* is high. The AI at 51% correctly placed the case in the review queue, since any medically reasonable $\tau^* > 0.51$ would route it to the radiologist.

Interpretation. The model was not “almost right”—it was maximally uncertain ($0.51 \approx$ coin flip). This is the correct output when the model genuinely cannot distinguish two diagnoses that look similar in its learned representation. Perfect calibration under epistemic uncertainty.

2A.6 Exercises

Exercise 2A.1 — Calibration Measurement

Implement Expected Calibration Error (ECE) and reliability diagram construction for a pretrained image classifier. Compare ECE before and after temperature scaling. Then compute ECE broken down by input subgroup (e.g., demographic, image condition). Where is calibration worst?

Exercise 2A.2 — Aleatoric vs. Epistemic Decomposition

Using MC Dropout:

1. Generate two overlapping Gaussian clusters (in-distribution, high aleatoric at boundary) and one separated Gaussian cluster (OOD, high epistemic).
2. Train a dropout classifier.
3. Run MC Dropout. Plot aleatoric and epistemic uncertainty separately.
4. Show that aleatoric peaks at the in-distribution decision boundary; epistemic peaks in the OOD region.

Exercise 2A.3 — HITL Routing Optimization

Given a binary classifier with calibrated confidence scores:

1. Set costs: $C_{\text{FN}} = 100$, $C_{\text{FP}} = 10$, $C_{\text{human}} = 2$.
2. Sweep threshold $\tau \in [0.5, 0.99]$ and compute expected total cost.
3. Find τ^* .
4. Compare accuracy at τ^* vs. full automation vs. full human review.

Exercise 2A.4 — OOD Detection Comparison

Compare three OOD detection methods (max softmax probability, MC Dropout predictive entropy, Mahalanobis distance) on a test set with both in-distribution and OOD samples. Report AUROC for each method. Explain which performs best and why.

2A.7 Recommended Reading

Foundational

- Kendall and Gal (2017) — What uncertainties do we need in Bayesian deep learning? *NeurIPS*.
- Guo et al. (2017) — On calibration of modern neural networks. *ICML*.
- Lakshminarayanan et al. (2017) — Simple and scalable predictive uncertainty estimation. *NeurIPS*.

Methods

- Gal and Ghahramani (2016) — Dropout as a Bayesian approximation. *ICML*.
- K. Lee et al. (2018) — A simple unified framework for detecting OOD samples. *NeurIPS*.
- Angelopoulos and Bates (2023) — A gentle introduction to conformal prediction. *arXiv:2107.07511*.

Applications

- Begoli et al. (2019) — Uncertainty quantification in machine-assisted medical decision making. *Nature Machine Intelligence*.
- Koenecke et al. (2020) — Racial disparities in automated speech recognition. *PNAS*.

This appendix provides the mathematical foundations for uncertainty quantification, calibration, and OOD detection as introduced in Chapter 2. The worked exercises connect the theory to the practical HITL scenarios described in the main chapter.

3 How Computers Learn (The Simple Version)

The trouble—the interesting trouble, the trouble that creates the natural home for human-in-the-loop systems—happens in step three.

— on machine learning

3.1 A Teacher Who Only Teaches Cats

Imagine a new teacher arrives at a school. Before their first day, someone hands them a thick folder: ten thousand photographs, each labeled with a word. Cat. Cat. Dog. Cat. Dog. Dog. Cat.

The teacher studies the folder carefully. After a few days, they’ve noticed patterns: pointy ears, vertical pupils, retractable claws—these go with “cat.” Floppy ears, round pupils, wagging tails—these go with “dog.” The teacher hasn’t been told any rules. They’ve just looked at enough examples to build a mental model of the difference.

Then a student shows up with a photograph of a Maine Coon. The teacher says: “That looks like a cat—big pointy ears, long fur, narrow face. I’d say cat.”

Another student shows up with a photo of a Dachshund. “That’s clearly a dog—low to the ground, floppy ears. Dog.”

Then a student shows up with a photograph of a Siberian Forest Cat sitting at an unusual angle, backlit, in partial shadow. The teacher pauses. “It...probably has pointed ears? The shape looks cat-like. I’ll say cat—but I’m not certain.”

This is machine learning. Exactly this, minus the teacher and the children and the folder, and plus a lot of mathematics.

3.2 What Machine Learning Actually Is

The dominant approach to modern AI—the technology behind image recognition, speech assistants, spam filters, fraud detection, and large language models—is called machine learning. Despite the name, it’s not magic, and it’s not mysterious. At its core, it is a statistical process.

Here is the three-step summary:

Step 1: Gather labeled examples. For a spam filter: collect thousands of emails humans have marked as “spam” or “not spam.” For an image recognizer: collect millions of images with labels attached—“dog,” “cat,” “car,” “tree.” For a language model: collect billions of sentences from the internet with no explicit labels, just the raw sequence of text. The examples are the teacher. The labels (when present) are the lessons.

Step 2: Learn patterns. Feed the examples through a mathematical model—often a neural network—that adjusts its internal parameters to match the patterns in the data. The model is not learning rules you could write down. It is learning weights—numbers that describe how much different features contribute to different outputs. A spam filter learns that the phrase “you have won” in the subject line is a strong predictor of spam. An image classifier learns that diagonal edges combined with a certain color distribution predict “cat ear.”

Step 3: Generalize. Apply what the model learned to new, unseen examples. The test isn't "can the model correctly classify emails it was trained on?"—of course it can. The test is "can it correctly classify emails it has never seen before?" Generalization is the whole point.

The trouble—the interesting trouble, the trouble that creates the natural home for human-in-the-loop systems—happens in step 3.

3.3 The Generalization Gap

When you train a machine learning model, you use a set of examples. When you deploy it, you encounter examples that are not in that set. The question is: how well does the model's learned pattern extend to new territory?

The gap between training performance and real-world performance is called the **generalization gap**. It exists because the world is larger than any training dataset. No matter how many examples you collect, the real world will eventually show you something that wasn't there.

This gap has a structure. It is not random. The model will generalize well on examples that are similar to things it has seen and poorly on examples that are different. The boundary of good performance corresponds roughly to the boundary of the training data.

The Generalization Gap and HITL Design

The cases that most need human help are the cases farthest from the training distribution. These are, almost by definition, the cases the model is least equipped to handle. They are the students showing up with unusual photographs—the Siberian Forest Cat backlit in shadow.

A well-designed HITL system uses the model's confidence (or, more precisely, the model's distance from familiar training examples) to detect when it has wandered into unfamiliar territory—and routes those cases to human review.

3.4 Overfitting: The Expert Who Never Learned to Generalize

There is a classic failure mode in machine learning called **overfitting**. It is exactly what it sounds like: the model fits the training data too well.

Consider the teacher with the photograph folder. Suppose that in the training set, every photo of a cat happened to have a blue background. The teacher notices this pattern and adds it to their mental model: “cats have blue backgrounds.” This pattern has nothing to do with what makes something a cat—it's a coincidence of the training set. But if you test the teacher only on training-set photos, this rule seems to work perfectly.

Now send in a real-world cat photo. Black background. The teacher hesitates: “The ears are pointy but the background is wrong. I'm less sure.”

This is overfitting: the model has learned patterns that are specific to the training data rather than patterns that describe the underlying truth. It performs well on training examples and poorly on new examples.

Overfitting creates a particular kind of false confidence. The model may be very sure of its predictions—having fit the training data tightly—while being systematically wrong on the cases that don't match the training set's quirks. The expressed confidence is high;

the actual accuracy on new examples is low. This is calibration failure in its most consequential form.

The human-in-the-loop implication: an overfit model may not flag the cases it most needs help with, because it has learned to be confident in its (flawed) patterns. This is why uncertainty detection alone isn't enough—you also need calibration auditing to catch systematic overconfidence.

3.5 Underfitting: The Expert Who Never Learned Enough

The opposite failure mode is **underfitting**: the model has not learned the training data well enough to generalize at all. It produces generic outputs regardless of input.

An underfit model is like the teacher who studied the photo folder for only ten minutes before their first class and absorbed almost nothing. They have a vague sense that “cats are smaller than dogs” but can't reliably tell them apart. Their uncertainty signals are constant and uninformative—they're unsure about everything, including the easy cases.

Underfitting is actually easier to diagnose than overfitting, because it produces poor performance even on the training set. The model fails on examples it has seen. In a HITL system, an underfit model would route almost everything to human review—not because it has good calibration, but because it has learned nothing worth being confident about.

In practice, the interesting and dangerous case is overfitting, not underfitting. Modern neural networks are powerful enough to memorize training data if you let them. The engineering challenge is finding the right balance: enough fitting to learn real patterns, not so much fitting that the model learns the noise.

3.6 Why Even a Perfect Learner Needs Help at the Edges

Here is the insight that connects machine learning to human-in-the-loop systems in a way that no engineering progress can eliminate.

Suppose you had a perfect machine learning system. It fits the training data exactly right—not too much, not too little. It generalizes as well as any statistical approach possibly could. It has achieved the theoretical optimum for learning from the data it was given.

This perfect learner is still bounded by its training distribution. There will be examples in the real world that differ from the training data in ways that matter—new domains, new populations, new edge cases, situations that didn’t exist when the training data was collected. No amount of better training, more parameters, or cleverer architecture eliminates the boundary between the distribution the model learned from and the distribution it will encounter in deployment.

At those boundaries, the perfect learner does not have perfect performance. It has performance that degrades gracefully—hopefully—as inputs move further from what it knows. And at the edges, where the learning is thinnest, human judgment becomes most valuable.

The Permanent Edge Problem

This is not a temporary limitation waiting to be engineered away. It is a fundamental property of statistical learning:

- A system trained on human-labeled data from 2020 will encounter situations from 2025 that didn’t exist in 2020.
- A system trained on medical records from a teaching hospital will encounter patients from rural clinics whose presentations it has never seen.

- A system trained on English language text will encounter code-switching, evolving slang, and new cultural references that weren't in the training corpus.

The edges are permanent. The need for human help at the edges is permanent. The question is not whether to involve humans at the boundaries of machine learning—the question is how to do it well.

A fair objection hangs over this example, and it deserves a direct answer: this book's evidence is selected on failure. The cases that open its chapters—the chatbot that invented a refund policy, the car that followed GPS directions into a lake—are vivid because something went wrong. But selection on failure can show that loops break, not that they always break, and a skeptic is right to ask where the base rates are. So here is the claim in falsifiable form: if a domain accumulates years of autonomous operation at better-than-human error rates with acceptable tail risk—as spam filtering, search ranking, and route planning largely have—then a human in that loop is cost without benefit, and full automation with monitoring is the right architecture there. The thesis of this book is not that every system needs a human at the edges. It is that in domains where errors are expensive, the distribution keeps moving, and someone must answer for the outcome, the edges regenerate faster than engineering closes them. Where those conditions hold, the loop is not a workaround. Where they do not, it is.

3.7 The Train/Test Gap as HITL Metaphor

Machine learning practitioners split their data into training sets and test sets. The model is trained on the training set and evaluated on the test set. The test set represents “what will the model encounter in the real world?”

In a well-designed experiment, the test set is drawn from the same distribution as the training set, just held back so the model hasn't seen it. Performance on the test set is a proxy for performance in deployment.

But here is the thing practitioners know and users often don't: the real-world deployment set is not the same as the test set. The test set is carefully curated, drawn from the same distribution, cleaned of pathological edge cases. The real world has none of those properties.

The gap between test performance and deployment performance is one of the most consistent phenomena in applied machine learning. Models that look great on benchmark tests often underperform in deployment—not because the benchmark was badly constructed, but because the deployment distribution includes things the benchmark didn't.

The Deployment Gap

This gap is the formal expression of why human-in-the-loop systems exist. The model handles the cases that look like the training distribution. The cases that don't look like the training distribution—the cases that exist in deployment but not in evaluation—are strangers to the model—and a human should make the introduction.

3.8 Patterns, Not Rules

One of the most important practical consequences of how machine learning works is that the systems cannot always explain their decisions in terms humans can verify.

When a spam filter classifies an email, it is not applying a rule you could read. It is combining hundreds or thousands of weighted features in ways that reflect statistical patterns in its training data. The output “spam” means “the weighted sum of all these features, in the configuration

present in this email, matches the pattern I learned from spam examples.” The explanation is the weights—and the weights are not human-readable.

This creates a specific challenge for human oversight: you are being asked to review a decision whose reasoning you cannot fully inspect. The human reviewer sees the email and the classification—not the internal calculations that produced it.

This is not unique to ML. Doctors make diagnoses whose reasoning is partly intuitive. Judges make rulings whose logic is partly implicit. Experts routinely know things they cannot fully articulate. But in machine learning, the gap between the decision and the explanation is structural and systematic, not occasional.

Human-in-the-loop systems must account for this. When a reviewer is asked to confirm or correct an AI classification, they are not auditing the model’s reasoning. They are providing a second, independent judgment that uses different—often more contextual—information. The value of human review is not that it traces the model’s internal logic. It is that it provides an independent check from a different vantage point.

This is why the framing of “human oversight” sometimes misleads: the human is not watching the AI and catching its mistakes in real time. The human is providing an independent signal that the system incorporates alongside the AI’s signal.

Find three predictions an AI system has made about you this week—autocorrect suggestions, content recommendations, search result rankings. For each one, ask: what was the system trained on that led to this prediction? What might it be missing about you specifically? You don’t need to know the technical answer. The

practice of asking the question changes how you think about AI outputs.

Chapter Summary

Key Concepts

- Machine learning: three steps—gather labeled examples, learn patterns, generalize to new cases
- The generalization gap: performance degrades as inputs move away from training distribution
- Overfitting: too much fit to training data; model is confident but wrong on new examples
- Underfitting: too little fit; model fails even on familiar examples
- Even a perfect learner has permanent performance boundaries at the edges of its training distribution—not a limitation to be engineered away, but the reason HITL design is the correct architecture, not a workaround
- The train/test gap is a formal expression of why HITL systems are permanently needed
- ML systems learn patterns, not rules; human review provides independent judgment, not audit of reasoning

Key Examples

The Maine Coon and Dachshund Easy generalization within training distribution—clear cases far from the decision boundary

The Siberian Forest Cat in shadow Edge case requiring epistemic HITL—backlit, unusual angle, uncertain territory

Overfit model with “blue background = cat” rule Confident but wrong on real-world examples; calibration failure in its most consequential form

Real-world deployment vs. test benchmark performance gap The consistent phenomenon that motivates HITL system design

Five Dimensions Check

- *Uncertainty Detection* (Dimension 1): Learning theory explains why uncertainty concentrates at distribution edges
- *Timing* (Dimension 3): Cases far from training distribution are the natural trigger for human escalation

In the next chapter, we ask how a system turns uncertainty into a decision. The apparently simple step of choosing an action from a probability distribution turns out to involve deep questions about value and stakes.

Technical Appendix 3A: Machine Learning Foundations for HITL Design

Chapter 3 introduced machine learning as a three-step process of gathering examples, learning patterns, and generalizing. This appendix provides the mathematical framework underlying those concepts: empirical risk minimization, the bias-variance tradeoff, generalization bounds, and distribution shift.

— Formal foundations for Chapter 3

3A.1 Empirical Risk Minimization

The core framework for supervised machine learning is **empirical risk minimization** (ERM). Given training data $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$, we find parameters θ that minimize:

$$\hat{R}(\theta) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f_{\theta}(x_i), y_i)$$

where $\mathcal{L}$ is a loss function (e.g., cross-entropy for classification, mean squared error for regression) and f_{θ} is the model.

The **true risk** we actually care about is:

$$R(\theta) = \mathbb{E}_{(x,y) \sim P}[\mathcal{L}(f_{\theta}(x), y)]$$

The **generalization gap** is:

$$\text{gap}(\theta) = R(\theta) - \hat{R}(\theta)$$

For parameters θ fitted by ERM, this gap is non-negative *in expectation*—the empirical risk is an optimistically biased estimate of the true risk, since the same data are used to fit the model and to score it—and its expected size tends to grow with model complexity relative to dataset size. (For a single *fixed* θ chosen independently of the data, the empirical risk is unbiased and the gap can fall on either side of zero.)

3A.2 The Bias-Variance Tradeoff

For a regression model, the expected prediction error decomposes as:

$$\mathbb{E}[(y - \hat{f}(x))^2] = \underbrace{\text{Bias}^2(\hat{f}(x))}_{\text{Systematic error}} + \underbrace{\text{Var}(\hat{f}(x))}_{\text{Sensitivity to training data}} + \underbrace{\sigma^2}_{\text{Irreducible noise}}$$

- **Bias:** Error from the model not capturing the true relationship (underfitting)
- **Variance:** Error from sensitivity to fluctuations in training data (overfitting)
- σ^2 : Aleatoric uncertainty — cannot be reduced

The bias-variance tradeoff: reducing bias (using a more complex model) typically increases variance (model fits noise), and vice versa.

HITL Design Implication of Bias vs. Variance

- **High-variance models** benefit most from human-in-the-loop review because their errors are unpredictable — concentrated at unusual inputs with no systematic pattern.
- **High-bias models** benefit most from systematic data collection for the classes of inputs where bias is highest.

3A.3 Generalization Bounds

3A.3.1 VC Dimension Bound

For a hypothesis class $\mathcal{H}$ with VC dimension d_{VC} , with probability $1 - \delta$ (Vapnik & Chervonenkis, 1971):

$$R(\theta) \leq \hat{R}(\theta) + \sqrt{\frac{8}{n} \left(d_{VC} \cdot \left(\ln \frac{2n}{d_{VC}} + 1 \right) + \ln \frac{4}{\delta} \right)}$$

The generalization gap shrinks at rate $\mathcal{O}(\sqrt{d_{VC}/n})$. More complex models (higher d_{VC}) need more data to close the gap.

3A.3.2 PAC-Bayes Bound

A tighter bound for stochastic classifiers (McAllester, 1999):

$$\mathbb{E}_{\theta \sim Q}[R(\theta)] \leq \mathbb{E}_{\theta \sim Q}[\hat{R}(\theta)] + \sqrt{\frac{\text{KL}(Q\|P) + \ln(n/\delta)}{2n}}$$

where P is the prior and Q is the posterior distribution over parameters.

3A.4 Distribution Shift

3A.4.1 Types of Shift

Covariate shift. $P(x)$ changes but $P(y | x)$ stays the same:

$$P_{\text{train}}(x) \neq P_{\text{test}}(x), \quad \text{but} \quad P_{\text{train}}(y | x) = P_{\text{test}}(y | x)$$

Example: a spam filter trained on emails from 2020, deployed in 2025 — email styles changed but spam is still spam.

Label shift. $P(y)$ changes but $P(x | y)$ stays the same:

$$P_{\text{train}}(y) \neq P_{\text{test}}(y), \quad \text{but} \quad P_{\text{train}}(x | y) = P_{\text{test}}(x | y)$$

Example: fraud rate changes seasonally; the appearance of fraud is similar but its base rate changes.

Concept drift. The relationship $P(y | x)$ changes:

$$P_{\text{train}}(y | x) \neq P_{\text{test}}(y | x)$$

Example: a hiring model trained before a policy change; the definition of “qualified” shifts.

Wildcard. A completely new feature space:

$$\exists x \in \mathcal{X}_{\text{test}} : p_{\text{train}}(x) \approx 0$$

Example: a vision model encountering medical images after being trained on natural photographs.

3A.4.2 Detecting Distribution Shift

Maximum Mean Discrepancy (MMD). A kernel-based two-sample test for whether two samples are drawn from the same distribution (Gretton et al., 2012):

$$\text{MMD}^2(P, Q) = \|\mathbb{E}_{x \sim P}[\phi(x)] - \mathbb{E}_{x \sim Q}[\phi(x)]\|_{\mathcal{H}}^2$$

For an RBF kernel $k(x, x') = \exp(-\|x - x'\|^2 / 2\sigma^2)$:

```

1 def compute_mmd(X_train, X_test, sigma=1.0):
2     """Estimate MMD between training and test
3       distributions."""
4     def rbf_kernel(X, Y):
5         dists = cdist(X, Y, 'sqeuclidean')
6         return np.exp(-dists / (2 * sigma**2))
7
8     K_tt = rbf_kernel(X_train, X_train)
9     K_ee = rbf_kernel(X_test, X_test)
10    K_te = rbf_kernel(X_train, X_test)
11
12    return K_tt.mean() - 2*K_te.mean() + K_ee.mean()

```

Listing 3.1: MMD estimation between training and test distributions

Simple monitoring approach for production HITL.

```

1 class DistributionMonitor:
2     def __init__(self, reference_predictions,
3       window_size=1000):
4         self.reference = np.array(reference_predictions)
5         self.window_size = window_size
6         self.recent = deque(maxlen=window_size)

```

```

7     def update(self, new_prediction):
8         self.recent.append(new_prediction)
9
10    def drift_score(self):
11        if len(self.recent) < self.window_size:
12            return None
13        recent = np.array(self.recent)
14        # KS test p-value (low p = significant drift)
15        stat, p_value = ks_2samp(self.reference, recent)
16        return stat, p_value

```

Listing 3.2: Distribution monitor with KS-test drift detection

3A.5 The PAC Learning Framework

Probably Approximately Correct (PAC) learning, introduced by Valiant, formalizes when learning is feasible (Valiant, 1984).

A concept class $\mathcal{C}$ is PAC learnable if there exists an algorithm $\mathcal{A}$ such that for any $\varepsilon, \delta > 0$ and any distribution P over $\mathcal{X}$, given $m \geq \text{poly}(1/\varepsilon, 1/\delta, d, \text{size}(c))$ examples (with d the dimension of the instance space), with probability $\geq 1 - \delta$:

$$R(\mathcal{A}(\mathcal{D}_m)) \leq \varepsilon$$

provided the labels are generated by some concept $c \in \mathcal{C}$ (the *realizable* case). In the *agnostic* case, where P is an arbitrary distribution over $\mathcal{X} \times \mathcal{Y}$, no algorithm can guarantee error below ε in general; the guarantee weakens to $R(\mathcal{A}(\mathcal{D}_m)) \leq \min_{c \in \mathcal{C}} R(c) + \varepsilon$, with sample complexity scaling as $1/\varepsilon^2$ rather than $1/\varepsilon$.

PAC Theory and Rare-Category HITL

Learning requires a finite number of labeled examples proportional to $1/\epsilon$ in the realizable case ($1/\epsilon^2$ in the agnostic case). For very rare categories (low prior probability), the sample required to learn them grows rapidly. Rare-category detection is a natural HITL application: human review of uncertain cases on rare categories is often the only practical way to achieve acceptable error rates.

3A.6 Exercises

Exercise 3A.1 — Distribution Shift Diagnosis

A spam filter achieves 99.2% accuracy on its test set but 96.1% accuracy in deployment. The test set was drawn from the same email corpus as training. Identify: (a) what type of distribution shift is most likely; (b) how you would diagnose it; (c) what HITL intervention would help most.

Exercise 3A.2 — Bias-Variance and Fine-Tuned Models

Using the bias-variance decomposition, explain why a large pre-trained language model fine-tuned on a small domain-specific dataset might be riskier in HITL deployment than a simpler model trained on a larger domain-specific dataset. What monitoring would you add?

Exercise 3A.3 — MMD Monitoring

Implement a simple version of maximum mean discrepancy monitoring in Python. Use it to detect when a batch of new predictions has drifted from a reference batch. What threshold for MMD would you use to trigger a human review audit?

Exercise 3A.4 — Irreducible Error and HITL

The irreducible error σ^2 in the bias-variance decomposition is analogous to aleatoric uncertainty. Explain in technical terms why human-in-the-loop intervention cannot reduce this component of error, and what a HITL system should do when it detects a case with high aleatoric uncertainty.

3A.7 Recommended Reading

Foundational

- Vapnik, V. (1998). *Statistical Learning Theory*. Wiley.
- Shalev-Shwartz, S., & Ben-David, S. (2014). *Understanding Machine Learning: From Theory to Algorithms*. Cambridge University Press.

Distribution Shift

- Quinonero-Candela, J., et al. (Eds.). (2009). *Dataset Shift in Machine Learning*. MIT Press.
- Sugiyama, M., & Kawanabe, M. (2012). *Machine Learning in Non-Stationary Environments*. MIT Press.

- Gretton, A., et al. (2012). A kernel two-sample test. *Journal of Machine Learning Research*, 13, 723–773.

PAC Learning

- McAllester, D. (1999). PAC-Bayesian model averaging. *COLT*.

This appendix provides the mathematical foundations for empirical risk minimization, the bias-variance tradeoff, generalization bounds, and distribution shift as introduced in Chapter 3. The worked exercises connect theory to practical HITL monitoring and intervention design.

4 From Confusion to Decision

The threshold is not the end of the story. It is the beginning of the conversation.

— on thresholds and values

4.1 The Paramedic's Rule

In emergency medicine, there is a principle called the “treat first, ask later” rule. When a paramedic arrives at the scene of an unconscious patient, they don’t wait for a confirmed diagnosis before beginning resuscitation. The cost of not treating someone who needs it is so high—death—that they accept a significant rate of treating people who didn’t need the intervention.

The inverse principle exists in surgery: you don’t operate unless you’re confident it’s necessary. An unnecessary surgery carries real risks—infection, anesthetic complications, recovery time. So the threshold for surgery is set higher than the threshold for CPR.

Both of these are threshold decisions. They define the point at which uncertainty becomes action. And they are calibrated—deliberately—to the stakes of being wrong in each direction.

This is not a medical insight. It is a design principle. And it applies to every AI system that has to turn a probability into a choice.

4.2 The Probability Isn't the Decision

When a machine learning model processes an input, its output is a probability. The spam filter doesn't output "spam" or "not spam" directly. It outputs something like: "There is a 0.73 probability that this email is spam."

Then something has to turn 0.73 into an action.

That something is a **threshold**. A line is drawn somewhere on the probability scale—often at 0.5, but sometimes much higher or lower—and every input above that line goes in one bucket, everything below it in another. The model's job is to produce probabilities. The threshold's job is to produce decisions.

This seems like a small technical detail. It is not. The choice of threshold is one of the most consequential decisions in the design of any AI system—and it is fundamentally not a technical question. It is a question about values.

To see why, consider what happens when the threshold is wrong.

4.3 The Two Mistakes

Every classification system can make exactly two kinds of mistakes:

Type 1 error (False Positive): The system says "yes" when the true answer is "no." The spam filter moves a legitimate email to the spam folder. The fraud detection system blocks a valid purchase. The cancer screening algorithm flags healthy tissue for biopsy.

Type 2 error (False Negative): The system says "no" when the true answer is "yes." The spam filter lets a phishing email into the inbox. The fraud detection system approves a stolen card transaction. The cancer screening algorithm misses a tumor.

The Irreducible Tradeoff

You cannot minimize both types of error simultaneously. The tradeoff is real. If you lower the threshold, you flag more cases as positive—catching more true positives (better!) but also flagging more false positives (worse!). If you raise the threshold, you let more cases through as negative—missing more true positives (worse!) but also avoiding more false alarms (better!).

The threshold is a dial that trades one type of error for the other. And the question of where to set that dial is not answerable by looking at the data alone. It requires knowing how much each type of error costs.

4.4 The Cost Asymmetry

In some domains, the two types of error are roughly equivalent in cost. If a movie recommendation system fails in either direction—recommending a movie you hate, or failing to recommend one you'd love—the cost is roughly “a bit of time wasted.” Either error is about as bad as the other, and a threshold of 0.5 makes reasonable sense.

In other domains, the two types of error have wildly different costs. A cancer screening algorithm that misses a malignant tumor (Type 2 error) condemns a patient to delayed diagnosis. The cost is potentially measured in years of life. The same algorithm that flags healthy tissue for unnecessary biopsy (Type 1 error) costs the patient an uncomfortable procedure, some anxiety, and a medical bill. Painful—but recoverable.

Given this asymmetry, the cancer screening algorithm should be set with a threshold biased toward sensitivity: flag more positive cases, accept more false alarms, minimize the chance of missing the real thing. The cost structure demands it.

Child Protective Services: When Both Errors Are Serious

Consider a child protective services AI that screens reports of suspected abuse. A false negative—failing to flag a genuine case—means a child remains in danger. A false positive—flagging a family incorrectly—means a traumatic and potentially family-disrupting investigation. Both errors are genuinely serious.

The question of which is worse is not answerable technically. It is a question of values, priorities, and the society we want to live in. Different reasonable people give different answers, and those differences translate directly into different threshold settings. This is why setting the threshold is a human decision, not a model decision.

4.5 Decision Boundaries and Edge Cases

The threshold creates a **decision boundary**: a line in the probability space that separates one action from another. Understanding this boundary geometrically helps clarify where human-in-the-loop systems add the most value.

Consider a simple two-class problem: emails that are spam and emails that are not. In the high-dimensional space of email features (words, sender, format, etc.), the spam emails cluster in certain regions and the legitimate emails cluster in other regions. The classifier's job is to draw a boundary between these regions.

In the interior of each cluster, the classifier is confident. Emails that look very much like spam are far from the boundary, deep inside the spam region. Emails that look very much like legitimate communication are far on the other side. The classifier handles these with high confidence and high accuracy.

Near the boundary is where the interesting—and difficult—cases live. These are the emails that contain features from both regions: sales pitches that use legitimate business language, phishing attempts that mimic real vendors, marketing emails that share formatting with scam templates. The classifier is uncertain here. Its probability estimates land near 0.5, in the ambiguous zone.

The Boundary Zone as HITL Territory

The boundary region is the natural home of human-in-the-loop intervention. Not because every borderline case is important—most aren't. But because the rate of error is highest near the boundary, and the cases that are important (the sophisticated phishing attempt, the false alarm that costs a relationship) are disproportionately concentrated in this region.

A HITL system that routes only the near-boundary cases to human review—the cases where the classifier's confidence is between, say, 0.3 and 0.7—is efficiently targeting human attention where it has the highest marginal value (Figure 4.1).




Figure 4.1: The two-threshold HITL band. Scores below τ_{low} are decided automatically as negatives and scores above τ_{high} as positives; the uncertain middle band is routed to human review.

4.6 The Threshold Is Not Fixed

One mistake in HITL design is treating the threshold as a fixed engineering parameter, set once and left alone. In practice, the optimal threshold changes:

Over time. The base rate of positive cases changes. If fraud increases from 0.1% to 0.5% of transactions, the same model now flags far more transactions, and the review queue balloons. The threshold needs adjustment.

Across populations. A credit model may need different thresholds for different customer segments if the cost of errors differs across those segments.

With stakes. A high-stakes case—large transaction amount, high-profile account—warrants a more conservative threshold than a routine small transaction.

With system load. When human reviewers are overwhelmed, routing more cases to them is counterproductive. Alert fatigue sets in. The threshold should adapt to the capacity of the human-in-the-loop component, not just the performance of the model.

This last point is underappreciated. A HITL system is not just a model plus humans. It is a system whose behavior depends on the humans' current capacity and performance. A threshold setting appropriate for a team of 10 reviewers is often inappropriate for a

team of 3. And no threshold setting is appropriate for a team that is so overwhelmed it has stopped reading carefully.

4.7 Three Real Systems, Three Threshold Philosophies

Medical imaging triage uses a low threshold for flagging: it is far better to send a healthy patient for a follow-up scan than to miss a malignant finding. Radiologists operating in HITL systems regularly review flagged cases that turn out to be benign—and they accept this as the price of the system’s sensitivity. The threshold is set at a level that radiologists are comfortable with: they’d rather do ten unnecessary reviews than miss one cancer.

Banking fraud detection uses an asymmetric two-threshold system. Below the lower threshold: approve automatically. Above the upper threshold: decline automatically. Between them: route to human review or send an SMS confirmation to the customer. The band in the middle—the “ask the human” zone—is calibrated to the bank’s risk appetite and the capacity of its review team. Banks adjust these thresholds seasonally, during high-fraud periods, and when they detect coordinated attack patterns.

Social media content moderation uses category-specific thresholds. Content that is almost certainly violating (terrorist imagery, child exploitation material) is removed automatically—the confidence bar for removal is set deliberately low, because the cost of a false negative is extreme. Content that might violate community standards (graphic violence in ambiguous context, potentially coordinated harassment, political speech near the edge of rules) is routed to human reviewers. Content that clearly doesn’t violate

is passed through. Each category has its own threshold, reflecting the different costs of false positives (removing legitimate speech) and false negatives (allowing harmful content) in that category.

title

These are not arbitrary choices. They reflect deliberate judgments about the costs of each type of error in each domain—judgments that cannot be made by the algorithm.

4.8 The Human Value Judgment That Hides in the Algorithm

When engineers set a threshold for an automated system, they are making a value judgment. This is true even if nobody calls it that. When a credit scoring algorithm sets its rejection threshold, whoever sets that number is making a judgment about the acceptable rate of denying legitimate applicants versus the acceptable rate of approving fraudulent ones. That judgment has redistributive consequences—it affects who gets access to credit.

The pretense that the threshold is just a “technical parameter” obscures the nature of the choice. It allows consequential value judgments to be made without accountability, without deliberation, and without the participation of the people who will bear the costs of being wrong.

Human-in-the-loop design, at its best, forces the threshold question into the open. By designating cases near the threshold for human review, it creates a visible space where value tradeoffs are exercised explicitly—where a human decision-maker faces a real case and has to commit to an answer. That process generates information: about where the model is uncertain, about where human judgment differs from model output, and about whether the threshold is set in the right place.

The threshold is not the end of the story. It is the beginning of the conversation.

Think of a decision you made recently that involved a threshold—a cutoff between “do it” and “don’t do it.” Maybe it was whether to send an email, whether to raise a concern at work, whether to go to a doctor. What were the costs of the two types of error in that situation? Did your implicit threshold reflect those costs accurately? Could you describe, in plain language, what your threshold was?

Chapter Summary

Key Concepts

- Thresholds convert probabilities to actions—a necessary step that is not a technically neutral choice
- Two types of error: false positives (saying yes when no) and false negatives (saying no when yes)
- The optimal threshold minimizes expected cost, which depends on the cost asymmetry of each error type
- Decision boundaries define where confusion concentrates; near-boundary cases are the natural home of HITL
- The threshold should adapt to domain, population, stakes, and reviewer capacity

- Threshold-setting embeds value judgments that should be made explicitly, not hidden in technical parameters
- The two-threshold HITL band—automating high-confidence and cautious-direction cases, routing the middle to humans—is the core design pattern this book returns to throughout: not humans or machines, but each doing what they do best

Key Examples

Emergency medicine threshold asymmetry CPR vs. surgery: the same principle of cost-calibrated thresholds, in the highest-stakes domain

Cancer screening High sensitivity required—accept false alarms, minimize false negatives

Credit scoring Threshold embeds judgment about who bears the cost of each error type

Banking fraud detection Two-threshold system with human-in-the-loop band, dynamically adjusted

Five Dimensions Check

- *Stakes Calibration* (Dimension 4): This chapter explores it directly—the threshold is the mechanism by which stakes enter the system design
- *Intervention Design* (Dimension 2): The “ask the human” band between thresholds is the intervention zone

In the next chapter, we move from knowing that we need to ask to how to ask well. The design of HITL interventions turns out to be a surprisingly deep question—one that draws on cognitive psychology, communication design, and the mechanics of how human attention works.

Technical Appendix 4A: Threshold Optimization and Cost-Sensitive Learning

Chapter 4 introduced the threshold as the mechanism that turns a probability into an action, and argued that optimal threshold-setting requires knowing the cost structure of each error type. This appendix formalizes that argument through ROC analysis, precision-recall tradeoffs, cost-sensitive learning, and multi-threshold HITL band design.

— Formal foundations for Chapter 4

4A.1 The ROC Curve and AUC

4A.1.1 Definition

For a binary classifier with scores $s(x) \in [0, 1]$, sweeping the threshold $\tau \in [0, 1]$ traces the Receiver Operating Characteristic (ROC) curve (Fawcett, 2006):

$$\text{TPR}(\tau) = P(s(x) \geq \tau \mid y = 1) \quad [\text{True Positive Rate / Sensitivity}]$$

$$\text{FPR}(\tau) = P(s(x) \geq \tau \mid y = 0) \quad [\text{False Positive Rate}]$$

The Area Under the ROC Curve (AUC) equals the probability that a random positive is scored higher than a random negative:

$$\text{AUC} = P(s(x^+) > s(x^-))$$

AUC measures ranking quality, independent of threshold choice.

AUC vs. Threshold: Separate Questions

AUC tells you how separable your classes are. The threshold tells you where to make the cut given your cost structure. These are separate questions. A model with high AUC can still have a badly chosen threshold; a model with modest AUC but a carefully chosen threshold can outperform a high-AUC model with a naive threshold.

4A.1.2 Optimal Threshold Selection

The optimal threshold τ^* minimizes expected cost:

$$\mathbb{E}[\text{Cost}(\tau)] = C_{\text{FN}} \cdot \text{FNR}(\tau) \cdot P(y = 1) + C_{\text{FP}} \cdot \text{FPR}(\tau) \cdot P(y = 0)$$

where C_{FN} is the cost of a false negative (miss) and C_{FP} is the cost of a false positive (false alarm).

For a calibrated score $s(x) = P(y = 1 | x)$, setting $d\mathbb{E}[\text{Cost}(\tau)]/d\tau = 0$ gives the optimal threshold:

$$\tau^* = \frac{C_{\text{FP}}}{C_{\text{FP}} + C_{\text{FN}}}$$

The class priors $P(y = 0)$ and $P(y = 1)$ do not appear: a calibrated posterior already incorporates them, so the threshold depends only on the cost ratio. (This is why `optimal_decision` below thresholds on $c_{\text{fp}}/(c_{\text{fp}} + c_{\text{fn}})$ directly.)

- For equal costs ($C_{FN} = C_{FP}$): $\tau^* = 0.5$
- For $C_{FN} \gg C_{FP}$ (missing a positive is very costly, e.g., cancer screening): $\tau^* \ll 0.5$

4A.2 Precision-Recall Tradeoff

4A.2.1 When AUC Is Misleading

AUC is dominated by performance on the majority class. For imbalanced problems (rare positive class), precision-recall analysis is more informative (Davis & Goadrich, 2006); Figure 4.2 contrasts the two views of the same classifier.

$$\text{Precision}(\tau) = \frac{\text{TP}}{\text{TP} + \text{FP}} = P(y = 1 \mid s(x) \geq \tau)$$

$$\text{Recall}(\tau) = \frac{\text{TP}}{\text{TP} + \text{FN}} = P(s(x) \geq \tau \mid y = 1)$$

The F_β score balances precision and recall with weight β :

$$F_\beta = \frac{(1 + \beta^2) \cdot \text{Precision} \cdot \text{Recall}}{\beta^2 \cdot \text{Precision} + \text{Recall}}$$

- $\beta = 1$: Equal weight on precision and recall (F1 score)
- $\beta = 2$: Recall twice as important (use when false negatives are more costly)
- $\beta = 0.5$: Precision twice as important (use when false positives are more costly)




Figure 4.2: ROC and precision–recall views of the same classifier on a test set with 25% positive prevalence. The diagonal on the ROC plot is chance performance; the marked operating point is the threshold chosen for deployment. On an imbalanced problem the precision–recall curve exposes the precision cost of raising recall that the ROC view can mask.

4A.2.2 HITL Band as Precision-Recall Bridge

A two-threshold HITL system $(\tau_{\text{low}}, \tau_{\text{high}})$ with human review for scores in $(\tau_{\text{low}}, \tau_{\text{high}})$:

$$\text{Automated TPR} = P(s(x) \geq \tau_{\text{high}} \mid y = 1)$$

$$\text{Automated FPR} = P(s(x) \geq \tau_{\text{high}} \mid y = 0)$$

$$\text{Human Review Rate} = P(\tau_{\text{low}} \leq s(x) < \tau_{\text{high}})$$

This system achieves high precision in the automated band (only scores above τ_{high} are automatically positive) at the cost of sending the uncertain band to humans. Optimizing the band width trades reviewer workload against precision/recall tradeoffs.

4A.3 Cost-Sensitive Learning

4A.3.1 Class Weighting

Most ML frameworks allow cost-sensitive training via class weights (Elkan, 2001). For a binary classifier with $C_{FN} \neq C_{FP}$:

```
1 # PyTorch
2 weight = torch.tensor([C_FP, C_FN])
3 loss_fn = nn.CrossEntropyLoss(weight=weight)
4
5 # scikit-learn
6 from sklearn.linear_model import LogisticRegression
7 model = LogisticRegression(class_weight={0: C_FP, 1: C_FN})
```

Listing 4.1: Cost-sensitive training in PyTorch and scikit-learn

This biases the model to minimize the more costly error during training, effectively shifting the optimal threshold without post-hoc adjustment.

4A.3.2 Calibrated Decision Rules

Given a well-calibrated model with output $p(x) = P(y = 1 \mid x)$:

```
1 def optimal_decision(score, c_fp, c_fn, margin=0.1):
2     """
3     Make optimal decision given cost structure.
4
5     Returns: 'positive', 'negative', or 'uncertain' (for
6             HITL routing)
7     """
8     # Optimal threshold: score where cost_positive =
9         cost_negative
10    tau_optimal = c_fp / (c_fp + c_fn)
```

```

9
10  # HITL band: add uncertainty margin around optimal
    threshold.
11  # Note: margin is a design parameter, not derived from
    the cost ratio -
12  # it trades reviewer workload against the precision of
    automated decisions.
13  if score >= tau_optimal + margin:
14      return 'positive'
15  elif score <= tau_optimal - margin:
16      return 'negative'
17  else:
18      return 'uncertain'  # Route to human review

```

Listing 4.2: Optimal decision rule with HITL routing band

4A.4 Multi-Threshold HITL Systems

4A.4.1 Three-Way Classification

A HITL band system partitions the score space into three regions:

Score range	Action
$[0, \tau_L)$	Automatic NEGATIVE (high-confidence negative)
$[\tau_L, \tau_H]$	HUMAN REVIEW (uncertain)
$(\tau_H, 1]$	Automatic POSITIVE (high-confidence positive)

Optimization. Find τ_L and τ_H to minimize total system cost:

$$\begin{aligned}
\text{Cost}(\tau_L, \tau_H) = & C_{\text{FN}} \cdot P(s(x) < \tau_L \mid y = 1) \cdot P(y = 1) \quad [\text{missed positives}] \\
& + C_{\text{FP}} \cdot P(s(x) > \tau_H \mid y = 0) \cdot P(y = 0) \quad [\text{false alarms, automated}] \\
& + C_H \cdot P(\tau_L \leq s(x) \leq \tau_H) \quad [\text{human review cost}] \\
& + C_{H,\text{FN}} \cdot P(\text{human misses } y = 1 \text{ in band}) \\
& + C_{H,\text{FP}} \cdot P(\text{human false alarms in band})
\end{aligned}$$

where C_H is the per-case review cost and $C_{H,\text{FN}}$, $C_{H,\text{FP}}$ account for human errors in the review band.

4A.4.2 Capacity Constraints

When human review has a capacity constraint (max reviews per unit time $= R_{\text{max}}$):

$$P(\tau_L \leq s(x) \leq \tau_H) \cdot \text{arrival_rate} \leq R_{\text{max}}$$

Threshold and Reviewer Capacity

As arrival rate increases (busy periods), the band $(\tau_H - \tau_L)$ must narrow — fewer cases go to human review, more are handled autonomously. This is the mathematical basis for the observation in Chapter 4 that threshold settings must adapt to reviewer capacity. Ignoring this constraint leads to alert fatigue, which degrades the quality of every human decision in the review queue.

4A.5 Exercises

Exercise 4A.1 — Optimal Threshold for Fraud Detection

A fraud detection system achieves $AUC = 0.94$. The base rate of fraud is 0.2%, $C_{FN} = \$500$ (average fraud amount), $C_{FP} = \$2$ (customer service cost of false decline). Compute the optimal threshold τ^* using the formula from Section 4A.1.2. Then explain why the false positive rate at τ^* cannot be determined from AUC alone, and what additional information you would need to estimate it. What does this imply for the operational feasibility of deploying at τ^* ?

Exercise 4A.2 — Dynamic Threshold Under Queue Pressure

A content moderation system has a human review capacity of 1000 cases/hour. The arrival rate of flagged content is 800/hour normally, but spikes to 2000/hour during breaking news events. Design a dynamic threshold system that adjusts τ_L and τ_H based on current review queue length. What are the consequences for each error type during a high-traffic period?

Exercise 4A.3 — HITL Band Simulation

Implement the three-way classification system from Section 4A.4.1 in Python. Create a simulation where you vary the margin parameter and observe how the review rate, false negative rate, and false positive rate change. Find the margin that minimizes total cost for a given cost structure.

Exercise 4A.4 – ROC vs. Precision-Recall on Imbalanced Data

Draw the ROC curve and precision-recall curve for the same model on an imbalanced dataset (1% positive class). Identify the operating point on each curve that corresponds to the optimal threshold from Section 4A.1.2. Why does the ROC curve look more favorable than the PR curve for this scenario?

4A.6 Recommended Reading

Foundational

- Elkan, C. (2001). The foundations of cost-sensitive learning. *IJCAI*, 17(1), 973–978.
- Fawcett, T. (2006). An introduction to ROC analysis. *Pattern Recognition Letters*, 27(8), 861–874.

Calibration and Decision Rules

- Zadrozny, B., & Elkan, C. (2001). Obtaining calibrated probability estimates from decision trees and naive Bayesian classifiers. *ICML*.
- Davis, J., & Goadrich, M. (2006). The relationship between precision-recall and ROC curves. *ICML*.

Unified Treatment

- Hernández-Orallo, J., Flach, P., & Ferri, C. (2012). A unified view of performance metrics: Translating threshold choice into expected classification loss. *Journal of Machine Learning Research*, 13, 2813–2869.

This appendix provides the mathematical foundations for threshold optimization, ROC and precision-recall analysis, cost-sensitive learning, and multi-threshold HITL band design as introduced in Chapter 4. The exercises connect the theory to concrete design decisions in fraud detection, content moderation, and other high-stakes HITL systems.

5 The Art of Asking for Help (Without Being Annoying)

The art of asking for help. Not whether to ask. How.

— on the art of asking

5.1 A Better Question

At the start of a life-or-death call, a 911 dispatcher faces a deceptively simple task: find out whether the patient is breathing. It sounds binary—either someone is breathing or they are not. But cardiac arrest does not always look the way people expect. In the first minutes after the heart stops, many patients produce **agonal breathing**: slow, gasping, reflexive breaths that are not real breathing at all, yet look and sound enough like it to fool a frightened bystander.

For years this created a deadly ambiguity. When a dispatcher asked, “Is he breathing?”—the natural, obvious question—a caller watching a loved one gasp would answer “yes.” The dispatcher, hearing “yes,” would not recognize a cardiac arrest, would not begin telephone CPR instructions, and would not get a bystander pushing on the chest in the seconds that matter most. The single most important intervention a dispatcher can trigger was being lost to one word in a question.

Researchers studying out-of-hospital cardiac arrest in Dallas introduced a revised dispatch protocol built around better asks. Instead of

letting a caller’s “yes” settle the question, dispatchers were given concrete assessment steps: have the caller count the patient’s breaths, ask them to describe what the breathing sounds like—in effect, to judge whether it sounds *normal*—and have them hold the phone next to the patient so the dispatcher can hear for themselves (Roppolo et al., 2009).

Each of these steps does the same work. It forces the caller to judge the *quality* of the breathing rather than answer a yes/no question they had already half-answered in their head. A gasping, agonal breath is many things, but it is not normal breathing—and once a caller is asked to count it, describe it, or hold the phone close enough for the dispatcher to hear it, they hesitate, they report what is actually happening, and the dispatcher can recognize the arrest.

After the protocol changed, the share of cardiac-arrest patients mistakenly thought to be breathing normally fell, and dispatchers who had detected almost no agonal breathing in the months before the change were identifying it case after case afterward—each instance a chance to start CPR minutes earlier (Roppolo et al., 2009).

No new technology. A handful of small changes to the questions dispatchers ask hundreds of times a day changed how callers understood what they were being asked—and changed what information the dispatchers got back.

This is the art of asking for help. Not whether to ask. How.

5.2 The Ask Is a Design Problem

Everything we have established in the previous chapters points toward a moment: the moment when a system detects that it is uncertain, evaluates that the case is worth routing to a human, and then... asks. Or alerts. Or flags. Or notifies. Or interrupts.

That moment is where most HITL systems fail, even when everything before it has been done right.

Chapter 1 gave us the framework: five dimensions, three failure modes. Chapter 2 explained that confusion itself is useful information. Chapter 3 showed that learning systems produce predictable patterns of uncertainty at the edges of their training. Chapter 4 introduced the threshold problem and the human value judgments embedded in it.

All of that engineering, all of that calibration, all of that careful cost analysis—it delivers the system to the moment of the ask. And then a poorly designed question throws it all away.

The fraud alert that comes at 3 AM and wakes you from sleep to ask about a charge you made yourself that morning. The spam filter notification that flags the same newsletter for the hundredth time. The pop-up that asks for your permission to access your location before you have even understood what the app is for. The car alarm that fires whenever a large truck passes.

Each of these is a system that detected uncertainty correctly, made a reasonable threshold decision, and then asked badly. The ask was the failure.

5.3 Three Levers

The design space for an ask has three primary dimensions:

Timing. When does the ask happen? Is it synchronous—immediately interrupting what the person is doing—or asynchronous, deferred until a natural break? Is it too early (before the person has the information needed to answer) or too late (after the decision has already been made by default)?

Frequency. How often does the ask happen? Is it rare enough to command attention, or common enough to become noise? Does the system have a mechanism to prevent the same person from being asked the same type of question too many times?

Format. How is the ask phrased? What choices does it offer? How much cognitive work does it require? Does the phrasing bias the answer?

These are not independent. A well-timed, infrequent ask can survive imperfect format. A perfectly formatted ask at the wrong moment fails. The goal is to optimize all three together, and to understand how the human's context—what they are doing, what they are attending to, what they have already been asked—shapes every dimension of the optimal ask.

5.4 The Trust–Attention–Response Triangle

Human psychology and the design of alerts interact in a system that has its own dynamics. Call it the **trust–attention–response triangle**.

When a system asks for human input, three things are required: the human must **trust** that the ask is worth responding to; the human must have **attention** available to process it; and the human must be willing and able to **respond** effectively.

Each of these can fail independently.

Trust failure. A system that has cried wolf too many times loses the trust that makes its alerts actionable. The alerts still arrive. The human learns to dismiss them. This is the clinical **alert fatigue** problem described in Chapter 1, but it applies everywhere. The fire alarm that fires so frequently that residents stop evacuating. The fraud alert service that flags so many routine purchases that cardholders start ignoring the texts. The over-alerting system trains its own users to be non-users.

Attention failure. A system that asks at the worst possible moment gets ignored not because the human does not trust it but because the human literally cannot process the ask right now. The update prompt that fires during a video call. The security notification that arrives mid-presentation. Timing failures are attention failures.

Response failure. A system that asks a question the human cannot answer effectively gets useless or harmful responses. A fraud alert that asks “Was this transaction legitimate?” when you cannot remember whether you made it, at an amount that is not extraordinary—you will answer randomly or default to “yes” without thinking. The ask’s format must give the human the information they need to answer accurately.

The trust–attention–response triangle is not static. It changes over time, across users, and across contexts. A system that maintains all three is rare. A system that degrades all three is common.

5.5 The Lazy User Assumption

Here is a design assumption that underlies most bad alert design, usually implicitly: *users will not try*.

The premise behind asking questions with six items to review, or requiring users to navigate three menus to respond to a notification, or sending a fraud alert in a format that requires calling a phone number—all of these assume that the user is lazy and that the friction is acceptable because dedicated users will push through it.

This assumption is wrong in two ways.

First, it is empirically incorrect. Users in well-designed systems routinely engage with alerts, respond to notifications, and participate in feedback mechanisms—when the design makes it easy. The problem is not user

laziness. The problem is designers treating users as if they have infinite attention and zero competing demands.

Second, the assumption fails even on its own terms. A willing user given a four-step response format is not four times as likely to respond as one given a single step. The friction does not select for motivated users. It just loses them.

The Design Principle

Good HITL intervention design starts from the opposite assumption: *users will try if the ask is worth their time*. The design challenge is making the ask worthy.

“Worth their time” means: the ask is rare enough to be surprising; it arrives at a moment when the human can actually respond; the format makes the question clear and the response easy; and the human has reason to believe their response will actually matter.

5.6 Two-Factor Authentication: Lessons from the Field

Two-factor authentication (2FA) is perhaps the most widely deployed HITL mechanism in existence. Every time a website sends you a six-digit code and asks you to type it in, it is inserting a human into a loop that would otherwise be purely automated.

The security rationale is sound: even if an attacker has your password, they are unlikely to also have your phone. The human in the loop—you, confirming your identity by entering a code only you could receive—provides a verification layer that pure password authentication cannot.

But 2FA’s implementation landscape is a masterclass in how to design the ask both well and badly.

Poorly designed 2FA.

- Push-approval prompts where a single tap approves the login—so-called *MFA fatigue attacks* exploit users who approve reflexively just to make the prompts stop
- SMS codes sent to numbers that may be SIM-swappable
- Time pressure so aggressive that the window expires before the user finds their phone
- No context provided: “Your authentication code is 847291”—but for what? From where?
- Prompts sent at 3 AM that train users to approve requests without fully waking up

Well-designed 2FA.

- Provides context: “Someone is trying to sign into your account from a new device. If this was you, enter this code: [code]”
- Number matching: the login screen shows a number and the push notification asks “Did you see the number [X]?”—defeating reflexive approval
- Adaptive authentication: only prompts for 2FA when the login context is unusual (new device, new location, unusual time)
- Clear, low-friction response options with a prominent “I did not initiate this login” option

The difference between these two designs is not the technology. It is the application of the trust–attention–response triangle. Good 2FA provides the context the user needs (trust), arrives when the user is actively trying

to log in (attention), and makes both “approve” and “deny” equally easy (response).

5.7 The Fraud Alert as Design Canvas

Bank fraud alerts are another rich case study, because the same underlying technology—“this transaction looks unusual, confirm it was you”—can be implemented in ways that range from excellent to actively harmful.

An Excellent Fraud Alert

- Sent within seconds of the transaction, while you are still in the situation
- Contains specific, identifiable information: the merchant name, the exact amount, the last four digits of the card
- Asks a simple binary question: “Did you make this \$47.23 purchase at Costa Coffee? Reply YES or NO”
- Takes one action (a single text reply) to confirm
- If you reply NO: immediate card freeze plus one-tap call option to the fraud line
- If you do not reply within 10 minutes: gentle follow-up, then escalation protocol

- Sent six hours after the transaction, when context is gone

- Contains: “Your account has been flagged for unusual activity”
- Requires you to call a number, navigate an IVR system, verify your identity, explain the situation
- If the transaction was legitimate: 20 minutes of your life gone
- Teaches you that fraud alerts are time-consuming, training you to ignore the next one

The first design respects the trust–attention–response triangle at every step. The second violates all three. And the second is, unfortunately, still common.

5.8 Content Moderation and the Interface of Judgment

Content moderation is a domain that makes the design of the ask visible in ways that other HITL applications do not, because the “human” in the loop is often a professional reviewer whose job is entirely about responding to system-generated asks.

Professional content moderators at major platforms spend their days responding to a stream of flagged content: videos, images, posts, comments—each one a system-generated ask requiring a human response. The design of the interface they use shapes the quality of every response they give (Roberts, 2019).

Research into content moderation interfaces has documented several consistent findings.

Context improves accuracy. Moderators shown content with its original context (the thread, the account history, the apparent intent) make more accurate decisions than those shown the same content in isolation. But showing context slows throughput. The ask’s format creates a tradeoff between speed and quality.

Category granularity matters. Interfaces with many specific categories (“graphic violence: real vs. artistic context,” “hate speech: individual vs. systemic”) produce more consistent and useful decisions than interfaces with coarse categories (“violent,” “non-violent”). But more categories require more cognitive work per decision.

Throughput pressure wears reviewers down. The literature on content moderators links heavy workloads and sustained exposure to disturbing material to degraded well-being—fatigue and strain that accumulate across a shift and that consistent, careful judgment can ill afford (Steiger et al., 2021). The ask format cannot compensate for structural workload problems.

Explanation requests improve downstream model quality. When moderators are asked not just “is this violating?” but “what makes this violating?” (even briefly), the resulting labels are more useful for model training (Zaidan et al., 2007). The ask format shapes what information is captured, not just what decision is made.

Each of these findings is a design lesson. The ask format is not neutral. It actively shapes the information the system gets back.

5.9 The 911 Dispatcher, Revisited

Return to the dispatcher's questions. A handful of small changes to the ask—count the breaths, describe the breathing, hold the phone close—changed how often a cardiac arrest was recognized in time.

Why did these changes matter so much?

The original question (“Is he breathing?”) asked the caller for a yes/no fact they thought they already knew. A gasping patient *is* moving air, after all; to a panicked bystander, that reads as breathing. The framing let them confirm what they wanted to believe.

The revised asks demanded a judgment rather than a fact. Describing the breathing, counting it, holding the phone close enough to hear it—each made the caller compare what was in front of them against what breathing is supposed to look and sound like, and agonal gasping fails that comparison. The new prompts turned a reflexive “yes” into an act of observation.

These are not just clarity improvements. They are changes to what information the question extracts. The question's format shapes the answer's content.

The Core Insight

The question you ask determines the information you get. A vague question gets vague information. A question framed to invite confirmation gets confirmation. A question framed to invite nuance gets nuance.

And in a HITL system, the quality of the information the human provides determines how much the system improves. A question designed to extract maximum useful information from a minimum human effort is not just a UX nicety—it is the core engineering of the feedback loop.

5.10 Feedback Integration: Does It Learn?

There is a fifth dimension in our framework that often gets overlooked in the design of individual asks: **Feedback Integration**—whether the system actually learns from the human’s response.

The ask is useless if the response disappears into a void. If a fraud alert’s “NO” response does not update the model’s representation of the user’s spending patterns, the same type of transaction will keep getting flagged. If a content moderator’s “not violating” decision on a borderline case does not inform how similar cases are classified in the future, the system is simply routing work rather than learning.

The design of the ask and the design of the feedback loop are not separate problems. The ask’s format determines what information is captured. The feedback loop determines what happens with that information.

Well-designed HITL systems close this loop explicitly:

- The user’s response is labeled and stored with the full context of the ask
- Cases where human decisions systematically differ from model predictions are flagged for calibration review
- The frequency of human corrections in a particular category is used to prioritize model improvement
- The human reviewer occasionally receives confirmation that their feedback mattered: “Your corrections in this category helped improve accuracy by X%”

That last point—feedback to the feedback provider—may seem like a courtesy. It is actually a mechanism for maintaining engagement. A human reviewer who never sees evidence that their responses matter eventually

stops trying. The trust–attention–response triangle degrades from the trust side.

For the next week, pay attention to every digital alert or notification that asks you to take an action. For each one, ask three questions: (1) Is this well-timed—are you in a position to respond effectively right now? (2) Is this well-framed—do you understand exactly what is being asked and have what you need to answer? (3) Does your response matter—does it feel like your answer will change anything? Keep a log of your answers. You will quickly develop a sense for which systems are designed to get useful human input, and which are designed to appear to involve humans without actually extracting useful judgment.

Chapter Summary

Key Concepts

- The design of the ask is as important as the decision to ask—a technically correct routing decision can be wasted by a poorly designed intervention
- Three primary levers: **timing, frequency, format**
- The trust–attention–response triangle: all three must be maintained for HITL to work
- The Lazy User Assumption is wrong: users will engage with asks that respect their time and attention

- The question's format determines the information captured—not just what decision is made
- Feedback Integration closes the loop: the ask is wasted if the response does not improve the system

Key Examples

911 dispatch (cardiac arrest) Asking “Is he breathing *normally*?” instead of “Is he breathing?” helped dispatchers catch agonal breathing and recognize cardiac arrest in time

Two-factor authentication The same security mechanism implemented well vs. poorly—context and number matching vs. fatigue-inducing push approval

Bank fraud alerts Specific, timely, binary ask vs. generic, delayed, high-friction request

Content moderation interfaces Context, category granularity, and throughput pressure all shape decision quality

Key Principles

- Timing must align with the human's attention window—ask when they can answer
- Frequency must be low enough to preserve the signal value of each ask
- Format must give the human what they need to answer accurately in minimum steps

- Good design assumes users will try; bad design assumes they will not
- The ask and the feedback loop are one system, not two separate problems

Five Dimensions Check

- *Intervention Design* (Dimension 2): The central subject of this chapter—the design of the ask itself, not merely the decision to ask
- *Timing* (Dimension 3): One of the three primary levers; the ask must fall within the human’s attention window
- *Feedback Integration* (Dimension 5): The ask and the response are one system—the ask is wasted unless the answer improves the model

This chapter closes the first part of our journey: from understanding that smart systems need to ask for help, through understanding what confusion is and where it comes from, to understanding how to turn uncertainty into decisions, and finally to understanding how to make the ask itself effective. The next part turns to specific domains where HITL design matters most.

Technical Appendix 5A: Optimal Ask Design

Chapter 5 introduced three design levers for human-in-the-loop asks: timing, frequency, and format. This appendix provides mathematical frameworks for optimizing each, drawing on information theory, cognitive psychology, and reinforcement learning.

— Formal foundations for Chapter 5

5A.1 Information-Theoretic Foundations of Ask Design

5A.1.1 Information Value of a Human Response

The value of asking a human is the expected reduction in uncertainty about the correct action. Let Y be the correct decision (unknown to the system), H the human's response, and X the system's current information. The **information gain** from asking is:

$$\text{IG}(H; Y | X) = \mathbb{H}(Y | X) - \mathbb{H}(Y | H, X)$$

where $\mathbb{H}(\cdot)$ denotes entropy. A high information gain means the human's response substantially reduces uncertainty about the right action.

Optimal ask criterion. Ask a human when the expected value of the information gained exceeds the cost of asking:

$$V_{\text{decision}} \cdot \text{IG}(H; Y \mid X) > C_{\text{ask}}$$

where C_{ask} is the cost (time, attention, trust) and V_{decision} is the value of resolving uncertainty about the decision, per bit. Since IG is measured in bits, both sides are in the same units; equivalently, ask when $\text{IG}(H; Y \mid X) > C_{\text{ask}}/V_{\text{decision}}$.

5A.1.2 Query Optimisation

The **optimal query** Q^* maximizes information gain subject to cognitive load constraints:

$$Q^* = \arg \max_{Q \in \mathcal{Q}} [\text{IG}(H_Q; Y \mid X) - \lambda \cdot \text{CogLoad}(Q)]$$

where $\text{CogLoad}(Q)$ is the cognitive effort required to answer query Q and λ is a cost–effort tradeoff parameter.

Simplicity vs. Information

A simpler query (lower cognitive load) is only preferable if its information loss is proportionally small. The 911 example from Chapter 5 illustrates the flip side: a one-word change (“breathing” to “breathing *normally*”) added almost nothing to the caller’s effort, yet sharply increased the information the question extracted—a negligible cost for a large gain.

5A.1.3 Human–AI Complementarity

The value of the full human–AI system depends on the *complementarity* between human and AI knowledge:

$$I(H_{\text{human}}; Y \mid H_{\text{AI}}, X) = \mathbb{H}(Y \mid H_{\text{AI}}, X) - \mathbb{H}(Y \mid H_{\text{human}}, H_{\text{AI}}, X)$$

This is the additional uncertainty reduction from the human, given the AI has already processed X . High complementarity means the human knows things the AI does not—this is when HITL intervention is most valuable.

5A.2 Alert Fatigue: A Dynamic Model

5A.2.1 Trust as a State Variable

Let $\tau_t \in [0, 1]$ represent the human's trust in the system's alerts at time t . Trust evolves dynamically:

$$\tau_{t+1} = \tau_t + \alpha \cdot \mathbf{1}[\text{alert}_t \text{ was valid}] - \beta \cdot \mathbf{1}[\text{alert}_t \text{ was false}]$$

where α is the trust gain from a true alert and β is the trust loss from a false alarm. Empirically $\beta > \alpha$: trust erodes faster than it builds (J. D. Lee & See, 2004). Effective alert value therefore decays with sustained false-positive rates:

$$V_{\text{effective}}(t) = \tau_t \cdot V_{\text{intrinsic}}$$

5A.2.2 Precision–Fatigue Tradeoff

Define r as the alert rate, $p(r)$ as precision (decreasing in r), $\tau(r)$ as equilibrium trust (also decreasing in r), and $\rho(\tau)$ as response probability. The effective information extracted per hour:

$$I_{\text{effective}}(r) = r \cdot p(r) \cdot \tau(r) \cdot \rho(\tau(r)) \cdot V_{\text{per alert}}$$

```

1 from scipy.optimize import minimize_scalar
2
3 def effective_information(r, precision_decay=0.1,
4                           trust_decay=0.05,
5                           base_value=1.0):
6     precision = 1 / (1 + precision_decay * r)
7     false_alarm_rate = r * (1 - precision)
8     trust = 1 / (1 + trust_decay * false_alarm_rate)
9     response_prob = trust ** 0.5
10    return r * precision * trust * response_prob *
11        base_value
12
13 result = minimize_scalar(
14     lambda r: -effective_information(r),
15     bounds=(0.1, 100), method='bounded'
16 )
17 print(f"Optimal alert rate: {result.x:.2f} alerts/hour")

```

Listing 5.1: Effective information vs. alert rate

5A.2.3 Habituation and Recovery

Consecutive alerts induce faster degradation:

$$\tau_t = \tau_{\text{baseline}} \cdot e^{-\gamma \cdot n_{\text{recent}}}$$

where n_{recent} is the number of alerts in the recent window. Recovery is slower:

$$\tau_{t+T} = \tau_t + (1 - \tau_t)(1 - e^{-\delta T})$$

Spacing alerts recovers attention; clustering alerts degrades it faster than average-rate models predict. A well-designed HITL interface enforces a minimum inter-alert gap, not just a maximum average rate.

5A.3 Optimal Timing Models

5A.3.1 Attention Windows

Let $a(t) \in [0, 1]$ be the human's attention availability. The effective response quality for an alert sent at time t :

$$Q_{\text{response}}(t) = Q_{\text{max}} \cdot a(t) \cdot \tau(t)$$

Optimal timing: $t^* = \arg \max_t Q_{\text{response}}(t)$. Interruption at natural task breakpoints consistently produces higher response quality than mid-task interruption.

5A.3.2 The Window of Relevance Problem

Context availability decays after an event:

$$C(\Delta t) = e^{-\kappa \Delta t}$$

for decay rate κ depending on event type. Optimal alert delay balances context availability against attention window:

$$t^* = \arg \max_t a(t + t_{\text{event}}) \cdot C(t)$$

For financial transactions κ is large, so $t^* \approx 0$. Asking about a charge 24 hours later yields near-chance accuracy.

5A.4 Format Design: Cognitive Load and Response Accuracy

5A.4.1 Cognitive Load Model

Following Sweller’s cognitive load theory (Sweller, 1988), response accuracy degrades beyond working memory capacity W :

$$A(c, W) = \begin{cases} A_{\max} & \text{if } c \leq W \\ A_{\max} \cdot (W/c)^\alpha & \text{if } c > W \end{cases}$$

5A.4.2 Optimal Query Granularity

A k -choice query captures $\log_2 k$ bits at $O(k)$ cognitive cost:

$$k^* = \arg \max_k \frac{\log_2 k}{k^\beta \cdot C_{\text{per-choice}}}$$

For most practical settings $k^* \in [2, 5]$, depending on the complexity exponent β . At the $\beta \approx 0.6\text{--}0.8$ typical of simple perceptual tasks, 4–5 choices represent the practical optimum. At higher β (e.g., $\beta = 1.5$ as in the exercise below), $k^* = 2$ – binary questions dominate.

5A.5 Feedback Integration: Closing the Learning Loop

5A.5.1 Active Learning Formulation

Each human response is a labeled data point. Select the next query to maximize expected model improvement:

$$Q_{\text{next}}^* = \arg \max_Q \mathbb{E}[\Delta \text{ModelPerf} \mid H_Q, \mathcal{D}_{\text{current}}]$$

In practice this is approximated by querying cases near the decision boundary, in underrepresented categories, or where human and model systematically disagree.

5A.5.2 RL Formulation of the Full HITL Loop

The ask–learn loop formalised as an MDP (Christiano et al., 2017):

- **State** s_t : model weights, trust τ_t , queue depth, context
- **Action** a_t : ask with format Q , or act autonomously
- **Reward**: $r_t = \Delta \text{Acc}(a_t) - C_{\text{ask}} \cdot \mathbf{1}[a_t = \text{ask}] - \lambda \cdot \Delta \tau_t$

Optimal policy π^* maximizes $\mathbb{E}[\sum_t \gamma^t r_t]$, capturing the tension between short-run accuracy and long-run trust.

5A.6 Exercises

Exercise 5A.1 — Information Value Calculation

A spam filter has confidence 0.65 on an email. A human reviewer has 90% accuracy on boundary cases.

1. Calculate the expected information gain from asking the reviewer.
2. If the cost of asking is 30 seconds of reviewer time (valued at \$0.01/second), at what minimum email value must human review be worthwhile?
3. How does information gain change if filter confidence was 0.95?

Exercise 5A.2 — Alert Fatigue Simulation

Using the dynamic trust model:

```

1 import numpy as np
2
3 def simulate_alert_fatigue(n_alerts=300,
4                             precision=0.7,
5                             alpha=0.05, beta=0.15):
6     trust, history = 1.0, [1.0]
7     for _ in range(n_alerts):
8         is_valid = np.random.random() < precision
9         trust = min(1.0, trust + alpha) if is_valid \
10                  else max(0.0, trust - beta)
11         history.append(trust)
12     return history

```

Listing 5.2: Alert fatigue simulator

Simulate: (a) 10/hour at 70% precision; (b) 3/hour at 90% precision; (c) 20/hour at 90% precision. Compare equilibrium trust levels.

Exercise 5A.3 — Optimal Query Granularity

For a content moderation interface with 5 labels vs. 2 labels:

1. Calculate maximum information per response for each.
2. Assuming $\beta = 1.5$, compute k^* .
3. Design an experiment measuring actual information captured using inter-annotator agreement as a proxy.

Exercise 5A.4 — HITL Policy Optimisation

Implement a Q-learning agent learning when to ask. State: confidence (binned 0–10), trust (high/medium/low). Actions: auto-decide, binary query, detailed query. Reward: +10 correct, −1 ask cost, $-0.5 \times (1 - \text{trust})$ incorrect auto-decision. Simulate 10,000 episodes. At what confidence level does the optimal policy switch from asking to auto-deciding?

5A.7 Recommended Reading

Alert Design

- EEMUA (2013) — Engineering Equipment and Materials Users' Association: alarm system management guidelines.
- Drew et al. (2013) — Attention and awareness in simulated radiological search. *Psychological Science*.

Active Learning and HITL Feedback

- Settles (2009) — Active learning literature survey.
- Christiano et al. (2017) — Deep reinforcement learning from human preferences. *NeurIPS*.

This appendix provides the formal foundations for ask design, alert fatigue, timing, format, and feedback integration as introduced in Chapter 5.

Part II

The Human Side of the Loop

A HITL system is only as good as its worst interface.

— Chapter 6

6 Interfaces That Make Humans Smarter

A HITL system is only as good as its worst interface—the design of a screen can make a doctor brilliant, or helpless.

— on interfaces

6.1 The Judge and the Number

Imagine a judge in Broward County, Florida, opening a pre-sentencing report. Near the top is a score. The number was produced by COMPAS—Correctional Offender Management Profiling for Alternative Sanctions—an algorithmic risk assessment tool used in dozens of American jurisdictions. The score is supposed to tell the judge how likely the defendant is to reoffend. High risk or low risk. A number to guide the sentence.

What data produced it? Which factors drove the score up or down? Was it the defendant's age? Their prior record? Their neighborhood? The tool's vendor, Northpointe (now Equivant), considered those inputs proprietary. The score arrived without explanation, like a verdict without a trial.

The judge did what many judges do: they used the number. Research by ProPublica, published in 2016, would later show that Black defendants who did not reoffend were nearly twice as likely as comparable white defendants to be falsely classified as high risk (Angwin et al., 2016). A separate study

by Dressel and Farid, published in *Science Advances* in 2018, added an equally troubling finding: untrained laypeople—random people recruited online and shown only the same basic information—predicted recidivism about as accurately as the algorithm. COMPAS was no better at the task than a person on the street (Dressel & Farid, 2018). The algorithm had a race problem and a competence problem, both obscured by a single clean-looking integer. The ProPublica analysis was itself contested: a reanalysis of the same Broward County data argued that the tool’s predictive accuracy was comparable across groups, and that the fairness verdict depended on which metric—calibration or error rates—one chose to privilege (Flores et al., 2016). That dispute is the impossibility theorem this book returns to in Chapter 14, playing out in public: both sides measured fairness correctly, and differently.

But here is the part of the story that gets less attention: the judge was not irrational. They were responding, predictably and almost inevitably, to the way the information was presented. A number with no explanation exerts a strange psychological gravity. It feels more objective than a probation officer’s narrative. It feels more certain than it is. The interface between the algorithm and the human had not made the judge smarter. It had quietly replaced their judgment.

This is the subject of this chapter: not whether AI systems should assist human decision-makers, but *how*. The interface—the screen, the dashboard, the number, the recommendation—is not a neutral transmission channel. It shapes what humans notice, what they ignore, and ultimately what they decide (Norman, 2013). A badly designed interface does not just fail to help. It actively makes its human partner worse.

6.2 The Invisible Architecture of Attention

In 2017, a Stanford team led by Andrew Ng unveiled CheXNet, a deep learning system that could detect pneumonia from a chest X-ray about as well as a board-certified radiologist (Rajpurkar et al., 2017). The headlines celebrated the accuracy. But the more consequential discovery came later, from researchers studying how radiologists actually *used* systems like it—and the answer was not what anyone expected (Gaube et al., 2021).

In a preregistered experiment, Susanne Gaube, Harini Suresh, and their colleagues gave physicians—radiologists and emergency-medicine doctors—chest X-rays accompanied by diagnostic advice labeled as coming either from an AI system or from an experienced radiologist. Some of the advice was deliberately wrong. The purported source turned out to be nearly irrelevant: what the advice was *called* had no measurable effect on anyone’s diagnostic accuracy. What mattered was whether the advice was *correct*. When it was wrong, physicians followed it far too often regardless of the label—more than a quarter of the radiologists went along with the inaccurate advice every time they received it.

The label did move one thing: what the experts *said*. The radiologists rated the identical advice significantly lower in quality when it was attributed to the AI—a clear case of algorithmic aversion in their stated judgments—and then relied on that same advice just as much as on the human-labeled kind. Their expressed skepticism and their actual behavior came apart.

The Say-Do Gap

Ask experts whether they trust AI advice and they say they trust it less than a colleague’s. Watch what they *do*, and the gap closes: bad advice is followed just as readily whatever its purported source. Stated attitudes are not revealed reliance. Interfaces must be de-

signed for what humans actually do—not for what they say they will do.

This **say–do gap** between expressed attitudes and revealed reliance is what makes the architecture of attention so important. What arrives as a recommendation, what arrives as a probability, what requires a human to click to see the evidence behind a recommendation—each of these choices shapes what the human notices, and designers cannot rely on the humans involved to report how those choices are affecting them.

6.3 Automation Bias and the Seduction of the Algorithm

In 2023, a team led by Thomas Dratsch published a study with an uncomfortable finding. Twenty-seven radiologists read mammograms and assigned each a BI-RADS category with the help of what they were told was an AI system. For some images the AI's suggestion was correct; for others it was deliberately wrong. The question was whether the radiologists would catch the AI's mistakes.

When the AI was correct, it helped. When the AI was wrong, however, a distressing number of radiologists followed it anyway. Accuracy collapsed under an incorrect suggestion—even experienced readers fell from correctly categorizing about 82% of cases to roughly 46%, and less experienced readers dropped from nearly 80% to under 20%. The AI's suggestion was overriding the radiologists' own judgment (Dratsch et al., 2023).

This is **automation bias**: the tendency to over-rely on automated recommendations, particularly when the human is uncertain or fatigued. It is not a character flaw. It is a documented cognitive response to the presence

of a system that presents itself as authoritative. The more confident-looking the recommendation, the more powerful the bias.

Automation bias was first systematically documented in aviation, where it was called “automation complacency.” Pilots in highly automated cockpits were found to be less vigilant, less likely to notice anomalies, more likely to accept erroneous automated suggestions (Parasuraman & Manzey, 2010). The same pattern appears in air traffic control, nuclear power plant monitoring, and now, everywhere AI recommendations are shown to human reviewers.

The unsettling implication is that AI assistance can make expert humans worse. Not in every case, and not in every domain, but under the right conditions—a fatigued reviewer, a plausible-looking error, a confident-seeming display—the AI’s recommendation can function as a kind of anchor that pulls human judgment toward it even when judgment should pull away.

One partial defense is not to hide the AI recommendation. It is to design the interface so that the human forms their own judgment before seeing the AI’s, or is required to articulate a reason when they accept the AI’s recommendation over their own initial assessment. Research on automation bias in simulated high-tech cockpits found that operators given accountability instructions—told they would have to justify their decisions—double-checked the automation more often, though the effect on error rates was mixed (Skitka et al., 1999). How the human is asked to engage—accountably, and on their own judgment first—turns out to matter as much as what the interface shows.

6.4 Death by a Thousand Clicks

In her ethnography of commercial content moderation, Sarah Roberts documented the working conditions at firms contracted by the largest social media platforms (Roberts, 2019). The moderators she observed reviewed hundreds of posts per shift—posts flagged by platform AI as potentially violating community guidelines. Their interfaces presented complex, multi-level violation taxonomies with dozens of categories and sub-categories. Flagging a post for “coordinated inauthentic behavior” could require navigating to a separate tool, entering a case ID, and then returning to the original queue. Every additional click consumed time that was not spent on judgment.

Average review times were measured in seconds per post. The measurement itself was a design decision: it pushed moderators toward speed over scrutiny.

The problem is not that the time per post is too short in some abstract sense. The problem is that a large fraction of that time is consumed by interface mechanics—clicking, navigating, confirming, entering IDs—rather than by actual judgment about the content. The throughput pressure is real: moderators on contracts paying by volume face an incentive to process quickly, and the interface adds its own tax on top of that pressure (Stureborg et al., 2023).

The specific cognitive cost of operating a complex interface under time pressure is well-documented in vigilance research: sustained attention on repetitive classification tasks degrades over time, and every unnecessary click is a small withdrawal from a limited attention budget.

The Interface Changes the Labels

In a controlled study of crowdsourced text annotation, Stureborg and colleagues tested six different interface designs for the same hierarchical labeling task, collecting more than eighteen thousand individual annotations. The design alone changed the quality of the labels: interfaces that grouped related concepts and made the label hierarchy visible improved workers' accuracy substantially, with the largest gains on the hardest examples (Stureborg et al., 2023). The annotators were drawn from the same pool, looking at the same content, applying the same guidelines. The interface had made them better at their job.

The fix is rarely glamorous. It is usually keyboard shortcuts, sensible defaults, streamlined workflows, and a ruthless audit of which clicks actually serve the quality of the human decision and which merely serve the recordkeeping needs of the system. But it requires asking, seriously and specifically: what cognitive work do we actually want the human to do here? And then, for every interface element: does this support that work, or impede it?

6.5 What Explainable AI Actually Looks Like

The COMPAS story prompted a wave of interest in what researchers call **Explainable AI**, or XAI. The basic idea is simple: AI systems should be able to explain why they produced a particular output. A judge should be able to see not just a risk score but the factors that drove it. A doctor should see not just a diagnosis probability but the regions of the scan that drove the model's confidence.

The reality of XAI in practice is messier.

The most common approach is **LIME**—Local Interpretable Model-Agnostic Explanations, developed by Marco Ribeiro and colleagues at the University of Washington (Ribeiro et al., 2016). LIME works by generating a simpler model around a specific prediction that approximates the complex model’s behavior in that local region.

The problem is that LIME explanations can be unstable. Two inputs that are nearly identical—two scans a radiologist would call the same—can receive noticeably different explanations of which features mattered. Researchers have shown that small perturbations of the input can shift an explanation enough that a human reviewer comparing the two outputs might come away with a different picture of what the model was doing (Alvarez-Melis & Jaakkola, 2018).

A competing approach, **SHAP** (SHapley Additive exPlanations), provides more mathematically grounded explanations based on game-theoretic principles (Lundberg & Lee, 2017). SHAP is more consistent but more computationally expensive. Neither LIME nor SHAP tells you whether the model is actually reasoning correctly—they show you which features the model is using, not whether those features are the right ones to use.

For radiology AI, the most common XAI approach is the attention map or **Grad-CAM**—a visualization showing which pixels in an image most influenced the model’s prediction (Selvaraju et al., 2017). Studies by Zech and colleagues at Mount Sinai showed that early pneumonia models were sometimes attending to text tokens burned into the image by portable X-ray scanners—a proxy for scans taken on inpatient wards and in the emergency department, where pneumonia cases are concentrated (Zech et al., 2018). The model had learned a spurious correlation, and Grad-CAM made it visible.

An explanation is only useful if the human receiving it can actually act on it. A heatmap that a radiologist cannot interpret provides no safety benefit. A SHAP plot that a loan officer cannot read does not protect the borrower. Explanation design is interface design, and all the same rules apply: what cognitive work do we actually want the human to do?

6.6 Interfaces That Support Judgment vs. Interfaces That Replace It

Imagine a spectrum. At one end: an interface that shows the human only the AI's final recommendation—"Approve," "Deny," "High Risk," "Low Risk." The human has information about the conclusion but not the reasoning. This is the COMPAS interface.

At the other end: an interface that shows the human all the raw information the AI has access to, all the uncertainty estimates, all the feature contributions, a full audit trail. This is maximal transparency, and it is also practically useless. No human can process all that information in the time available for a real decision.

Between these extremes is a design space that has been surprisingly little explored until recently. The best interfaces in this space share a few properties.

They separate the AI's job from the human's job. The AI handles pattern recognition and data synthesis; the human handles judgment about edge cases, context, and stakes. The interface makes this division visible: here is what the AI found, here is where the AI is confident, here is where the AI is not.

They present uncertainty as information, not as failure. A confidence level of 62% is not the AI apologizing—it is the AI directing the human’s attention to a case that genuinely requires human judgment. An orange flag on uncertain cases and a green flag on confident ones tells the human not “the AI isn’t sure,” but “the AI is inviting your expertise here.”

They preserve the human’s capacity for independent judgment. Some interface designs are so anchored around the AI recommendation that the human reviewer finds themselves responding to the AI rather than to the underlying case. Consider legal document review: a reviewer shown only a document summary, without the original, is more likely to follow that summary even when it is wrong than a reviewer who sees both—the interface has quietly substituted the AI’s reading for the human’s.

They make disagreement easy. If a reviewer wants to flag disagreement with the AI, the interface should make that disagreement easy to record and route to appropriate review. Many current HITL interfaces make agreement a single click and disagreement a three-menu navigation. This asymmetry is a thumb on the scale.

6.7 The Radiology Interface Wars

Few fields have experimented more intensively with AI-assisted human decision-making interfaces than radiology.

Consider how a confidence display can backfire. A natural design choice is to show the AI’s confidence as a horizontal bar from 0% to 100%. It looks informative. But research on clinical decision support has repeatedly found that the *form* in which an AI’s output is presented shapes how clinicians use it—often in ways the designers never intended (Bussone et al., 2015).

A confidence bar is a case in point. It can be read not as “here is how sure the model is” but as “here is how much attention this case deserves.” High-confidence cases then draw long scrutiny, as the radiologist tries to understand why the AI is so sure; low-confidence cases get waved through, because the low marker makes the case feel benign. That is the opposite of what a confidence signal is supposed to do—and a reviewer’s time ends up allocated by the display rather than by the difficulty of the case.

The fix is rarely more precision. A simpler signal—“AI found a region of interest” versus “AI found no region of interest”—often serves the human’s judgment better than a precise-looking number that invites misinterpretation.

Sequencing and Explanation

A simulation reader study by Pesapane and colleagues had radiologists interpret mammograms in three phases: unassisted, then with AI assistance, then with AI plus explanation—with some AI suggestions deliberately wrong. The unassisted reads established each reader’s own judgment; seeing a wrong AI suggestion afterward produced automation bias in roughly a third of manipulated cases—readers followed the error—and an anchoring bias pulled readers into revising correct answers at a comparable rate. Neither failure was rare. What helped most was explanation: pairing the recommendation with a saliency heatmap that showed *where* the AI was looking cut both biases roughly in half (Pesapane et al., 2026).

The lesson is not that one sequence wins. It is that an interface which simply hands over a verdict—whenever it arrives—leaves the human exposed to a predictable bias, while one that shows its reasoning gives the human something to push back against.

6.8 The Five Dimensions Through the Interface Lens

The five-dimension framework maps cleanly onto interface design choices.

Uncertainty Detection must be made visible. An AI's internal uncertainty estimate is worthless if the interface does not communicate it to the human reviewer. The COMPAS system showed a score without uncertainty bounds. A poorly designed confidence bar can be misread as a cue for how much attention a case deserves. A highlighted region paired with a plain-language description of what drove the AI's concern is more effective at directing human attention appropriately.

Intervention Design is the interface's core problem. What is the human being asked to do? The best interventions are designed with a specific cognitive task in mind: the human is being asked to do something they are genuinely better at than the AI.

Timing in interface terms is about sequencing. Does the human see the AI recommendation before forming their own view, after, or simultaneously? Each order carries its own bias, so the sequence is a deliberate design choice, not an afterthought. The COMPAS score was delivered at the moment of sentencing in a format that made it feel like a fact rather than a recommendation.

Stakes Calibration should shape the visual weight of recommendations. A high-stakes flagged case should look different from a routine flagged case. Many annotation interfaces treat all flags equally, presenting a false confidence to human reviewers that every flag deserves the same attention.

Feedback Integration requires the interface to make the human's feedback visible and learnable. If a radiologist consistently disagrees with the AI about a particular type of finding, that pattern should be detectable—to improve the model, to route those cases to this radiologist, or to trigger a review of the AI's training data.

6.9 Why Design Is a Safety Issue

The same AI model can help at one site and hurt at another. A multi-site evaluation of a single machine-learning system across different hospitals found that its performance did not transfer cleanly: strong where the local data resembled what it was trained on, materially degraded elsewhere (J. Yang et al., 2022). Part of that variation is about data—the populations and scanners differ from one hospital to the next. But part of it is about how the model is put in front of people.

Whether the AI's finding is shown to the clinician before or after they form their own judgment changes whether it reduces errors or introduces them (Pesapane et al., 2026). The same model can be a safe second opinion under one interface and a dangerous first impression under another.

The AI did not make those clinicians worse. The deployment did.

Interface Design as Safety Engineering

A HITL system whose interface makes humans worse is not a HITL system that is working well. It is a HITL system that has quietly removed the loop from human-in-the-loop. Interface design is a safety issue with the same practical significance as model accuracy.

The next time you use a system that gives you a recommendation—a navigation app, a streaming service, a shopping site—notice the order. Did you form your own preference before you saw the recommendation, or did the recommendation arrive first? Pay attention to whether having seen the recommendation makes it harder to decide differently, even if you genuinely want to. That slight resistance is automation bias in action. It is real, it is measurable, and it is built into the interface whether the designers intended it or not.

Chapter Summary

Key Concepts

- Interfaces are not neutral transmission channels—they shape what humans notice, how they reason, and what they decide
- **Automation bias** is the documented tendency to over-rely on automated recommendations; it degrades with fatigue and scales with displayed AI confidence
- XAI tools (LIME, SHAP, Grad-CAM) have real value but also real failure modes when outputs are misinterpreted or poorly presented
- The best interfaces separate the AI's job from the human's job, present uncertainty as information, preserve independent judgment, and make disagreement easy
- Interface design is a safety issue with the same practical significance as model accuracy

Key Examples

COMPAS risk scoring A score without explanation replaced judicial reasoning and embedded racial bias in a format that looked like objectivity

CheXNet-era radiology AI (Gaube et al.) Radiologists rated AI-labeled advice lower in quality than identical human-labeled advice, yet relied on both equally—stated distrust never became actual skepticism

Mammography automation bias (Dratsch et al.) A wrong AI BI-RADS suggestion caused even experienced radiologists to miss findings they would otherwise have caught—automation bias in expert practice

Confidence displays A precise-looking confidence bar can be misread as a cue for how much attention a case deserves; a simpler flag often serves judgment better

Sequencing and explanation (Pesapane et al.) Showing AI first invites automation bias and showing it second invites anchoring bias; a saliency heatmap that shows where the AI looked cuts both

Interface design changes labels (Stureborg et al.) With the same annotators and the same task, interface design alone substantially changed annotation quality

Five Dimensions Check

- *Uncertainty Detection* (Dimension 1): An uncertainty estimate the interface never communicates might as well not exist—how confidence is displayed determines whether the human can act on it

- *Intervention Design* (Dimension 2): The interface is the intervention—what the human is asked to do, in what order, with what evidence, decides whether review supports judgment or replaces it
- *Timing* (Dimension 3): Sequencing is a timing decision at the scale of seconds—AI-first invites automation bias; human-first invites anchoring on the reader’s own initial read. Neither order is free; explanation helps both

In the next chapter, we turn to the other side of the interface: not what the AI shows the human, but what the human provides to the AI. Labels are the foundation of supervised learning—but they are not neutral, and extracting good ones from human annotators requires its own design discipline.

Technical Appendix 6A: Interface Design for Human-in-the-Loop Systems

The interfaces described in Chapter 6 can be analyzed using cognitive load theory, decision-theoretic models of automation bias, and formal metrics for explanation quality.

— Formal foundations for the examples in Chapter 6

6A.1 Cognitive Load Theory and Interface Design

6A.1.1 The Three-Component Model

Following Sweller (1988) and Paas and Van Merriënboer (1994), we identify three components of cognitive load relevant to HITL interface design:

Intrinsic Load (IL). The inherent complexity of the decision task.

$$IL = f(\text{task complexity, element interactivity})$$

Extraneous Load (EL). Cognitive overhead imposed by interface design that does not contribute to the decision.

$$EL = f(\text{navigation steps, visual noise, context switches})$$

Germane Load (GL). Cognitive effort directed toward actual judgment.

$$GL = \text{Total Working Memory} - IL - EL$$

The design goal for HITL interfaces: minimize *EL*, so *GL* can be maximized within working memory limits.

6A.1.2 Formalizing the Annotation Click Problem

Let *T* be total session time, *N* be the number of items to review, and define:

- *t_j* = time on genuine judgment per item
- *t_m* = time on interface mechanics per item
- *q₀* = baseline judgment quality (early session)
- $\delta(t_m)$ = quality degradation as a function of accumulated mechanical overhead

Then:

$$\text{Session Quality} = q_0 - \int_0^N \delta(t_m \cdot n) \, dn$$

Table 6.1: Annotation quality vs. interface design (illustrative values)

Interface Design	Avg. <i>t_m</i> per item	Session kappa at item 400
Legacy (menu-based)	14 sec	0.61
Optimized (keyboard shortcuts)	5 sec	0.79
Optimal (single-action)	2 sec	0.84

6A.2 Automation Bias: Formal Model

6A.2.1 The Anchor-Adjustment Framework

Automation bias can be modeled as an anchoring process (Tversky & Kahneman, 1974). When a human reviewer h is shown an AI recommendation r_{AI} before forming their independent judgment:

$$r_h = r_{AI} + \alpha \cdot (r_{\text{independent}} - r_{AI})$$

where $\alpha \in [0, 1]$ is the adjustment coefficient. When $\alpha = 1$, the human makes a fully independent judgment. When $\alpha = 0$, the human simply adopts the AI recommendation.

The adjustment coefficient degrades with fatigue and AI confidence display:

$$\alpha(f, c) = \alpha_0 \cdot e^{-\lambda_f f} \cdot e^{-\lambda_c c}$$

where f is fatigue (session hours elapsed), c is displayed AI confidence, and λ_f, λ_c are empirically estimated decay constants.

As an illustrative range (in the spirit of the tables above), values of $\alpha \in [0.3, 0.7]$ are plausible under standard conditions; the decay constants λ_f, λ_c must be estimated for the deployment setting.

6A.2.2 Interface Interventions That Increase α

Pre-commitment approach. Require the human to record their initial assessment before revealing the AI recommendation:

$$\alpha_{\text{pre-commit}} \approx 1.0 \quad (\text{by construction})$$

Studies of cockpit automation report the same direction of effect: operators prompted to commit to an independent, accountable judgment before consulting the automation verify it more often and commit fewer automation-bias errors than those shown the recommendation first (Mosier et al., 1998).

Disagree-facilitation. If $C_{\text{agree}} = 1$ click and $C_{\text{disagree}} = k$ clicks:

$$P(\text{disagree}) \propto \frac{1}{k}$$

Setting $k = 1$ significantly increases the rate of warranted independent overrides.

6A.3 Explainability Output Quality Metrics

6A.3.1 Faithfulness and Stability

An explanation method E that maps model f and input x to a feature importance vector is evaluated on:

Faithfulness.

$$\text{Faithfulness}(E, f, x) = \text{corr}(\Delta f(x, \text{mask}), E(f, x))$$

Stability.

$$\text{Stability}(E, f, x) = \exp\left(-\mathbb{E}_{\epsilon}[\|E(f, x + \epsilon) - E(f, x)\|_2]\right)$$

The exponential kernel keeps the metric in $(0, 1]$: identical explanations score 1, and increasingly unstable ones approach 0.

Illustrative magnitudes: LIME stability scores around 0.4–0.6; SHAP stability scores around 0.7–0.85 for the same models.

6A.3.2 Human Interpretability Metric

$$\begin{aligned} \text{Utility}(E, h) = & P(\text{correct override} \mid E \text{ shown to } h) \\ & - P(\text{correct override} \mid E \text{ not shown to } h) \end{aligned}$$

Table 6.2: Human utility of explanation types (illustrative values)

Explanation Type	Human Utility (mean)
Probability score only	−0.04
Probability + highlighted region	+0.12
Probability + region + plain-language rationale	+0.21
Full SHAP feature plot (no training)	−0.07

6A.4 Interface Sequence Models

6A.4.1 Information Ordering and Decision Quality

Under ideal Bayesian updating, the order of evidence presentation does not matter. Humans are not ideal Bayesian updaters. Empirically:

$$\Delta_{\text{order}} = P(\text{correct} \mid \text{Human-first}) - P(\text{correct} \mid \text{AI-first})$$

As an illustrative magnitude, $\Delta_{\text{order}} \in [0.05, 0.15]$ in favor of human-first sequencing is a plausible range, particularly when the AI recommendation is wrong; the actual effect size must be measured in the deployment setting.

6A.4.2 Optimal Interface Sequencing Algorithm

```
1 def optimal_review_sequence(item, ai_model, human_reviewer,
2                             HIGH_CONF_THRESHOLD=0.92):
3     # Phase 1: Establish human prior (independent)
4     initial_judgment = human_reviewer.assess(item)
5     record(initial_judgment, timestamp=now())
6
7     # Phase 2: Compute AI recommendation
8     ai_pred, ai_conf, ai_expl = ai_model.predict(item)
```

```

9
10 # Phase 3: Conditional display based on uncertainty
11 if ai_conf < HIGH_CONF_THRESHOLD:
12     show_to_human(ai_pred, ai_conf, ai_expl)
13     final_judgment = human_reviewer.review(
14         item,
15         own_judgment=initial_judgment,
16         ai_input=ai_pred
17     )
18 else:
19     if initial_judgment == ai_pred:
20         final_judgment = initial_judgment
21         log(agreement=True, reviewed=False)
22     else:
23         # Disagreement on high-confidence case: escalate
24         final_judgment = escalate_to_senior(item)
25
26 record(final_judgment, ai_pred, ai_conf)
27 return final_judgment

```

Listing 6.1: Optimal review sequencing for HITL interfaces

6A.5 Attention Map Quality Assessment

6A.5.1 Grad-CAM Formal Definition

For a convolutional network with class activation weights α_k^c and feature map activations A^k :

$$L_{\text{Grad-CAM}}^c = \text{ReLU}\left(\sum_k \alpha_k^c A^k\right)$$

where:

$$\alpha_k^c = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}$$

6A.5.2 Radiologist Utility Conditions

Grad-CAM is useful to radiologists when:

1. **Localization is correct:** IoU between Grad-CAM highlight and ground-truth annotation > 0.5
2. **Not dominated by artifacts:** Attention mass on non-anatomical features $< 5\%$
3. **Uncertainty not masked:** Diffuse low-magnitude Grad-CAM vs. tight high-magnitude should be visually distinguished

```
1 def assess_gradcam_quality(gradcam_map, ground_truth_mask,
2   artifact_mask):
3     intersection = np.logical_and(gradcam_map > 0.5,
4       ground_truth_mask)
5     union = np.logical_or(gradcam_map > 0.5,
6       ground_truth_mask)
7     iou = intersection.sum() / union.sum()
8
9     artifact_fraction = (gradcam_map * artifact_mask).sum()
10    / gradcam_map.sum()
11    p = gradcam_map / gradcam_map.sum()
12    entropy = -np.sum(p * np.log(p + 1e-8))
13
14    return {
15      'localization_iou': iou,
16      'artifact_fraction': artifact_fraction,
17      'attention_entropy': entropy,
18      'display_recommended': iou > 0.5 and
19        artifact_fraction < 0.05
20    }
```

Listing 6.2: Grad-CAM quality assessment

6A.6 The Review Queue as a Cognitive System

6A.6.1 Queue Design and Cognitive Fatigue

Reviewer quality $q(n)$ as a function of items reviewed n :

$$q(n) = (q_0 - q_{\text{floor}}) \cdot e^{-\lambda_q n} + q_{\text{floor}}$$

where q_0 is the fresh-start quality (so $q(0) = q_0$), q_{floor} is the asymptotic floor reached under fatigue, and λ_q is the decay rate.

Break scheduling to maintain quality above threshold q_{min} :

$$n_{\text{break}} = \frac{1}{\lambda_q} \ln \left(\frac{q_0 - q_{\text{floor}}}{q_{\text{min}} - q_{\text{floor}}} \right)$$

For $q_0 = 1.0$, $q_{\text{floor}} = 0.7$, $q_{\text{min}} = 0.85$, $\lambda_q = 0.005$:

$$n_{\text{break}} = 200 \cdot \ln(2) \approx 139 \text{ items}$$

For high-stakes annotation, breaks should be enforced approximately every 140 items.

6A.6.2 Priority Routing in Review Queues

$$\text{Priority}(x) = w_1 \cdot U_{\text{model}}(x) + w_2 \cdot S(x) + w_3 \cdot T(x)$$

where:

- $U_{\text{model}}(x)$ = model uncertainty for item x
- $S(x)$ = estimated stakes of a wrong decision

- $T(x)$ = time pressure (queue age of item)
- w_1, w_2, w_3 = domain-specific weighting parameters

Recommended starting weights:

- Content moderation (safety-critical): $w_1 = 0.3, w_2 = 0.5, w_3 = 0.2$
- Medical imaging triage: $w_1 = 0.4, w_2 = 0.5, w_3 = 0.1$
- General annotation: $w_1 = 0.5, w_2 = 0.3, w_3 = 0.2$

6A.7 Exercises and Assignments

Exercise 6A.1 — Interface Audit

Take any annotation or AI-assisted review interface. Count the number of clicks required to: (a) confirm the AI recommendation, (b) override the AI recommendation, (c) escalate to senior review. Calculate $C_{\text{disagree}}/C_{\text{agree}}$ (the ratio k from the disagree-facilitation model above). What does this ratio imply about the interface's design assumptions regarding override frequency?

Exercise 6A.2 — Automation Bias Measurement

Design a controlled experiment to measure the automation bias coefficient α for a binary classification review task. Define your Control, AI-First, and Human-First conditions. Describe how you estimate α from the resulting data. Compute the sample size needed to detect a difference of 0.05 in mean accuracy between conditions (power = 0.80, $\alpha = 0.05$).

Exercise 6A.3 — Explanation Utility

Implement a simple binary classifier. For 50 test cases, have three evaluators assess each case under three conditions: (a) no explanation, (b) LIME explanation, (c) SHAP explanation. Calculate Human Utility for each explanation type. Which performs best? Are there systematic differences based on evaluator data literacy?

Exercise 6A.4 — Cognitive Load Reduction

Choose an annotation task requiring at least 5 clicks per item. Redesign the interface to reduce clicks to 2 or fewer while maintaining all necessary data capture. Using the cognitive fatigue model, estimate the expected change in the session quality threshold crossing point n_{break} for your redesign.

Exercise 6A.5 — Priority Queue Design

A content moderation queue contains 10,000 items. You have 5 reviewers and 8 hours. Using the $\text{Priority}(x)$ formula, design a routing system. How do you estimate $U_{\text{model}}(x)$ and $S(x)$ in practice? How do you choose weights? Conduct a capacity analysis: how many items can be reviewed, and what fraction of the highest-priority items will receive human review?

This appendix provides the formal foundations for the interface design principles introduced in Chapter 6. The cognitive load model, automation bias

framework, and explanation utility metric together constitute a quantitative toolkit for evaluating and improving HITL interfaces in production systems.

7 Getting Good Information from People

Human labels are the foundation of supervised learning—but labels are not neutral. Every annotation is an act of interpretation, shaped by who is labeling, what instructions they received, and what world they live in.

— on labels

7.1 The Photograph That Changed How We See Data

In 2015, Stanford researcher Fei-Fei Li gave a TED talk in Vancouver describing what she and her collaborators had been building since 2007: ImageNet (Li, 2015). The project, assembled with the help of Amazon Mechanical Turk, had grown to 15 million photographs organized into 22,000 categories. It was the largest labeled image dataset ever assembled, and it had already become the training ground for the deep learning revolution (J. Deng et al., 2009).

By 2019, a series of studies had identified something troubling in the data. The category “cheerleader” contained photographs that were overwhelmingly of women. The category “guard” was overwhelmingly of men. The annotations were accurate in a narrow sense—the photographs did show what they were labeled as—but the dataset had encoded cul-

tural stereotypes so thoroughly that models trained on it absorbed and reproduced those gendered associations (K. Yang et al., 2020).

The labels were not wrong. They were incomplete descriptions of a biased sample, faithfully transcribed into a format that the model interpreted as ground truth.

The ImageNet case illustrates something that the whole field of machine learning spent the 2010s slowly learning: **labels are not neutral**. Every label is an act of interpretation, shaped by who is doing the labeling, what instructions they were given, what cultural context they bring, and how tired they are. The data annotation problem is not a logistical problem that can be solved by deploying enough workers. It is an epistemological problem—a problem about how we know what we know, and what it means to reduce the complexity of a photograph, a sentence, or a human interaction to a single categorical label.

7.2 The Instruction Problem

Imagine you are about to annotate a set of online comments for toxicity. You have been given a single instruction: “Label each comment as TOXIC or NOT TOXIC.”

You see this comment: *“I can’t believe you people still think this is acceptable. You’re all completely delusional.”*

Is it toxic? Probably yes, you think.

Now imagine the same comment, but a second annotator is working with a slightly expanded instruction: “Label each comment as TOXIC or NOT TOXIC. Toxicity means explicit slurs, direct threats, or content that would make a member of a protected group feel physically unsafe.”

The second annotator looks at the same comment. Contemptuous? Yes. An explicit slur, a direct threat, or content likely to make a protected group feel physically unsafe? Not exactly. They label it NOT TOXIC.

Same comment. Different instructions. Different labels.

This is what annotation researchers call the **instruction problem**, and it is pervasive. Röttger and colleagues showed that even the broad *paradigm* a task adopts changes its labels: a *descriptive* approach, which invites annotators to apply their own subjective judgment, and a *prescriptive* one, which asks them to apply a single specified standard regardless of their personal view, produce systematically different label distributions on the same data—and neither is simply “correct,” because they are answering different questions (Röttger et al., 2022). The same gap opens with a single reworded question. “Is this comment hateful?” is not the same task as “Would a reasonable person find this comment toxic?” The second framing invokes a hypothetical external standard that moderates individual variation—it asks the annotator to model a composite person rather than report their own reaction.

The instruction problem compounds when guidelines are long. Professional annotation companies routinely produce annotation guidelines that run to 50, 80, even 120 pages for complex labeling tasks. Annotators working from a 100-page guideline develop a mental model of the guideline that is partly shared and partly individual. The shared part produces consistent labels. The individual part produces the kind of systematic variation that shows up in low inter-annotator agreement scores.

7.3 Measuring Agreement: Cohen's Kappa

When two annotators label the same set of items, their agreement is measurable. The simplest measure is raw agreement—the proportion of items they label the same way. But raw agreement is inflated by chance agreement: if 90% of items are NOT TOXIC and each annotator independently labeled 90% of items NOT TOXIC at random, they would agree 82% of

the time by chance alone—the sum of both-TOXIC and both-NOT-TOXIC chance cells—while making no genuine classification decisions.

The standard correction is **Cohen’s kappa**, developed by psychologist Jacob Cohen in 1960 (Cohen, 1960):

$$\kappa = \frac{P_o - P_e}{1 - P_e}$$

where P_o is the observed agreement and P_e is the expected agreement by chance. A kappa of 1.0 means perfect agreement; 0 means agreement no better than chance; negative values mean systematic disagreement.

The Landis-Koch Interpretation Scale

- $\kappa < 0.20$: Slight agreement
- $0.20 \leq \kappa < 0.40$: Fair agreement
- $0.40 \leq \kappa < 0.60$: Moderate agreement
- $0.60 \leq \kappa < 0.80$: Substantial agreement
- $\kappa \geq 0.80$: Almost perfect agreement

These thresholds, proposed by Landis and Koch in 1977 (Landis & Koch, 1977), are conventional, not theoretically grounded. Always interpret kappa in the context of the specific task.

When the toxicity detection field began systematically measuring annotator agreement, the results were sobering. A landmark study of the OLID benchmark by Zampieri and colleagues in 2019 found a Fleiss’ kappa of 0.83 in a small trial with five expert annotators labeling 21 tweets (Zampieri et al., 2019)—Fleiss’ kappa being the multi-rater counterpart to the two-rater Cohen’s kappa derived above, and 0.83 being “almost perfect” on

the Landis-Koch scale. Yet subsequent annotation projects in the toxicity domain with larger, less-controlled annotator pools have reported substantially lower values, in the 0.45–0.60 range. That gap—between what experts achieve under controlled conditions and what larger pools achieve in practice—is the sobering result, illustrating how sensitive inter-annotator agreement is to annotator composition and task framing. Kappa is chance-corrected, so it does not map directly onto a disagreement rate: with an even split of toxic and benign comments, $\kappa = 0.45$ means two annotators would disagree on roughly a quarter of items; with a skewed label distribution, the same kappa can sit atop much higher raw agreement. A model trained on such data is not learning to detect toxicity; it is learning to predict the label distribution of a specific group of annotators working under specific conditions.

7.4 The Annotator Is Not a Recording Device

There is a useful fantasy in machine learning that annotators are passive transcribers—that they observe some objective feature of the data and record it, and that the label accurately reflects something real in the world. The fantasy is comfortable because it preserves the idea that labels are ground truth.

It is false.

An annotator is a person with a history, a worldview, a cultural context, and a present state that varies by hour of day, level of fatigue, number of items reviewed, and recent sequence of items seen.

Sequence effects. An annotator who has just reviewed 20 highly toxic comments will have a different calibration threshold for the next borderline comment than an annotator who has just reviewed 20 benign ones. Studies of machine translation quality annotation found that annotators

systematically rated the same translations differently depending on the quality of the surrounding translations in their queue (Mathur et al., 2017). The translation had not changed. The annotator's reference point had.

Demographic effects. Research led by Sap and colleagues at the University of Washington found that tweets containing African American English (AAE) were up to twice as likely to be labeled offensive compared to similar tweets in Mainstream American English, even when the content was neutral or positive. When annotators were primed with information about AAE dialect features before labeling, the racial bias in labeling decreased (Sap et al., 2019). Subsequent work by Sap et al. ("Annotators with Attitudes," 2022) found the same pattern at the annotator level: Black and white annotators differed substantially in their toxicity judgments (Sap et al., 2022). The practical consequence is that a model trained on annotations from a predominantly white annotator pool will have a systematically higher false toxicity rate for AAE.

This is not a problem that better annotation guidelines can fully fix, because it is embedded in who the annotators are and what cultural context they bring to the task. Diversifying the annotator pool is a practical necessity, not merely a fairness gesture.

7.5 The Economics of Getting Labels

The ImageNet dataset was completed with the help of Amazon Mechanical Turk, a crowdsourcing platform that allows requesters to post small tasks and have them completed by a globally distributed workforce for small payments. A typical toxicity labeling task on Mechanical Turk pays on the order of \$0.02 to \$0.05 per item. In a 2018 study of all work on the

platform, Hara and colleagues found a median effective hourly wage of approximately \$2/hr once unpaid time—searching for tasks, waiting on requesters, switching between jobs—was counted (Hara et al., 2018).

The economic logic of crowdsourcing annotation is compelling. A dataset that would cost hundreds of thousands of dollars to label through professional linguists can be labeled for tens of thousands of dollars through Mechanical Turk, at higher speed.

The quality control problem is also real. Without oversight, crowd-sourced annotation produces noise: workers who click randomly, who misread instructions, whose cultural context does not match the task requirements. The standard approach—multiple annotators per item plus majority voting—reduces individual noise. But majority voting on disagreement can enshrine the biases of the majority as ground truth.

The distinction between **noise disagreement** and **genuine ambiguity** matters enormously. If two annotators disagree because one was not paying attention, the disagreement is noise and should be resolved by majority vote. If two annotators disagree because the item is genuinely ambiguous, the disagreement is information. Collapsing it to a single label loses the information.

Label Distribution Training

Some modern machine learning practice has begun to work with label distributions rather than single labels—training on the full distribution of annotator judgments rather than the majority vote. Research by Davani and colleagues demonstrated that models trained on distributions were better calibrated: they expressed more uncertainty on genuinely ambiguous items, and more confidence on items where annotators had agreed. This is a more honest representation of what the data actually contains (Davani et al., 2022).

7.6 Framing Effects: How the Question Changes the Answer

Framing effects describe the phenomenon that how a question is posed shapes the answer, independently of what the answer “truly” is. A classic study found that a medical treatment described as having a “90% survival rate” was viewed more favorably than the same treatment described as having a “10% mortality rate.” Same statistic. Very different responses (McNeil et al., 1982).

The same dynamic appears in annotation. Consider the difference between:

Version A: “Rate each response on a scale of 1–5 for helpfulness.”

Version B: “Consider: would a patient receiving this information be better equipped to make decisions about their health? Rate 1–5.”

Version B anchors the annotator on a concrete downstream beneficiary and asks them to simulate that person’s experience. This produces more consistent labels across annotators and labels that better predict actual user satisfaction. The question is not whether Version B is more “correct.” It is that Version B more precisely specifies what the annotation is supposed to measure.

The most consequential framing decisions in annotation are often invisible because they are embedded in what is *not* included in the instructions. Medical image annotation typically asks annotators to identify the presence or absence of a finding, not to classify its clinical significance. A model trained on such annotations learns to detect findings, not to assess their relevance to patient management—a critical difference when the model is deployed to assist diagnosis.

7.7 Calibration Sessions and the Iterative Guideline

The annotation community has developed practices to address the instruction problem, the most important of which is the **calibration session**.

In a calibration session, a group of annotators reviews and discusses a set of carefully selected items together. The items are chosen to cover the range of cases that arise in the task, including clear cases, borderline cases, and cases that expose ambiguities in the annotation guidelines. Annotators record their labels independently, then compare and discuss disagreements.

The output is twofold: an updated annotation guideline that addresses the exposed ambiguities, and a shared mental model that enables more consistent application going forward. Companies like Scale AI and Appen conduct formal calibration sessions at the start of annotation projects and regularly thereafter, particularly when the task involves subjective judgment or cultural nuance.

OpenAI's work on RLHF (Reinforcement Learning from Human Feedback) is explicit about this iterative process. The human feedback that shapes RLHF training is not collected once and used permanently; it is continuously re-evaluated, guidelines are updated, and new annotation rounds are conducted to capture changing standards (Ouyang et al., 2022). The annotation is not a solved problem but an ongoing conversation between the model's behavior and human preferences.

7.8 What Low Agreement Tells You

When inter-annotator agreement on a task is low—say, kappa below 0.40—the conventional response is to interpret this as a quality problem. The solution is to revise the guidelines and improve the training.

This interpretation is sometimes correct. But low agreement can also be a signal about the *nature of the task itself*. Some annotation tasks ask annotators to make distinctions that are genuinely difficult, genuinely subjective, or genuinely contested. A task asking annotators to distinguish between “mildly offensive” and “moderately offensive” is asking for a discrimination that may not be reliably present in human judgment, not because the annotators are failing but because the categories do not map onto stable psychological reality.

When low agreement persists despite careful guidelines, calibrated training, and expert annotators, the right response is often not to force agreement but to interrogate the task. Is the distinction the annotation is trying to capture actually important for the downstream use case? Could the model perform adequately with coarser categories where agreement is higher? Is the genuine ambiguity in the data actually carrying information that should be preserved?

A Low Kappa Score Is Telling You Something Real

In the most honest framing: a low kappa score is telling you something real about the data, about the task, and about the limits of reducing human judgment to a scalar label. The answer to that signal is not always to try harder. Sometimes it is to listen.

Take any three subjective opinions you hold—about a film, a piece of writing, a public figure’s behavior. For each one, write down the implicit criteria you are using to form the opinion. Now try to write those criteria as annotation guidelines, clearly enough that a stranger could apply them reliably. Notice how quickly “good”

requires unpacking, how “harmful” requires definition, and how much of your judgment is tacit knowledge that resists specification. This is the annotation problem, and it does not have a clean solution.

Chapter Summary

Key Concepts

- Labels are not neutral—every annotation is an act of interpretation shaped by who is labeling, what instructions they received, and what cultural context they bring
- The instruction problem: even small differences in guideline wording produce systematically different label distributions
- **Cohen’s kappa** measures inter-annotator agreement corrected for chance; low kappa can signal poor guidelines, genuine task ambiguity, or both
- Framing effects show that the question shape changes the answer, independently of what is being classified
- The annotator is not a recording device—sequence effects, demographic background, and fatigue all systematically shape labels

Key Examples

ImageNet (2007–2019) Cultural stereotypes baked into labels by an annotation process that faithfully transcribed annotator worldview as ground truth

AAE toxicity labeling Demographically homogeneous annotator pools produce systematically biased labels for minority language registers

RLHF human feedback (OpenAI) Iterative guideline revision as a necessary feature of annotation for AI training

Scale AI and Appen calibration sessions The standard industry practice for reducing annotation inconsistency

Label distribution training Training on the full annotation distribution rather than majority-vote labels produces better-calibrated models

Five Dimensions Check

- *Feedback Integration* (Dimension 5): Labels are the training feedback that drives the loop; biased or noisy annotation corrupts everything downstream
- *Intervention Design* (Dimension 2): Guideline wording and question framing are interface choices that shape what the annotator produces
- *Uncertainty Detection* (Dimension 1): Low inter-annotator agreement is itself a signal—of an ambiguous task, unclear guidelines, or both

In the next chapter, we zoom out from individual interactions—interfaces, labels, human judgments—and look at the HITL pipeline as a whole: the data flows, the engineering challenges, and what happens when a system that worked beautifully in the lab breaks down in production.

Technical Appendix 7A: Annotation Metrics and Label Quality

The annotation challenges described in Chapter 7 can be formalized using agreement statistics, Bayesian aggregation models, and signal detection theory.

— Formal foundations for the examples in Chapter 7

7A.1 Cohen’s Kappa: Derivation and Properties

7A.1.1 Basic Formula

For two annotators labeling N items into k categories, let n_{ij} be the number of items where annotator 1 assigned category i and annotator 2 assigned category j .

$$P_o = \frac{\sum_i n_{ii}}{N} \quad (\text{observed agreement})$$

$$P_e = \sum_i \frac{n_{i\cdot}}{N} \cdot \frac{n_{\cdot i}}{N} \quad (\text{expected chance agreement})$$

$$\kappa = \frac{P_o - P_e}{1 - P_e}$$

where $n_{i\cdot} = \sum_j n_{ij}$ is the row marginal and $n_{\cdot j} = \sum_i n_{ji}$ is the column marginal.

7A.1.2 Variance and Confidence Intervals

$$\text{Var}(\kappa) \approx \frac{P_o(1 - P_o)}{N(1 - P_e)^2}$$

Approximate 95% confidence interval: $\kappa \pm 1.96\sqrt{\text{Var}(\kappa)}$

This variance formula is a simplified approximation that ignores variance from estimating P_e itself, producing confidence intervals that are slightly too narrow—especially when P_e is large or N is small. For publication-quality intervals use the exact formula from Fleiss (1971) or a bootstrap estimate.

As an absolute minimum, require $N \geq 30$ items per annotator pair for the estimate to be computable and roughly unbiased. Confidence-interval quality at that sample size is poor (at $P_o = 0.8$, $P_e = 0.5$, the 95% half-width is about ± 0.29); for tighter margins use the sample-size formula in §7A.5.1 below.

7A.1.3 Weighted Kappa for Ordinal Categories

With linear weights $w_{ij} = 1 - |i - j|/(k - 1)$ or quadratic weights $w_{ij} = 1 - [(i - j)/(k - 1)]^2$, define weighted observed and expected agreement:

$$P_o^w = \sum_{ij} \frac{w_{ij} n_{ij}}{N}, \quad P_e^w = \sum_{ij} w_{ij} P_{e,ij}, \quad P_{e,ij} = \frac{n_{i\cdot}}{N} \cdot \frac{n_{\cdot j}}{N}$$

Weighted kappa follows the same structure as unweighted kappa (Cohen, 1968):

$$\kappa_w = \frac{P_o^w - P_e^w}{1 - P_e^w}$$

Quadratic weights are standard for medical rating scales; linear weights for quality assessments.

7A.2 Fleiss' Kappa: Extension to Multiple Annotators

For M annotators labeling N items into k categories, let n_{ij} be the number of annotators who assigned item i to category j :

$$\begin{aligned} \bar{P} &= \frac{1}{NM(M-1)} \sum_i \sum_j n_{ij}(n_{ij} - 1) & \bar{P}_e &= \sum_j \left(\frac{\sum_i n_{ij}}{NM} \right)^2 \\ \kappa_F &= \frac{\bar{P} - \bar{P}_e}{1 - \bar{P}_e} \end{aligned}$$

Table 7.1: Agreement benchmarks by annotation domain

Domain	Typical Range	Target
Binary toxicity (general)	0.45–0.65	≥ 0.60
Medical image (radiology)	0.55–0.75	≥ 0.75
Sentiment analysis (clear)	0.70–0.85	≥ 0.70
Named entity recognition	0.80–0.95	≥ 0.85
Legal document classification	0.50–0.70	≥ 0.65
RLHF instruction quality	0.55–0.75	≥ 0.65
AAVE toxicity (diverse pool)	0.40–0.55	≥ 0.60

7A.3 Annotation as a Bayesian Inference Problem: Dawid-Skene

7A.3.1 Generative Model

The Dawid-Skene model (Dawid & Skene, 1979) treats annotator labels as noisy observations of a true latent label, estimating both true labels and annotator error rates jointly.

For item i , annotator j , and observed label y_{ij} :

$$P(y_{ij} = l \mid z_i = k) = \pi_{jkl}$$

where z_i is the true label for item i and π_{jkl} is annotator j 's probability of labeling as l when the true label is k .

7A.3.2 EM Algorithm

E-step. Compute posterior probability of true label:

$$P(z_i = k \mid \mathbf{y}_i) \propto p_k \prod_{j \in \mathcal{J}_i} \prod_l \pi_{jkl}^{\mathbf{1}[y_{ij}=l]}$$

where $\mathcal{J}_i$ is the set of annotators who labeled item i .

M-step. Update annotator error rates and class priors:

$$\hat{\pi}_{jkl} = \frac{\sum_{i \in \mathcal{J}_j} P(z_i = k) \cdot \mathbf{1}[y_{ij} = l]}{\sum_{i \in \mathcal{J}_j} P(z_i = k)} \quad \hat{p}_k = \frac{1}{N} \sum_i P(z_i = k)$$

where $\mathcal{J}_j$ is the set of items labeled by annotator j , matching the restricted sums in the code below.

```

1 def dawid_skene_em(annotations, n_classes, n_iters=50):
2     """
3     annotations: dict {item_id: {annotator_id: label}}
4     Returns: estimated true label posteriors T, annotator
           conf. matrices pi
5     """
6     N, K = len(annotations), n_classes
7     items = list(annotations.keys())
8     annotators = list(set(a for item in
           annotations.values() for a in item))
9     M = len(annotators)
10
11     # Initialize T with majority vote (ties → uniform; EM
           resolves them)
12     T = np.zeros((N, K))
13     for i, item in enumerate(items):
14         for v in annotations[item].values():
15             T[i, v] += 1
16             T[i] /= T[i].sum() # one-hot for single-annotation
           items; a row with no votes would be NaN
17
18     pi = np.ones((M, K, K)) / K
19
20     for _ in range(n_iters):
21         # M-step: update pi
22         for j, ann in enumerate(annotators):
23             for k in range(K):
24                 for l in range(K):
25                     num = sum(T[i,k] for i,item in
                           enumerate(items)
26                               if ann in annotations[item]
27                               and annotations[item][ann] ==
                           l)

```

```

28         den = sum(T[i,k] for i,item in
29                     enumerate(items)
30                     if ann in annotations[item])
31         pi[j,k,l] = num / (den + 1e-10)
32
33     prior = T.mean(axis=0)
34
35     # E-step: update T
36     log_T = np.log(prior[np.newaxis,:] + 1e-10) *
37             np.ones((N,K))
38     for i, item in enumerate(items):
39         for j, ann in enumerate(annotators):
40             if ann in annotations[item]:
41                 l = annotations[item][ann]
42                 log_T[i] += np.log(pi[j,:,l] + 1e-10)
43
44     log_T -= log_T.max(axis=1, keepdims=True)
45     T = np.exp(log_T)
46     T /= T.sum(axis=1, keepdims=True)
47
48     return T, pi

```

Listing 7.1: Dawid-Skene EM (simplified)

7A.4 Framing Effects: Signal Detection Theory Framework

Model annotation as a signal detection problem. The annotator observes item x with latent “harmfulness” signal $s(x) \in \mathbb{R}$, and classifies by comparing to threshold θ .

$$\hat{y} = \begin{cases} 1 & s(x) > \theta \\ 0 & s(x) \leq \theta \end{cases}$$

The annotation guideline framing shifts θ for individual annotators:

$$\theta_A = \theta_0 + \Delta_A$$

Inter-annotator agreement is a function of the variance in individual thresholds:

$$\kappa \approx 1 - f(\sigma_\theta^2)$$

Framing effects that reduce σ_θ^2 (by more precisely specifying the threshold) directly increase κ . The “reasonable person” standard reduces idiosyncratic threshold variation at the cost of introducing shared cultural blind spots.

7A.5 Sample Size Planning for Annotation Projects

7A.5.1 Statistical Power for Kappa Estimation

To estimate kappa with margin of error e at confidence level α :

$$N \geq \frac{z_{\alpha/2}^2 \cdot P_o(1 - P_o)}{e^2(1 - P_e)^2}$$

For typical annotation settings ($P_o \approx 0.80$, $P_e \approx 0.50$, $e = 0.05$, $\alpha = 0.05$):

$$N \geq \frac{3.84 \times 0.16}{0.0025 \times 0.25} \approx 983 \text{ items}$$

7A.5.2 Redundancy and Quality Trade-off

Expected kappa improvement from n_r annotators per item with majority-vote aggregation:

$$\kappa(n_r) \approx \kappa_1 + (1 - \kappa_1) \cdot (1 - e^{-\lambda(n_r-1)})$$

where κ_1 is single-annotator kappa and $\lambda \approx 0.3$. Increasing redundancy beyond $n_r = 5$ yields diminishing returns; under this model the marginal quality gain from a 6th annotator is $(1 - \kappa_1)(e^{-1.2} - e^{-1.5}) \approx 0.08(1 - \kappa_1)$, i.e., roughly 0.02–0.04 kappa for typical $\kappa_1 \in [0.5, 0.75]$ (dropping below 0.01 only when $\kappa_1 \geq 0.87$).

The value $\lambda \approx 0.3$ is an empirical heuristic derived from annotation studies across multiple NLP and CV tasks; it is not derived from first principles. In practice, λ varies by task difficulty and annotator pool. Treat the formula as a planning heuristic, not a precise prediction, and calibrate against pilot annotation data for any specific project.

7A.6 Exercises and Assignments

Exercise 7A.1 — Cohen's Kappa Computation

Two annotators classify 100 social media posts:

	Ann. 2: T	Ann. 2: N
Ann. 1: T	20	15
Ann. 1: N	10	55

Compute P_o , P_e , κ , and a 95% confidence interval. Interpret using the Landis-Koch scale and recommend whether to proceed with model training.

Exercise 7A.2 — Fleiss' Kappa Properties

Three annotators each label 50 documents. Compute Fleiss' kappa and compare to average pairwise Cohen's kappa. Under what conditions do they diverge most significantly? Construct a 3-annotator example where average pairwise kappa is misleadingly optimistic.

Exercise 7A.3 — Dawid-Skene Simulation

Simulate annotation data for 100 items with two expert annotators (error rate 5%) and eight novice annotators (error rate 30%). Compare accuracy of majority voting vs. Dawid-Skene EM vs. expert-only majority vote. How much does including novice annotators help or hurt under each method?

Exercise 7A.4 — Framing Effect Experiment

Design an experiment to measure framing effects on a toxicity annotation task. Specify: two guideline versions, number of items, number of annotators per condition, statistical tests for kappa difference and positive label rate difference, and expected results based on the signal detection model.

Exercise 7A.5 — Annotation Budget Optimization

Budget: \$5,000. Task: binary annotation, minimum 2,000 items. Expert: \$2.00/label, $\kappa_1 = 0.75$. Crowdworkers: \$0.05/label, $\kappa_1 = 0.50$. Design three annotation strategies. For each: estimate expected quality (using the redundancy formula), coverage, and total cost. Which strategy do you recommend, and under what conditions would you choose differently?

This appendix provides the formal tools for quantifying annotation quality and diagnosing label problems described in Chapter 7. The kappa family of statistics, the Dawid-Skene model, and the signal detection framework together constitute the annotation quality toolkit for applied HITL system development.

8 The Blueprint for Smart Helper Systems

Most systems that claim to be “human-in-the-loop” are actually pipelines.

— on architecture

8.1 The Alarm Nobody Heard

In 2018, an Uber self-driving test vehicle struck and killed Elaine Herzberg in Tempe, Arizona. She was walking her bicycle across a multi-lane road at night. The onboard systems detected her 5.6 seconds before impact. But the software never settled on what it was seeing: it classified her first as a vehicle, then as an unknown object, then as a bicyclist—and each time the classification flipped, the system discarded its tracking history and started over. It never classified her as a pedestrian, and it never predicted her path (National Transportation Safety Board, 2019).

About 1.2 seconds before impact, the system finally recognized an imminent collision. It did not brake. Emergency braking had been deliberately disabled while the computer was driving, on the theory that a human safety driver was watching the road and would intervene. The system’s design response was to sound an alert and hand the moment to the human. By then, it was too late.

The backup safety driver, meanwhile, had not been watching the road.

Here is what makes this a *system architecture failure* rather than a sensor failure or a human failure: both the algorithm and the human were in the loop. Both failed. And the system's architecture—the way the two were connected, the expectations placed on each, the escalation rules that governed handoffs—was designed without sufficient attention to how human and algorithmic components interact under real operating conditions.

The human was supposed to be in the loop. She was sitting in the driver's seat. But the system had been designed in a way that made meaningful human oversight essentially impossible: the safety driver was expected to monitor continuous operation but had no clear trigger for when to engage. The uncertainty signal that the software was generating was not communicated to her in any actionable form. The “human in the loop” was a checkbox, not a design.

This is the architecture problem.

8.2 The Pipeline Is Not the Loop

Before we discuss what HITL system architecture should look like, it is worth naming the failure mode it is designed to avoid: treating a human-in-the-loop system as a **pipeline** with a human at one end.

A pipeline is a sequence of steps. Data goes in, processing happens, output comes out. In a pipeline with a human step, the human reviews something and passes it along. The human's role is fixed, their inputs are fixed, their outputs are fixed. There is no feedback. There is no learning.

A loop is different. In a loop, the outputs from the human step flow back into the system. The human's responses change how the system behaves in the future. The system is designed to improve based on human feedback. And the human is designed to be supported—with good interfaces, good context, appropriate caseload—so their feedback is worth incorporating.

Most systems that claim to be “human-in-the-loop” are actually pipelines. The human step exists but contributes nothing to system learning. The model is static. The human is a quality control check, not a contributor to the system’s intelligence.

Real HITL architecture requires several components that a simple pipeline does not have.

8.3 The Components of Real HITL Architecture

The Prediction Layer. The machine learning model that processes inputs and produces outputs with uncertainty estimates. The key architectural requirement is that it produces not just predictions but **confidence signals**—calibrated probabilities that the downstream routing layer can act on.

The Uncertainty Quantification Layer. A separate component (or integrated capability) that computes calibrated uncertainty for each prediction. This may be as simple as temperature-scaled softmax outputs, or as complex as a separate ensemble model for uncertainty estimation (Lakshminarayanan et al., 2017). The key is that the output is a calibrated confidence score, not a raw logit.

The Routing Layer. The decision logic that determines what happens to each prediction. Above the high-confidence threshold: automatic action. Below the low-confidence threshold: automatic action in the cautious direction. In between: route to human review. This is where the threshold band design from Chapter 4 is implemented.

The Interface Layer. The system through which human reviewers receive and respond to cases. This is more than a display—it is the context

provision, the question framing, the response format, and the ergonomics that determine whether human feedback is useful or garbage. Chapter 6 covers this in depth. Architecturally, the key requirement is that it can present information appropriately for different case types and capture structured responses that can be fed back into training.

The Feedback Integration Layer. The mechanism by which human responses improve the system. This may be a retraining pipeline, an active learning module, a rule update system, or a calibration-adjustment component. Without this layer, the “loop” in HITL is not a loop—it is an open-ended pipeline.

The Monitoring Layer. Ongoing tracking of system behavior—prediction distributions, human response patterns, accuracy metrics, throughput, reviewer load, calibration drift. This is the component that keeps the system honest over time, detecting when things are drifting in ways that require intervention (Beyer et al., 2016).

The Audit Layer. Logging of decisions, escalations, human responses, and outcomes. This is what makes the system accountable: a record of what was decided, why, by whom (human or machine), and what happened next. In regulated industries, this is often legally required. In all industries, it is practically required for debugging and improvement.

8.4 The Feedback Loop: From Human Response to Model

The most important and most underbuilt component is **feedback integration**. This is where the loop actually closes—where the system gets smarter over time because of what humans tell it.

Several architectures exist, ranging from simple to sophisticated.

Batch retraining. Human labels are accumulated over a period (day, week, month), quality-checked, and added to the training dataset. The model is retrained periodically on the augmented dataset. Simple, robust, but slow—the model does not improve from this morning’s reviews until next month’s retraining.

Online learning. Human labels are incorporated immediately into model updates. Faster learning, but also faster propagation of errors if human quality drops. Requires careful quality monitoring of the incoming labels.

Active learning queues. The routing layer preferentially sends to human review the cases where the model would benefit most from a label—not just the uncertain cases, but the uncertain cases in underrepresented regions of the training distribution. This is an intelligent version of batch retraining where human effort is directed where the learning value is highest. Chapter 17 covers active learning in more depth.

RLHF. Reinforcement Learning from Human Feedback (RLHF) is a more sophisticated architecture used in large language models where human preferences (expressed as comparisons between outputs) are used to train a reward model, which is then used to fine-tune the generative model (Christiano et al., 2017). This is the architecture behind ChatGPT and similar systems (Ouyang et al., 2022). The architectural novelty is that humans do not label correct answers—they express preferences between alternatives, which is often easier and cheaper.

8.5 Logging, Monitoring, and Drift Detection

One of the most consistent failure modes in production HITL systems is drift without detection. The system performs well at launch, degrades slowly over the following months, and the degradation is not noticed until a visible failure—a bad press story, a complaint escalation, an audit finding—prompts a review.

Good HITL architecture prevents this with systematic monitoring.

Prediction monitoring. Track the distribution of confidence scores, label predictions, and model outputs over time. A shift in any of these distributions is a signal that something has changed—either the model’s behavior or the input distribution.

Human response monitoring. Track human reviewer agreement rates, time-per-review, and the rate at which human reviewers override automated decisions. Declining agreement rate or drastically faster reviews often signal reviewer fatigue or training distribution change. Rising override rates suggest the model is performing worse, or that the reviewers have developed new judgment that the model has not yet learned.

Outcome monitoring. Where outcomes are observable (fraud actually occurred or did not, content actually caused reported harm or did not), track whether the system’s predictions corresponded to reality. This is the ground truth check. It requires outcome labels, which are often delayed and incomplete—but even partial outcome monitoring is valuable.

Calibration monitoring. Periodically recompute the Expected Calibration Error (ECE) on held-out samples with known outcomes. If calibration is drifting, the routing rules need to be recalibrated.

8.6 Where Things Break in Production

Real HITL systems break in predictable ways.

Reviewer bottleneck The volume of cases routed to human review exceeds reviewer capacity. Reviewers speed up, quality drops, the feedback signal degrades. Downstream: model does not improve, error rates stay high.

Label quality degradation As a task becomes routine, reviewer attention decreases. Labels become less careful. Training data quality drops. Model performance plateaus or regresses.

Feedback lag Human labels from today do not reach the model for weeks or months. By the time the model learns from today's edge cases, those edge cases are no longer edge cases—the distribution has moved.

Distribution shift without detection The real-world data changes faster than the monitoring system detects. The model is making decisions on a distribution it was not trained on, but the monitoring dashboard still shows “all green” because it is tracking the wrong metrics.

The ghost reviewer A reviewer position that exists on paper but is chronically understaffed or filled by someone with insufficient training. The system reports that human oversight is in place. The reality is that oversight is nominal.

Specification gaming The model learns to game the metrics it is being evaluated on rather than learning the true task. If

the feedback loop is measuring the wrong thing, the model will optimize the wrong thing.

8.7 The Lifecycle: From Prototype to Deployed System

HITL systems evolve over time, and their architecture needs to evolve with them.

Phase 1—Human-heavy prototype. Almost all cases go to human review. The model exists but has thin training. This is the data collection phase—the human labels are the product, not the model outputs.

Phase 2—Hybrid operation. The model handles high-confidence cases automatically. Human review is concentrated on uncertain cases. The model is updated regularly on reviewed cases. This is the steady state for most production HITL systems.

Phase 3—Model-dominant. As training data accumulates, the model becomes more capable. Fewer cases require review. Human review concentrates on genuinely hard cases and edge cases. The human role becomes more expert and less routine.

Phase 4—Monitoring-dominant. The model handles most cases well. Human involvement shifts to monitoring, auditing, and handling edge cases and appeals. The architecture focuses on drift detection and outcome monitoring rather than case-by-case review.

Each phase has different architectural requirements. A system optimized for Phase 1 data collection will need significant changes to serve

well in Phase 3 model-dominant operation. Planning for these transitions is part of good HITL architecture design.

Sketch the HITL architecture for a system you interact with regularly—your email spam filter, your bank’s fraud detection, a recommendation system. Where is the prediction layer? Where is the routing layer? Where is the interface? And most importantly: can you identify a feedback integration layer? If you cannot find the feedback loop, what does that tell you about whether the system is actually designed to improve?

Chapter Summary

Key Concepts

- HITL architecture requires a loop, not just a pipeline—human responses must flow back into the system
- Seven components: prediction, uncertainty quantification, routing, interface, feedback integration, monitoring, and audit layers
- Feedback integration can be batch retraining, online learning, active learning, or RLHF—each with different latency and quality tradeoffs
- Production HITL systems break in predictable ways: reviewer bottleneck, label quality degradation, feedback lag, distribution shift without detection

- The system lifecycle has distinct phases (human-heavy → hybrid → model-dominant → monitoring-dominant), each requiring different architecture

Key Examples

Uber Tempe incident (2018) A human-in-the-loop that was not—a pipeline with a human placeholder; uncertainty was detected but not communicated to the safety driver

Active learning queues Intelligent routing for maximum learning value, not just uncertainty-based routing

RLHF in language models Preference-based feedback as an alternative to label-based feedback; the architecture behind ChatGPT and similar systems

Five Dimensions Check

- *Feedback Integration* (Dimension 5): The core subject of this chapter
- *Uncertainty Detection* (Dimension 1): Required for the routing layer to function
- *Stakes Calibration* (Dimension 4): Informs the routing thresholds

Next chapter: a practical guide to the tools anyone can use—from annotation platforms to cloud-based HITL services to open-source options for getting started with real data.

Technical Appendix 8A: HITL System Architecture Patterns

Chapter 8 introduced the blueprint for smart helper systems: routing, queuing, monitoring, and feedback. This appendix provides formal definitions, complete reference implementations, and production considerations for each component.

— Formal foundations for Chapter 8

8A.1 The Pipeline vs. Loop Distinction

A **HITL Pipeline** is a directed acyclic graph $G = (V, E)$ where: V includes exactly one human review node H ; H receives input from a model node M ; H produces output consumed by a downstream process; and no path exists from any downstream node back to M .

A **HITL Loop** is a directed graph $G = (V, E)$ where: V includes at least one human review node H ; and a path exists from H back to some component that influences future model behavior.

The loop property is necessary but not sufficient for a learning HITL system. The path from H back to the model must also satisfy:

1. Human labels are available for integration into training.
2. Labels are quality-filtered before integration.

3. The training process is actually triggered periodically.

A common failure mode is a nominally “looping” HITL system whose feedback queue is never processed. The system collects human labels, stores them, and retrains on schedule—except the schedule was never activated in production. Verify that the retraining job is running, not just that the queue is filling.

8A.2 The Three-Zone Router

The core routing component partitions incoming cases into three zones: auto-positive, auto-negative, and human review, with dynamic threshold adjustment based on reviewer capacity.

```

1  from dataclasses import dataclass
2  from enum import Enum
3  from typing import Optional, Callable
4
5  class Disposition(Enum):
6      AUTO_POSITIVE = "auto_positive"
7      AUTO_NEGATIVE = "auto_negative"
8      HUMAN_REVIEW = "human_review"
9
10 @dataclass
11 class RoutingConfig:
12     tau_low: float # Below this: auto-negative
13     tau_high: float # Above this: auto-positive
14     capacity_fn: Callable # Returns current reviewer
15         capacity [0,1]
16     overflow_narrowing: float = 0.05

```



```
45         self.tau_high = max(self.tau_high -
46                             self.config.overflow_narrowing,
47                             self.tau_low + min_band)
48     elif capacity > 0.8: # Underutilised: widen the band
49         base_low = self.config.tau_low
50         base_high = self.config.tau_high
51         self.tau_low = max(self.tau_low -
52                             self.config.underflow_widening,
53                             base_low - 0.15)
54         self.tau_high = min(self.tau_high +
55                             self.config.underflow_widening,
56                             base_high + 0.15)
```

Listing 8.1: Three-zone router with capacity adaptation

Stakes Override

The `stakes` parameter implements a high-stakes override: when a case is flagged as high-stakes (e.g., a medical decision, a financial transaction above a threshold), the review band widens, so that even moderately confident scores fall into the human-review zone. This pushes high-consequence cases toward human review regardless of reviewer capacity.

8A.3 Active Learning Queue Implementation

Cases routed for human review should be prioritized not by arrival order but by *learning value*: the expected improvement to the model from labeling this particular case. The implementation below combines uncertainty (proximity to the decision boundary) with diversity (dissimilarity from already-queued cases).

```

1 from dataclasses import dataclass, field
2 from typing import List
3 import numpy as np
4 from heapq import heappush, heappop, heapify, nsmallest
5
6 @dataclass(order=True)
7 class LearningQueueItem:
8     learning_value: float
9     score: float = field(compare=False)
10    case_id: str = field(compare=False)
11    features: np.ndarray = field(compare=False)
12
13    def __post_init__(self):
14        object.__setattr__(self, 'learning_value',
15                           -abs(self.learning_value)) #
16                                                    max-heap
17
18 class ActiveLearningQueue:
19     def __init__(self, uncertainty_weight=0.7,
20                 diversity_weight=0.3, max_size=1000):
21         self.heap = []
22         self.uw = uncertainty_weight
23         self.dw = diversity_weight
24         self.max_size = max_size
25         self._labeled = []
26
27     def _uncertainty(self, score: float) -> float:
28         return 1 - abs(2 * score - 1) # peaks at 0.5
29
30     def _diversity(self, features: np.ndarray) -> float:
31         if not self._labeled:
32             return 1.0
33         labeled = np.array(self._labeled)

```

```

33     min_dist = np.linalg.norm(labeled - features,
34                               axis=1).min()
35     return min(min_dist / (min_dist + 1.0), 1.0)
36
37 def add(self, case_id: str, score: float, features:
38       np.ndarray):
39     lv = self.uw * self._uncertainty(score) + \
40         self.dw * self._diversity(features)
41     heappush(self.heap, LearningQueueItem(lv, score,
42                                           case_id,
43                                           features))
44
45     if len(self.heap) > self.max_size:
46         self.heap = nsmallest(self.max_size, self.heap)
47         heapify(self.heap)
48
49 def pop_top_k(self, k: int) -> List[LearningQueueItem]:
50     results = []
51     for _ in range(min(k, len(self.heap))):
52         item = heappop(self.heap)
53         results.append(item)
54         self._labeled.append(item.features)
55     return results

```

Listing 8.2: Active learning priority queue

8A.4 Production Monitoring

A production HITL system requires continuous monitoring of three signals: *confidence distribution drift* (is the model seeing different inputs than it was trained on?), *human override rate* (are reviewers frequently disagreeing with automated decisions?), and *review rate* (is the fraction of cases reaching human review unexpectedly high?).


```
1 from collections import deque
2 from datetime import datetime
3 from typing import Dict, Optional
4 import numpy as np
5 from scipy import stats
6
7 class HITLMonitor:
8     def __init__(self, reference_confidence_dist:
9         np.ndarray,
10         window_size: int = 1000,
11         review_band: tuple = (0.3, 0.7),
12         alert_thresholds: Optional[Dict] = None):
13         self.reference = reference_confidence_dist
14         self.conf_window = deque(maxlen=window_size)
15         self.override_window = deque(maxlen=window_size)
16         self.review_low, self.review_high = review_band
17         self.thresholds = alert_thresholds or {
18             'drift_pvalue': 0.01,
19             'override_rate_high': 0.30,
20             'review_rate_high': 0.60,
21         }
22
23     def record(self, score: float, disposition: str,
24         human_override: Optional[bool] = None):
25         self.conf_window.append(score)
26         if human_override is not None:
27             self.override_window.append(int(human_override))
28
29     def check_health(self) -> Dict:
30         if len(self.conf_window) < self.conf_window.maxlen
31             // 2:
32             return {'status': 'insufficient_data'}
33         current = np.array(self.conf_window)
```

```

32     ks_stat, p_value = stats.ks_2samp(self.reference,
33                                     current)
34     issues = {}
35     if p_value < self.thresholds['drift_pvalue']:
36         issues['distribution_drift'] = {
37             'severity': 'HIGH',
38             'ks_stat': round(ks_stat, 4),
39             'p_value': round(p_value, 6),
40         }
41     if len(self.override_window) >= 50:
42         override_rate =
43             np.mean(list(self.override_window))
44         if override_rate >
45             self.thresholds['override_rate_high']:
46             issues['high_override_rate'] = {
47                 'severity': 'MEDIUM',
48                 'override_rate': round(override_rate,
49                                     3),
50             }
51     review_rate = sum(self.review_low <= s <=
52                     self.review_high
53                     for s in self.conf_window) \
54                     / len(self.conf_window)
55     if review_rate >
56         self.thresholds['review_rate_high']:
57         issues['high_review_rate'] = {
58             'severity': 'LOW',
59             'review_rate': round(review_rate, 3),
60         }
61     return {'status': 'issues_detected' if issues else
62           'healthy',
63           'issues': issues}

```

Listing 8.3: HITL system health monitor

The Kolmogorov–Smirnov test used above detects distributional drift in one-dimensional confidence scores. It does not detect drift in the high-dimensional feature space that produces those scores. A model may show stable confidence distributions while operating on inputs that have shifted significantly from its training distribution. Complement confidence monitoring with feature-space monitoring using dimensionality reduction.

8A.5 Exercises

Exercise 8A.1 — Feedback Integration Layer

Implement a feedback integration layer that: ingests human labels from the review queue; quality-filters them using inter-annotator agreement scores; and triggers retraining when accumulated new labels exceed a threshold. What safeguards prevent bad labels from degrading the model?

Exercise 8A.2 — Time-of-Day Threshold Adjustment

Extend the `ThreeZoneRouter` to implement time-of-day threshold adjustment: lower the review band during business hours (more capacity) and narrow it during nights and weekends. How would you test whether this adjustment improves overall system performance?

Exercise 8A.3 — Ghost Reviewer Detection

Implement a detector for the “ghost reviewer” failure mode: a review queue that has been accumulating items without being processed. What metrics would indicate this? How would you design an alert that distinguishes reviewer absence from legitimate queue growth?

Exercise 8A.4 — High-Dimensional Drift Detection

The monitor checks for distribution drift using a KS test on confidence scores. What are the limitations of this approach? Propose an alternative using dimensionality reduction of input features (e.g., UMAP or t-SNE on penultimate-layer activations). How would you set alert thresholds?

This appendix provides reference implementations for the core components of a production HITL architecture: routing, active learning queues, and system health monitoring, as introduced in Chapter 8.

9 Tools Anyone Can Use

The tools are not magic. But the landscape is rich enough now that the question of what tool to use has real answers.

— on tools

9.1 The Label Studio Moment

Open-source annotation tools like Label Studio did something their creators did not fully anticipate: they put real labeling capacity in the hands of organizations that could never afford a commercial vendor. Built by a small startup and released as open source, Label Studio was aimed at machine learning teams at well-funded labs—but it was taken up by a very different audience: non-profits labeling satellite images of deforestation, hospital systems annotating medical records, small newsrooms classifying articles for studies of media coverage.

None of these organizations had data science teams in the traditional sense. None of them could afford enterprise annotation platforms. But they had data, they had a task, and they had people willing to do the labeling work. What they needed was infrastructure.

This is the story of the past decade in HITL tooling. The infrastructure—the platforms, the pipelines, the APIs, the hosted services—has become accessible to organizations that would have been entirely locked out five years earlier. The engineering sophistication required to build a human-in-the-loop system from scratch has dropped dramatically. And with it,

the range of people and institutions who can participate in building AI systems that use human judgment has expanded.

The tools are not magic. They have real limitations, and choosing the wrong tool for a task can cost more in rework than building the right thing from scratch. But the landscape is rich enough now that the question “what tool should I use?” has real answers.

9.2 What Annotation Tools Actually Do

Before surveying the landscape, a quick reminder of what these tools are for.

An annotation tool provides:

- A way to present data items (images, text, audio, video, structured records) to reviewers
- A way for reviewers to record their judgments in structured form
- A way to store those judgments and export them for use in training
- Usually: some way to manage the workflow—who reviews what, when, quality control

That is the core. Advanced tools add:

- Agreement tracking and quality metrics
- Active learning queues (routing cases by model uncertainty)
- Model-assisted annotation (the model pre-labels, humans correct)
- Analytics dashboards
- Multi-reviewer reconciliation

- Team management and access control

The choice of tool depends on your data type, your task complexity, your team size, your budget, and how much customization you need.

9.3 Self-Hosted and Open-Source Options

Label Studio is among the most widely used open-source annotation platforms. It supports all major data types: text, images, audio, video, time series, and HTML. The labeling interface is configured with a simple XML-style tag language, which keeps it relatively accessible. It has a Python SDK and REST API, active community development, and a commercial cloud offering for teams that want hosted infrastructure.

Best for: research projects, small-to-medium teams with technical capacity, organizations with privacy requirements that prevent cloud hosting.

Limitations: setup requires some engineering; advanced features (active learning integration, enterprise workflow management) require more work.

Argilla (formerly Rubrix) focuses specifically on NLP annotation—text classification, named entity recognition, question answering, and evaluation of language model outputs. It integrates well with the Hugging Face ecosystem and has strong support for RLHF-style feedback collection. It is notably good for LLM evaluation use cases.

Best for: NLP teams, LLM evaluation, RAG pipeline improvement.

CVAT (Computer Vision Annotation Tool) is purpose-built for computer vision tasks. Originally developed by Intel, it supports bounding boxes, polygons, keypoints, and semantic segmentation. It is the tool of choice for many computer vision annotation projects.

Best for: image and video annotation, object detection, segmentation tasks.

Docanno and Brat are simple open-source text annotation tools requiring minimal technical setup. They are effective for single annotators or small teams on text classification tasks, but do not scale well.

Prodigy is a commercially licensed but developer-friendly tool from the creators of spaCy. It is designed for fast, scriptable annotation workflows with a particular emphasis on active learning integration. Prodigy's model-in-the-loop approach, where the model pre-annotates and humans correct, can dramatically reduce annotation time on NLP tasks.

Best for: NLP tasks where speed is critical and you have some technical capacity.

9.4 Cloud Platform Approaches

For organizations that want managed infrastructure rather than self-hosted tooling, several cloud platforms offer HITL services.

Scale AI and **Surge AI**, along with similar enterprise platforms, offer managed annotation workforces, quality controls, domain-specific pipelines (medical, legal, autonomous vehicle), and APIs for integrating annotation into ML pipelines. These managed workforces handle the recruitment, training, and quality management of annotators—a significant operational burden for organizations that do not want to manage their own annotation team. One caution on Scale in particular: after Meta's June 2025 purchase of a 49% stake in the company, several of Scale's largest customers pulled back, and buyers now weigh its independence as a vendor more carefully than they once did.

Best for: large-scale annotation projects, organizations without in-house annotation capacity, high-stakes domains requiring managed quality control.

Amazon SageMaker Ground Truth integrates annotation directly into AWS's ML toolchain. It offers both Amazon Mechanical Turk (crowd-sourced) and private workforce (your own team) options. Its automated data labeling feature provides a built-in form of active learning, automatically sending uncertain cases to human reviewers and using confirmed labels to improve the model. One significant caveat: AWS has announced that it will close Ground Truth to new customers in July 2026 and does not plan to introduce new features, so organizations choosing a platform for a new project should weigh the alternatives in this chapter.

Best for: existing AWS customers with Ground Truth deployments; new projects should look elsewhere.

Google Vertex AI Data Labeling is Google Cloud's managed annotation service. Similar in positioning to SageMaker Ground Truth but within GCP. Tight integration with Vertex AI for active learning loops.

Labelbox is a mid-market platform with strong support for complex multi-step annotation projects, quality assurance workflows, and a notable focus on video annotation.

Best for: medium-to-large organizations with complex annotation workflows, teams that need strong quality assurance tooling.

9.5 Crowdsourcing vs. Expert Labeling

This is one of the most consistently misunderstood trade-offs in annotation work.

Crowdsourced annotation (Amazon Mechanical Turk, Scale AI’s managed workforce, similar platforms) provides large-scale, fast, relatively low-cost annotation. The workforce is diverse and large. The limitation: for any task with significant domain requirements—medical text, legal documents, specialized imagery, culturally specific content—crowdworkers without that domain expertise produce noisy labels. Noise in training data is not random; it clusters around the hard cases that require domain judgment. These are often the most important cases.

Expert annotation (domain specialists—doctors, lawyers, linguists, security analysts) produces high-quality labels on hard cases. The limitation: expensive, slow, difficult to scale, and hard to find. Expert annotators are also not immune to bias—their expertise brings systematic viewpoints that may not reflect the diversity of real-world cases (Davani et al., 2022).

The practical resolution for most organizations: use the task to determine who labels it. Objectively determinable tasks (does this image contain a dog?) can use crowdsourcing successfully. Subjectively determinable tasks requiring domain judgment (does this symptom presentation suggest a particular diagnosis?) require experts. For tasks in between, consider a two-stage approach: crowd for initial filtering, experts for ambiguous cases.

Chapter 7 covers inter-annotator agreement in depth. Here, the practical tooling note: most annotation platforms support measuring inter-annotator agreement automatically (Krippendorff’s α , Cohen’s κ), and using agreement scores to route low-agreement items to additional review.

9.6 No-Code Options and Their Limits

Several platforms offer annotation workflows without any coding required.

Roboflow provides a no-code interface for computer vision annotation with integrated training and deployment. For image classification and object detection projects with small-to-medium scale, it is a complete pipeline with minimal technical friction.

Datasaur offers a no-code text annotation platform with good NLP support and collaboration features.

The honest assessment of no-code options: they work well for small, well-defined tasks. They hit ceilings quickly when you need active learning integration or custom label schemas, or when you have privacy requirements or annotation volumes that rule out the managed services. They are a useful starting point, not a long-term architecture.

9.7 When to Build vs. Buy

The build vs. buy decision in HITL tooling is more nuanced than in most software contexts because the annotation workflow is coupled to the ML pipeline in ways that create vendor lock-in risks.

Buy when:

- Your task fits a standard annotation template (object detection, text classification, NER)
- You need scale quickly and do not have engineering capacity
- Data privacy is not a concern that rules out cloud hosting
- Your workflow requirements are stable (you know what you are building)

Build when:

- Your annotation task is non-standard or highly domain-specific
- You have strong privacy requirements that rule out external processing
- Your HITL workflow needs deep integration with your ML pipeline
- You have long-term product vision that requires full workflow control

Most organizations end up using an off-the-shelf annotation tool (open-source or commercial) as the user interface layer, while building custom integration with their ML infrastructure for routing, feedback, and monitoring. The annotation UX is not where the architectural moat is. The feedback integration layer is.

9.8 Getting Started Practically

A realistic path to a working HITL system for a solo developer or small team:

Week 1: Define the task precisely. What is the input? What label is a reviewer assigning? What does “correct” look like? Write the annotation guidelines before you pick the tool. If you cannot write clear annotation guidelines, the task is not defined enough to build a system for.

Week 2: Collect a small sample and try labeling it yourself. 100–200 examples. Where do you hesitate? Those are your hard cases. How long does each label take? That determines your reviewer throughput estimate.

Week 3: Set up Label Studio or Argilla. Get a working annotation interface. Have a second person try to use it without help. Fix the parts they cannot figure out.

Week 4: Train a baseline model on your first 200 labels. Even a simple model. Use it to get uncertainty estimates. Route the bottom third by confidence to your review queue.

Month 2: Iterate. Review 200 more cases. Retrain. Check whether the low-confidence cases the model routed were actually harder. If they were—your uncertainty detection is working. If not—debug your calibration.

This is not glamorous. It is what building a real HITL system looks like.

Visit the Label Studio GitHub repository and browse the example templates. Pick one task that matches something you know about. Spend 15 minutes imagining what annotation guidelines you would write for that task. What is easy to write? What is hard? What would you need to know before you could write clear instructions? This exercise—writing annotation guidelines before touching any code—is the most important diagnostic for whether a HITL project is ready to build.

Chapter Summary

Key Concepts

- The annotation tooling landscape ranges from open-source (Label Studio, Argilla, CVAT) to cloud platforms (Scale AI, SageMaker Ground Truth) to no-code options
- Crowdsourced vs. expert annotation: match the annotator expertise level to the task's domain requirements
- No-code tools hit ceilings quickly; the right long-term architecture often combines off-the-shelf annotation UX with custom feedback integration
- Build vs. buy depends on task novelty, privacy requirements, and long-term pipeline coupling
- The most important first step is writing annotation guidelines—before choosing a tool

Key Examples

Label Studio The tool that democratized annotation for non-ML organizations—used by non-profits, hospitals, and news organizations

RLHF-focused tools like Argilla Specialized for LLM evaluation and preference collection

Active learning in SageMaker Ground Truth Architectural integration of uncertainty routing in a managed cloud service—and a lesson in vendor churn, with the service closing to new customers in July 2026

Five Dimensions Check

- *Feedback Integration* (Dimension 5): Tools are the infrastructure for closing the loop
- *Intervention Design* (Dimension 2): The annotation interface is the intervention design implementation

Next chapter: language is the hardest thing to teach a computer, and it is also where the most interesting HITL work is happening right now.

Technical Appendix 9A:

Annotation Pipeline Tooling

Chapter 9 introduced the annotation economy and the quality control mechanisms that keep it from producing noise. This appendix provides Python reference implementations for production annotation pipelines using Label Studio, inter-annotator agreement measurement, and multi-stage quality routing.

— Formal foundations for Chapter 9

9A.1 Label Studio API Integration

Label Studio provides an open-source annotation platform with a Python SDK supporting programmatic project management, task import, and export.

```
1 from label_studio_sdk import Client
2
3 ls = Client(url='http://localhost:8080',
4             api_key='YOUR_API_KEY')
5 ls.check_connection()
6
7 # Text classification project
8 project = ls.start_project()
```

```
8     title='Document Classification',
9     label_config=''
10    <View>
11        <Text name="text" value="$text"/>
12        <Choices name="label" toName="text" choice="single">
13            <Choice value="relevant"/>
14            <Choice value="irrelevant"/>
15            <Choice value="uncertain"/>
16        </Choices>
17    </View>
18    '''
19 )
```

Listing 9.1: Label Studio connection and project creation

9A.1.1 Task Import with Pre-Annotations

```
1  # Import with pre-annotations (model-assisted annotation)
2  pre_annotated_tasks = [
3      {
4          "text": "Document text here.",
5          "predictions": [{
6              "result": [{
7                  "from_name": "label",
8                  "to_name": "text",
9                  "type": "choices",
10                 "value": {"choices": ["relevant"]},
11                 "score": 0.87
12             }]
13         }]
14     }
15 ]
16 project.import_tasks(pre_annotated_tasks)
```

Listing 9.2: Importing tasks with model pre-annotations

9A.1.2 Export and Parsing

```
1 import pandas as pd
2
3 annotations = project.export_tasks(export_type='JSON_MIN')
4 results = []
5 for task in annotations:
6     if task.get('annotations'):
7         ann = task['annotations'][0]
8         label = ann['result'][0]['value']['choices'][0]
9         results.append({
10             'id': task['id'],
11             'text': task['data']['text'],
12             'label': label,
13             'annotator': ann.get('completed_by',
14                                 {}).get('id'),
15             'created_at': ann.get('created_at'),
16         })
17 df = pd.DataFrame(results)
```

Listing 9.3: Annotation export and label parsing

9A.2 Complete HITLAnnotationPipeline

The following class integrates quality control, inter-annotator agreement measurement, and expert routing into a single pipeline object.

```

1 import numpy as np
2 from label_studio_sdk import Client
3 from sklearn.metrics import cohen_kappa_score
4 from collections import defaultdict
5 from typing import Dict, List, Tuple
6 import logging
7
8 logger = logging.getLogger(__name__)
9
10 class HITLAnnotationPipeline:
11     """
12     Production annotation pipeline with quality control.
13
14     Features:
15     - Automatic overlap sampling for IAA measurement
16     - Per-annotator kappa computation with outlier detection
17     - Multi-stage routing: fast pass -> expert review
18     """
19
20     def __init__(self, ls_url: str, api_key: str,
21                  project_id: int,
22                      min_kappa: float = 0.60,
23                      overlap_fraction: float = 0.10,
24                      expert_threshold: float = 0.45):
25         self.client = Client(url=ls_url, api_key=api_key)
26         self.project = self.client.get_project(project_id)
27         self.min_kappa = min_kappa
28         self.overlap_fraction = overlap_fraction
29         self.expert_threshold = expert_threshold
30
31     def sample_overlap_items(self, items: List[dict]
32                             ) -> Tuple[List[dict],
33                                         List[dict]]:
```

```

33     n_overlap = max(30, int(n * self.overlap_fraction))
34     np.random.shuffle(items)
35     return items[:n_overlap], items[n_overlap:]
36
37     def compute_project_kappa(self) -> Tuple[float,
38         List[str]]:
39         tasks =
40             self.project.export_tasks(export_type='JSON')
41         overlap = [t for t in tasks
42             if len(t.get('annotations', [])) >= 2]
43         if not overlap:
44             return float('nan'), []
45
46         per_task: Dict[int, Dict[str, str]] = {}
47         for task in overlap:
48             per_task[task['id']] = {}
49             for ann in task['annotations']:
50                 aid = str(ann.get('completed_by',
51                     {}).get('id', 'unknown'))
52                 if ann.get('result'):
53                     res = ann['result'][0]['value']
54                     label = res['choices'][0]
55                     per_task[task['id']][aid] = label
56
57         all_ann = list({a for labels in per_task.values()
58             for a in labels})
59         pair_kappas = {}
60         for i in range(len(all_ann)):
61             for j in range(i + 1, len(all_ann)):
62                 a, b = all_ann[i], all_ann[j]
63                 la, lb = [], []
64                 for labels in per_task.values():
65                     if a in labels and b in labels:

```

```

63         la.append(labels[a]);
           lb.append(labels[b])
64         if len(la) >= 30: # Appendix 7A: require N
           >= 30 items per pair
65             pair_kappas[(a, b)] =
               cohen_kappa_score(la, lb)
66
67         if not pair_kappas:
68             return float('nan'), []
69
70         # Aggregate as the mean of per-pair kappas: pooling
           labels
71         # across different pairs would mix rater marginals,
           and the
72         # result would be no pair's kappa (nor Fleiss'
           kappa).
73         aggregate =
           float(np.mean(list(pair_kappas.values()))))
74
75         avg_kappa = defaultdict(list)
76         for (a, b), k in pair_kappas.items():
77             avg_kappa[a].append(k); avg_kappa[b].append(k)
78         low_quality = [aid for aid, ks in avg_kappa.items()
79                       if np.mean(ks) <
                           self.expert_threshold]
80
81         return aggregate, low_quality
82
83     def route_for_expert_review(
84         self, completed: List[dict]
85     ) -> Tuple[List[dict], List[dict]]:
86         _, low_quality = self.compute_project_kappa()
87         accepted, expert_queue = [], []
88         for task in completed:

```



```
116                                     / total_anns,  
117                                }
```

Listing 9.4: Production annotation pipeline with quality control

9A.3 Tool Comparison

9A.4 Exercises

Exercise 9A.1 — Label Studio Setup

Using the Label Studio SDK, set up a binary text classification project. Import 50 sample items. Annotate 20 items manually. Export the results and compute the distribution of labels.

Exercise 9A.2 — Majority Vote Resolution

Extend the HITLAnnotationPipeline class to support configurable majority vote resolution for overlap items. When 3 annotators label an item and 2 agree, accept the majority. When all 3 disagree on a 3-class task, route to expert review.

Exercise 9A.3 — Real-World IAA Analysis

Using the SNLI or MultiNLI datasets from HuggingFace (which include per-annotator labels), compute pairwise kappa for all annotator pairs. Identify the 10% most contested items. What is the kappa improvement from removing them?

Exercise 9A.4 — Cost Estimation

Using the `cost_estimate` function, calculate the total cost for a project of 10,000 items at \$0.10/annotation with the default 10% overlap fraction. Then extend the function with `expert_fraction` and `expert_cost_usd` parameters modeling expert re-annotation at \$1.00/annotation, and compute the total cost when 20% of items receive expert re-annotation.

This appendix provides production-ready Python tooling for annotation pipelines, inter-annotator agreement measurement, and quality-controlled expert routing, as introduced in Chapter 9.

Table 9.1: Annotation platform comparison

Tool	Data types	Managed	Self-host	Best use case
Label Studio	All	No	Yes	Research, privacy-sensitive
Argilla	NLP, LLM eval	No	Yes	RLHF preference collection
CVAT	Image, video	No	Yes	Object detection, segmentation
Prodigy	NLP	No	License	Fast NLP with active learning
Scale AI	All	Yes	No	High-volume commercial
SageMaker GT	All	Yes	AWS	AWS-integrated workflows
Labelbox	All	Yes	No	Enterprise, collaborative

Part III

Teaching Machines to Perceive

Language is uniquely hard for computers—and the difficulty is a gift, because it keeps the human in the loop.

— Chapter 10

10 Teaching Computers to Understand Language

Language is the most human thing we do. That computers can appear to do it is either our greatest technical achievement—or our most elaborate illusion.

— on language and machine

10.1 The Lawyer Who Trusted the Machine

On June 8, 2023, a federal judge in New York held an unusual sanctions hearing. Two attorneys—Steven Schwartz and Peter LoDuca of the firm Levidow, Levidow & Oberman—sat before Judge P. Kevin Castel in the Southern District of New York. They had filed a legal brief in a personal injury case against Avianca airlines. The brief cited half a dozen case precedents. The judge had noticed something strange: he couldn't find them.

He couldn't find them because they didn't exist.

ChatGPT, the AI assistant attorney Schwartz had used to research the case, had fabricated six court decisions—complete with plausible-sounding case names, docket numbers, courts of jurisdiction, and summarized holdings. *Martinez v. Delta Air Lines*. *Varghese v. China Southern Airlines*. *Shaboon v. Egypt Air*. None were real. When opposing counsel filed a brief pointing this out, Schwartz asked ChatGPT to verify the cases. The AI

doubled down, insisting they were “real and can be found in reputable legal databases.” Schwartz attached printed ChatGPT transcripts to a court filing, apparently as evidence that he had done his due diligence.

Judge Castel imposed sanctions on both attorneys in a written decision issued two weeks later. In the order that had first flagged the fabrications, he had already written something that should be printed on the login screen of every AI assistant: “The Court is presented with an unprecedented circumstance. A submission filed by plaintiff’s counsel in opposition to a motion to dismiss is replete with citations to non-existent cases.” (“Mata v. Avianca, Inc.” 2023)

This story reveals something profound about how language models work—and how they fail. These are systems that can discuss Aristotle, write sonnets, and explain quantum mechanics, yet they will confidently cite cases that never happened, with the same untroubled certainty they’d use to tell you the capital of France. To understand why, we need to understand what language models actually are—and what “understanding” means when you’re made of math.

10.2 Why Language Is Hard

Before computers, before neural networks, before anyone dreamed of machines that could write—long before all of that—there was the problem of language itself.

Language is, in a word, slippery.

Consider the sentence: “I saw the man with the telescope.” Did you use the telescope to see him? Or was he carrying the telescope? Both readings are grammatically valid. English doesn’t tell you.

Now try: “The chicken is ready to eat.” Is the chicken about to eat something? Or has it been cooked and is awaiting consumption? Again—both are correct. Context resolves the ambiguity, but context is everything: it’s the restaurant menu, the farmyard, the smell coming from the kitchen.

Or consider irony: “Oh, great. Another Monday.” A computer reading that sentence, unaware that Monday often brings dread, might classify it as a positive statement. Sentiment analysis tools still stumble on sarcasm, because understanding irony requires a model of what the speaker wanted you to *think they meant* as opposed to what they *literally said*—a layered theory of mind that took humans hundreds of thousands of years to evolve.

Language also carries cultural baggage. “That’s sick” can be a compliment among teenagers and an insult at the dinner table. “Wicked” means evil in standard English and excellent in Boston. References to “the shot heard round the world” mean something entirely different to a baseball fan than to a history teacher.

And then there’s implication. “Do you know what time it is?” is not a question about your epistemological relationship to clock-reading. It’s a request for the time. Humans automatically perform this translation. Computers, for decades, could not.

These challenges—ambiguity, context, irony, cultural reference, implication—are not edge cases. They are the *substance* of language. Every sentence ever uttered is woven through with them. For five decades, researchers tried to solve this with rules. Write down enough grammar rules. Build enough ontologies. It mostly failed. Language was too vast, too fluid, too context-dependent for any rulebook.

Then came a different idea: what if instead of teaching computers the rules of language, you let them *learn* language from the sheer volume of text that humans had already written?

10.3 Words as Coordinates in Meaning-Space

The first breakthrough in modern language AI came in 2013, from a small team at Google. Tomáš Mikolov and colleagues published a technique called **Word2Vec** (Mikolov et al., 2013). The idea sounds almost too simple: train a neural network to predict which words appear near other words in text, and the process of learning those predictions will encode meaning.

Here's the intuition. Words that appear in similar contexts tend to mean similar things. "Dog" and "cat" both appear near "pet," "owner," "veterinarian," "food bowl." "King" and "queen" both appear near "throne," "crown," "royal," "palace."

Word2Vec turned these contextual patterns into coordinates—**vectors** in a high-dimensional space, where each dimension captures some aspect of meaning. The result was startling: not only did similar words cluster near each other geometrically, but the *relationships* between words were encoded in the directions between them.

The famous example: the vector for "king" minus the vector for "man," plus the vector for "woman," gave something very close to the vector for "queen." The geometry of meaning-space mirrored the geometry of human concepts.

But word embeddings had a fundamental limitation. They gave each word a single fixed location in meaning-space, regardless of context. The word "bank" has the same vector whether you're talking about a river bank or a financial institution. Real understanding requires knowing which *sense* of a word is meant in a particular sentence—for that, you need something that reads the whole sentence to determine meaning.

10.4 The Attention Mechanism: Reading While Noticing

In 2017, a team at Google published a paper with a deceptively casual title: “Attention Is All You Need” (Vaswani et al., 2017). It was, by any measure, one of the most consequential papers in the history of artificial intelligence.

The key idea was a mechanism called **attention**, which can be understood through an analogy.

Imagine you’re reading a mystery novel. You’re on page 300, and you’ve just reached a scene where the detective notices a red coat. To understand why the red coat is significant, your brain doesn’t process just the current sentence. Without conscious effort, your attention drifts back to page 47, where a witness mentioned a figure in red. And to page 112, where the suspect claimed to own no bright clothing. Your mind holds the current word and simultaneously asks: *which other words, anywhere in what I’ve read, are relevant to understanding this one?*

That is attention.

The **transformer** architecture—the “T” in ChatGPT, the engine inside nearly every modern language AI—does exactly this, but for every word simultaneously.

When a transformer processes the sentence “The animal didn’t cross the street because it was too tired,” it asks, for the word “it”: which other words in this sentence should I weight heavily when interpreting this one? “Animal” scores highly. “Street” scores low. “Tired” scores very high, because the attribution of tiredness resolves the pronoun reference. The word “it” is processed *in relation* to all other words at once.

This happens not once but many times simultaneously, in parallel “heads” of attention, each head learning to track different kinds of relationships—syntactic ones, semantic ones, coreference chains, discourse

structure. When all these heads are combined and processed through layers upon layers of the same mechanism, something remarkable emerges: representations of text that capture context at every level of abstraction.

The transformer architecture unleashed a scaling law that no one fully expected (Kaplan et al., 2020). The more text you fed it, the more parameters you used, the more compute you threw at the problem—the better the results, in a way that showed no sign of plateauing. Researchers fed transformers first millions of documents, then billions of web pages, then essentially the entire written output of human civilization. The models got better and better at predicting text, which—it turned out—was a surprisingly good proxy for getting better at understanding it.

This is how we arrived at **large language models**: GPT-4, Claude, Gemini, LLaMA, Mistral. Vast statistical engines with hundreds of billions of learned parameters.

10.5 What a Language Model “Knows”

Here is where we need to be careful—because the word “knows” is doing a lot of work, and it’s often doing it dishonestly.

When a language model appears to “know” something, what’s actually happening is this: during training, the model processed enormous amounts of text in which certain patterns co-occurred. It learned, in a statistical sense, that “Paris” and “capital of France” appear together constantly. That “Abraham Lincoln” and “1865” and “assassinated” appear together. That “Avianca” and “airline” appear together, and that “Martinez v. Delta Air Lines”—well, that’s where things get interesting.

The model didn’t learn “Martinez v. Delta Air Lines” from any case reporter. It learned the *pattern* of what court citations look like. And it learned what kinds of phrases appear near “airline injury lawsuits.” When asked to produce precedents for airline injury cases, it did what it

always does: it generated text that was statistically plausible given the prompt. It produced strings of text that *looked like* court citations without any mechanism to distinguish between citing something that existed and generating something plausible-sounding.

The model has no memory in the human sense. It has no index pointing from case names to actual decisions. It has something more like a vast, compressed statistical portrait of what human writing looks like—and it uses that portrait to generate text that continues naturally from whatever you’ve given it.

This is not the same as knowing. It is, at its best, an extraordinary mimicry of knowing. The distinction matters enormously—and it’s a distinction the models themselves cannot make.

10.6 The Hallucination Problem

The AI research community has a name for what happened to attorney Schwartz’s legal brief: **hallucination**.

A language model hallucinates when it produces text that is confident, fluent, plausible in form—and factually wrong. The term captures something real: the wrongness isn’t random noise. It’s confabulation—the generation of plausible-seeming content that fills a gap where actual knowledge should be.

Why does hallucination happen? Several reasons, each illuminating.

Training objective. The model is trained to predict what comes next in text. Real text sometimes contains errors, propaganda, myths, and mistakes. The model learned from all of it. When it generates text, it generates what’s statistically likely given its training—not what’s true.

No “I don’t know” token. The model has no mechanism for expressing genuine uncertainty about a specific fact. Every query gets a response. The model has learned to produce fluent text, and silence is not a token in its vocabulary.

Confidence is stylistic, not calibrated. The model’s confidence is encoded in the fluency and assertiveness of its output style, not in any internal calibration. It writes “The case was decided on April 14, 1992, by the Second Circuit” with the same stylistic assurance whether the case is real or fabricated.

Research into hallucination rates is sobering. A 2024 Stanford study found that general-purpose language models hallucinated on 58% to 88% of specific, verifiable questions about federal court cases (Dahl et al., 2024). A companion study of *purpose-built* legal research tools—systems explicitly marketed as nearly hallucination-free—found that even these still fabricated information between roughly 17% and 33% of the time (Magesh et al., 2025). And in medical AI, the uncertainty behind diagnostic outputs is rarely communicated in ways that clinicians can interpret or act on (Kompa et al., 2021). Even retrieval-augmented systems—models that can look things up—hallucinate at measurable rates. These figures were measured on early 2024 systems; subsequent model generations have reduced, though not eliminated, these errors.

The Calibration Core of Hallucination

The hallucination problem is, at its core, a calibration problem. The model isn’t uncertain where it should be uncertain. Its expressed confidence doesn’t match its actual accuracy. This is the same problem this book has been tracking all along—in GPS systems, in

medical AI, in the Air Canada chatbot. The system doesn't know what it doesn't know.

10.7 Three Shapes of Human-in-the-Loop in Language AI

Language AI has not one but three very distinct shapes of HITL deployment, each demanding very different design thinking.

10.7.1 The Chatbot: Real-Time Assistance

The first shape is the one most people encounter: the AI assistant deployed to answer questions, help write emails, or provide customer service in real time.

Here, the HITL challenge is primarily one of **uncertainty communication and escalation**. When should the system answer directly? When should it hedge? When should it say “I recommend you check this with a human expert”? When should it refuse to answer entirely?

The Air Canada chatbot failed this test catastrophically (Civil Resolution Tribunal of British Columbia, 2024). It had no mechanism to distinguish confident knowledge from confident confabulation. It had no escalation path to a human agent when uncertainty was high. It behaved, in effect, as though every question were equally within its competence.

Well-designed chatbot systems do the opposite. They are trained to recognize categories of query where the model's knowledge is unreliable—time-sensitive information, jurisdiction-specific legal matters, medical diagnosis, anything requiring personal context the model doesn't have. For these categories, they hedge explicitly: “I can give you general information, but for your specific situation, you should consult a licensed professional.”

This hedging is not a limitation—it is a feature, exactly the kind of *Stakes Calibration* we described in earlier chapters. The stakes for getting medical advice wrong are not the same as the stakes for getting a recipe wrong.

10.7.2 The Content Moderator: Harmful Content Detection

The second shape of HITL in language AI is less visible to most users but arguably more consequential: content moderation at scale.

Language AI does the first pass. Human reviewers handle the escalations. This sounds elegant, but the practice is bruising. The language of harm is deeply contextual. “I want to kill my boss” is, in virtually all instances, idiomatic frustration, not a genuine threat. Satire of a political figure can look exactly like earnest endorsement of the views being lampooned. A slur reclaimed by the community it targets means something different from the same slur used as an attack.

Automated systems miss these distinctions constantly. They over-remove legitimate content—often disproportionately affecting minority communities whose speech patterns fall outside the training distribution—and under-remove harmful content that’s been camouflaged through deliberate adversarial manipulation.

The human reviewers who handle escalations are some of the least-discussed people in the AI ecosystem. Research has documented significant psychological harm to content moderators exposed to graphic violence, child abuse material, and extremist propaganda (Steiger et al., 2021). The feedback they provide trains the next version of the automated system.

This is HITL in its most difficult form: humans in the loop not as occasional consultants but as the cognitive infrastructure that makes the whole system function.

10.7.3 The Alignment Trainer: Preference Feedback

The third shape of HITL in language AI—using human feedback to align language model behavior with human values—has received the most attention from researchers and is the domain of RLHF, DPO, and Constitutional AI.

10.8 Teaching Values: RLHF, DPO, and Constitutional AI

10.8.1 Reinforcement Learning from Human Feedback

Imagine you’ve trained a language model on billions of words. It can generate fluent text. But “fluent” is not the same as “good.” The model might generate text that’s plausible but harmful, factually wrong, or simply unhelpful.

Reinforcement Learning from Human Feedback (RLHF)—developed through work at OpenAI, DeepMind, and Anthropic (Ouyang et al., 2022)—works roughly like this:

Generate a lot of model outputs for a variety of prompts. Show pairs of outputs to human raters and ask: which is better? The raters compare responses on dimensions like helpfulness, honesty, harmlessness, and accuracy. These judgments train a **reward model**—an AI that learns to predict which outputs humans prefer. The original language model is then fine-tuned using reinforcement learning, nudging it toward outputs the reward model would rate highly (Christiano et al., 2017).

The genius of RLHF is that it converts the fuzzy notion of “good output” into something a model can optimize for. Instead of trying to write rules for every possible scenario, you hire humans to express their preferences, train a model on those preferences, and use that as a training signal.

The limitation of RLHF is that it's expensive, slow, and the humans in the loop introduce their own biases. There's also a well-documented problem of **reward hacking**—the model finding ways to score highly on the reward model without actually being helpful, a bit like a student who learns to write essays that sound good to automated graders without engaging with the substance.

10.8.2 Direct Preference Optimization

Direct Preference Optimization (DPO) (Rafailov et al., 2023) is a more recent approach that sidesteps the need for a separate reward model. DPO directly adjusts the language model's weights based on preference comparisons, using a mathematical insight: the reinforcement learning problem RLHF is trying to solve has a closed-form solution computable directly from preference data. This makes training more stable, more efficient, and less prone to some forms of reward hacking.

For our purposes, the important point is the same: humans are in the loop, providing preference judgments that shape what the model learns to do.

10.8.3 Constitutional AI

A different approach was pioneered by Anthropic under the name **Constitutional AI (CAI)** (Bai et al., 2022). Rather than relying purely on human raters to judge millions of individual outputs, CAI starts with a set of written principles—a “constitution”—describing how the AI should behave: be helpful, be honest, avoid harmful outputs, respect human autonomy.

The AI is then asked to critique and revise its own outputs based on these principles, in a process that can be largely automated. This drastically reduces the need for human feedback on individual examples, while still

allowing humans to exert significant influence—through the design of the constitution itself.

Constitutional AI represents an interesting reconfiguration of the HITL relationship. The human is in the loop not at the level of individual judgments (“this output is better than that one”) but at the level of values (“these are the principles we want the system to embody”). It’s a shift from line-editing to authorship—from correcting individual mistakes to writing the ethical framework the system will use to correct itself.

None of these approaches is perfect. They all depend, ultimately, on humans—on the quality of the raters’ judgments, the wisdom of the constitution’s authors, the diversity and representativeness of the human feedback. That dependency is not a weakness to be engineered away. It is the point.

10.9 The Five Dimensions, Applied to Language AI

Uncertainty Detection. The calibration picture for language models is nuanced: base models can be reasonably well calibrated on multiple-choice formats, but the free-form verbalized confidence users actually encounter often is not (Kadavath et al., 2022). They express uncertainty through hedging language (“I think,” “I believe,” “you may want to verify”), but this hedging is learned stylistic behavior, not a reflection of actual probability estimates. The challenge: the model generates output token by token with no natural point where genuine uncertainty is computed and communicated.

Intervention Design. The best implementations use a tiered approach: answer directly when confidence is high and stakes are low; hedge and suggest verification when confidence is low or ambiguous; refuse and escalate when stakes are high and confidence is insufficient. The chatbot

that says “For specific legal advice, please consult a licensed attorney” is practicing good intervention design.

Timing. In real-time language AI, timing of human intervention is particularly constrained. A customer service chatbot that routes to a human after ten minutes of unhelpful circular conversation has failed the Timing dimension. Intervention needed to happen at the first signal of inadequate confidence.

Stakes Calibration. Language AI systems need domain-aware stakes calibration. Getting a recipe wrong has different consequences than getting a drug interaction wrong. The best systems have been trained with category-specific behavior: more hedging in high-stakes domains, clearer escalation for medical, legal, and safety-critical queries.

Feedback Integration. RLHF and its descendants are, at bottom, systematic feedback integration—the most elaborate version of this dimension yet designed. Every preference judgment from a human rater is a piece of feedback that shapes future behavior. The challenge is ensuring this feedback loop is representative, unbiased, and fast enough to keep pace with the ways people use the model.

10.10 Does It Understand? A Philosophical Aside

We have been using the word “understand” carefully throughout this chapter. Let’s address it directly, because it matters more than it might seem.

When a transformer model processes “The animal didn’t cross the street because it was too tired” and correctly identifies that “it” refers to the animal rather than the street—has it *understood* that sentence?

One view says: yes. It produced the correct output. It tracked the relevant relationships. What more could understanding require?

Another view says: no. Understanding implies something beyond correct output. It implies a model of the world inside—representations of animals and streets and tiredness that connect to a broader understanding of physical reality. The transformer has learned statistical patterns over text. It has, perhaps, a kind of textual shadow of these concepts. Whether that shadow constitutes genuine comprehension is unclear (Bender et al., 2021; Searle, 1980).

This is not an idle philosophical question. It has direct practical implications for where human oversight is most needed.

A system that genuinely understands what it's saying will, presumably, notice when what it's saying is false. A system that is producing statistically plausible text has no such internal check. If the distinction matters, it tells us something important: language AI systems need human oversight not as a workaround for a temporary technical limitation, but as a structural feature of what these systems are.

The attorney's problem was not that ChatGPT was a flawed prototype that future versions would fix. It was that ChatGPT is, at its core, a very sophisticated text prediction machine—and text prediction machines, however sophisticated, do not have a relationship to truth. They have a relationship to plausibility. Newer models have narrowed that gap, but closing it requires humans in the loop.

The next time you use an AI assistant for something that matters—a health question, a legal question, a factual claim you'll repeat to someone else—ask it to express its confidence. Then verify the answer through another source. How often does it hedge appropri-

ately? How often is it confidently wrong? What does this tell you about when AI tools need human supervision?

Chapter Summary

Key Concepts

- Language is intrinsically hard for computers because of ambiguity, context-dependence, irony, cultural reference, and implication
- Word embeddings gave words coordinates in “meaning-space”—similar words cluster together
- Transformer models use attention to process words in context of all other words simultaneously
- Language models “know” things statistically, not factually—they have no internal check on truth
- Hallucination is a calibration failure: models express high confidence when they should express uncertainty
- Three HITL shapes in language AI: real-time chatbot assistance, content moderation, and alignment training
- RLHF, DPO, and Constitutional AI are systematic approaches to incorporating human values through preference feedback

Key Examples

Mata v. Avianca (2023) Attorneys sanctioned for submitting AI-fabricated legal citations; demonstrates the hallucination problem's real-world consequences

Air Canada chatbot (2024) No escalation path for high-stakes queries; confidently wrong about policy

Content moderation reviewers The unseen human cognitive infrastructure behind automated moderation

RLHF alignment training Human preference judgments as the signal that makes AI helpful rather than merely fluent

Constitutional AI Human values encoded at the constitutional level, not the individual judgment level

Five Dimensions Check

- *Uncertainty Detection* (Dimension 1): Hallucination is confident wrongness—a language model that cannot signal what it does not know fails this dimension by construction
- *Stakes Calibration* (Dimension 4): The Avianca brief shows what happens when a high-stakes domain inherits a system calibrated for low-stakes fluency
- *Feedback Integration* (Dimension 5): RLHF and its successors are feedback integration at industrial scale—human preferences, not just labels, shaping what the model says

In Chapter 11, we turn from language to vision—and discover that the way machines learn to see reveals something just as strange about the nature of machine intelligence.

Technical Appendix 10A: The Mathematics of Language Models

The surprising thing is not that attention works — it's that it works so well that we still aren't sure why.

— on the mathematics of meaning

10A.1 Token Probability Distributions and Sampling

Language models don't generate words — they generate **tokens**. A token is a sub-word unit, typically a word or a common word-fragment produced by Byte Pair Encoding (BPE) (Sennrich et al., 2016). At each step, the model produces a probability distribution over the entire vocabulary — typically on the order of 100,000–250,000 tokens in current models (e.g., Llama 3: 128K, GPT-4o: ~200K, Gemma: 256K); earlier models used 32,000–100,000.

Formally, given a sequence of previous tokens $x_1, x_2, \dots, x_t$, the model produces:

$$P(x_{t+1} \mid x_1, x_2, \dots, x_t) = \text{softmax}(\mathbf{W}_o \cdot \mathbf{h}_t) \quad (10.1)$$

where $\mathbf{h}_t$ is the final hidden-state vector and $\mathbf{W}_o$ is the output projection matrix. The softmax ensures probabilities sum to 1:

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}} \quad (10.2)$$

10A.1.1 Sampling Strategies

Temperature sampling divides logits by temperature T before softmax:

$$P_T(x_{t+1}) = \text{softmax}\left(\frac{\mathbf{z}}{T}\right) \quad (10.3)$$

At $T < 1$, high-probability tokens dominate (conservative output); at $T > 1$, distribution flattens (more varied output).

Top- p (nucleus) sampling (Holtzman et al., 2020) selects the smallest set of tokens whose cumulative probability exceeds threshold p :

$$\mathcal{N}_p = \min S \subseteq V \text{ such that } \sum_{x \in S} P(x) \geq p \quad (10.4)$$

This adapts dynamically: a peaked distribution yields a small nucleus (confident model); a flat distribution yields a large nucleus (uncertain model).

A model's hallucination tendency is partly a sampling artifact. When the model is genuinely uncertain — the correct token has, say, 15% probability and the second-best has 12% — temperature and nucleus parameters determine whether this uncertainty is amplified or suppressed. The model has no internal mechanism to communicate “I am sampling from a very flat distribution here; my output should be hedged.” This is why calibration research matters.

10A.2 The Attention Mechanism

The transformer's core operation is **scaled dot-product attention** (Vaswani et al., 2017):

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V \quad (10.5)$$

where:

- $Q \in \mathbb{R}^{n \times d_k}$ — **query** matrix (what we're looking for)
- $K \in \mathbb{R}^{m \times d_k}$ — **key** matrix (what's available to be looked up)
- $V \in \mathbb{R}^{m \times d_v}$ — **value** matrix (what gets returned)
- $\sqrt{d_k}$ — scaling factor preventing dot products from saturating the softmax

In **self-attention**, queries, keys, and values are all derived from the same input sequence $X \in \mathbb{R}^{n \times d_{\text{model}}}$:

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V \quad (10.6)$$

where $W_Q, W_K \in \mathbb{R}^{d_{\text{model}} \times d_k}$ and $W_V \in \mathbb{R}^{d_{\text{model}} \times d_v}$ are learned projection matrices. The output $\text{softmax}(QK^\top / \sqrt{d_k})$ is the **attention weight matrix** $A \in \mathbb{R}^{n \times m}$, where A_{ij} represents how strongly position i attends to position j .

10A.2.1 Multi-Head Attention

Multi-head attention runs h attention operations in parallel with different learned projections:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W_O \quad (10.7)$$

$$\text{head}_i = \text{Attention}(QW_{Q_i}, KW_{K_i}, VW_{V_i}) \quad (10.8)$$

Different heads learn to track different relationship types: coreference resolution, syntactic dependencies, semantic roles. The combination captures language structure at multiple levels simultaneously.

10A.2.2 Transformer Layer

Each transformer layer combines multi-head attention with a feed-forward network:

$$\mathbf{h}' = \text{LayerNorm}(\mathbf{x} + \text{MultiHead}(\mathbf{x})) \quad (10.9)$$

$$\mathbf{h}'' = \text{LayerNorm}(\mathbf{h}' + \text{FFN}(\mathbf{h}')) \quad (10.10)$$

where $\text{FFN}(\mathbf{x}) = \max(0, \mathbf{x}W_1 + b_1)W_2 + b_2$ is a position-wise MLP with ReLU activation. Modern LLMs stack tens to hundreds of such layers; the deepest layers capture the most abstract semantic properties.

10A.3 RLHF Pipeline

10A.3.1 Stage 1: Supervised Fine-Tuning (SFT)

Starting from a pretrained language model π_{PT} , collect high-quality human-written responses y^* to prompts x , and fine-tune via maximum likelihood:

$$\mathcal{L}_{\text{SFT}} = -\mathbb{E}_{(x, y^*) \sim \mathcal{D}_{\text{SFT}}} [\log \pi_{\theta}(y^* | x)] \quad (10.11)$$

10A.3.2 Stage 2: Reward Model Training

Collect preference data: for pairs of outputs (y_w, y_l) shown to human raters, record $y_w \succ y_l$. Train reward model $r_{\phi}(x, y)$ using the Bradley-Terry pairwise preference model (Christiano et al., 2017):

$$P(y_w \succ y_l) = \sigma(r_\phi(x, y_w) - r_\phi(x, y_l)) \quad (10.12)$$

Minimize negative log-likelihood of observed preferences:

$$\mathcal{L}_{\text{RM}} = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}_{\text{pref}}} \left[\log \sigma(r_\phi(x, y_w) - r_\phi(x, y_l)) \right] \quad (10.13)$$

10A.3.3 Stage 3: RL Fine-Tuning via PPO

Fine-tune π_{SFT} against the learned reward model using Proximal Policy Optimization (Ouyang et al., 2022):

$$\max_{\theta} \mathbb{E}_{x \sim \mathcal{D}} \left[\mathbb{E}_{y \sim \pi_\theta(\cdot | x)} [r_\phi(x, y)] - \beta \cdot \text{KL}[\pi_\theta(\cdot | x) \parallel \pi_{\text{SFT}}(\cdot | x)] \right] \quad (10.14)$$

The KL divergence term $\beta \cdot \text{KL}[\pi_\theta \parallel \pi_{\text{SFT}}]$ is a crucial regularizer that prevents the fine-tuned model from drifting into degenerate behavior that exploits the reward model. The PPO clipped update is:

$$\mathcal{L}_{\text{PPO}} = \mathbb{E}_t \left[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon) \hat{A}_t) \right] \quad (10.15)$$

where $r_t(\theta) = \pi_\theta(y_t | x, y_{<t}) / \pi_{\text{old}}(y_t | x, y_{<t})$ is the probability ratio and $\hat{A}_t$ is the estimated advantage.

10A.4 Direct Preference Optimization (DPO)

DPO (Rafailov et al., 2023) achieves the same alignment goal as RLHF without a separate reward model.

Key insight: The RLHF objective (10.14) has a closed-form optimal solution:

$$\pi^*(y | x) = \frac{1}{Z(x)} \pi_{\text{ref}}(y | x) \exp\left(\frac{r(x, y)}{\beta}\right) \quad (10.16)$$

Rearranging, the reward is expressible in terms of the optimal policy:

$$r(x, y) = \beta \log \frac{\pi^*(y | x)}{\pi_{\text{ref}}(y | x)} + \beta \log Z(x) \quad (10.17)$$

Substituting into the Bradley-Terry preference probability (noting $\log Z(x)$ cancels):

$$P(y_w \succ y_l | x) = \sigma \left(\beta \log \frac{\pi_{\theta}(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \beta \log \frac{\pi_{\theta}(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right) \quad (10.18)$$

The DPO training objective directly maximizes the likelihood of observed preferences:

$$\mathcal{L}_{\text{DPO}}(\theta) = -\mathbb{E}_{(x, y_w, y_l)} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \beta \log \frac{\pi_{\theta}(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right) \right] \quad (10.19)$$

Gradient interpretation: The DPO gradient simultaneously increases $\pi_{\theta}(y_w | x)$ and decreases $\pi_{\theta}(y_l | x)$, weighted by how much the current model disagrees with the preference. The reference model π_{ref} keeps the optimization anchored, preventing degenerate solutions where the model assigns near-zero probability to everything except winning responses.

10A.5 Constitutional AI: Two-Phase Process

Constitutional AI (CAI) (Bai et al., 2022) reduces dependence on human-labeled comparisons by using AI feedback guided by a written constitution of principles.

10A.5.1 Phase 1: Supervised Learning for Constitutional AI (SL-CAI)

1. **Generate:** Sample outputs $y \sim \pi_0(\cdot | x)$ for adversarial “red-team” prompts x designed to elicit harmful responses.
2. **Critique:** Prompt the model to critique its output against constitutional principle c_i :

$$\text{critique} = \pi_0(\cdot | x, y, c_i) \quad (10.20)$$

3. **Revise:** Prompt the model to revise based on the critique:

$$y' = \pi_0(\cdot | x, y, \text{critique}) \quad (10.21)$$

4. **Fine-tune:** Train π_0 on (x, y') pairs to get $\pi_{\text{SL-CAI}}$.

10A.5.2 Phase 2: Reinforcement Learning from AI Feedback (RLAIF)

1. For each prompt, generate two outputs from $\pi_{\text{SL-CAI}}$.
2. Obtain AI preference labels: present the pair with constitutional principle c_i to the model and elicit a preference:

$$P_{\text{AI}}(y_w \succ y_l | x) = \pi_{\text{SL-CAI}}(\text{“A is better”} | x, y_1, y_2, c_i) \quad (10.22)$$

3. Train reward model r_ϕ on AI preference labels using the Bradley-Terry objective (10.13).
4. Fine-tune $\pi_{\text{SL-CAI}}$ via PPO against r_ϕ using the objective (10.14).

The HITL Reconfiguration in CAI

In standard RLHF, humans are in the loop at the level of individual output comparisons. In CAI, humans are in the loop at the level of constitutional design. This shifts the human contribution from millions of individual judgments to a small number of carefully considered principles. The quality of the resulting model depends critically on the quality, completeness, and internal consistency of the constitution.

10A.6 Calibration Metrics for Language Models

A model is **calibrated** if its expressed confidence matches its empirical accuracy. For classification outputs, the canonical measure is **Expected Calibration Error (ECE)** (Guo et al., 2017):

$$\text{ECE} = \sum_{b=1}^B \frac{|B_b|}{n} |\text{acc}(B_b) - \text{conf}(B_b)| \quad (10.23)$$

where predictions are grouped into B bins by confidence level, $|B_b|$ is the number of predictions in bin b , $\text{acc}(B_b)$ is the fraction of correct predictions in that bin, and $\text{conf}(B_b)$ is the mean confidence.

A perfectly calibrated model has $\text{ECE} = 0$: every confidence bin's accuracy exactly matches its stated confidence.

For generative language models, calibration is harder to measure because outputs are sequences, not class probabilities. Active research areas include:

- **Verbalized calibration:** Asking the model to state its confidence in percentage terms; measuring ECE on those verbalized estimates (Kadavath et al., 2022)

- **Token-level probability calibration:** Measuring whether the probability assigned to the correct completion correlates with actual correctness frequency
- **Semantic calibration:** Measuring calibration at the level of factual claims, not individual tokens

Research consistently shows that large language models are overconfident — their verbalized confidence significantly exceeds their actual accuracy on factual queries, particularly in specialized domains like law, medicine, and science (Magesh et al., 2025).

These technical foundations underpin the practical HITL design questions in Chapter 10: when to trust a language model’s outputs, when to require human verification, and how to build systems that communicate their limitations honestly.

11 Teaching Computers to See

We assume that things that look the same to us look the same to everything that looks. The adversarial example is a reminder that this assumption is wrong.

— on machine vision

11.1 The Stop Sign That Wasn't

In 2017, a team of researchers from the University of Washington, UC Berkeley, and other institutions published a paper with an unsettling finding (Eykholt et al., 2018). Using nothing more exotic than a color printer and some stickers, they could reliably make a stop sign invisible to autonomous vehicle vision systems—while remaining completely recognizable to every human who looked at it.

The attack was elegant. The researchers printed small black-and-white sticker patterns—some designed to look like graffiti, some like ordinary wear-and-tear—and affixed them to real stop signs. To any passing driver or pedestrian, the signs still read “STOP” in large white letters on the familiar red octagon. To the machine vision system they tested, the signs were classified as something else entirely: “Speed Limit 45” was the most common misclassification.

Not occasionally. Reliably. In multiple weather conditions. From multiple viewing angles. At vehicle speeds.

If machines see differently from humans in ways this dramatic and this exploitable, what does “teaching a computer to see” actually mean? The answer is stranger and more profound than it might initially seem—and understanding it reveals why human oversight of machine vision is not a temporary accommodation on the road to full automation. It is a structural requirement, rooted in the fundamental nature of what machine vision is.

11.2 Pixels as Numbers, Patterns as Everything

To a digital camera, an image is nothing more than a rectangular grid of numbers. Each position in the grid is a pixel. Each pixel is a triplet of numbers between 0 and 255 representing the intensity of red, green, and blue light. A 640×480 image is a block of $640 \times 480 \times 3 = 921,600$ numbers. No shapes. No objects. No meaning. Just numbers arranged in a grid.

Early computer vision tried to extract meaning from these numbers using hand-crafted rules: look for sharp transitions in pixel values (edges), then look for edges that form closed curves (contours), then match those contours against stored templates. This worked tolerably well for simple, controlled scenarios and catastrophically badly for anything more complex.

The revolution came in 2012, when Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton submitted an entry to the ImageNet Large Scale Visual Recognition Challenge (Krizhevsky et al., 2012). Their entry—AlexNet—used a **convolutional neural network** trained on two GPUs for about a week. It crushed the competition, achieving a top-5 error rate almost 11 percentage points lower than the next best system.

The deep learning era in computer vision had begun.

11.3 What a Convolutional Neural Network Actually Does

The intuition behind convolutional neural networks is one of the most beautiful ideas in modern AI.

Imagine systematically scanning an image for a face. You might first look for local patterns: gradual smooth regions (like a cheek), abrupt curved boundaries (like a hairline), specific arrangements of dark and light (like eyes against a forehead). Then you'd look for combinations of these patterns at the right scales and spatial relationships.

A **convolutional neural network** (CNN) does something structurally similar, learned from data rather than programmed by hand.

The **lowest layer** learns to detect the simplest patterns: edges—places where pixel values change sharply in a particular direction.

The **next layer** combines edge detections into slightly more complex patterns: corners, curves, short line segments.

The **layer after that** combines corners and curves into textures, simple shapes, partial outlines.

And so on, up through the network. By the time you reach the highest layers, the network is detecting things like “wheel-like circular shapes” or “skin-tone blobs flanking a dark vertical stripe”—components of concepts approaching the human-level categories we'd use to describe what we see.

The crucial insight: the network *learns* all of this from examples. You show it a million images labeled “cat” and a million labeled “not cat,” and it adjusts its internal weights until it can reliably distinguish them. The features it learns are not ones a human would have specified. They're whatever works, extracted from the statistics of the training images.

And this is precisely the source of both the power and the vulnerability.

11.4 What the Model Actually Sees—and What It Doesn't

The features a convolutional neural network learns are not the features humans use to understand images.

We know this because of **adversarial examples**.

Ian Goodfellow and colleagues at Google published a landmark 2015 paper demonstrating that you could take any image a neural network classifies correctly, add a carefully calculated pattern of noise—often invisible to the human eye—and cause the network to classify the image as anything you want, with near-absolute confidence (Goodfellow et al., 2015).

The canonical example: take a photograph of a panda. The network correctly identifies it with 57.7% confidence. Add a very specific pattern of pixel perturbations—perturbations so small they don't visibly change the image—and the network now classifies it as a gibbon with 99.3% confidence. The image looks identical to a human. The network sees something completely different.

This isn't a trick enabled by cheating or system access. The adversarial pattern exploits something fundamental about how convolutional networks process images: they are extremely sensitive to specific high-frequency patterns in the input that humans don't perceive and don't use for recognition.

Think about what this means. The network and the human are processing the same numbers and arriving at completely different interpretations. The stop sign with stickers is a physical version of the same phenomenon: the sticker pattern creates specific local features that cascade up through the architecture into a confident wrong classification, while human perception processes the sign at the level of semantic content—reading the word “STOP,” recognizing the shape, understanding the context—in ways that are robust to such local perturbations.

The Root of Machine Vision's Vulnerability

The adversarial example vulnerability is not a bug that will be fixed in the next software version. It is an inherent feature of how these systems work, rooted in the fact that they learned from pixel statistics rather than from the conceptual structure of the visual world. A system with genuine conceptual understanding of what a stop sign is—what it's for, how it's designed, how it differs from other signs—would be robust to sticker attacks, just as human vision is.

11.5 Three HITL Domains in Computer Vision

This fundamental difference in how machines see creates different challenges in different high-stakes applications. Let's examine three in detail.

11.5.1 Medical Imaging: The Augmented Radiologist

Medicine was among the first high-stakes domains where AI vision found genuine traction—and where the HITL question has been most carefully and formally studied.

Modern radiology AI systems can detect early-stage lung cancer in chest CTs and flag pneumonia in chest X-rays with accuracy competitive with experienced radiologists on the specific tasks they've been trained for (Ardila et al., 2019; Rajpurkar et al., 2017).

The key phrase is “specific tasks they've been trained for.” A radiologist looking at a chest CT is not only looking for the specific pathology her AI assistant was trained to detect. She is looking at the whole image, bringing decades of clinical intuition, noticing unexpected findings, integrating the scan with the patient's history, symptoms, and prior imaging. She is doing something far broader and more contextual than image classification.

This asymmetry has shaped how medical AI is deployed. The dominant paradigm is not replacement but **augmentation**: the AI examines the image, generates findings, assigns confidence scores, and the radiologist reviews both the raw image and the AI's assessment.

Research on this workflow has produced a genuinely important finding: radiologists who see the AI's confidence score systematically change their behavior in ways that are not always improvements.

When the AI confidently flags an abnormality, radiologists are more likely to agree with the AI's characterization. And when the advice itself is wrong, it drags the human down with it: in a controlled study, radiologists given deliberately inaccurate advice made significantly worse diagnoses than those given accurate advice—regardless of whether the advice was labeled as coming from an AI or from a human expert (Gaube et al., 2021). The AI's confident output acts as a cognitive anchor.

This is **automation bias** operating in its most consequential form: expressed AI confidence suppresses human vigilance precisely in the cases where human vigilance is most needed.

The design response has been nuanced. Some systems present AI findings *after* radiologists have done their initial read, so that human perception is not anchored by the AI's prior. Others present AI findings only as flagged regions, not diagnoses—prompting the radiologist to look without telling them what to conclude. The **Intervention Design** dimension turns out to be as consequential as the accuracy of the underlying algorithm.

The FDA has developed a specific regulatory pathway for AI/ML-based medical devices (U.S. Food and Drug Administration, 2021). Over 1,000 AI-enabled medical devices had received FDA authorization as of 2025—the large majority in radiology—with the number growing rapidly each year. This regulatory architecture is itself a form of institutionalized HITL: before any medical vision AI can be used on actual patients, it must demonstrate

performance to a human oversight body with the authority to require changes.

11.5.2 Autonomous Vehicles: Human in the Loop at 60 MPH

The physics of driving establish a constraint that makes the HITL design problem in autonomous vehicles uniquely brutal.

At 60 miles per hour, a car travels 88 feet per second. Typical human reaction time to an unexpected event is approximately 1.5 seconds. In the 1.5 seconds between when an autonomous system recognizes it cannot handle the current situation and when a human driver could realistically respond, the vehicle has already traveled 132 feet—nearly half a football field.

This makes the naive version of “human in the loop”—AI drives until uncertain, then asks human to take over—deeply inadequate as a safety model. By the time the human has processed the alert, looked up, assessed the situation, and responded, the situation has evolved far beyond the moment when the AI first recognized its uncertainty.

The autonomous vehicle industry has responded to this constraint by developing radically different approaches to where the human fits in the loop.

Tesla’s Full Self-Driving (Supervised) operates in a “Level 2” framework: the human driver is always responsible, hands near the wheel, required to be ready to take over at any moment. Tesla has essentially exited the “Timing” problem by requiring continuous human engagement. The downside: drivers habituate to the AI’s competence and stop paying genuine attention—precisely the automation bias problem observed in radiology, now at highway speeds.

Waymo’s approach is the opposite extreme. Waymo’s commercial vehicles operate without a human driver in the vehicle at all. The HITL happens remotely: human operators can observe vehicles that encounter uncertain situations and provide guidance. This solves the reaction-time problem (remote operators provide guidance that the vehicle’s AI then executes, not real-time physical control) but requires sufficient operational design confidence to drive without a physical fallback.

Cruise’s approach —before its service was suspended following a 2023 incident in San Francisco, and GM ultimately wound the robotaxi business down entirely in late 2024—was a kind of middle path: vehicles with safety drivers in some conditions, without in others, with remote human monitoring in all conditions.

What this divergence reveals is that “human in the loop for autonomous vehicles” is not a single design pattern but a design space, with fundamentally different architectures making fundamentally different tradeoffs between human reaction time, automation bias, regulatory acceptability, and operational cost. The **Timing** dimension of the Five Dimensions framework is at its most urgent and most difficult here.

11.5.3 Content Moderation at Scale: Hashing, Perception, and the Adversary

Image-based content moderation faces a challenge that language moderation does not: the stakes are higher for certain categories and the adversaries more resourceful.

For child sexual abuse material (CSAM), the dominant technical approach is **perceptual hashing**. A perceptual hash is a compact fingerprint of an image, computed in a way that’s deliberately similar for visually similar images. The PhotoDNA system, developed by Microsoft, computes a hash for every uploaded image and compares it against a database of

known CSAM hashes maintained by the National Center for Missing and Exploited Children (Farid, 2021).

The design is elegant: known harmful images can be detected and blocked automatically, without any human having to review the content.

But perceptual hashing systems have a fundamental vulnerability: they can only identify *known* harmful content—images that have already been identified, hashed, and added to the database. They cannot identify novel content: new images produced after the database was compiled.

And adversaries learn. Once a bad actor understands that image hashing is the detection mechanism, they can modify images in ways that change the hash while preserving the harmful content for human viewers. The arms race between hashing systems and adversarial evasion is active and ongoing (Jain et al., 2022).

This is where human updating of the hash database becomes critical: every new image of harm that is identified by a human reviewer, removed, and hashed adds to the detection capability. The human is in the loop not just as an adjudicator of edge cases but as the source of new ground truth in an adversarially evolving system.

The second challenge for image moderation is context. A photograph of a human body might be art, medical documentation, evidence of abuse, or exploitation. No hash-based system can distinguish context from pixels. Human judgment on these contextual distinctions is not a temporary workaround—it is the only mechanism available for evaluations that require understanding of intent, context, and meaning.

11.6 What “Seeing” Means When You’re Made of Math

Human vision is not just image processing. It’s perception—active, contextual, laden with meaning and expectation. When a human sees a stop sign,

they understand that it is embedded in a physical environment, connected to rules they've internalized, part of a situation with other actors following those rules.

When a convolutional neural network classifies an image, it is detecting statistical patterns that correlate with the presence of things labeled a certain way in training data. It has no model of the environment, no understanding of why things are labeled as they are, no connection between what it "sees" and any broader understanding of the world.

This is why adversarial examples are possible. The vulnerability exists precisely because the network's classification is based on statistical pattern matching, not conceptual understanding (Goodfellow et al., 2015).

The practical implication is the same either way: machine vision and human vision are complementary, with different strengths, different failure modes, and different robustness profiles. The radiologist's eyes catch the unexpected finding the AI wasn't trained to detect. The AI catches the early-stage nodule the radiologist would have dismissed as artifact. The human in the loop is not compensating for a deficiency—they're contributing a genuinely different capability.

11.7 The Five Dimensions in Visual AI

Uncertainty Detection. A well-calibrated vision system expresses low confidence on images outside its training distribution. In practice, most classification networks are overconfident, producing high-confidence predictions on inputs where it should be uncertain (Guo et al., 2017). Techniques like MC Dropout (Gal & Ghahramani, 2016) and deep ensembles (Lakshminarayanan et al., 2017) improve calibration but remain imperfect.

Intervention Design. In medical imaging, the specific format in which AI findings are presented to radiologists—highlighted regions, confidence

scores, differential diagnoses—significantly affects the quality of radiologist decisions. Research has found that the format of AI findings shapes human attention as much as the AI's accuracy (Gaube et al., 2021).

Timing. In autonomous vehicles, the gap between when the system detects uncertainty and when a human can meaningfully respond is a physics problem, not a software problem. Architectural choices (whether humans are physically present, remotely monitored, or not involved until post-hoc review) represent fundamentally different responses to this timing constraint.

Stakes Calibration. In image moderation, the intervention threshold should be set differently for different content categories, with human review calibrated to the stakes of each error type. Missing CSAM has different consequences than over-removing legitimate content. Precision-recall tradeoffs must be set with the asymmetric cost structure in mind.

Feedback Integration. When the training distribution changes—new camera hardware, new environments, new types of adversarial attack—the model's performance silently degrades. Human monitoring and feedback integration is how systems detect and respond to this drift. In content moderation, the hash database is a form of feedback integration: human reviewers contribute new ground truth that updates the automated system.

11.8 What's Next

The next four chapters of this book turn from the modalities—language and vision—to the system around them: measurement, human psychology, fairness, and failure. In each domain, the five dimensions take on specific texture. The timing constraint of a fraud detection queue looks nothing

like the timing constraint of a medical imaging review. The stakes of getting a hiring recommendation wrong are not the stakes of getting a moderation decision wrong.

But the underlying structure is the same: a system with genuine capabilities and genuine limitations, working alongside humans who have complementary capabilities and complementary limitations, with the quality of their collaboration determined by how well the system is designed to recognize its uncertainties and communicate them.

Machine vision will become more capable. But the stop sign with stickers is not a temporary limitation of 2017-era deep learning. It is a window into what machine vision fundamentally is: a powerful pattern detector that sees differently from us, fails differently from us, and will always need us most precisely at the moments when what it sees diverges most sharply from what is actually there.

Find an image recognition tool online—Google Lens, Apple’s Visual Look Up, or any image classifier. Photograph the same object under three different conditions: normal lighting, an unusual angle, and with some small modification (adding a sticker, changing the background, partial occlusion). How does the classifier’s confidence change? What does this tell you about what it’s actually detecting?

Chapter Summary

Key Concepts

- Images are grids of numbers; CNNs learn hierarchical pattern detectors from edge detections to abstract concepts
- The features CNNs learn differ fundamentally from the features humans use—adversarial examples make this concrete
- Machine vision and human vision are complementary, not interchangeable: different failure modes, different strengths
- Three HITL domains: medical imaging (augmented radiologists, automation bias), autonomous vehicles (the escalation timing problem), and content moderation (hashing, adversarial evasion, contextual judgment)
- Automation bias—AI confidence suppressing human vigilance—is as important as AI accuracy in medical AI design
- Machine vision's distinctive failure modes—adversarial sensitivity, distribution shift, and automation bias—are not temporary limitations of current technology; they are the permanent reasons why human-AI collaboration in visual domains is not a workaround but the correct architecture

Key Examples

AlexNet (2012) The deep learning breakthrough that launched modern computer vision

Adversarial examples (Goodfellow et al., 2015) Invisible pixel perturbations that fool neural networks with near-absolute confidence

Physical adversarial stop sign attack (2017) Stickers that fool autonomous vehicle vision while remaining obvious to humans

FDA AI device authorization pathway Institutional HITL for medical vision AI—over 1,000 authorized devices as of 2025

Waymo vs. Tesla vs. Cruise Three fundamentally different architectures for human oversight in autonomous vehicles—two survivors and a middle path GM shut down in late 2024

PhotoDNA Perceptual hashing for CSAM detection, and its structural vulnerability to adversarial evasion

Key Principles

- Machine vision failure modes are structural, not accidental—they arise from what deep learning is
- The HITL design challenge in autonomous vehicles is a physics problem (reaction time), not just a software problem
- In content moderation, the human is in the loop as ground-truth generator, not just edge-case adjudicator
- Automation bias means that how AI findings are presented is as important as the AI's accuracy

Five Dimensions Check

- *Uncertainty Detection* (Dimension 1): Adversarial examples and distribution shift mean a vision model's confidence can be highest

exactly where it is most wrong—external checks, not self-reports, must carry the signal

- *Timing* (Dimension 3): The autonomous-vehicle handoff shows timing as physics—a human alerted 1.5 seconds before impact is not meaningfully in the loop
- *Stakes Calibration* (Dimension 4): The FDA authorization pathway is stakes calibration institutionalized—medical vision AI must prove itself to a human oversight body before touching a patient

In the chapters ahead, we apply these insights to specific domains—and discover that knowing when to hand off to a human is the deepest form of machine intelligence.

Technical Appendix 11A: The Mathematics of Computer Vision

Chapter 11 introduced how machines learn to see. This appendix provides the mathematical treatment of convolutional networks, the standard metrics for object detection, uncertainty estimation methods, and active learning strategies for reducing annotation cost.

— Formal foundations for Chapter 11

11A.1 Convolutional Neural Networks

11A.1.1 The Convolution Operation

A convolutional layer applies a set of learned **filters** to an input feature map. For a 2D input $X \in \mathbb{R}^{H \times W}$ and a filter $K \in \mathbb{R}^{k \times k}$, the output at position (i, j) is:

$$Z[i, j] = \sum_{m=0}^{k-1} \sum_{n=0}^{k-1} X[i + m, j + n] \cdot K[m, n] + b$$

For color images with C channels and F filters, the full operation produces an output volume $Z \in \mathbb{R}^{H' \times W' \times F}$:

$$Z[i, j, f] = \sum_{c=1}^C \sum_{m=0}^{k-1} \sum_{n=0}^{k-1} X[i + m, j + n, c] \cdot K[m, n, c, f] + b_f$$

The filter $K[m, n, c, f]$ learns a feature detector: for edge detection it responds to intensity gradients; for complex patterns it encodes whatever local structure discriminates between training classes.

11A.1.2 Pooling and Spatial Hierarchy

Max pooling reduces spatial resolution while preserving the strongest local activations:

$$\text{MaxPool}(Z)[i, j] = \max_{m \in [0, k), n \in [0, k)} Z[i \cdot s + m, j \cdot s + n]$$

where s is the stride. Stacking convolution–activation–pooling layers creates a **spatial hierarchy**: lower layers detect edges and textures; higher layers detect object parts and whole objects.

11A.1.3 Residual Connections

He et al. (2016) showed that very deep networks suffer from vanishing gradients and introduced **residual connections** that allow gradients to flow directly across layers:

$$\mathbf{h}^{(l+1)} = F(\mathbf{h}^{(l)}) + \mathbf{h}^{(l)}$$

where $F(\mathbf{h}^{(l)})$ is the learned residual. This enables networks 50–152 layers deep without training instability.

11A.2 Object Detection: IoU and mAP

11A.2.1 Intersection over Union

Object detection requires localizing objects via bounding boxes. The standard metric for localization quality is **Intersection over Union** (IoU):

$$\text{IoU}(B_{\text{pred}}, B_{\text{gt}}) = \frac{\text{Area}(B_{\text{pred}} \cap B_{\text{gt}})}{\text{Area}(B_{\text{pred}} \cup B_{\text{gt}})}$$

For axis-aligned rectangles:

$$\begin{aligned} x_{\text{inter}} &= \max(0, \min(x_2^p, x_2^g) - \max(x_1^p, x_1^g)) \\ \text{IoU} &= \frac{x_{\text{inter}} \cdot y_{\text{inter}}}{\text{Area}(B_{\text{pred}}) + \text{Area}(B_{\text{gt}}) - x_{\text{inter}} \cdot y_{\text{inter}}} \end{aligned}$$

$\text{IoU} \in [0, 1]$. The standard threshold for a “correct detection” is $\text{IoU} \geq 0.5$.

11A.2.2 Mean Average Precision

For a detection system evaluated across N classes, the standard metric is **mean Average Precision** (mAP):

$$\text{mAP} = \frac{1}{N} \sum_{c=1}^N \text{AP}_c$$

where AP_c is the area under the precision–recall curve for class c . A detected box is a *true positive* if $\text{IoU} \geq \tau$ with a ground truth box of that class; otherwise a *false positive*. Ground truth boxes not matched by any detection are *false negatives*.

In real datasets, class distribution is heavily skewed. Standard mAP aggregates equally across all classes, masking poor performance on rare classes. For safety-critical systems, report per-class AP separately and monitor performance on underrepresented conditions.

11A.3 Uncertainty Estimation in Vision

11A.3.1 Monte Carlo Dropout

Gal and Ghahramani (2016) showed that applying dropout at inference time approximates Bayesian inference over network weights. Making T stochastic forward passes yields an ensemble of predictions:

```
1 import torch
2 import torch.nn.functional as F
3
4 def mc_dropout_predict(model, image, T=50):
5     model.train() # Keep dropout active at inference
6     predictions = []
7     for _ in range(T):
8         with torch.no_grad():
9             logits = model(image)
10            probs = F.softmax(logits, dim=-1)
11            predictions.append(probs)
12    predictions = torch.stack(predictions) # (T,
13        num_classes)
14    mean_pred = predictions.mean(dim=0)
15    uncertainty = predictions.std(dim=0)
16    return mean_pred, uncertainty
```

Listing 11.1: MC Dropout uncertainty estimation for vision

Decomposing uncertainty.

$$\hat{\sigma}_{\text{epistemic}}^2 = \frac{1}{T-1} \sum_{t=1}^T (f_t(x) - \bar{f}(x))^2 \quad \hat{\sigma}_{\text{aleatoric}}^2 = \frac{1}{T} \sum_{t=1}^T \sigma_t^2(x)$$

In medical imaging: epistemic uncertainty means the model has seen few training examples like this image (actionable — request a radiologist). Aleatoric uncertainty reflects inherent image ambiguity such as poor contrast (also worth flagging, but for different reasons).

11A.3.2 Deep Ensembles

Train M models with different random initialisations (Lakshminarayanan et al., 2017):

$$\bar{p}(y | x) = \frac{1}{M} \sum_{m=1}^M p_m(y | x) \quad \hat{\sigma}^2 = \frac{1}{M} \sum_{m=1}^M (p_m(y | x) - \bar{p}(y | x))^2$$

Deep ensembles produce better-calibrated uncertainties than MC Dropout but require M times the training compute and storage.¹

11A.4 Confusion Matrix for Object Detection

For object detection, the confusion matrix must account for missed detections (false negatives — objects present but not detected) and spurious detections (false positives — detections with no corresponding ground truth). An extended matrix for N classes plus a “background” class has shape $(N + 1) \times (N + 1)$:

¹The MC-dropout estimator above uses the unbiased $1/(T - 1)$ sample variance because the T passes are stochastic draws from an approximate posterior; the ensemble uses $1/M$, treating the M trained models as the full population rather than a sample.

- TP_c : objects of class c correctly detected with correct label and $IoU \geq \tau$
- $Conf(i \rightarrow j)$: objects of class i detected as class j (localization correct, label wrong)
- FN_c : objects of class c present but not detected
- FP_c : detections labeled as class c with no corresponding ground truth

In medical imaging, FN (missed findings) carry far higher consequences than FP (unnecessary follow-up). This asymmetry justifies operating the system at a threshold that minimizes FN rate, even at the cost of a higher FP rate.

11A.5 Active Learning for Image Annotation

11A.5.1 Uncertainty Sampling

Given a pool of unlabeled images $\mathcal{U}$, select the image with highest predictive entropy:

$$x^* = \arg \max_{x \in \mathcal{U}} \mathbb{H}[p(y | x)] = - \sum_c p(y = c | x) \log p(y = c | x)$$

Or using BALD (Bayesian Active Learning by Disagreement) (Houlsby et al., 2011), which selects images where uncertainty is epistemic (reducible with more data) rather than aleatoric:

$$x^* = \arg \max_{x \in \mathcal{U}} \mathbb{I}[y; \theta | x, \mathcal{D}] = \mathbb{H}[p(y | x)] - \mathbb{E}_\theta[\mathbb{H}[p(y | x, \theta)]]$$

11A.5.2 Core-Set Selection

Select a set of images $S \subseteq \mathcal{U}$ that maximally covers the feature space:

$$S^* = \arg \min_{|S|=n} \max_{x \in \mathcal{U}} \min_{s \in S} d(x, s)$$

where $d(x, s)$ is cosine distance in penultimate-layer feature space. This ensures the annotation budget covers the diversity of the unlabeled pool rather than clustering around already-well-represented examples.

11A.5.3 Active Learning Loop

```

1 def active_learning_loop(model, labeled_pool,
2   unlabeled_pool,
3   annotation_budget,
4   query_strategy):
5     while annotation_budget > 0:
6       model.fit(labeled_pool)
7       scores = query_strategy.score(model,
8         unlabeled_pool)
9       batch_size = min(BATCH_SIZE, annotation_budget)
10      to_annotate = select_top_k(unlabeled_pool, scores,
11        batch_size)
12
13      # The HITL intervention: human annotation
14      new_labels = human_annotator.annotate(to_annotate)
15
16      labeled_pool.add(to_annotate, new_labels)
17      unlabeled_pool.remove(to_annotate)
18      annotation_budget -= batch_size
19      log_performance(model.evaluate(validation_set),
20        len(labeled_pool))
21    return model, labeled_pool

```

Listing 11.2: Active learning loop for image annotation

The `human_annotator.annotate()` call represents the HITL intervention. The quality of the final model depends on both the query strategy (selecting informative images) and the quality of the human annotation (correctly labeling them).

11A.6 Exercises

Exercise 11A.1 — IoU Implementation

Implement IoU from scratch for axis-aligned bounding boxes. Test it on a case of perfect overlap ($\text{IoU} = 1$), no overlap ($\text{IoU} = 0$), and a partial case you compute by hand. Then implement mAP for a 3-class detector using a synthetic dataset.

Exercise 11A.2 — MC Dropout Uncertainty Decomposition

Using a pretrained image classifier with dropout layers:

1. Generate in-distribution test images and OOD images (from a different dataset).
2. Run MC Dropout ($T = 50$ passes). Compute epistemic and aleatoric uncertainty separately.
3. Show that epistemic uncertainty is higher for OOD images.
4. Plot uncertainty maps as a function of model confidence.

Exercise 11A.3 — Active Learning for CIFAR-10

Implement an active learning loop on CIFAR-10:

1. Start with 500 randomly labeled examples. Train a baseline CNN.
2. Compare random sampling, entropy sampling, and BALD as query strategies.
3. Plot model accuracy as a function of the number of labeled examples.
4. How many labeled examples does the best active learning strategy need to match the accuracy of 5,000 random labels?

This appendix provides the mathematical foundations for understanding how machines see, how we measure their seeing, and how humans remain essential in the loop of teaching machines to see better, as introduced in Chapter 11.

12 Measuring Success

When a measure becomes a target, it ceases to be a good measure.
— on the measurement trap

12.1 The 400-Reviews-Per-Hour Problem

Picture a mid-sized e-commerce platform—the case that follows is a composite, but every element of it is drawn from documented moderation failures. The platform deployed a content moderation system to screen product listings for policy violations. The algorithm handled roughly 85% of listings automatically. The rest—borderline cases involving possibly counterfeit goods, prohibited items, or ambiguous descriptions—got routed to a team of human reviewers.

Management wanted to know: is this working?

They chose a simple metric to answer that question: reviews completed per hour. Reviewers who processed more cases were rated higher. Reviewers who lagged got performance improvement conversations. Within six weeks, average throughput climbed from about 40 reviews per hour to over 400.

Impressive. And catastrophically wrong.

What had actually happened was that reviewers, under pressure to move faster, had learned to approve most borderline cases with a glance. The algorithm was routing them uncertain cases—exactly the cases that *needed* careful human attention. But “careful attention” doesn’t score well

when the metric is speed. So reviewers pattern-matched the easy ones, skimmed past the hard ones, and clicked approve.

The model, which was designed to learn from reviewer decisions, dutifully incorporated these rubber-stamp approvals as ground truth. Over several months, its threshold for flagging violations drifted upward. More borderline cases were auto-approved. Fewer were sent to humans. The feedback loop tightened.

By the time anyone noticed—when a journalist published a story about the platform’s counterfeit problem—the model was worse than it had been before the HITL system was introduced. The reviewers, optimizing for the metric they were actually measured on, had broken the system entirely.

Goodhart’s Law

In human-AI collaboration, the consequences are particularly severe because feedback loops are tight, mistakes are invisible, and the model gets smarter—in exactly the wrong direction.

12.2 Why Accuracy Alone Is Never Enough

The most natural question to ask of any AI system is: how often is it right? Accuracy is necessary. But it is not sufficient—and in HITL systems, it can be actively misleading.

Consider a fraud detection system applied to a dataset where 1% of transactions are fraudulent. A model that labels every single transaction “legitimate” achieves 99% accuracy. It also catches zero fraud. An accuracy-first evaluation would call this a good model. Anyone whose credit card gets stolen would disagree.

This is why real measurement requires looking at the full picture: **precision** (when the model flags a case, how often is it actually a problem?), **recall** (of all the real problems, how many did the model catch?), and the **F1**

score that balances the two. For HITL systems specifically, these metrics must be tracked separately for auto-processed cases and human-reviewed cases.

A system where auto-processed cases achieve 98% accuracy and human-reviewed cases achieve 94% accuracy might look fine in aggregate. But it is telling you something important: the cases humans review are harder for humans too. The question is whether human review is adding value, or whether the system is routing the wrong cases, or whether reviewers are fatigued and not giving hard cases the attention they deserve.

12.3 A Complete Measurement Framework

Good HITL measurement covers seven distinct dimensions. Each illuminates something the others miss.

Model accuracy and calibration. Accuracy captures raw performance. **Calibration** captures something subtler: does the model's expressed confidence correspond to its actual reliability? A well-calibrated model that says it is 80% confident should be right about 80% of the time. A poorly calibrated model misroutes cases in both directions: confident-but-wrong cases sail through unreviewed, while cases it could have decided correctly clog the human queue.

This matters enormously for HITL design because routing decisions depend on confidence scores. If those scores don't mean what you think they mean, the whole intervention design breaks down.

Human workload. How many cases per hour is each reviewer handling? How does this change over a shift? Research on human performance consistently shows a fatigue curve: in vigilance studies, detection performance degrades measurably within the first 20–30 minutes of a monitoring

task (Warm et al., 2008). A HITL system routing 500 cases per hour to a team of five implicitly assumes each reviewer can sustain 100 per hour with consistent quality. That assumption is almost certainly wrong for anything requiring genuine judgment.

Time-to-decision. How long from flagging until a decision is made? In fraud detection, this matters in seconds. In medical imaging, hours may be acceptable. Time-to-decision is a system-level metric that captures bottlenecks no amount of model accuracy data will reveal.

Cost per decision. Human review is expensive. This needs to be measured explicitly, because cost shapes every design tradeoff. A system routing 40% of cases to human review when 5% would achieve similar quality wastes resources that could be spent on better interventions elsewhere.

Error rates by pathway. Track error rates separately for auto-processed and human-reviewed cases, and watch whether they move in the right direction over time. If reviewers are gaming the throughput metric, error rates may look stable while underlying quality deteriorates.

Reviewer agreement rate. When the same uncertain case is reviewed by two different reviewers, how often do they agree? Tracking agreement rates over time reveals distribution shifts, fatigue, or diverging interpretation of decision criteria. A sudden drop in agreement is a warning signal.

Feedback loop latency. How long until a human decision improves the model? Short latency means faster learning, but also faster propagation of reviewer errors. The tradeoff depends on the quality control processes applied before incorporating feedback into training.

12.4 The A/B Testing Mindset

Here is something many HITL practitioners overlook: you can run a controlled experiment on your HITL design.

Want to know whether your current confidence threshold for routing is optimal? Randomly assign incoming cases to two groups—one gets the current threshold, one gets a slightly lower threshold. Compare outcomes across every metric in your framework. The group whose threshold produces better outcomes at acceptable cost tells you which direction to move.

Most organizations set routing thresholds once, at launch, based on intuition or informal testing, and never revisit them. Meanwhile, the distribution of incoming cases shifts, reviewer expertise changes, and model calibration drifts—all of which can make the original threshold obsolete.

The Experimental Approach

A/B testing in HITL systems requires adaptations from standard experimental design: reviewers must not know they are in an experiment (and that concealment carries its own ethics burden—covert experiments on workers need the same consent and disclosure review this book asks of every other deployment); cases must be assigned in a statistically sound way (not arm A on Monday, arm B on Friday); and a held-out test set evaluated against stable ground truth is essential. The payoff is continuous, evidence-driven improvement rather than intuition-driven stagnation.

12.5 The Divergence Trap

There is a failure mode in HITL measurement that deserves its own name: **the divergence trap**—optimizing short-term accuracy at the cost of long-term learning and coverage.

Imagine a medical imaging HITL system where the model flags uncertain scans for radiologist review. The radiologists handle hard cases well. Short-term accuracy looks excellent. But the model is learning from their decisions on the cases it was most uncertain about—which tend to be unusual presentations, rare variants, edge cases.

Over time, the model incorporates these corrections and becomes better calibrated on the common cases. It grows more confident about them and routes fewer to humans. The unusual ones continue flowing to the human queue.

Apparent accuracy improves. But the model's effective coverage has narrowed: it handles more of the easy cases autonomously and pushes the unusual ones—the cases with the most to teach it—into the human review queue. If that queue becomes overwhelmed, or if reviewers lose familiarity with the common cases they never review anymore, the system becomes brittle when the case distribution shifts.

The divergence trap is visible only if you track long-term calibration, not just short-term accuracy. One practical countermeasure: periodically inject known-easy cases into the human review queue, both to keep reviewers calibrated and to monitor whether the model's confidence on common cases is drifting.

12.6 Goodhart's Law and the Metric You Choose

The British economist Charles Goodhart made a related observation in 1975; the crisp formulation usually quoted—“when a measure becomes a target, it ceases to be a good measure”—is actually the anthropologist Marilyn Strathern's later restatement of his idea (Strathern, 1997). The principle has been rediscovered repeatedly across healthcare, education, software development, and now AI—because it describes something fundamental about how people respond to measurement.

In HITL systems, Goodhart's law operates on three levels simultaneously.

Reviewer-level Goodhart: Reviewers optimize for whatever metric they are evaluated on. If it is throughput, they go fast. If it is agreement with the algorithm, they defer to the algorithm. If it is accuracy on a held-out test set they know about, they learn the test set.

Organization-level Goodhart: Organizations optimize for whatever metric they report upward. If that metric is “percentage of cases auto-processed,” there is organizational pressure to lower the routing threshold without quality controls.

Model-level Goodhart: Models train on reviewer behavior, silently absorbing whatever biases or gaming that behavior contains.

The only real protection is to measure multiple things simultaneously and to hold some things sacred: a periodic human audit of a random sample of *all* decisions—not just the flagged ones—evaluated by reviewers who do not know the original decisions. This is operationally awkward. It is also the only rigorous way to detect system-wide degradation before it becomes a public problem.

12.7 What Good Measurement Looks Like

The 400-reviews-per-hour story had a preventable outcome. The organization measured exactly one thing and nothing else. A complete measurement framework would have caught the problem in weeks rather than months.

In practice, this means: a dashboard tracking accuracy, calibration error, throughput, and reviewer agreement at weekly intervals; a separate track comparing error rates by pathway, watching for them to diverge; an automatic alert when reviewer agreement on borderline cases drops below the historical baseline; and a quarterly audit where a sample of cases from all pathways—including auto-approved—is reviewed by senior reviewers who do not know the original decisions.

And critically: a metric measuring *model improvement over time*. If the model is not getting better from human feedback, either the feedback loop is broken, the feedback is not good enough to learn from, or the cases being reviewed are too hard to yield useful signal. That last possibility is an argument for routing some easy, clearly-resolved cases into the review queue to provide the model with strong training signal.

Measurement in HITL systems is not surveillance. It is the visibility necessary to improve.

Think of a decision-making system you interact with regularly—a spam filter, a content recommendation engine, a loan approval process. What single metric would you guess the organization tracks to evaluate it? Now think of two other metrics that single measure might be hiding. What would you need to see to know whether the system was actually working well?

Chapter Summary

Key Concepts

- Accuracy is necessary but not sufficient—a model labeling everything “legitimate” on a 99:1 imbalanced dataset achieves 99% accuracy while catching zero fraud
- A complete HITL measurement framework covers seven dimensions: model accuracy and calibration, human workload, time-to-decision, cost per decision, error rates by pathway, reviewer agreement, and feedback loop latency
- Goodhart’s Law operates at three levels simultaneously: reviewer, organization, and model
- The divergence trap: optimizing short-term accuracy narrows long-term learning by routing unusual cases—the most informative ones—into the human queue
- A/B testing applied to HITL design enables continuous, evidence-driven improvement of routing thresholds

Key Examples

The 400-reviews-per-hour case Throughput metric incentivized rubber-stamping; model trained on gamed labels; system degraded over months

Class imbalance and accuracy A 99%-accurate fraud detector that labels everything “legitimate” catches zero fraud

Feedback loop latency The lag between human correction and model improvement shapes both learning speed and error propagation risk

Quarterly pathway audits Random samples of auto-processed cases, reviewed by senior staff without knowledge of original decisions, are required to detect drift

Key Principles

- Measure multiple things simultaneously; single metrics invite gaming
- Track metrics separately by pathway: auto-processed versus human-reviewed
- Hold periodic audits on cases the system handles autonomously, not only the cases it flags
- Treat measurement design as an intervention-design problem—what you measure shapes what reviewers and organizations optimize for

Five Dimensions Check

- *Feedback Integration* (Dimension 5): Measurement is how the loop proves it is learning—feedback-loop latency and pathway error rates are the core signals
- *Stakes Calibration* (Dimension 4): Tracking error rates by pathway and auditing autonomous cases keeps the system honest about where mistakes carry weight
- *Intervention Design* (Dimension 2): What you choose to measure is itself an intervention—it shapes what reviewers and organizations optimize for

The next chapter turns from measuring the system to the human inside it: the cognitive biases, expertise effects, and fatigue that shape how people actually behave in the loop.

Technical Appendix 12A:

Measuring HITL System Performance

Measurement theory applied to human-in-the-loop systems: calibration metrics, pathway decomposition, inter-rater reliability, and the statistics of controlled experiments.

— Formal foundations for Chapter 12

12A.1 Classification Metrics: Formal Definitions

12A.1.1 The Confusion Matrix

For a binary classifier, outcomes fall into four cells:

Table 12.1: Binary classifier confusion matrix

	Predicted Positive	Predicted Negative
Actual Positive	True Positive (TP)	False Negative (FN)
Actual Negative	False Positive (FP)	True Negative (TN)

12A.1.2 Core Metrics

Accuracy.

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Misleading under class imbalance.

Precision (Positive Predictive Value).

$$\text{Precision} = \frac{TP}{TP + FP}$$

Of all cases flagged as positive, how many were correct? A precision-focused system minimizes false alarms.

Recall (Sensitivity, True Positive Rate).

$$\text{Recall} = \frac{TP}{TP + FN}$$

Of all actually positive cases, how many were caught? A recall-focused system minimizes missed detections.

F1 Score.

$$F_1 = 2 \cdot \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2 \cdot TP}{2 \cdot TP + FP + FN}$$

The harmonic mean of precision and recall. The harmonic mean penalizes extreme imbalances between precision and recall more severely than the arithmetic mean.

F_β Score.

$$F_\beta = (1 + \beta^2) \cdot \frac{\text{Precision} \times \text{Recall}}{(\beta^2 \cdot \text{Precision}) + \text{Recall}}$$

When $\beta > 1$, recall is weighted more heavily (missing a positive is more costly). When $\beta < 1$, precision is weighted more heavily.

12A.1.3 HITL-Specific Pathway Decomposition

Let $\mathcal{A}$ denote the set of auto-processed cases and $\mathcal{H}$ the set of human-reviewed cases, with $|\mathcal{A}| + |\mathcal{H}| = N$.

Pathway Error Rates.

$$\text{Error}_{\mathcal{A}} = \frac{FP_{\mathcal{A}} + FN_{\mathcal{A}}}{|\mathcal{A}|}, \quad \text{Error}_{\mathcal{H}} = \frac{FP_{\mathcal{H}} + FN_{\mathcal{H}}}{|\mathcal{H}|}$$

System-Level Error Rate.

$$\text{Error}_{\text{system}} = \frac{|\mathcal{A}|}{N} \cdot \text{Error}_{\mathcal{A}} + \frac{|\mathcal{H}|}{N} \cdot \text{Error}_{\mathcal{H}}$$

A system can show improving aggregate error while $\text{Error}_{\mathcal{A}}$ increases, if the routing threshold is raised (more cases auto-processed). Tracking pathway-specific error rates separately detects this divergence.

Human Value-Added.

$$\Delta_H = \text{Error}_{\mathcal{H}}^{(\text{auto-only})} - \text{Error}_{\mathcal{H}}^{(\text{human-reviewed})}$$

Estimated counterfactual: how much higher would the error rate on human-reviewed cases have been if the model had decided them autonomously?

12A.2 Calibration Metrics

12A.2.1 Why Calibration Matters for HITL

A HITL system routes cases to humans based on model confidence. If confidence is miscalibrated—if an 80% confidence score corresponds to 60% accuracy—the routing threshold is miscalibrated too.

Definition: Calibration. A probabilistic classifier f is perfectly calibrated if, for all $p \in [0, 1]$:

$$P(\hat{y} = y \mid f(x) = p) = p$$

12A.2.2 Expected Calibration Error (ECE)

Partition predictions into M equally-spaced confidence bins. Let B_m denote the set of cases in bin m , with $|B_m|$ elements:

$$\begin{aligned} \text{acc}(B_m) &= \frac{1}{|B_m|} \sum_{i \in B_m} \mathbf{1}[\hat{y}_i = y_i], & \text{conf}(B_m) &= \frac{1}{|B_m|} \sum_{i \in B_m} \hat{p}_i \\ \text{ECE} &= \sum_{m=1}^M \frac{|B_m|}{n} |\text{acc}(B_m) - \text{conf}(B_m)| \end{aligned}$$

A perfectly calibrated model has $\text{ECE} = 0$. Values below 0.02 are considered good; above 0.05 warrant intervention (Guo et al., 2017).

12A.2.3 Maximum Calibration Error (MCE)

ECE is an average; a model can have low ECE while severely miscalibrated in high-confidence bins where routing decisions are made:

$$\text{MCE} = \max_m |\text{acc}(B_m) - \text{conf}(B_m)|$$

For HITL systems, MCE in the high-confidence bins should be tracked separately from the overall ECE.

12A.2.4 Calibration Methods

Temperature Scaling (Guo et al., 2017). A single learned parameter $T > 0$ scales the logits before the softmax:

$$\hat{p}_i = \text{softmax}(z_i/T)$$

When $T > 1$, confidence is reduced. T is optimized on a held-out validation set by minimizing negative log-likelihood.

Platt Scaling. Fits a logistic regression $P(y = 1 \mid f) = \sigma(af + b)$ on a validation set, where (a, b) are learned parameters.

Isotonic Regression. Non-parametric monotonic calibration. More flexible than temperature scaling but requires a larger validation set to avoid overfitting. Preferred when calibration errors are concentrated in specific confidence regions.

12A.3 Confidence Intervals for HITL Metrics

12A.3.1 Wilson Score Interval for Proportions

For an observed proportion $\hat{p} = k/n$, the Wilson score interval:

$$\left[\frac{\hat{p} + z^2/(2n) \pm z\sqrt{\hat{p}(1-\hat{p})/n + z^2/(4n^2)}}{1 + z^2/n} \right]$$

where $z = 1.96$ for a 95% interval. Preferred over the Wald interval when $\hat{p}$ is near 0 or 1.

12A.3.2 Bootstrap Confidence Intervals

F1, ECE, and inter-rater agreement are nonlinear functions of the data. Closed-form intervals are unreliable. The bootstrap:

1. Draw $B = 10,000$ bootstrap samples (with replacement) from n observations.
2. Compute the metric on each: $\theta_1^*, \dots, \theta_B^*$.
3. The 95% interval is the 2.5th and 97.5th percentiles of $\{\theta_b^*\}$.

12A.3.3 Minimum Detectable Effect

Required sample size to detect error-rate difference δ with power $1 - \beta$ at significance α :

$$n = \frac{(z_{\alpha/2} + z_{\beta})^2 \cdot 2\bar{p}(1 - \bar{p})}{\delta^2}$$

where $\bar{p}$ is the pooled error rate. For a 5% baseline and 1 percentage point improvement, 80% power requires $n \approx 6,746$ cases per arm.

12A.4 Inter-Rater Reliability

12A.4.1 Cohen's Kappa

Raw agreement is inflated by chance. Kappa corrects for this:

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

where p_o is observed agreement and $p_e = \sum_k (n_{k\cdot}/n)(n_{\cdot k}/n)$ is expected chance agreement.

Benchmarks (Landis & Koch, 1977): $\kappa < 0.20$ slight; 0.21–0.40 fair; 0.41–0.60 moderate; 0.61–0.80 substantial; >0.80 almost perfect.

12A.4.2 Fleiss' Kappa for Multiple Reviewers

For $k > 2$ reviewers, Fleiss' generalization:

$$\kappa_F = \frac{\bar{P} - \bar{P}_e}{1 - \bar{P}_e}$$

12A.4.3 Intraclass Correlation Coefficient (ICC)

For continuous annotations (e.g., severity scores), two-way mixed ICC:

$$\text{ICC}(2, 1) = \frac{MS_R - MS_E}{MS_R + (k - 1)MS_E + k(MS_C - MS_E)/n}$$

where MS_R , MS_C , MS_E are mean squares from a two-way ANOVA decomposition.

12A.5 Statistical Significance in A/B Tests

12A.5.1 Two-Proportion Z-Test

Null hypothesis $H_0 : p_A = p_B$. Test statistic:

$$z = \frac{\hat{p}_A - \hat{p}_B}{\sqrt{\hat{p}(1 - \hat{p})(1/n_A + 1/n_B)}}$$

where $\hat{p} = (k_A + k_B)/(n_A + n_B)$. Reject H_0 at $\alpha = 0.05$ when $|z| > 1.96$.

12A.5.2 Complications in HITL Experiments

Three HITL-Specific Complications

- **Non-independence:** Cases reviewed by the same reviewer are correlated. Use cluster-robust variance estimation with reviewers as clusters.
- **SUTVA violations:** Reviewer behavior can change based on queue length, which is affected by arm assignment. Minimize by balancing queue lengths across arms.
- **Learning effects:** If the model updates during the experiment, freeze model weights for the experimental period or use pre-registered stopping rules.

12A.5.3 CUPED: Variance Reduction Using Pre-Experiment Data

CUPED reduces variance by regressing out baseline covariates (A. Deng et al., 2013):

$$Y_{\text{CUPED}} = Y - \theta(X - \mathbb{E}[X])$$

where $\theta = \text{Cov}(Y, X)/\text{Var}(X)$ is estimated from pre-experiment data. Can reduce variance by 30–50%, cutting required sample size by the same 30–50%.

12A.6 Feedback Loop Latency Measurement

12A.6.1 Defining Latency

Let t_{decision} be the time a human reviewer finalizes a decision, and $t_{\text{incorporated}}$ be the time that decision is reflected in the model's predictions:

$$L_i = t_{\text{incorporated},i} - t_{\text{decision},i}, \quad \bar{L} = \frac{1}{|\mathcal{H}|} \sum_{i \in \mathcal{H}} L_i$$

12A.6.2 Measuring Feedback Value

The value of the feedback loop is the improvement attributable to human corrections:

$$V_{\text{feedback}} = \text{Error}_t^{(\text{without feedback})} - \text{Error}_{t+\Delta}^{(\text{with feedback})}$$

Estimating the counterfactual requires freezing model snapshots at multiple time points and evaluating each against new ground truth—operationally complex but the only rigorous approach.

12A.7 Exercises

Exercise 12A.1 — Calibration Assessment

A model has predictions in 10 confidence bins with accuracy values: {0.05, 0.11, 0.18, 0.30, 0.42, 0.55, 0.63, 0.74, 0.82, 0.90} (equal numbers per bin, $n = 1,000$ total).

- (a) Compute the ECE.
- (b) Describe the reliability diagram shape.
- (c) Assess calibration quality; is temperature scaling appropriate?
- (d) What does this calibration imply for HITL routing at a threshold of 0.80?

Exercise 12A.2 — Sample Size Determination

A HITL system has a current error rate of 8%. A proposed redesign is expected to reduce this to 6%.

- (a) Compute the required sample size per arm at 80% power, $\alpha = 0.05$.
- (b) Recompute for 90% power.
- (c) At a throughput of 500 human-reviewed cases per day, how long would each experiment take? What practical challenges arise at this timescale?

Exercise 12A.3 — Goodhart's Law Simulation

A model routes the bottom 20% of confidence cases to human review. Original reviewer accuracy: 92%. Gaming increases throughput 3× but reduces accuracy to 65%. System auto-processes the remaining 80% at a 3% error rate.

- (a) Compute the system-level error rate before and after gaming using the pathway decomposition from Section 12.5.
- (b) Explain why aggregate accuracy might appear stable to management while the system degrades.
- (c) What single monitoring metric would most reliably detect gaming within two weeks?

Exercise 12A.4 — Inter-Rater Reliability

Two reviewers independently label 100 borderline cases. Reviewer A: 60 positive, 40 negative. Reviewer B: 55 positive, 45 negative. Agreement: 50 both-positive, 35 both-negative.

- (a) Compute Cohen's kappa.
- (b) Interpret the level of agreement.
- (c) What trend in kappa over time would constitute a warning signal? What might cause it?

Exercise 12A.5 — Feedback Loop Quality Control

Design a quality control procedure for incoming reviewer decisions before model training. The system processes 10,000 cases per day, routes 15% to human review, and targets 24-hour feedback loop latency.

- (a) Specify four statistics to compute on each incoming daily batch.
- (b) Define alert thresholds for each statistic.
- (c) Describe the decision rule: when to proceed, when to hold for audit, when to reject.
- (d) What is the cost of a false hold (rejecting a good batch)? Of a false pass (accepting a bad batch)?

This appendix provides the mathematical foundations for Chapter 12's measurement framework. The exercises are calibrated so that Exercises 12A.1–12A.4 are accessible to readers with introductory statistics, while Exercise 12A.5 requires systems-level thinking appropriate for practitioners and advanced students.

Part IV

The Human Factor

The human is not a neutral input device.

— Chapter 13

13 The Psychology of Human-AI Teams

The better the automation, the less prepared the human for the moment it fails.

— on the irony of reliable automation

13.1 The Night the Computers Went Silent

At 2:10 in the morning on June 1, 2009, Air France Flight 447's autopilot disconnected over the South Atlantic Ocean. Four minutes later the aircraft hit the ocean. All 228 people aboard were killed. It remains the deadliest aviation accident in Air France's history (Bureau d'Enquêtes et d'Analyses pour la Sécurité de l'Aviation Civile, 2012).

The plane—an Airbus A330 equipped with some of the most sophisticated autopilot systems in the world—had flown into an area of severe turbulence over the Intertropical Convergence Zone. Ice crystals formed in the pitot tubes, the sensors that measure airspeed. The autopilot, suddenly unable to trust its airspeed data, did exactly what it was designed to do: it disengaged. It handed control to the pilots.

The crew had, for the previous several hours, been managing a comfortable, mostly automated flight. Then, in the space of a few seconds, they were thrust into a rapidly deteriorating emergency with faulty and

conflicting sensor data—a situation requiring immediate, accurate action on the aircraft’s manual controls.

What the flight data recorders revealed was devastating: the pilots were confused. One pilot pulled back on the stick, increasing the plane’s angle of attack even as the aircraft stalled. The confusion lasted four and a half minutes, from autopilot disengagement to impact with the ocean.

The BEA’s conclusion was precise: the crew had not been sufficiently trained in high-altitude manual flight, and they had lost situational awareness during the period of automation. When the system handed control back to them, they were not ready for it.

Air France 447 is not a story about bad pilots or bad technology. It is a story about what happens when the human is in the loop for so long without actually needing to exercise skill that the skill quietly atrophies—and then suddenly, catastrophically, the skill is required.

13.2 The Paradox of Reliable Systems

There is something counterintuitive at the heart of automation design: the better the automation, the more dangerous the gap when it fails.

A poor autopilot that disengages frequently keeps pilots sharp. An excellent autopilot that handles 99.97% of situations without intervention produces pilots who have far less practice with the remaining 0.03%—precisely the moments when judgment is most critical.

This is called **automation complacency**, and it has been documented in aviation, nuclear power, industrial process control, and AI-assisted medical diagnosis. The pattern is consistent: as system reliability increases, human operators become less vigilant, less practiced, and less prepared for the moments when the system fails (Parasuraman & Manzey, 2010).

The paradox runs deeper than practice alone. When a system is highly reliable, there is a rational case for trusting it. If an AI flagging system

has 98% precision, a reasonable reviewer should doubt their own judgment when it conflicts with the AI's recommendation—that's sensible Bayesian updating. The problem is that this rational deference, taken too far, becomes **automation bias**.

Bainbridge's Irony of Automation (Bainbridge, 1983)

Designers assume that the human will reliably oversee the automated system. But automation removes the very practice that makes reliable oversight possible. The designers are correct that human operators are unreliable—but automation is part of the reason why.

13.3 Automation Bias: The Research Record

The term was formalized by Linda Skitka and colleagues in 1999. Pilots shown automated recommendations became statistically less likely to notice when the automation was wrong—even when they had been explicitly told to apply independent judgment, and even when the errors were visible to an unbiased observer (Skitka et al., 1999).

The finding has been replicated across dozens of studies and multiple domains. Medical residents shown an AI-generated differential diagnosis are less likely to consider diagnoses the AI omitted, even when their own history-taking revealed relevant clues. Radiologists reviewing AI-flagged scans spend more time on flagged areas and proportionally less time on unflagged areas—increasing the likelihood of missing pathology the AI didn't catch.

What makes automation bias particularly insidious in HITL systems is that it works against the purpose of having a human in the loop. The whole point of human review is to catch cases the model gets wrong. Automation bias systematically reduces the probability that humans catch exactly those cases—the ones where the model is confident but incorrect.

The mechanism: automation bias is not primarily about deference to authority, or cognitive laziness. Research by Parasuraman and Riley, who produced the foundational taxonomy of human-automation interaction, suggests the core mechanism is **attentional allocation**—the recommendation satisfies the brain’s drive toward closure, reducing effortful processing of the case itself (Parasuraman & Manzey, 2010; Parasuraman & Riley, 1997). This means “apply independent judgment” instructions, without interface redesign, accomplish very little.

13.4 Overtrust and Undertrust: Both Cost You

Automation bias is a form of **overtrust**. But the psychological literature is equally clear that **undertrust** is a problem, and often an underappreciated one.

Undertrust manifests as systematic distrust of algorithmic recommendations even when the algorithms are accurate: clinicians who treat AI flags as “suggestions” and override them at high rates even when the flags were correct; judges reluctant to “let an algorithm decide” who override risk predictions at rates statistical accuracy would not justify.

Undertrust has real costs. If a screening algorithm correctly identifies high-risk cases and clinicians consistently override toward lower-risk assessments, the system’s effective recall drops—undermining the whole purpose of AI augmentation.

The optimal human-AI collaboration requires **calibrated trust**: trust that accurately tracks the system’s actual performance domain by domain. Neither systematic over-reliance nor systematic under-reliance, but flexible, evidence-sensitive deference.

Calibrated trust is hard to achieve and harder to maintain. It requires that reviewers have accurate mental models of what the AI is good at and what it isn’t. It requires feedback loops that let reviewers learn when

they're right to override and when they're wrong. And it requires an organizational culture that neither rewards blind deference nor treats algorithm skepticism as a virtue (Dietvorst et al., 2015).

Algorithm Aversion

Dietvorst and colleagues (2015) found that seeing an algorithm err *once* causes people to systematically avoid using it in the future—even when the algorithm still outperforms human judgment. The research label is “algorithm aversion.” In HITL contexts, reviewers regularly see the model’s errors (that is why the cases are in the queue). Without deliberate feedback design, algorithm aversion is a predictable outcome.

13.5 The Expertise Effect: Better in Some Ways, Worse in Others

Domain experts make better HITL reviewers in several respects: they catch errors novices miss, recognize unusual presentations, and provide harder-to-fool audits of AI outputs.

But expertise introduces its own vulnerabilities.

Experts are more susceptible to **expectation confirmation**—the **consistency trap**: strong prior beliefs about what they will find, and strong expectations about what a decision “should” look like, lead them to process new evidence in light of those priors and to dismiss AI outputs that violate those expectations—even when the outputs are correct. An experienced radiologist who “knows” that a certain pattern correlates with benign disease may be less willing to register a finding that contradicts that expectation, even when an AI correctly flags it.

Highly experienced reviewers have longer histories with their own intuitions. They have been right a lot. When AI recommendations conflict

with expert intuition, the expert has more to defend and more history of being right. The resistance to AI suggestions that contradict expertise can be fierce.

This is not an argument against expert reviewers. It is an argument for designing HITL systems that account for expert psychology—presenting AI recommendations in ways that invite genuine evaluation rather than triggering confirmation or resistance.

13.6 Fatigue and the Time-of-Day Effect

Human judgment is not constant. It varies with time of day, length of shift, number of prior decisions, and emotional state—factors that most HITL evaluation frameworks entirely ignore.

A 2011 study by Danziger and colleagues examined 1,112 judicial hearings. The probability of a favorable parole decision dropped from roughly 65% at the start of each session to nearly zero toward the end, before recovering after breaks (Danziger et al., 2011). To the extent the effect is real, it operated below the judges' awareness.

The causal mechanism behind the Danziger finding—whether it reflects decision fatigue, glucose depletion, or a statistical artifact of case ordering—has been contested in subsequent methodological work. The precise effect size should be treated cautiously. The directional finding (decision quality degrades over a session) is, however, well-supported by the broader human performance literature independent of this specific study.

The finding points to something well-established in human performance research: decision quality declines under sustained cognitive load. Annotation quality, medical diagnosis accuracy, content moderation deci-

sions, and financial review accuracy all follow this pattern. All are better early in a work session and worse late. The most rigorous documentation comes from medicine: in one study, physicians' antibiotic prescribing decisions degraded measurably over the course of a clinic day (Linder et al., 2014).

A HITL system that is measured, calibrated, and staffed without accounting for time-of-day effects is ignoring a significant and predictable source of variance in its human component's performance. Routing the most difficult cases to reviewers at hour four of a shift is not equivalent to routing them at hour one. The two populations—same person, different time—can have measurably different error rates.

Practical implications: randomize case routing to prevent clustering of hard cases at shift end; build mandatory break schedules into the review workflow; use intra-batch agreement rates as a near-real-time proxy for reviewer fatigue; consider two-pass review for the highest-stakes decisions falling late in long shifts.

13.7 Team Dynamics Without a Human Team

Traditional organizational research on team performance was developed to understand groups of humans. Almost none of it applies cleanly when one team member is an algorithm.

An algorithm does not experience psychological safety. It does not feel threatened by dissent. It does not have a bad day. It does not grow resentful of overrides. It does not build an informal theory of which reviewers are worth listening to.

What this means in practice is that the “team dynamics” of a human-AI system are almost entirely determined by how the human relates to the AI—and those dynamics are heavily influenced by system design choices that organizations rarely treat as team management decisions.

How confident does the AI appear? Systems presenting recommendations with high apparent confidence elicit more deference than systems that explicitly communicate uncertainty. The same model, the same accuracy, different framing, different behavior.

Does the AI explain itself? Explainable AI can reduce automation bias by giving reviewers something to evaluate. But it can also increase automation bias if the explanations are plausible-sounding but wrong—providing post-hoc narrative that convinces reviewers the AI must be right.

Does the reviewer get feedback on their overrides? Without feedback, there is no mechanism to calibrate trust. With feedback, reviewers can learn which domains their judgment beats the AI’s and which it doesn’t. This feedback is among the most powerful interventions available for producing calibrated trust.

13.8 What HITL System Design Can Do About Human Psychology

These psychological findings are not arguments for resignation. They are arguments for designing systems that support the kind of human engagement that actually produces the outcomes we want.

Interface design for independent judgment. Present the AI recommendation in a way that invites evaluation rather than closure. One tested approach: have reviewers complete their own initial assessment before

the AI recommendation is revealed. This forces genuine processing uncorrupted by automation bias.

Explicit uncertainty communication. A system that displays “85% confidence” gives reviewers something to reason about. A system that displays “RECOMMENDED: APPROVE” encourages closure. Calibrated uncertainty communication, introduced in Chapter 1 and formalized in Chapter 12, shapes human psychology as much as it does routing.

Structured disagreement protocols. For high-stakes decisions, require reviewers to document their reasoning when they override the AI, and require a second review when reviewer and AI disagree. This creates a context in which the human is genuinely doing the work the HITL design requires.

Feedback on overrides. If you want calibrated trust, give reviewers information about the accuracy of their overrides. Without feedback, trust cannot calibrate. With it, even experienced reviewers become more accurate about which overrides are worth making.

Fatigue management. Build mandatory break structures and randomized case routing into the system design. Do not let the hardest cases cluster at shift end. Track intra-batch agreement rates as an operational fatigue signal.

The next time you use an AI-powered suggestion—autocomplete, a navigation recommendation, a grammar suggestion—pay attention to your own psychological response. Do you evaluate the

suggestion independently before accepting it, or does the presence of a recommendation change your process? What would you have decided if no recommendation had appeared? This is automation bias in your own cognition, in real time.

Chapter Summary

Key Concepts

- **Automation complacency:** The more reliable the system, the less prepared the human for the moment it fails
- **Automation bias:** Humans shown an AI recommendation become less likely to catch the AI's errors—even when instructed to apply independent judgment
- Overtrust and undertrust are both costly; calibrated trust requires feedback loops as a deliberate design element
- The expertise effect: domain experts catch more errors but resist disconfirming AI recommendations more strongly
- Fatigue and time-of-day effects on decision quality are predictable and large; HITL designs that ignore them implicitly accept a major variance source
- When one team member is an algorithm, human-AI dynamics are largely shaped by interface and system design choices

Key Examples

Air France 447 (2009) Highly reliable autopilot produced pilots unpracticed in manual flight; sensor failure disengaged automation in crisis; skill atrophy produced fatal confusion

Skitka et al. (1999) Experimental demonstration that automation bias persists under explicit independence instructions; the mechanism is pre-conscious attention allocation

Radiologist AI studies Attention shifts toward AI-flagged areas; proportionally less time on unflagged areas; higher miss rate for pathology AI didn't flag

Danziger et al. (2011) Judicial parole decisions fall from 65% favorable early in session to nearly zero at end; recovers after breaks

Key Principles

- The human in the loop is not a neutral input device—psychology shapes performance in predictable, designable ways
- Interface design choices (when the recommendation is revealed, how confidence is displayed, whether explanations are provided) dramatically affect automation bias
- Calibrated trust is the goal; it requires feedback loops and cannot be achieved by instruction alone
- HITL system design must incorporate the literature on human performance under automation, including fatigue effects and the paradox of reliable systems

Five Dimensions Check

- *Intervention Design* (Dimension 2): When the recommendation is revealed, how confidence is displayed, and whether explanations are shown all shape automation bias
- *Feedback Integration* (Dimension 5): Calibrated trust is built through feedback loops over time, not through instruction alone
- *Timing* (Dimension 3): Fatigue and time-of-day effects mean the same reviewer performs differently depending on when the case arrives

The next chapter turns from individual psychology to collective consequences: how HITL systems become fair or unfair, and why human oversight alone does not guarantee ethical outcomes.

Technical Appendix 13A: Signal Detection Theory and Human Reviewer Behavior

Signal Detection Theory provides a formal decomposition of human reviewer performance into discriminability and response bias, predicting when and how automation affects each.

— Formal foundations for Chapter 13

13A.1 Signal Detection Theory: The Formal Framework

13A.1.1 Why SDT?

Standard accuracy metrics treat reviewer errors as random noise. SDT provides a richer decomposition separating two fundamentally different aspects of performance:

1. **Discriminability** (d'): How well can the reviewer distinguish positive (signal) from negative (noise) cases? This reflects cognitive ability.
2. **Response bias** (c or β): Where has the reviewer set their threshold? This reflects implicit weighting of false alarms versus misses, shaped by instructions, incentives, and automation feedback.

A reviewer with high d' and miscalibrated c produces different errors from a reviewer with moderate d' and well-calibrated c . Standard accuracy metrics cannot distinguish them. SDT can.

13A.1.2 The Basic Model

In SDT, each reviewed case generates a decision variable X . For positive (signal) cases, X is drawn from $\mathcal{N}(\mu_S, \sigma^2)$; for negative (noise) cases, from $\mathcal{N}(0, \sigma^2)$.

By convention, $\mu_N = 0$, $\sigma = 1$, and $\mu_S = d' > 0$. We write Φ for the standard normal CDF.

The reviewer compares X to a criterion location k :

$$\text{Response} = \begin{cases} \text{“positive”} & \text{if } X > k \\ \text{“negative”} & \text{if } X \leq k \end{cases}$$

13A.1.3 The Four Outcomes and Their Probabilities

Table 13.1: SDT outcome matrix

	“Positive” Response	“Negative” Response
Signal present	Hit (TP): $P(H) = \Phi(d' - k)$	Miss (FN): $1 - P(H)$
Signal absent	False Alarm (FP): $P(FA) = \Phi(-k)$	Correct Rejection: $1 - P(FA)$

13A.1.4 Computing d' and c

$$d' = z(P(H)) - z(P(FA))$$

where $z(\cdot) = \Phi^{-1}(\cdot)$ is the inverse standard normal CDF.

$$c = -\frac{1}{2}[z(P(H)) + z(P(FA))]$$

The bias c and the criterion location k are related by $c = k - d'/2$, so $c = 0$ places the criterion exactly midway between the noise and signal means ($k = d'/2$, unbiased), whatever the value of d' .

Content Moderation Example

A moderator correctly flags 80% of policy violations ($P(H) = 0.80$) and incorrectly flags 15% of compliant posts ($P(FA) = 0.15$):

$$d' = z(0.80) - z(0.15) = 0.842 - (-1.036) = 1.878$$

$$c = -\frac{1}{2}[0.842 + (-1.036)] = 0.097$$

Slight conservative bias ($c > 0$).

Interpreting c .

- $c = 0$: Unbiased criterion
- $c > 0$: Conservative bias—reviewer requires strong evidence to flag
- $c < 0$: Liberal bias—reviewer flags at low evidence

13A.2 SDT Formally Predicts Automation Bias

13A.2.1 The Effect of an Automated Recommendation on Criterion

When a reviewer is shown an automated recommendation, it shifts their effective criterion c :

$$c_{\text{with-auto}} = c_0 - \delta_{\text{auto}}$$

where $\delta_{\text{auto}} > 0$ is the criterion shift induced by the recommendation. The reviewer becomes more liberal when automation says “flag” and more conservative when automation says “pass.”

This criterion shift produces automation bias: hit rate increases when automation correctly recommends “flag” (beneficial), but miss rate also increases when automation incorrectly recommends “pass” on a true positive (harmful). Automation bias is primarily a *criterion effect*, not a discriminability effect.

13A.2.2 Empirical Support

Studies by Skitka et al. (1999) and Goddard et al. (2012) found that hit rates for signals the automation correctly flagged increased substantially with automation, while miss rates for signals the automation incorrectly passed also increased substantially. d' estimates changed modestly—consistent with criterion shift, not discriminability loss, as the primary mechanism (Goddard et al., 2012; Skitka et al., 1999).

13A.2.3 Optimal Criterion for HITL Routing

Let p^* be the prevalence of positive cases in the human review queue. The cost-optimal criterion location:

$$k^* = \frac{d'}{2} - \frac{1}{d'} \ln \left(\frac{p^* \cdot C_{FN}}{(1 - p^*) \cdot C_{FP}} \right)$$

where C_{FN} and C_{FP} are the costs of misses and false alarms.

In a HITL system, routing has already changed p^* relative to the population base rate. If the base rate is 5% and the model routes uncertain cases, the queue prevalence may be 20–40%. Reviewers who implicitly use the base-

rate prior rather than the queue-enriched prior will have miscalibrated criteria. Explicitly informing reviewers of the queue base rate is a low-cost, high-impact intervention.

13A.3 The ROC Curve and HITL Design

13A.3.1 The Receiver Operating Characteristic

The ROC curve traces (false alarm rate, hit rate) pairs as the criterion k varies. For the Gaussian equal-variance model:

$$P(H) = \Phi(d' + \Phi^{-1}(P(FA)))$$

The area under the ROC curve (AUC) is a criterion-independent discriminability summary:

$$AUC = \Phi(d' / \sqrt{2})$$

13A.3.2 Optimal Operating Point

The optimal operating point maximizes:

$$\text{Benefit} = p^* \cdot C_{FN} \cdot P(H) - (1 - p^*) \cdot C_{FP} \cdot P(FA)$$

The slope of the ROC curve at the optimal point:

$$\text{slope}^* = \frac{(1 - p^*) \cdot C_{FP}}{p^* \cdot C_{FN}}$$

Automation bias shifts the operating point away from slope^{*} toward the automation's recommended response, degrading overall system benefit whenever the automation is wrong.

13A.4 Vigilance Decrement Models

13A.4.1 The Vigilance Problem

Vigilance—sustained attention to detect infrequent signals—declines over time. This is among the most robust findings in human performance psychology, established by Mackworth (1948) and replicated hundreds of times.

Three parameters govern vigilance decrement:

1. **Event rate:** More frequent signals maintain vigilance better
2. **Signal-noise ratio:** Harder discriminations decrement faster
3. **Knowledge of results:** Feedback maintains vigilance

13A.4.2 SDT Representation of Vigilance Decrement

Vigilance decrement manifests in SDT as a decline in d' over time—the opposite of automation bias, which is primarily a criterion effect.

$$d'(t) = d'_0 \cdot e^{-\lambda t}$$

Empirical estimates: λ ranges from 0.01 to 0.05 per minute for demanding tasks, implying meaningful discriminability loss within 20–40 minutes.

13A.4.3 Combined Effects

Automation bias and vigilance decrement compound. A reviewer 90 minutes into a shift, reviewing AI-recommended decisions, experiences:

- **Reduced d'** (vigilance decrement): Less capable of distinguishing signal from noise

- **Criterion shift toward automation** (automation bias): More reliant on AI recommendation

Both effects increase miss rate on cases where the automation is wrong. Their joint effect on miss rate is approximately multiplicative.

13A.5 Formal Model of Calibrated Trust

13A.5.1 The Optimal Trust Model

Let $\tau \in [0, 1]$ be the reviewer's trust level. The reviewer's effective decision:

$$D_{\text{reviewer}} = (1 - \tau) D_{\text{human}} + \tau D_{\text{automation}}$$

Expected accuracy of the human-AI team:

$$P(\text{correct}) = \tau \cdot P_A + (1 - \tau) \cdot P_H + f(\tau) \cdot P_C$$

where P_A = automation accuracy, P_H = human accuracy, P_C = complementarity (probability exactly one is correct), and $f(\tau) = \tau(1 - \tau)$ captures the team synergy term—highest when both parties contribute to the decision, vanishing under full deference ($\tau = 1$) or full dismissal ($\tau = 0$).

13A.5.2 Optimal Trust Level

Setting the derivative of $P(\text{correct})$ to zero, $P_A - P_H + P_C(1 - 2\tau) = 0$, gives the optimal trust level:

$$\tau^* = \frac{P_A - P_H + P_C}{2 P_C}$$

(clipped to $[0, 1]$ if the parameters fall outside the interior-solution range).

When Is Deference Optimal?

When $P_A > P_H$, the optimal trust level $\tau^* > 0.5$ (lean toward automation). Some automation bias is *optimal* given cost structure. The problem is not deference per se, but deference uncalibrated to actual relative accuracy—which requires feedback to achieve and maintain.

13A.6 Measurement Implications for HITL Practitioners

13A.6.1 Extracting SDT Parameters from Production Data

```

1 import numpy as np
2 from scipy.stats import norm
3
4 def compute_sdt_params(hits, total_signal, false_alarms,
5     total_noise):
6     """
7     Compute d' and criterion c from production HITL data.
8
9     Parameters:
10     -----
11     hits: int - number of signal cases correctly flagged
12     total_signal: int - total actual positive cases
13     false_alarms: int - noise cases incorrectly flagged
14     total_noise: int - total actual negative cases
15     """
16     # Apply log-linear correction for extreme rates
17     pH = (hits + 0.5) / (total_signal + 1)
18     pFA = (false_alarms + 0.5) / (total_noise + 1)

```

```

19     d_prime = norm.ppf(pH) - norm.ppf(pFA)
20     c = -0.5 * (norm.ppf(pH) + norm.ppf(pFA))
21     auc = norm.cdf(d_prime / np.sqrt(2))
22
23     return {'d_prime': d_prime, 'c': c, 'auc': auc,
24            'hit_rate': pH, 'fa_rate': pFA}

```

Listing 13.1: SDT parameter extraction from HITL logs

13A.6.2 Detecting Automation Bias via SDT

Compare c estimates in two conditions:

- **Condition A:** AI recommendation visible (standard production)
- **Condition B:** AI recommendation hidden (blind review sample)

Automation bias is operationally defined as $c_A \neq c_B$. Specifically: $c_A < c_B$ on AI-flagged cases (criterion shift toward liberal), $c_A > c_B$ on AI-passed cases (criterion shift toward conservative).

```

1  def detect_automation_bias(condition_a_data,
2      condition_b_data):
3      """
4      Compare SDT parameters between AI-visible and
5      blind-review conditions.
6
7      Returns bias magnitude (delta_c) and significance test.
8      """
9      params_a = compute_sdt_params(**condition_a_data)
10     params_b = compute_sdt_params(**condition_b_data)
11
12     delta_c = params_a['c'] - params_b['c']
13     delta_dprime = params_a['d_prime'] - params_b['d_prime']

```

```

12
13     return {
14         'automation_bias_criterion_shift': delta_c,
15         'discriminability_change': delta_dprime,
16         'bias_direction': 'liberal' if delta_c < 0 else
17             'conservative',
18         'mechanism': 'criterion_shift' if abs(delta_c) >
19             abs(delta_dprime)
20             else 'discriminability_change'
21     }

```

Listing 13.2: Automated bias monitoring in production

13A.6.3 Sequential Design to Counteract Automation Bias

SDT provides a prescription: anchor criterion independently before the automation recommendation induces a shift.

```

1  # Sequential design: human-first assessment
2  def sequential_review_protocol(case):
3      """
4      Implement sequential design to reduce automation bias.
5      """
6      # Step 1: Present case, collect human assessment
7      human_assessment = present_case_get_assessment(case)
8      record_initial_decision(human_assessment)
9
10     # Step 2: Reveal AI recommendation
11     ai_recommendation = get_ai_recommendation(case)
12     display_ai_recommendation(ai_recommendation,
13                               human_assessment)
14
15     # Step 3: Allow override with documented rationale

```

```
15     if human_assessment != ai_recommendation:
16         rationale = require_disagreement_rationale()
17         final_decision = present_with_rationale_requirement(
18             human_assessment, ai_recommendation, rationale
19         )
20     else:
21         final_decision = human_assessment
22
23     return final_decision
```

Listing 13.3: Sequential design decision flow

13A.7 Exercises

Exercise 13A.1 — Computing d' and c

A medical AI review system records: hit rate = 0.88, false alarm rate = 0.22.

- Compute d' and c .
- Is the reviewer biased conservative or liberal? Interpret in medical imaging context.
- If the cost of a missed anomaly is 10× the cost of a false alarm and queue base rate is 5%, compute the optimal c^* .
- Compare observed c to optimal c^* . What does this tell us about how to communicate the base rate to reviewers?

Exercise 13A.2 — Automation Bias via SDT

Baseline (no AI): $P(H) = 0.75$, $P(FA) = 0.18$. AI-assisted: hit rate = 0.85 when AI flags; miss rate increases to 0.32 when AI passes.

- (a) Compute d' and c for the baseline condition.
- (b) Characterize how criterion c has changed in the AI-assisted condition.
- (c) Is d' the primary driver of the change, or is criterion shift? What does this imply for intervention design?

Exercise 13A.3 — Vigilance Decrement

A reviewer starts a 4-hour shift with $d'_0 = 2.1$ and $\lambda = 0.012$ per minute.

- (a) Compute d' at $t = 30, 60, 90, 120$ minutes.
- (b) Compute AUC at each time point.
- (c) If misses cost 5× more than false alarms and queue prevalence is 12%, compute optimal c^* at each time point. How does the optimal operating point shift?
- (d) Design a production monitoring protocol that would detect this decay in near real-time.

Exercise 13A.4 — Optimal Trust Level

A loan review system has: automation accuracy 87%, human accuracy 78%, human is correct 60% of the time when they disagree.

- (a) Estimate $P(\text{complementary})$, taking it here to be the probability the human is correct conditional on disagreeing with the automation.
- (b) Compute τ^* .
- (c) Interpret: should reviewers defer to automation, use their own judgment, or use a mixed strategy? At what level?
- (d) Propose two interface design interventions that would help reviewers achieve τ^* .

This appendix provides the formal SDT framework underlying Chapter 13's psychological discussion. The key insight—that automation bias is a criterion effect, not a discriminability effect—has direct design implications: interventions must target criterion anchoring before automation recommendations are presented, not simply instruct reviewers to apply independent judgment.

14 Building Fair and Ethical Systems

Human oversight does not automatically produce human fairness.
— on the myth of neutral oversight

14.1 The Proxy Problem

In 2018, Reuters reported that Amazon had disbanded an experimental recruiting tool it had been developing since 2014. The AI was designed to screen résumés and surface top candidates for software engineering positions. Amazon’s machine learning team had trained it on ten years of hiring data—the résumés of everyone the company had actually hired.

The problem: Amazon’s engineering workforce had historically been predominantly male. The model learned this pattern. It taught itself that women’s résumés were less suitable, even though gender was not an explicit input variable. It found proxies. The model learned to penalize résumés that included the word “women’s”—as in “women’s chess club” or “women’s college.” It downgraded graduates of two all-women’s colleges.

Here is the part of the story that matters most. Amazon’s engineers discovered the gender bias themselves, roughly a year into development. They tried to fix it. They edited the model to neutralize the gender-correlated terms—but that was no guarantee the model would not devise other ways of sorting candidates that could prove discriminatory. And the tool’s results

were underwhelming in their own right, its recommendations little better than random, sometimes surfacing unqualified candidates. Executives lost hope, and the project was disbanded before it was ever deployed as a decision system. Amazon stated that its recruiters looked at the recommendations the tool generated when searching for new hires, but never relied solely on its rankings to make hiring decisions (Dastin, 2018). No applicant was ever rejected by this tool; the harm it could have done stayed hypothetical, caught in development rather than discovered in production—which is exactly what makes the case worth studying.

The Amazon recruiting case is instructive not as a story of an oversight failure that went undetected for years, but as a demonstration of a deeper structural problem: bias is not a bug you can patch. It is a signal the model correctly learned from biased data. The model was doing exactly what it was trained to do. The data reflected a biased world. The model faithfully encoded it. And the fix—removing proxies one by one—was unwinnable, because the entire training distribution was saturated with structural gender patterns.

14.2 The Layers Where Bias Enters

There is no single place where bias enters an AI system. It enters everywhere.

Collection bias is the most fundamental: whose data is collected? Early facial recognition datasets, used to train and evaluate models for decades, were overwhelmingly composed of light-skinned male faces. Models trained on these datasets were well-calibrated for some users and dramatically worse for others. A 2018 study by Joy Buolamwini and Timnit Gebru—the **Gender Shades** project—found commercial gender-classification systems from major technology companies had error rates

up to 34 percentage points higher for darker-skinned women than for lighter-skinned men (Buolamwini & Gebru, 2018). The systems had not been designed to fail for these groups. They had simply been trained primarily on data that did not include them.

The same pattern holds geographically. In 2017, researchers at Google showed that image classifiers trained on the field's standard datasets—collected overwhelmingly from North America and Europe—performed far worse on images from the rest of the world: wedding photographs from Ethiopia and Pakistan drew the lowest recognition confidence, bridegrooms mislabeled with guesses like “chain mail” and “academic gown” (Shankar et al., 2017). A *Nature* commentary the following year supplied the pattern's most famous illustration: a bride in a white Western dress labeled *bride* and *wedding*; a bride in Indian wedding clothing labeled *costume* and *performance art* (Zou & Schiebinger, 2018). The classifiers had not learned what a wedding is; they had learned what weddings look like in the places the data came from.

Annotation bias is the next layer: whose judgment decides what is correct? Training data requires human labels—someone has to say which résumés were “good,” which social media posts are “harmful,” which recidivism predictions turned out to be accurate. Those labels reflect the values, biases, and social positions of the people who made them. Research on content moderation annotation has found that annotators from different demographic groups make systematically different judgments about the same content, particularly on questions involving race, political content, and hate speech (Sap et al., 2019). There is no neutral label. Every annotation embeds a perspective (Bender et al., 2021).

Model bias is what happens when a machine learning algorithm amplifies the patterns in the data. A model trained on historical criminal justice

data will learn that certain neighborhoods, or certain demographic groups, have higher rates of conviction—because that is what the data shows. It will then apply this learned association to new cases, without encoding any understanding of whether the historical disparities were themselves the product of unequal policing, unequal prosecution, or unequal sentencing. The model reproduces structural inequities as if they were natural facts.

Deployment bias is the mismatch between the population on which a system was developed and the population on which it is used. A medical diagnostic AI trained on data from academic medical centers in the United States performs differently when deployed in rural clinics, in other countries, or on populations underrepresented in the training data. A fraud detection model calibrated for typical spending patterns in one economic context may generate systematically higher false positive rates for users from different economic contexts—essentially treating unfamiliar behavior as suspicious.

The Normal Condition

None of these biases are obscure edge cases. They are the normal condition of AI systems built from historical data about an unequal world. The question is not whether they exist but what to do about them.

14.3 The Feedback Loop Problem

Perhaps the most insidious form of algorithmic bias is the one that makes itself worse over time.

Consider a predictive policing system that uses historical crime data to predict where future crimes are most likely to occur. The model says: “Deploy officers to these neighborhoods.” Officers are deployed there.

More arrests happen there—as a mathematical certainty, because there are more officers there to make arrests. The arrest data flows back into the training dataset. The model learns, with even higher confidence, that those neighborhoods are where crime happens. It directs even more police resources there. The cycle intensifies.

The humans reviewing the system's outputs see increased arrests and interpret them as validation of the model's predictions. The model, in their experience, is “working.” But what the system is actually measuring is not underlying crime; it is the combination of underlying crime plus policing intensity plus historical data artifacts—and it is making policing intensity more unequal over time.

This feedback loop problem is not unique to predictive policing. It applies wherever an AI system's outputs become the inputs for future training data, and where the system's deployment affects the phenomenon it is trying to measure. Credit scoring systems that affect access to credit affect future creditworthiness. Hiring algorithms that affect who gets interviews affect who has professional experience. Recidivism prediction tools that affect sentencing affect post-release conditions that themselves influence recidivism.

In each case, HITL oversight does not break the loop. Human reviewers validate individual decisions, not the systemic effect of those decisions aggregated over time. The loop runs beneath the level where individual human review operates.

14.4 COMPAS: A Case Study in Algorithmic Justice

In May 2016, ProPublica published an investigation that changed the conversation about AI in criminal justice. Julia Angwin, Jeff Larson, and colleagues had obtained two years of COMPAS recidivism scores for de-

fendants in Broward County, Florida, and compared those scores to actual re-arrest data over the following two years (Angwin et al., 2016).

COMPAS—Correctional Offender Management Profiling for Alternative Sanctions—is a proprietary risk assessment tool widely used in American courts to predict the likelihood that a defendant will reoffend. Its scores influence pretrial detention decisions, sentencing recommendations, and parole determinations. It is, in short, an algorithm that affects how much time people spend incarcerated.

ProPublica’s analysis found that the system’s errors were not distributed equally. Black defendants who did not actually reoffend were nearly twice as likely as white defendants who did not reoffend to be labeled as high-risk—a false positive rate disparity. White defendants who did reoffend were more likely to have been labeled low-risk—a false negative rate disparity.

The company behind COMPAS, Northpointe, responded that their tool was calibrated—that among people assigned a given risk score, the proportion who actually reoffended was approximately equal for Black and white defendants. They were right about this. And ProPublica was also right about the disparate false positive and false negative rates. Both claims were mathematically true simultaneously.

This is not a contradiction. It is the **fairness impossibility theorem**.

14.5 When Fairness Definitions Conflict

Researchers have formally identified at least three distinct mathematical definitions of algorithmic fairness, and in 2017, Alexandra Chouldechova published a proof that they cannot all be satisfied simultaneously when base rates differ across groups (Chouldechova, 2017).

Demographic parity (also called statistical parity): The proportion of individuals receiving a positive prediction is equal across groups. A hiring algorithm satisfies demographic parity if it recommends candidates for interviews at the same rate regardless of race.

Calibration (predictive rate parity): Among people assigned a given risk score, the actual outcome rate is the same across groups. A recidivism tool satisfies calibration if, among everyone scored 7/10, the same percentage of Black and white defendants actually reoffend.

Equalized odds : False positive rates and false negative rates are equal across groups. A recidivism tool satisfies equalized odds if the rate of falsely labeling a non-reoffending person as high-risk is the same for all demographic groups.

Chouldechova (2017) proved that when the base rate of the predicted outcome differs across groups—whatever the reason for the difference—it is mathematically impossible for a model to simultaneously satisfy calibration and equal false positive rates.

This is not a solvable engineering problem. It is a logical impossibility. When base rates differ, satisfying one fairness criterion necessarily violates another. The choice between fairness definitions is a normative, political question about whose errors we are most concerned about, whose false positives are most costly, what justice requires—not a technical question with a correct answer. The COMPAS case was not, at bottom, a story about a poorly built algorithm. It was a story about a society choosing, perhaps without realizing it was choosing, which fairness definition it preferred to

apply to decisions about human liberty. The algorithm made that choice visible. It did not make the choice.

14.6 The Limits of Human Oversight as an Ethics Solution

There is a tempting but inadequate response to all of this: “Add more human review.” If the model is biased, have humans check more decisions. If the algorithm is unfair, make sure people are looking over its shoulder.

This response fails for three reasons.

First, humans reviewing AI outputs can absorb and launder the AI’s biases rather than catching them. The Amazon case is the exception that shows what catching bias looks like: the bias surfaced because a dedicated engineering team went looking for it during development, not because routine review flagged it in production. When reviewers trust the system, when the system appears to be working, when the organizational culture treats the algorithm as a tool rather than as an autonomous agent, human review provides accountability theater rather than genuine oversight.

Second, human reviewers carry their own biases. Research on human hiring decisions, judicial sentencing, medical diagnosis, and loan approval has repeatedly found that unaided human judgment is also racially disparate, gender-disparate, and socioeconomically disparate. Adding more humans to a biased process does not neutralize the bias; it may entrench the majority view of what “normal” looks like.

Third, human oversight is most valuable for individual case review and largely blind to systemic effects. A human can evaluate whether the reasoning for a specific decision sounds right. A human cannot, while reviewing individual decisions at production speed, track whether the

system is producing disparate outcomes across protected groups at the aggregate level. That requires **statistical auditing**, not case review.

This is not an argument that human oversight is worthless. It is an argument that human oversight, without intentional fairness design, statistical auditing, diverse reviewer teams, and explicit contestability mechanisms, does not constitute meaningful ethical accountability.

14.7 The Accountability Gap

When a HITL system makes a harmful decision, who is responsible?

This question has no clean answer in current legal frameworks, and the ethical answers are contested. Consider a sequence: a defendant is scored high-risk by COMPAS; a judge who knows the score sends the defendant to prison rather than probation; the score was the product of a model trained on racially disparate historical data; the model was built by a private company; the score was used by a public agency; the algorithm is proprietary and cannot be reviewed in open court.

Who bears moral and legal responsibility? The model's developers? The company that sold it? The judge who used it? The agency that deployed it? The political system that produced the racially disparate data?

Current frameworks tend to locate responsibility with the human decision-maker who used the AI output—the judge, the loan officer, the hiring manager. This is not unreasonable: humans should remain accountable for consequential decisions. But it creates a perverse incentive structure when those decision-makers are given little ability to interrogate the AI's reasoning, little visibility into its aggregate effects, and institutional pressure to trust its outputs.

Explainability. Understanding why a system made a specific prediction. If a loan application is rejected and the applicant asks why, they should receive an explanation that is genuinely informative. “The algorithm decided” is not an explanation; it is an evasion. The EU’s General Data Protection Regulation (Articles 13–15) (European Parliament and Council of the European Union, 2016) grants individuals a right to “meaningful information about the logic involved” in automated decisions, and the EU Court of Justice—in Case C-203/22, *Dun & Bradstreet Austria* (2025) (Court of Justice of the European Union, 2025)—has read this as requiring a concise, intelligible explanation of the procedure and principles actually applied to the individual—not a disclosure of the algorithm itself, but not an evasion either.

Auditability. The ability to independently review a system’s behavior at scale. Individual explanations reveal specific decisions; audits reveal patterns across many decisions. Audits can detect disparate impact that no individual case review would ever uncover. Mandatory independent auditing of high-stakes AI systems is an increasingly common policy proposal across jurisdictions.

Contestability. The ability to meaningfully challenge a decision. A person who receives an unfavorable algorithmic decision and is told they can “appeal” to a process that simply re-runs the same algorithm has not been given meaningful contestability. True contestability requires human review by a decision-maker who has both the authority and the information to override the algorithm, and who is trained to consider whether the algorithmic input was appropriate in the specific case.

14.8 Designing for Fairness: What It Actually Requires

Building a fair HITL system requires more than good intentions and diverse training data, though both help.

It requires deciding, explicitly and transparently, which fairness definition matters most for the specific application—and being honest about the tradeoffs that choice entails. Different applications warrant different choices. A hiring algorithm and a medical screening tool and a recidivism prediction model involve different cost structures around false positives and false negatives, different legal frameworks, and different stakeholder communities.

It requires auditing at the aggregate level, on an ongoing basis, across protected characteristics relevant to the deployment context. Not a one-time pre-launch fairness assessment, but regular monitoring that tracks whether disparate impacts are appearing as the model's outputs change and as deployment context shifts.

It requires diverse teams—not diversity as a compliance gesture, but because the failure modes of AI systems are often visible to people with relevant lived experience and invisible to those without it. The facial recognition systems that failed for darker-skinned faces were built and tested primarily by teams that did not include many people with darker skin. Homogeneous teams have systematic blind spots.

It requires authentic contestability: clear pathways for people to challenge unfavorable decisions, implemented by people who understand those pathways are meant to be used, not minimized.

Fairness Is Not a Property

Fairness is not a property a system achieves once and possesses permanently. It is a relationship between a system, a deployment context, and a society—one that requires ongoing attention, ongoing audit, and ongoing willingness to change.

Find a decision that affected you recently and was, in whole or in part, made or influenced by an algorithm—a credit decision, a job application screening, a content moderation decision, a recommendation, a price quoted to you. Ask: Do you know why the decision was made? Do you have any mechanism to challenge it? Would you even know if the algorithm systematically disadvantaged people like you? The gap between what you can answer and what would constitute genuine accountability is the space this chapter has tried to describe.

Chapter Summary

Key Concepts

- Bias enters AI systems at multiple stages: collection bias (whose data?), annotation bias (whose judgment?), model bias (what patterns did it amplify?), deployment bias (who does it affect?), and feedback loop amplification
- HITL does not automatically produce fair systems: human reviewers can absorb and legitimize bias rather than catching it; aggregate effects are invisible to individual case review

- The COMPAS case illustrates that different mathematical definitions of fairness can simultaneously be true and contradict each other
- The fairness impossibility theorem (Chouldechova, 2017): when base rates differ across groups, calibration and equal false positive rates cannot both be satisfied—the choice is normative, not technical
- Explainability, auditability, and authentic contestability are the structural elements of accountable AI deployment

Key Examples

Amazon recruiting AI (2018) Model learned gender proxies; engineers caught the bias roughly a year into development, but each fix produced new proxies; project killed before deployment as unfixable

Gender Shades (Buolamwini & Gebru, 2018) Commercial facial recognition error rates up to 34 percentage points higher for darker-skinned women; collection bias in training datasets

COMPAS (ProPublica, 2016) Both Northpointe's calibration defense and ProPublica's disparate error rate finding were mathematically true; the fairness impossibility theorem in a real case

Predictive policing feedback loops Arrests in over-policed neighborhoods increase, feeding back into the model, intensifying predictions, requiring more resources to the same neighborhoods

Key Principles

- Choose fairness definitions explicitly and transparently, knowing that the choice involves normative tradeoffs
- Audit at the aggregate level, regularly, across protected characteristics
- Design for explainability, auditability, and authentic contestability
- Recognize that human oversight without intentional fairness infrastructure provides accountability theater, not accountability
- Fairness is not a property achieved once—it is an ongoing relationship between a system, a deployment context, and an affected community

Five Dimensions Check

- *Stakes Calibration* (Dimension 4): Fairness is a stakes question—who bears the cost when the system is wrong, and across which groups
 - *Intervention Design* (Dimension 2): Human oversight is a design choice with real limits; without fairness infrastructure it becomes accountability theater
 - *Feedback Integration* (Dimension 5): Feedback loops can amplify bias—predictive-policing arrests feed back into the model and intensify the original disparity
-

The next chapter turns from fairness to failure: what goes wrong when HITL systems break down in practice, how to trace the breakdown to its source, and what good remediation looks like.

Technical Appendix 14A: Fairness Metrics, Bias Auditing, and Algorithmic Accountability Tools

Fairness in machine learning is not a single property—it is a family of distinct mathematical requirements, some mutually incompatible. This appendix formalizes the major definitions, proves the impossibility theorem, and provides implementation tools for bias auditing and mitigation.

— Mathematical foundations and Python implementations for building and auditing fair AI systems

14A.1 Formal Fairness Definitions

Let:

- $\hat{Y}$ = model prediction (binary: 0 or 1)
- Y = true outcome
- A = sensitive attribute (e.g., race, gender, age group)
- X = all other features

Demographic Parity (Statistical Parity).

$$P(\hat{Y} = 1 \mid A = 0) = P(\hat{Y} = 1 \mid A = 1)$$

The probability of a positive prediction is equal across groups. A hiring algorithm satisfies demographic parity if it recommends candidates at equal rates regardless of protected attribute.

Equalized Odds (Hardt et al., 2016).

$$P(\hat{Y} = 1 \mid Y = y, A = 0) = P(\hat{Y} = 1 \mid Y = y, A = 1) \quad \text{for } y \in \{0, 1\}$$

Both the true positive rate (TPR) and false positive rate (FPR) are equal across groups. A medical screening tool satisfies equalized odds if it correctly identifies sick patients—and incorrectly flags healthy ones—at equal rates regardless of demographic group.

Equal Opportunity (weaker form of equalized odds): only the TPR constraint:

$$P(\hat{Y} = 1 \mid Y = 1, A = 0) = P(\hat{Y} = 1 \mid Y = 1, A = 1)$$

Predictive Parity (predictive rate parity—equal positive predictive value across groups; not to be confused with *calibration* in the score sense, which the impossibility theorem below concerns).

$$P(Y = 1 \mid \hat{Y} = 1, A = 0) = P(Y = 1 \mid \hat{Y} = 1, A = 1)$$

Among all individuals predicted positive, the fraction who are actually positive is equal across groups.

14A.2 The Fairness Impossibility Theorem

Chouldechova (2017) proved that calibration and equal FPR cannot simultaneously hold when base rates differ across groups.

Setup. Let $b_A = P(Y = 1 \mid A)$ denote the base rate of the outcome for group A .

Theorem (Chouldechova, 2017). If $b_0 \neq b_1$ (base rates differ), then no score s can simultaneously satisfy:

1. **Calibration:** $P(Y = 1 \mid s, A) = s$ for all groups
2. **Equal FPR:** $P(\hat{Y} = 1 \mid Y = 0, A = 0) = P(\hat{Y} = 1 \mid Y = 0, A = 1)$
3. **Equal FNR:** $P(\hat{Y} = 0 \mid Y = 1, A = 0) = P(\hat{Y} = 0 \mid Y = 1, A = 1)$

Intuitive proof. For a calibrated model at threshold τ :

$$\text{FPR}_A = \frac{(1 - \text{PPV}_A) \cdot \text{PR}_A}{1 - b_A}$$

where PPV_A is the positive predictive value and PR_A is the prediction rate for group A . If $b_0 \neq b_1$ and PPV is equalized (predictive parity holds), then FPR equality requires PR_A to adjust in a way that inevitably violates FNR equality. All three cannot hold simultaneously.

```

1 import numpy as np
2 from sklearn.metrics import confusion_matrix
3
4 def compute_fairness_metrics(y_true, y_pred, groups):
5     results = {}
6     for group_val in [0, 1]:
7         mask = groups == group_val
8         yt, yp = y_true[mask], y_pred[mask]
9         tn, fp, fn, tp = confusion_matrix(yt, yp,
10                                         labels=[0, 1]).ravel()
11         n = len(yt)
12         results[f'group_{group_val}'] = {

```

```

12         'TPR':      tp / (tp + fn) if (tp + fn) > 0
           else np.nan,
13         'FPR':      fp / (fp + tn) if (fp + tn) > 0
           else np.nan,
14         'PPV':      tp / (tp + fp) if (tp + fp) > 0
           else np.nan,
15         'pred_rate': (tp + fp) / n,
16         'base_rate': (tp + fn) / n,
17     }
18     return results
19
20 def demonstrate_impossibility(base_rate_0=0.2,
21                               base_rate_1=0.4, n=2000):
22     np.random.seed(42)
23     groups = np.concatenate([np.zeros(n // 2), np.ones(n //
24                               2)])
25
26     # Calibrated, informative scores:  $P(Y = 1 \mid s, \text{group}) =$ 
27     #  $s$ .
28     # Beta(2, 8) has mean 0.2 = base_rate_0; Beta(4, 6)
29     # mean 0.4.
30     scores = np.where(
31         groups == 0,
32         np.random.beta(2, 8, n),
33         np.random.beta(4, 6, n)
34     )
35     y = np.random.binomial(1, scores)
36
37     y_pred = (scores > 0.3).astype(int)
38     metrics = compute_fairness_metrics(y, y_pred, groups)
39
40     for g, gv in [('group_0', 0), ('group_1', 1)]:
41         m = metrics[g]
42         mask = groups == gv

```

```

39         print(f"{g}: TPR={m['TPR']:.3f},
40               FPR={m['FPR']:.3f}, "
41               f"PPV={m['PPV']:.3f},
42               PredRate={m['pred_rate']:.3f}, "
43               f"MeanScore={scores[mask].mean():.3f}, "
44               f"BaseRate={m['base_rate']:.3f}")
45
46     fpr_gap = abs(metrics['group_0']['FPR'] -
47                  metrics['group_1']['FPR'])
48     ppv_gap = abs(metrics['group_0']['PPV'] -
49                  metrics['group_1']['PPV'])
50     print(f"\nFPR gap: {fpr_gap:.3f} | PPV gap
51           (calibration): {ppv_gap:.3f}")
52     print("Note: minimizing PPV gap (calibration) cannot
53           prevent FPR gap "
54           "when base rates differ.")

```

Listing 14.1: Fairness impossibility demonstration

14A.3 Bias Detection with Fairlearn

Microsoft's Fairlearn library provides audit dashboards and mitigation algorithms.

```

1  from fairlearn.metrics import (
2      MetricFrame, demographic_parity_difference,
3      equalized_odds_difference, selection_rate,
4  )
5  from sklearn.metrics import accuracy_score,
6      precision_score, recall_score
7
8  def full_fairness_audit(y_true, y_pred, sensitive_features,
9                          feature_name):

```

```
8     mf = MetricFrame(  
9         metrics={  
10             'accuracy':    accuracy_score,  
11             'precision':   precision_score,  
12             'recall':     recall_score,  
13             'selection_rate': selection_rate,  
14         },  
15         y_true=y_true,  
16         y_pred=y_pred,  
17         sensitive_features=sensitive_features,  
18     )  
19     dp_diff = demographic_parity_difference(  
20         y_true, y_pred,  
21         sensitive_features=sensitive_features  
22     )  
23     eo_diff = equalized_odds_difference(  
24         y_true, y_pred,  
25         sensitive_features=sensitive_features  
26     )  
27     print(f"Fairness Audit: {feature_name}")  
28     print(mf.by_group.to_string())  
29     print(f"Demographic parity diff:  {dp_diff:.4f}  (0 =  
30         fair)")  
31     print(f"Equalized odds diff:      {eo_diff:.4f}  (0 =  
32         fair)")  
33     if abs(dp_diff) > 0.1:  
34         print(f"WARNING: Demographic parity gap  
35             ({dp_diff:.3f}) exceeds 0.1")  
36     if abs(eo_diff) > 0.1:  
37         print(f"WARNING: Equalized odds gap ({eo_diff:.3f})  
38             exceeds 0.1")
```



```
36     return {'metric_frame': mf, 'dp_diff': dp_diff,  
            'eo_diff': eo_diff}
```

Listing 14.2: Full fairness audit with Fairlearn

14A.4 Bias Mitigation Techniques

Three categories of mitigation correspond to three pipeline stages: **pre-processing** (modify training data before training), **in-processing** (add fairness constraints to the learning objective), and **post-processing** (adjust predictions after training).

```
1  from fairlearn.reductions import ExponentiatedGradient,  
    DemographicParity  
2  from fairlearn.postprocessing import ThresholdOptimizer  
3  from sklearn.linear_model import LogisticRegression  
4  import numpy as np  
5  
6  # In-processing: Exponentiated Gradient  
7  def train_fair_inprocessing(X_train, y_train,  
    sensitive_train):  
8      mitigator = ExponentiatedGradient(  
9          estimator=LogisticRegression(max_iter=500),  
10         constraints=DemographicParity(),  
11         eps=0.01,  
12     )  
13     mitigator.fit(X_train, y_train,  
        sensitive_features=sensitive_train)  
14     return mitigator  
15  
16 # Post-processing: Threshold Optimizer  
17 def train_fair_postprocessing(estimator, X_train, y_train,  
    sensitive_train):
```

```
18     postprocess_est = ThresholdOptimizer(  
19         estimator=estimator,  
20         constraints='equalized_odds',  
21         predict_method='predict_proba',  
22         objective='accuracy_score',  
23     )  
24     postprocess_est.fit(X_train, y_train,  
25         sensitive_features=sensitive_train)  
26     return postprocess_est  
27 # Pre-processing: Reweighting (removes group-label  
28     correlation by instance weights)  
29     def compute_reweighting_weights(y_train, sensitive_train):  
30         n = len(y_train)  
31         weights = np.ones(n)  
32         for g in np.unique(sensitive_train):  
33             for l in np.unique(y_train):  
34                 mask = (sensitive_train == g) & (y_train == l)  
35                 n_gl = mask.sum()  
36                 if n_gl > 0:  
37                     expected = (sensitive_train == g).sum() /  
38                         n \  
39                         * (y_train == l).sum() / n * n  
40                     weights[mask] = expected / n_gl  
41         return weights
```

Listing 14.3: In-processing, post-processing, and pre-processing mitigation

14A.5 Disparate Impact Ratio (AIF360)

The disparate impact ratio (DIR) is the standard metric used in US employment discrimination law (the “4/5ths rule”):

$$\text{DIR} = \frac{P(\hat{Y} = 1 \mid A = \text{minority})}{P(\hat{Y} = 1 \mid A = \text{majority})}$$

A DIR below 0.8 is a legal threshold for disparate impact in employment contexts.

```
1 from aif360.datasets import BinaryLabelDataset
2 from aif360.metrics import BinaryLabelDatasetMetric
3
4 def compute_disparate_impact(df, label_col, sensitive_col,
5                             privileged_val,
6                             unprivileged_val):
7     privileged_groups = [{sensitive_col: privileged_val}]
8     unprivileged_groups = [{sensitive_col:
9                             unprivileged_val}]
10
11     dataset = BinaryLabelDataset(
12         df=df[[sensitive_col, label_col]],
13         label_names=[label_col],
14         protected_attribute_names=[sensitive_col],
15     )
16     metric = BinaryLabelDatasetMetric(
17         dataset,
18         unprivileged_groups=unprivileged_groups,
19         privileged_groups=privileged_groups,
20     )
21     dir_val = metric.disparate_impact()
22     mean_diff = metric.mean_difference()
```

```

22     print(f"Disparate Impact Ratio: {dir_val:.4f}")
23     print(f" (Legal threshold: >= 0.80; below = potential
        disparate impact)")
24     print(f"Mean Difference (stat parity): {mean_diff:.4f}")
25     if dir_val < 0.80:
26         print("ALERT: Disparate impact below 4/5ths rule
            threshold.")
27     return {'disparate_impact': dir_val, 'mean_difference':
        mean_diff}

```

Listing 14.4: Disparate impact ratio via AIF360

14A.6 Counterfactual Fairness

Counterfactual fairness (Kusner et al., 2017) asks: would the prediction change if the individual belonged to a different demographic group, holding all causally independent features constant?

$$P(\hat{Y}_{A \leftarrow a}(U) = y \mid X = x, A = a) = P(\hat{Y}_{A \leftarrow a'}(U) = y \mid X = x, A = a)$$

This requires a causal model of the data-generating process. The testing principle can be approximated by flipping the sensitive attribute and measuring prediction change.

```

1  def counterfactual_fairness_test(model, X_test,
    sensitive_col,
2                                     group_a_val, group_b_val):
3      X_a = X_test[X_test[sensitive_col] ==
        group_a_val].copy()
4      X_a_counter = X_a.copy()
5      X_a_counter[sensitive_col] = group_b_val
6

```

```
7      # The model must see the protected attribute for the
      # flip to register.
8      # (A full counterfactual would also update the
      # attribute's causal
9      # descendants; this approximate test holds the other
      # features fixed.)
10     preds_original = model.predict(X_a)
11     preds_counter  = model.predict(X_a_counter)
12
13     changed = (preds_original != preds_counter).mean()
14     print(f"Counterfactual flip rate: {changed:.4f}")
15     print(f" (0 = counterfactually fair; >0 = predictions
      # changed by group flip)")
16     return changed
```

Listing 14.5: Approximate counterfactual fairness test

14A.7 Feedback Loop Detection

Detecting whether a deployed model is amplifying disparities over time requires monitoring group-level prediction rates across retraining cycles.

```
1  import numpy as np
2  from scipy import stats
3
4  def detect_feedback_amplification(predictions_history,
      group_labels,
5                                     n_cycles=5):
6      disparity_over_time = []
7      for cycle in range(min(n_cycles,
          len(predictions_history))):
8          preds = np.array(predictions_history[cycle])
9          groups = np.array(group_labels[cycle])
```

```

10     rates = {g: preds[groups == g].mean() for g in
11              np.unique(groups)}
12     disparity_over_time.append(max(rates.values()) -
13                                min(rates.values()))
14
15     cycles = np.arange(len(disparity_over_time))
16     slope, _, r_val, p_val, _ = stats.linregress(cycles,
17                                                  disparity_over_time)
18
19     print(f"Disparities over cycles: {[f'{d:.4f}' for d in
20                                         disparity_over_time]}")
21     print(f"Trend slope: {slope:.4f}/cycle,
22           R²={r_val**2:.3f}, p={p_val:.4f}")
23     if slope > 0 and p_val < 0.05:
24         print("WARNING: Statistically significant disparity
25               amplification.")
26     elif slope > 0:
27         print("NOTE: Upward trend in disparity (not yet
28               statistically significant).")
29     else:
30         print("No evidence of feedback amplification
31               detected.")
32
33     return {'slope': slope, 'p_value': p_val,
34            'amplification_detected': slope > 0 and p_val <
35            0.05}

```

Listing 14.6: Feedback loop amplification test

14A.8 Pre-Deployment Audit Checklist

A structured audit checklist derived from NIST AI RMF (2023), US disparate-impact doctrine (the EEOC four-fifths rule), and common practice for EU AI Act conformity preparation:

Stage	Check	Tool	Threshold
Data	Demographic representation in training set	EDA + sampling stats	No group < 5% of N
Data	Label distribution by group	MetricFrame	Parity diff < 0.10
Model	Demographic parity at deployment threshold	Fairlearn	Diff < 0.10
Model	Equalized odds	Fairlearn	Diff < 0.10
Model	Calibration by group	ECE per group	Per-group ECE < 0.05
Model	Disparate impact ratio	AIF360	DIR $\geq$ 0.80
Deployment	Slice-based monitoring (monthly)	SliceFinder	Alert at 5% degradation
Deployment	Feedback loop correlation test	Custom	Slope $p < 0.05 \rightarrow$ halt
Process	Reviewer demographic diversity	Audit log	Document + review
Process	Contestability mechanism functional	Process audit	100% cases have pathway

14A.9 Exercises

Exercise 14A.1 — The Impossibility in Practice

Run `demonstrate_impossibility(base_rate_0=0.15, base_rate_1=0.35, n=2000)`. Report the false-positive-rate gap and the PPV (calibration) gap between the two groups. Using the impossibility theorem, explain why minimizing the calibration gap cannot also eliminate the FPR gap when the base rates differ.

Exercise 14A.2 — Disparate Impact Audit

Using the Adult/Census Income dataset with `sex` as the protected attribute, compute the disparate impact ratio and the mean-difference statistic with AIF360. Does the classifier satisfy the four-fifths (0.80) rule? Propose one mitigation strategy and predict its effect on overall accuracy.

Exercise 14A.3 — Feedback Loop Simulation

Simulate a predictive-policing feedback loop over several cycles: two neighborhoods with different base crime rates, where each cycle's arrests become the next cycle's training data. (a) Build the model and run the loop. (b) Measure whether the between-group disparity amplifies across cycles. (c) Propose an intervention that breaks the loop, and demonstrate its effect.

Exercise 14A.4 — Counterfactual Fairness Check

For a hiring model trained on historical data, apply the counterfactual fairness test. (a) Flip the protected (gender) attribute and measure the prediction flip rate. (b) Interpret the result. (c) Explain why a flip rate of zero does not by itself guarantee fairness, considering proxy features.

14A.10 Further Reading

- Chouldechova (2017) — Fair prediction with disparate impact. *Big Data*, 5(2).
- Hardt et al. (2016) — Equality of opportunity in supervised learning. *NeurIPS* 29.
- Kusner et al. (2017) — Counterfactual fairness. *NeurIPS* 30.
- Bird, S., et al. (2020). Fairlearn: A toolkit for assessing and improving fairness in AI. *Microsoft Research*.

15 When Things Go Wrong (And How to Fix Them)

The organizations that handle failures best are not the ones that never have them—they're the ones that built the infrastructure to find them quickly.

— on finding failure fast

15.1 The Day the Market Lost Its Mind

At 2:32 PM on May 6, 2010, a large mutual fund's algorithm began selling 75,000 E-mini futures contracts to hedge its portfolio. Ten minutes later, with the Dow Jones Industrial Average already down more than 300 points on a nervous day, the equity market began to fall vertically—roughly 600 points in about five minutes—bottoming out nearly 1,000 points down for the day by 2:47 PM, briefly erasing almost \$1 trillion in market value. Blue-chip stocks that had been trading at \$40 that morning suddenly printed trades at a penny. Procter & Gamble fell from \$60 to \$39. Accenture, a Fortune 500 company, traded at one cent per share.

And then, just as suddenly, it reversed. By 3:07 PM, the market had recovered most of its losses.

No single bomb had gone off. No company had collapsed. No foreign power had intervened. What had happened, the CFTC and SEC's joint post-event analysis would later determine, was that automated trading

algorithms had entered a catastrophic feedback loop. Other algorithms, designed to detect and follow large trades, began mimicking the sell program. This accelerated the decline, which triggered more algorithms to sell, which accelerated the decline further (Kirilenko et al., 2017).

The humans who should have been in the loop—the traders, the risk managers, the market regulators—were watching their screens, helpless to intervene. The circuit breakers that existed were insufficient: the market-wide halts never triggered, single-stock trading pauses did not yet exist, and firm-level cutoffs required manual action no one could take in time. Some trading firms simply pulled their bid prices from the market entirely rather than catch the falling knives, which made the cascade worse.

The Flash Crash is the canonical example of what happens when you remove the human from the loop entirely in a complex, high-stakes, real-time system. It wasn't that there were no humans nearby. There were hundreds of them. The problem was that none of them could intervene in time, because the system had been designed—whether deliberately or by architectural accident—to operate faster than human reaction time allowed.

The lesson wasn't only about trading. It was about what failure looks like when HITL design is absent, inadequate, or overwhelmed. And it raised the questions this chapter is about: When HITL systems fail, why do they fail? How do you find the failure? And how do you fix it without making it worse?

15.2 A Taxonomy of HITL Failures

Not all HITL failures look alike. Before you can fix a system, you need to understand which part of it is broken. Across domains—credit scoring, content moderation, medical imaging, fraud detection, hiring tools—HITL failures fall into four distinct categories.

15.2.1 Failure Mode 1: Model Failure

The model itself is producing incorrect predictions. This is the failure mode most people think of first, and it's actually the most straightforward to diagnose, because it leaves a statistical fingerprint. Model failures tend to be systematic: the model consistently underperforms on certain subpopulations, certain input distributions, or certain combinations of features.

The classic example is gender-classification systems that perform well on light-skinned male faces and poorly on dark-skinned female faces—a finding documented by MIT researcher Joy Buolamwini in her landmark Gender Shades study (Buolamwini & Gebru, 2018). This wasn't a random error. It was a systematic one, traceable directly to training data that overrepresented certain demographic groups.

Model failures manifest as:

- **Distributional shift:** The world has changed since the model was trained. A fraud detection model trained on pre-pandemic spending patterns will struggle with post-pandemic ones.
- **Coverage gaps:** The model was never trained on certain cases. A medical AI trained only on hospital data from one region will behave unpredictably on patients from another.
- **Spurious correlations:** The model learned to use a feature that happened to correlate with the outcome in training but doesn't causally relate to it.

15.2.2 Failure Mode 2: Annotation Failure

The training labels are wrong. This is the sneakiest failure mode, because the model has been doing exactly what it was trained to do—the problem is that it was trained on incorrect ground truth.

Annotation failures are common in any domain where human judgment is subjective, difficult, or fatiguing. In content moderation, annotators asked to label thousands of pieces of content per day for “hate speech” or “misinformation” face compounding problems: annotators disagree with each other; standards drift over time as fatigue sets in; annotators bring their own biases; and the labeling guidelines themselves may be ambiguous or culturally specific.

15.2.3 Failure Mode 3: Design Failure

The model may be accurate, the labels may be correct, but the threshold or routing logic is wrong. The system asks for human review at the wrong times—either too often (alert fatigue) or too rarely (overconfident silence).

Design failures are often invisible in standard evaluation metrics. A fraud detection system might have 99.5% accuracy—a number that looks excellent—while still being systematically wrong on a specific category of transactions that matters enormously.

A systematic review of clinical decision support found that clinicians dismissed between 46% and 96% of the alerts they received (Poly et al., 2020). When every signal is noise, the actual signals disappear.

15.2.4 Failure Mode 4: Reviewer Failure

The model is right, the labels are right, the threshold is right—but the humans reviewing the flagged cases are making poor decisions. Reviewer failures happen because of:

- **Reviewer burnout:** Humans reviewing hundreds of cases per day become cognitively exhausted. Decision quality degrades.
- **Automation bias:** Human reviewers who see the model's score before making their own judgment tend to anchor on it—even when the model is wrong. This undermines the entire premise of the human being in the loop.
- **Inconsistent standards:** Different reviewers apply different criteria, creating unreliable labels that corrupt the feedback loop.
- **Perverse incentives:** Reviewers measured on throughput will process faster, review less carefully, and agree with the model more often.

The Four Failure Mode Taxonomy

Model Failure Predictions are systematically wrong for a subpopulation

Annotation Failure Training labels are incorrect, skewed, or drifted

Design Failure Threshold or routing logic is miscalibrated

Reviewer Failure Human decision quality is compromised

The four modes are not mutually exclusive—the most serious HITL failures typically involve multiple modes simultaneously.

15.3 The Diagnostic Question

When a HITL system starts failing, the first task is diagnosis. Before you can fix anything, you need to answer four questions in sequence:

Is the model wrong? Look at the model's performance metrics, sliced by subpopulation, feature combination, and time period. Are there groups of cases where accuracy is systematically lower? Are there time trends suggesting distributional shift?

Are the labels wrong? Look at annotator agreement statistics. Is inter-annotator agreement high? Are there systematic patterns in disagreements? Do different annotation teams produce different labels for the same cases?

Is the threshold wrong? Look at how the system routes cases. What fraction reaches human review? What fraction is auto-decided? Of the cases reaching humans, what fraction are genuinely uncertain versus obviously one thing or the other?

Is the human component failing? Look at reviewer behavior. Is there evidence of automation bias? Are reviewer decisions consistent within reviewers over time? Between reviewers on the same cases?

Each answer points to a different remediation. Getting the diagnosis wrong means applying the wrong fix—which can make things worse.

15.4 Spot Failures and Systematic Failures

There is a critical distinction in failure analysis that determines the urgency and scope of response.

A **spot failure** is a single bad output on a single case. The model misclassified one loan application. The content moderation system incorrectly removed one post. These failures are annoying and should be fixed, but they don't necessarily indicate a broken system.

A **systematic failure** is a pattern of bad outputs across a population of cases. The model is misclassifying all loan applications from a particular ZIP code. The content moderation system is removing a particular category of political speech at ten times the rate of equivalent speech from other perspectives.

Systematic Failures Hide Behind Aggregate Metrics

A model with 95% accuracy is making errors on 5% of cases—but if those errors are concentrated in a specific subpopulation, the harm may be severe while the headline metric looks acceptable. Slice-based evaluation (see the Technical Appendix) is the tool that makes systematic failures visible.

15.5 Case Study: The Apple Card Investigation

In 2019, Apple and Goldman Sachs launched the Apple Card. Within months, users began reporting a pattern on social media: women were receiving strikingly lower credit limits than their husbands, despite shared finances. The most visible case involved the wife of a technology entrepreneur, whose credit limit was one-twentieth of her husband's—despite the couple filing joint tax returns and her holding the higher credit score. When she contacted customer service, representatives could neither explain nor override the algorithm's output; her limit was raised only after her husband's account of the episode went viral (Vigdor, 2019).

The New York Department of Financial Services (NYDFS) opened an investigation. Its report found no evidence of unlawful discrimination: a statistical analysis of underwriting data covering roughly 400,000 New York applicants did not show that gender, directly or through proxies, had driven the outcomes. The model did not use gender as an explicit input,

and the population-level data did not establish a systematically gendered outcome (New York Department of Financial Services, 2021).

But the visible cases were real, and the suspected mechanism was exactly the one fairness researchers warn about: proxy bias. A model trained on variables correlated with gender—income, credit history, spending patterns, themselves shaped by decades of structural wage gaps and credit-market disparities—can reproduce historical inequality without ever seeing a gender field. Whether or not that happened at population scale here, the bank could not demonstrate what had or had not driven any individual decision, because no one could explain individual decisions at all (New York Department of Financial Services, 2021).

What the NYDFS did fault was the human layer. Goldman Sachs's customer service processes—the human side of the loop—had no adequate mechanism to explain, challenge, or appeal an algorithmic credit decision. The investigation concluded in 2021 with a finding that Goldman Sachs had not violated fair-lending law, but the NYDFS made clear that algorithmic credit decisions fall squarely within its regulatory authority and that lenders bear the burden of demonstrating their models do not discriminate (New York Department of Financial Services, 2021).

The story did not end there. In October 2024, the Consumer Financial Protection Bureau ordered Goldman Sachs to pay a \$45 million civil penalty and at least \$19.8 million in consumer redress over Apple Card servicing failures—uninvestigated transaction disputes, delayed refunds, and inaccurate information furnished to credit reports (Consumer Financial Protection Bureau, 2024). The underwriting model had passed its fairness audit; the human and servicing layer around it had not.

The story has several HITL lessons. First, the human customer-service representatives were in the loop but had no diagnostic tools and no effective authority to override—the limit was changed only after media pressure, outside any process. Second, the absence of a protected attribute in the

feature list is not a fairness guarantee—and an aggregate statistical clean bill of health is not an answer to an individual customer asking *why*. Third, the problem was discovered not by the bank’s internal monitoring but by users comparing notes on social media—a pattern-detection mechanism no compliance department had planned for. The case illustrates that human oversight cannot function when humans are given outputs without evidence, and that fairness testing must look at outcomes by protected group, not merely at whether protected attributes appear in the feature list.

15.6 Case Study: Annotation Bias in Content Moderation

Facebook’s 2020 civil rights audit—conducted by independent auditors commissioned by the company—raised concerns about uneven enforcement across communities: certain communities experienced systematic over-enforcement of content policies, while policy-violating content targeting other communities went under-enforced (Murphy & Meta Platforms, Inc., 2020). The underlying challenge is persistent: content moderation systems trained on annotator judgments can encode the skew of the annotator pool into automated enforcement at scale.

The mechanism is now well-documented in the research literature. Sap and colleagues showed that tweets containing African American English were up to twice as likely to be labeled offensive as equivalent tweets in Mainstream American English—an effect they traced to annotators’ insensitivity to dialect, which dropped sharply when annotators were told a tweet’s dialect (Sap et al., 2019). The annotators were not acting in bad faith—they were applying the guidelines as they understood them, but their understanding was shaped by cultural distance from the dialect they were labeling. The model learned their biases.

The pattern is self-reinforcing. A model trained on a demographically skewed annotator pool over-moderates content from particular communities. The platform’s complaint data—appeals and user reports—reflects the AI’s enforcement patterns, not the underlying distribution of violations. When that complaint data is then used to retrain or fine-tune the model, the bias compounds.

The Annotation Feedback Loop

Annotation bias operates like the credit-model feedback loop described in the previous section: a biased training signal produces biased predictions, which shape the data collected next. The difference is that content moderation operates in a domain where “ground truth” is inherently contested—what counts as toxic, hateful, or harmful is subject to cultural and political disagreement. There is no credit-report ground truth to validate against.

The fix requires what the audit recommended: diverse annotator pools, regular fairness audits stratified by community, clear annotation guidelines that explicitly address dialect and cultural variation, and ongoing monitoring of enforcement-rate disparities. It also requires the architecture of the loop to register disagreement as a signal rather than noise—when annotator disagreement clusters along demographic lines, it is telling you something about the task, not just about the annotators.

15.7 The Audit Trail: You Can’t Fix What You Can’t Trace

Both case studies above depended on the existence of records. The NYDFS inquiry into Apple Card could proceed because Goldman Sachs maintained logs of credit decisions; the Meta civil rights audit could measure enforce-

ment disparities because the platform logged enforcement actions. The data existed—and that existence was the difference between an actionable finding and an unverifiable anecdote.

This is not how most HITL systems are built by default. In many organizations, the model runs in a black box. Predictions are made, decisions are executed, and no one asks for a receipt. When a failure is later identified, the team faces the worst possible diagnostic situation: trying to reconstruct what happened from incomplete records.

An audit trail serves three functions:

Detection. Systematic monitoring of logged predictions can surface patterns that would otherwise be invisible.

Diagnosis. When a failure is identified, the audit trail lets you trace it to its cause. Was this a model error, an annotation error, or a threshold error?

Remediation. When harm has occurred, the audit trail tells you who was affected and when. This is the prerequisite for notification, restitution, and making things right.

Building audit trails into HITL systems from the start is not expensive. Storing every prediction with its context costs a fraction of what failure investigation and remediation do.

15.8 Post-Mortems and the HITL Five Whys

When a HITL failure is identified and remediated, the work is not done. The most valuable thing an organization can do after fixing a failure is understand why it happened—not to assign blame, but to prevent recurrence.

Google's Site Reliability Engineering handbook formalized the concept of the **blameless post-mortem**: the principle that the goal of failure analysis is to understand system failures, not to punish individuals (Beyer et al., 2016). This matters especially in HITL systems, where annotators, reviewers, and data scientists are usually responding rationally to the incentive structures and information environments around them.

The HITL version of the **five whys**—the lean manufacturing technique for tracing failures to root causes—looks like this:

1. Why did the model make systematically incorrect predictions? *Because it was trained on biased labels.*
2. Why were the labels biased? *Because the annotation team was working under time pressure and used available news framing as implicit context.*
3. Why were they working under time pressure? *Because the annotation schedule was set by a product team that didn't understand annotation quality degrades under pressure.*
4. Why didn't the product team understand this? *Because there was no standardized process for communicating annotation quality constraints to product planning.*
5. Why wasn't that process in place? *Because the organization's HITL processes were designed primarily around speed, not quality—and no one owned cross-functional accountability for label quality.*

The fifth “why” reveals a structural or organizational root cause that no amount of technical improvement alone will fix.

15.9 The Framework Applied: Five Dimensions of HITL Failure

The Five Dimensions framework provides a useful lens for categorizing HITL failures:

- The Flash Crash was a **Timing** failure: the system operated on timescales faster than human intervention was possible.
- The Apple Card case was a **Feedback Integration** failure layered on an **Intervention Design** failure: the model reproduced proxy bias absorbed from its training data, and the human layer of customer service had no tools or authority to interrogate the decision.
- The annotation-bias case was an **Uncertainty Detection** failure at the annotation level combined with a monitoring design failure: annotator disagreement clustering along demographic lines was treated as noise, and enforcement rates were not tested by community.
- The Nest thermostat (Section 1.5) was an **Intervention Design** failure: the human wasn't meaningfully in the loop even when they nominally were.

When diagnosing a HITL failure, a useful first step is to map it onto the Five Dimensions and ask: which dimension failed first? The answer usually points to the most leverage for remediation.

15.10 Recovery: What Good Remediation Looks Like

Once a failure is diagnosed, remediation has three phases:

Stop the bleeding. Reduce ongoing harm while the deeper fix is being designed. Temporarily raise human review thresholds, pause automated decisions in a specific domain, or route a category of cases entirely to human reviewers.

Fix the root cause. Address the actual failure mode—retrain the model on corrected labels, revise the threshold, restructure the review process, or redesign the feedback loop.

Prevent recurrence. Implement monitoring that would have caught this failure earlier. Create a test that encodes this specific failure pattern and will alert if it reappears.

The Infrastructure of Recovery

The organizations that handle HITL failures best are not the ones that never have failures. They are the ones that built the diagnostic infrastructure—audit trails, slice-based monitoring, blameless post-mortems—to find failures quickly, remediate them decisively, and prevent recurrence systematically.

Think of a HITL system you interact with regularly—a content feed, a recommendation engine, a workplace tool, a fraud detection system. Has it ever done something that seemed *systematically* wrong (the same kind of error, repeatedly, on the same kind of case)

rather than randomly wrong? If so, which of the four failure modes does it most resemble? If you can't tell—what information would you need to diagnose it?

Chapter Summary

Key Concepts

- HITL failures fall into four categories: model failure, annotation failure, design failure, and reviewer failure. Each requires a different diagnosis and a different fix.
- Spot failures (single bad outputs) are different from systematic failures (patterns across a population). Systematic failures demand systematic responses.
- The audit trail—comprehensive logging of predictions, decisions, and system changes—is the prerequisite for diagnosis, remediation, and accountability.
- Post-mortems should be blameless: the goal is to understand system failures, not punish individuals.
- The Five Whys traces failures to structural root causes, which are usually organizational rather than technical.

Key Diagnostic Questions

Is the model wrong? Check slice-based performance metrics across sub-populations and time periods

Are the labels wrong? Check inter-annotator agreement and annotation history

Is the threshold wrong? Check routing statistics and human review rates

Is the human component failing? Check reviewer behavior, consistency, and automation bias indicators

Five Dimensions Check

- *Timing* (Dimension 3): The Flash Crash is the pure timing failure—a system operating on timescales where no human intervention was physically possible
- *Intervention Design* (Dimension 2): Apple Card’s customer service was a human layer with no tools and no authority—presence in the loop without power over the loop
- *Feedback Integration* (Dimension 5): Both case studies compounded through their own feedback loops—biased predictions shaped the next round of data, quietly amplifying the failure

In the next chapter, we extend HITL analysis beyond text and images—into audio, time-series data, scientific discovery, and the physical world of robots and prosthetics.

Technical Appendix 15A: Error Analysis Frameworks

15A.1 Calibration Analysis

Standard ML evaluation—accuracy, precision, recall, F1, AUC—measures aggregate model performance. In HITL systems, aggregate metrics are necessary but not sufficient. A model with 95% accuracy may have systematic failure modes invisible in aggregate statistics.

15A.1.1 Cost-Asymmetric Threshold Design

In HITL systems, the asymmetry between false positive and false negative costs is critical for threshold design. For a binary classifier, the optimal threshold under asymmetric costs is:

$$\theta^* = \frac{C_{FP}}{C_{FP} + C_{FN}} \quad (15.1)$$

where C_{FP} is the cost of a false positive, C_{FN} is the cost of a false negative, and $s(x) = P(y=1 | x)$ is the model's *calibrated posterior probability*.

This threshold applies when $s(x)$ is a well-calibrated posterior. If the model instead outputs a likelihood ratio or an uncalibrated score, multiply by the class prior ratio: $\tau^* = C_{FP} P(y=0) / [C_{FP} P(y=0) + C_{FN} P(y=1)]$. For

calibrated deep classifiers the two are equivalent after accounting for the prior.

For a cancer screening system with $C_{FP} = \$200$ (unnecessary biopsy) and $C_{FN} = \$50,000$ (missed diagnosis):

$$\theta^* = \frac{200}{200 + 50,000} \approx 0.004$$

This means the system should auto-decide negative only when it is more than 99.6% confident—and should route to human review otherwise. The aggressively low threshold reflects the severe asymmetry: missing a cancer costs 250× more than a false alarm.

15A.1.2 Expected Calibration Error

A well-calibrated model’s confidence scores match empirical accuracy: a model that says “90% confident” should be correct about 90% of the time. The standard scalar metric is **Expected Calibration Error** (ECE):

$$ECE = \sum_{b=1}^B \frac{|B_b|}{n} |\text{acc}(B_b) - \text{conf}(B_b)| \quad (15.2)$$

where B_b is a bucket of predictions grouped by confidence, $\text{acc}(B_b)$ is empirical accuracy in that bucket, and $\text{conf}(B_b)$ is mean confidence (Guo et al., 2017).

If a model is systematically overconfident, its uncertainty estimates cannot be trusted to trigger human review at the right times. $ECE < 0.05$ is generally acceptable for HITL routing; $ECE > 0.10$ requires recalibration (temperature scaling or isotonic regression) before the model’s confidence

can be used for routing decisions. (See also the Ch. 12 appendix for a stricter research-oriented rubric: < 0.02 good, > 0.05 intervene.)

15A.2 Slice-Based Evaluation

The most powerful technique for detecting systematic failures is **slice-based evaluation**: instead of computing one aggregate metric, compute metrics separately for meaningful subpopulations.

15A.2.1 Defining Slices

Slices should be defined based on domain knowledge:

- **Demographics:** Age, gender, race/ethnicity, geographic region, language
- **Domain-specific attributes:** In medical AI, disease severity, imaging equipment type, institution; in lending, transaction amount, merchant category, account tenure
- **Temporal attributes:** Time of day, day of week, season, year
- **Input characteristics:** Text length, image quality, transaction amount

15A.2.2 Automated Slice Discovery

Single-dimensional slicing misses intersectional failures. A model may perform well on each individual attribute while failing on combinations. **Slice finder** methods (Chung et al., 2019) and the Domino approach (Eyuboglu et al., 2022) use clustering and decision-tree-based methods to automatically identify subpopulations where model performance is poor.

```
1 def compute_slice_metrics(df, slice_col, target_col,
2   pred_col,
3   min_slice_size=50):
4   baseline_acc = accuracy_score(df[target_col],
5     df[pred_col])
6   results = []
7
8   for val in df[slice_col].unique():
9     slice_df = df[df[slice_col] == val]
10    if len(slice_df) < min_slice_size:
11      continue
12
13    acc = accuracy_score(slice_df[target_col],
14      slice_df[pred_col])
15
16    results.append({
17      'slice_val': str(val),
18      'n': len(slice_df),
19      'accuracy': acc,
20      'gap_from_baseline': acc - baseline_acc
21    })
22
23 df = pd.DataFrame(results)
24 return df.sort_values('gap_from_baseline')
```

Listing 15.1: Slice-based evaluation core implementation

15A.3 Behavioral Testing: The CheckList Methodology

The CheckList methodology (Ribeiro et al., 2020) adapts software testing principles to ML models. Instead of asking “what is this model’s aggregate accuracy?” it asks “does this model satisfy specific behavioral invariances?”

15A.3.1 Three Test Types

Minimum Functionality Tests (MFT). Does the model handle basic, obvious cases correctly?

Invariance Tests (INV). Does the model’s output remain stable when inputs are changed in ways that should not affect the output? Example for content moderation: *sentiment toward a topic should not change when the demographic group mentioned in the text is changed to an equivalent group.*

Directional Tests (DIR). Does the model’s output change appropriately when inputs are changed in ways that should affect the output? Example: *adding explicit negation (“I do not recommend this”) should decrease positive sentiment score.*

15A.3.2 CheckList for Annotation Quality Diagnosis

CheckList can be applied to test annotation consistency rather than model behavior. For the annotation failure hypothesis (political speech artifact):

```
1 def annotation_invariance_test(annotation_sample,
2   group_pairs):
3     """
4     Test whether annotations are consistent across
5     equivalent
```

```
4     content mentioning different demographic or political
      groups.
5
6     group_pairs: list of (group_a, group_b) strings to
      substitute
7     """
8     consistency_results = []
9
10    for text_template in annotation_sample.templates:
11        labels_by_group = {}
12
13        for group_a, group_b in group_pairs:
14            text_a = text_template.replace('[GROUP]',
15                                           group_a)
16            text_b = text_template.replace('[GROUP]',
17                                           group_b)
18
19            labels_a = annotation_sample.get_labels(text_a)
20            labels_b = annotation_sample.get_labels(text_b)
21
22            kappa = cohen_kappa_score(labels_a, labels_b)
23            consistency_results.append({
24                'template': text_template,
25                'group_a': group_a,
26                'group_b': group_b,
27                'consistency_kappa': kappa,
28                'flag': kappa < 0.6
29            })
30
31    return pd.DataFrame(consistency_results)
```

Listing 15.2: Annotation invariance testing

15A.4 Data Debugging with Influence Functions

When a systematic model failure is identified, **influence functions** (Koh & Liang, 2017) provide a principled method for tracing it to specific training examples.

15A.4.1 The Influence Function Formula

The influence function estimates the change in model parameters if a single training example were removed:

$$\mathcal{I}(z, z_{test}) = -\nabla_{\theta} \mathcal{L}(z_{test}, \hat{\theta})^T H_{\hat{\theta}}^{-1} \nabla_{\theta} \mathcal{L}(z, \hat{\theta}) \quad (15.3)$$

where z is a training example, z_{test} is the test example of interest, $\hat{\theta}$ are the trained parameters, and $H_{\hat{\theta}}$ is the Hessian of the training loss.

Positive influence ($\mathcal{I} > 0$) means upweighting this training example increases the model's loss on the test point—it is a *harmful* example that supports the model's current (potentially incorrect) prediction. Negative influence means upweighting it would reduce test loss, improving accuracy on similar cases.

15A.4.2 Practical Application: Label Noise Detection

A simpler application in annotation-failure scenarios is high-loss training example identification:

```

1 def detect_potential_mislabeled(model, X_train, y_train,
2                               threshold_percentile=95):
3     """
4     Training examples with anomalously high loss are
5     candidates
6     for mislabeling. Return for human re-review.
```

```

6      """
7      preds = model.predict_proba(X_train)
8
9      losses = -np.array([
10         y_train[i] * np.log(preds[i, 1] + 1e-10) +
11         (1 - y_train[i]) * np.log(preds[i, 0] + 1e-10)
12         for i in range(len(y_train))
13     ])
14
15     threshold = np.percentile(losses, threshold_percentile)
16     suspicious_indices = np.where(losses > threshold)[0]
17
18     return {
19         'suspicious_indices': suspicious_indices,
20         'n_suspicious': len(suspicious_indices),
21         'fraction_suspicious': len(suspicious_indices) /
22             len(y_train)
23     }

```

Listing 15.3: Label noise detection via training loss

15A.5 Inter-Annotator Agreement Metrics

For annotation failure diagnosis, the key measurement is inter-annotator agreement.

Cohen's Kappa (two annotators, categorical labels):

$$\kappa = \frac{p_o - p_e}{1 - p_e} \quad (15.4)$$

where p_o is observed agreement and p_e is expected agreement by chance.

Fleiss' Kappa (more than two annotators):

$$\kappa_F = \frac{\bar{P} - \bar{P}_e}{1 - \bar{P}_e} \quad (15.5)$$

Krippendorff's Alpha (most flexible; handles missing data and ordinal scales):

$$\alpha = 1 - \frac{D_o}{D_e} \quad (15.6)$$

Table 15.1: Kappa interpretation for HITL annotation quality. Agreement labels follow the Landis–Koch scale (Landis & Koch, 1977); the action column is the HITL-specific overlay.

Kappa Range	Agreement	HITL Action
> 0.80	Almost perfect	Proceed with training
0.61–0.80	Substantial	Review guidelines for ambiguous cases
0.41–0.60	Moderate	Revise guidelines; additional annotator training
≤ 0.40	Slight to fair	Halt training; redesign annotation process

15A.6 The HITL Debugging Workflow

The five tools above are most powerful when used in sequence:

1. **Calibration analysis** → Are the model's confidence scores trustworthy? If not, uncertainty detection is broken at the source.
2. **Slice-based evaluation** → Are there subpopulations with systematically lower performance? If so, model failure or annotation failure.

3. **CheckList behavioral tests** → Are there specific invariances that fail? If so, likely annotation failure with demographic or contextual bias.
4. **Influence functions / label noise detection** → Which specific training examples are responsible? Prioritize for relabeling.
5. **Inter-annotator agreement analysis** → Is the labeling process itself reliable? If not, the entire feedback loop is suspect.

This workflow should be run at initial deployment, after every model update, monthly as a routine check, and immediately upon any external complaint about systematic failures.

Part V

Beyond and Ahead

The future won't be about machines that never need help. It will be about machines that are smart enough to know exactly when they do.

— Chapter 1

16 Beyond Text and Images

When data is not static and consequences are physical, the principles of human-in-the-loop design remain the same—but everything about their application changes.

— on data in motion

16.1 The Voice That Doesn't Fit the Transcript

A crisis-line counselor works almost blind. There is no face to read, no posture, no trembling hands—everything arrives through the voice: the catch of breath between sentences, the half-second pause before a word like “okay” that turns it from an answer into a concession.

Researchers have tried to teach machines to hear these signals too. Models trained on the *acoustic* texture of hotline calls—pitch, rhythm, timing, not just the words—can now flag elevated suicide risk better than chance. But only modestly so, and the people who build them are careful to frame them as an aid to the counselor rather than a substitute (Su et al., 2025). The reason is the gap this chapter is about: a transcript can read as perfectly composed—measured, even grateful—while the voice it was lifted from is nothing of the kind. A call whose words say “I’m fine” can be a call from someone who is not.

Text-based AI can be tested on its text; image AI can be evaluated on its images. But distress in a voice lives in the prosody—the rhythm, pitch,

and timing of speech—not in the semantic content. The model is trained to read the words. The human is trained to hear the person.

16.2 Domain 1: Audio and Speech

Speech is deceptive. It presents itself as text-with-a-channel, as if the sound were just a delivery mechanism for the words. But decades of research in linguistics and psychoacoustics make clear that the channel carries enormous information: speaker identity, emotional state, health status, engagement level, confidence, deception, fatigue. Strip away the channel and you strip away much of what humans actually respond to.

16.2.1 The Transcription Feedback Loop

The most mature HITL application in speech is transcription correction: a speech recognition system produces a transcript, human reviewers correct errors, and those corrections feed back into model training.

The interesting complication in speech is that correction is bidirectional. When a human corrects “they’re gonna leave at three” to “they’re going to leave at three,” are they correcting an error or imposing a style preference? If the speaker said “gonna,” the correct transcription is “gonna.” A reviewer who “corrects” it to “going to” is not fixing the model—they are injecting a dialect bias that will cause the model to systematically misrepresent speakers who use informal registers—a bias of exactly the kind measured in commercial systems by Koenecke et al. (Koenecke et al., 2020).

In court transcription, medical dictation, and recorded business meetings, systematic dialect bias in transcription can have material consequences for the people whose speech was transcribed.

Building transcription HITL requires explicit guidelines that distinguish error correction from style correction, and quality checks that detect systematic correction patterns correlated with speaker demographics.

16.2.2 Speaker Diarization: The Attribution Problem

Speaker diarization—determining “who said what” in a multi-speaker recording—is one of the hardest problems in speech processing, and one of the most interesting HITL design problems.

A diarization model must segment an audio stream, cluster segments by speaker, and label each cluster with a speaker identity. These errors compound: a segmentation error can cause a clustering error, which causes an attribution error. In legal settings—depositions, recorded hearings, criminal investigations—correct attribution determines what each person said and when. Diarization systems used in legal contexts require HITL designs that are unusually careful about error propagation and that maintain an audit trail connecting each attributed segment to the model’s confidence in that attribution.

16.2.3 Crisis Detection and the Prosody Gap

The Prosody Gap

Modern NLP models are very good at analyzing the semantic content of speech. They are much worse at the prosodic and paralinguistic features—hesitations, flat affect, tremor—that experienced counselors learn to hear. In mental health applications, the model’s uncertainty is not the only relevant signal: there may be cases

where the model is confident (the words are fine) but the human expert would not be.

One design approach: a secondary model, trained on prosodic features, that operates in parallel with the text-based model. When the two models disagree—text says low-risk, prosody says high-risk—the case escalates to human review. The disagreement between models is itself a signal of uncertainty even when neither individual model is uncertain. This principle generalizes: in multimodal systems, model disagreement between modalities is often more informative than single-model confidence.

16.3 Domain 2: Time-Series and Sensor Data

Text and images are static. But much of the world's most important data is dynamic: streams of sensor readings, vital sign monitors, industrial process measurements, seismic instruments. Time-series HITL is structurally different from text/image HITL.

16.3.1 Industrial IoT: The 0.1% That Matters

Consider a pressure sensor in a continuous chemical processing plant. It reads every 100 milliseconds, 24 hours a day—approximately 6 million data points per week. Of those 6 million readings, the vast majority will be within normal operating range, and most of the remainder will trip routine alarms the control system absorbs on its own. But a handful—perhaps a dozen per month—will signal something genuinely anomalous: a valve malfunction, an unexpected exothermic reaction, a blocked line.

The human operator's job is not to watch 6 million numbers. It's to handle the 0.1% of alarms that represent genuine anomalies requiring human judgment. The AI system's job is triage: distinguishing between

the 99.9% of alarms the automatic control system can handle and the 0.1% that need a human.

This creates a specific design challenge. In complex chemical plants, a poorly calibrated alarm system can generate hundreds of alerts per operator per hour—a condition the industry calls **alarm flood**. Research on industrial control room operators shows that alarm flood is one of the most significant contributors to operator error and industrial incidents (EEMUA, 2013).

16.3.2 Earthquake Early Warning: Oversight, Not Review

Earthquake early warning systems, such as those used in Japan and Mexico, broadcast alerts when seismic P-waves are detected, giving citizens seconds of warning before the more destructive S-waves arrive. The human's role in this HITL system is not case-by-case review—the alert decision is automated. The human's role is system-level oversight: monitoring for false alerts, which damage public trust and cause secondary harms, and checking for system failures that would prevent a true alert from broadcasting.

This is a different kind of HITL—oversight and governance rather than triage and review. As systems become faster and more autonomous, this oversight role becomes increasingly important.

16.3.3 ICU Patient Monitoring: The False Alarm Crisis

In the modern intensive care unit, a patient may be connected to a cardiac monitor, pulse oximeter, blood pressure cuff, ventilator alarm, and medication infusion pump. Studies have found that ICU alarm rates can exceed 700 alarms per patient per day. The false alarm rate in cardiac monitor studies alone ranges from 72% to 99% (Cvach, 2012).

Nurses in high-alarm environments report the same phenomenon documented in clinical decision support systems: they stop hearing the alarms. When the one genuine alarm fires, it sounds like the other 699. This is a Stakes Calibration failure at the system level: the HITL design has failed to distinguish between alarms that require immediate intervention and alarms that are merely information.

The solution is not better sensors—the sensors are already excellent. The solution is smarter filtering: a model that integrates all sensor streams in context and produces a ranked patient state assessment rather than a list of individual parameter alarms.

16.4 Domain 3: Scientific Discovery

16.4.1 AlphaFold and the Validation Problem

In 2020, DeepMind's AlphaFold2 achieved a stunning result at the biennial CASP competition: it predicted protein structures with accuracy comparable to experimental methods (Jumper et al., 2021). For many proteins, AlphaFold's predictions were indistinguishable from structures determined by X-ray crystallography.

This was also a HITL design challenge. AlphaFold provides per-residue confidence scores (pLDDT scores) indicating how confident the model is about each part of the predicted structure. Regions with high scores (above 90) are generally reliable. Regions with low scores (below 50) may be genuinely disordered in the protein, or may be regions where the model is guessing. Scores in between warrant case-by-case judgment.

The pLDDT scores are the HITL interface—they tell the human researcher which parts of the prediction to trust and which to verify ex-

perimentally. This illustrates a broader principle in scientific HITL: the model's job is often to narrow the hypothesis space, not to provide final answers. The human's job is to select which hypotheses to test.

16.4.2 Bayesian Optimization in Materials Science

In materials science, a researcher might want to find the composition of a material that maximizes some desired property. The search space is often enormous. Bayesian optimization maintains a probabilistic model of the relationship between composition and property, suggesting at each iteration the composition most likely to produce new information (Shahriari et al., 2016).

The human researcher plays a specific role: they run the experiment, report the outcome, and—crucially—decide whether to trust the reported outcome. Experimental failures in materials science (contaminated precursors, equipment issues, measurement errors) can produce data points that would mislead the optimization if treated as reliable. The researcher's judgment about data quality is part of the feedback signal. The HITL design challenge is to provide the researcher with enough information about the model's current state that they can make informed decisions about when to override the model's recommendation.

16.4.3 Drug Interaction Prediction

Drug-drug interactions (DDIs) are a significant source of adverse events. Machine learning models use molecular graph representations to predict whether two drugs are likely to interact, scanning millions of pairs and flagging candidates for experimental validation. The HITL design involves tiered escalation: high-confidence predictions go into a clinical reference database; low-confidence predictions are routed to pharmacology experts who bring structural reasoning and mechanistic knowledge the

model cannot access. The experts' assessments become additional training labels—active learning in a domain where generating ground truth requires laboratory experiments or human expert judgment.

16.5 Domain 4: Robotics and Physical Systems

When AI acts in the physical world, errors have physical consequences. A misclassified email can be recovered. A robot arm's wrong move may not be. HITL in robotics is not just about improving AI performance—it's about managing the risks of AI acting in domains where mistakes can cause injury, damage, or death.

16.5.1 Teleoperation: The Latency Problem

In hazardous environments—nuclear facilities, deep-sea operations, bomb disposal, space exploration—robots extend the human's reach into places people cannot safely go. Modern teleoperation incorporates AI assistance: the robot handles low-level motor control (balance, obstacle avoidance), while the human handles high-level task planning (which object to pick up, where to put it). This is **shared autonomy**—the human and the robot each contribute to the execution of a task, with the boundary varying based on task difficulty, model confidence, and communication latency.

The HITL design challenge in teleoperation is latency. When a human on Earth is operating a robot on the Moon, communication delays are around one and a half seconds each way—nearly three seconds before the operator can see the effect of a command. The AI assistance layer needs to buffer the human's intentions, execute them smoothly at the robot's end, and alert the human when something unexpected happens that their command stream did not anticipate (Hauser, 2013).

16.5.2 Shared Autonomy: Uncertainty-Driven Handoff

Shared autonomy research, developed for powered wheelchairs and robotic manipulation for users with motor impairments, has produced a framework that applies directly to HITL in physical systems: the robot estimates the probability distribution over the human's intended goal, executes actions consistent with the most likely goal while keeping other plausible goals reachable, and asks for clarification when the distribution is too uncertain to act effectively (Javdani et al., 2015).

This is Uncertainty Detection, Intervention Design, and Stakes Calibration all at once. The system acts autonomously in the region of high confidence about the human's intent. It asks or slows in the region of uncertainty. It errs toward safety in regions where the cost of a wrong action is high.

16.5.3 Prosthetics and BCIs: The Most Intimate HITL Systems

At the furthest boundary of HITL in physical systems are prosthetics and brain-computer interfaces (BCIs). Here, the human is not simply in the loop—the human's nervous system is the loop.

A modern myoelectric prosthetic hand uses surface electrodes to read electrical signals from the user's remaining muscles. A ML model interprets those signals as intended hand movements. The user's brain sends a signal; muscles contract; electrodes read the contraction; the model predicts the intended movement; the prosthetic actuates (Farina et al., 2014).

The calibration problem is unusually personal. Every user's residual muscle patterns are different. A model trained on one user's data will perform poorly on another. Personalization—continuous adaptation to the individual user's patterns over time—is not optional; it's the core

design requirement. This is Feedback Integration as an intimate, embodied process.

Co-Adaptation

In prosthetics and BCIs, **co-adaptation** describes the mutual adaptation of human and AI system over time: the user adapts to the model, and the model adapts to the user. Over time, a skilled BCI user and a well-designed model converge on a shared language. This is the most intimate form of human-AI collaboration we have yet developed—and it is also the most visible instance of a process that occurs in all HITL systems, albeit more slowly and less obviously.

16.6 What Changes When Data Is Not Static

Across all four domains, a common thread emerges. When data is not static, the design of HITL systems must change in specific ways:

Streaming data requires continuous evaluation. A static text classification system can be evaluated on a held-out test set before deployment. A streaming anomaly detection system is always encountering new data. Performance must be monitored continuously, with alerts for drift and degradation.

Real-time decisions compress the intervention window. In earthquake early warning, in ICU alarms, in teleoperation, the window for human intervention is measured in seconds or less. HITL design must either slow the system to create a meaningful intervention window, or design the human's role as oversight and calibration rather than case-by-case review.

Physical-world consequences change the asymmetry. In text classification, the cost of an incorrect automated decision is usually recoverable. In robotics, in patient monitoring, in industrial control, incorrect automated decisions can have consequences that are not recoverable. This shifts the cost-asymmetry calculation dramatically toward caution.

Multimodal data surfaces disagreements between sensors. When a system integrates audio, video, and text—or pressure, temperature, and flow rate—disagreements between modalities often signal conditions outside the model’s training distribution. Model disagreement as an uncertainty signal is a powerful design principle in multimodal HITL systems.

Think of a device or system that monitors something over time—your smartwatch, a weather app, a fitness tracker, a home security system. How does it handle uncertainty in its sensor data? When does it alert you, and when does it act automatically? Apply the Five Dimensions framework: which dimension is handled best? Which is handled worst? What would a better HITL design look like?

Chapter Summary

Key Concepts

- HITL design takes structurally different forms across audio/speech, time-series/sensor data, scientific discovery, and robotics.
- In audio, the prosody gap—information in the voice that transcripts don't capture—requires human oversight that model confidence scores alone cannot substitute for.
- In time-series, the challenge is triage at scale: distinguishing the 0.1% of alarms that matter from the 99.9% that don't.
- In scientific discovery, AI narrows hypothesis spaces and suggests next experiments; humans provide experimental judgment and validate results.
- In robotics, physical-world consequences change the stakes calibration for autonomous action. Co-adaptation describes the mutual learning of human and system over time.

Five Dimensions Check

- *Uncertainty Detection* (Dimension 1): Each modality needs its own signal—the prosody a transcript misses, the one alarm that matters in a flood of sensor data
- *Stakes Calibration* (Dimension 4): Physical-world consequences in robotics raise the stakes of autonomous action well above those of a text classifier

- *Timing* (Dimension 3): With streaming and physical systems the data never stops, so when the human can intervene becomes a first-order constraint
- *Intervention Design* (Dimension 2) and *Feedback Integration* (Dimension 5) take new shapes in each domain but are not this chapter's focus

In the next chapter, we explore the frontier: systems that get better at knowing what they don't know—and the progression from passive HITL to proactive AI collaboration.

Technical Appendix 16A:

Domain-Specific HITL Metrics

16A.1 Audio and Speech HITL Metrics

16A.1.1 Word Error Rate with Demographic Slicing

Word Error Rate (WER) is the standard transcription metric:

$$WER = \frac{S + D + I}{N} \quad (16.1)$$

where S is substitutions, D is deletions, I is insertions, and N is the number of words in the reference transcript.

For HITL audit purposes, WER must be computed by speaker demographics (not just in aggregate). A flagging threshold of 30% worse in relative terms—that is, a slice WER more than 1.3× the overall WER—is a practical starting point; the appropriate threshold depends on the application’s sensitivity requirements.

Koenecke et al. (2020) found that five major commercial ASR systems had an average WER of approximately 35% for Black speakers versus 19% for white speakers—roughly a 2× gap. This represents a systematic HITL failure: annotation from these systems will have systematically different reliability across speaker demographics.

16A.1.2 Diarization Error Rate

Diarization Error Rate (DER) evaluates speaker attribution:

$$DER = \frac{\text{Missed Speech} + \text{False Alarm Speech} + \text{Speaker Confusion}}{\text{Total Reference Speech Time}} \quad (16.2)$$

For HITL routing in legal and medical contexts, DER by confidence bin is critical: if high-confidence attributions are not more accurate than low-confidence attributions, the confidence score cannot be used to route cases to human review.

16A.1.3 Multimodal Crisis Escalation Logic

The crisis detection escalation combines semantic and prosodic risk scores:

$$\text{Escalate} = \begin{cases} 1 & \text{if } r_{\text{combined}} > \theta_H \\ 1 & \text{if } |r_{\text{text}} - r_{\text{prosody}}| > \delta \\ 0 & \text{otherwise} \end{cases} \quad (16.3)$$

where $r_{\text{combined}} = \alpha r_{\text{text}} + (1 - \alpha) r_{\text{prosody}}$, θ_H is the high-risk threshold, and δ is the disagreement threshold. Setting δ appropriately ensures that exactly the kind of case this chapter opened with—text risk low, prosodic risk high—always triggers escalation.

16A.2 Time-Series HITL Metrics

16A.2.1 Alarm Quality Metrics

Four metrics characterize alarm system quality:

Alarm Precision. Of all alarms generated, what fraction were genuine anomalies?

$$\text{Precision} = \frac{TP}{TP + FP} \quad (16.4)$$

Alarm Burden. Average alarms per operator per hour. EEMUA Publication 191 gives steady-state guidance: fewer than 6 alarms per hour (one per ten minutes) is very likely acceptable, 12 per hour (one per five minutes) is manageable, 30 per hour (one per two minutes) is likely over-demanding, and sustained rates above 60 per hour are very likely unacceptable (EEMUA, 2013).

Actionability Rate. Fraction of reviewed alarms that resulted in a human intervention.

Critical Miss Rate. Fraction of genuine anomalies that the system failed to alarm on.

Table 16.1: Alarm Management Performance Thresholds (EEMUA 191)

Alarm Burden	Category	HITL Implication
< 6/hr	Very likely acceptable	Target range for manageable oversight
~ 12/hr	Manageable	Upper bound for sustained oversight
~ 30/hr	Likely over-demanding	Alarm fatigue; redesign filtering
> 60/hr	Very likely unacceptable	Immediate intervention required

Rates well below one alarm per hour may instead indicate under-detection. Note also that these are steady-state averages: EEMUA 191 judges more than roughly 20 alarms in the first ten minutes after a major upset hard to cope with—flood-rate guidance, not a steady-state threshold.

16A.2.2 Temporal Uncertainty Quantification

For time-series anomaly detection, uncertainty estimates must account for temporal context. Monte Carlo Dropout applied to a sequence model provides epistemic uncertainty estimates that can be augmented with changepoint proximity:

$$u_t^* = u_t \cdot (1 + \beta \cdot \text{CP}_t) \quad (16.5)$$

where u_t is the base uncertainty estimate, CP_t is a changepoint proximity score (0–1), and β is a scaling factor. Uncertainty is higher near recent changepoints in the input signal.

16A.3 Scientific Discovery HITL Metrics

16A.3.1 AlphaFold pLDDT Confidence Classification

For drug design applications, the critical metric is the fraction of **binding site** residues in low or very-low confidence categories—not the overall structure mean. A protein with mean pLDDT of 85 may still have an unreliable binding site if the active site residues cluster in the 50–70 range.

16A.3.2 Bayesian Optimization Loop Metrics

Convergence rate Number of iterations to reach 90% of the best observed objective value

Table 16.2: AlphaFold pLDDT Score Interpretation for Drug Design HITL

pLDDT Range	Interpretation	Structural Reliability	Validation Required
≥ 90	Very high	Comparable to experiment	Low
70–90	High	Backbone reliable	Medium (side chains)
50–70	Low	Uncertain	High
< 50	Very low	Likely disordered	Critical

Human override rate Fraction of model suggestions the researcher overrode

Human value-add Average objective improvement of overridden experiments minus average objective improvement of model-suggested experiments. Positive values indicate the researcher is adding value beyond the model; negative values may indicate the researcher is underconfident in the model.

Data quality flag rate Fraction of experimental outcomes flagged by the researcher as unreliable. High rates indicate an experimental process quality issue rather than a model quality issue.

16A.4 Robotics and Physical Systems HITL Metrics

16A.4.1 Shared Autonomy Calibration

A key metric for shared autonomy systems is the correlation between system uncertainty and human intervention rate:

$$\rho = \text{Corr}(\text{Uncertainty}_t, 1[\text{Human Intervened}_t]) \quad (16.6)$$

A high positive correlation ($\rho > 0.5$) indicates the system is correctly identifying when to ask for help. A low or negative correlation indicates that the system is asking for help on the wrong cases—a **Timing** failure in the Five Dimensions framework. Because most deployed systems trigger interventions by an uncertainty threshold, ρ is high partly by construction, so it must be read jointly with the uncertainty–error correlation $\text{Corr}(\text{Uncertainty}_t, 1[\text{Model Error}_t])$: high ρ with low uncertainty–error correlation means the system reliably asks for help on the *wrong* cases.

16A.4.2 Prosthetic Adaptation Metrics

Adaptation trajectory Accuracy slope per session (accuracy fraction per session). A positive slope indicates continuing improvement; slope < 0.005 suggests convergence.

Correction rate Fraction of actuations where the user re-attempted due to incorrect movement. Above 15% suggests recalibration is needed.

Per-movement accuracy Accuracy by intended movement class. Identifies specific movements that require targeted retraining.

Latency Time from muscle signal to actuation onset. Should remain stable or decrease as adaptation proceeds; increasing latency may indicate hardware issues.

16A.5 Universal Metrics for Non-Static HITL Systems

16A.5.1 Distribution Shift Detection

The Kolmogorov-Smirnov two-sample test detects whether the current data distribution has shifted from the training distribution:

$$D_{KS} = \sup_x |F_1(x) - F_2(x)| \quad (16.7)$$

where F_1 and F_2 are the empirical CDFs of the training and recent data samples. Applied feature-by-feature, this identifies which input dimensions have drifted—pointing toward which aspect of the model’s training distribution is no longer representative.

16A.5.2 HITL Health Dashboard Metrics

For continuous deployment monitoring:

- **Human-model agreement rate:** Fraction of cases where the human agrees with the model’s recommendation. Below 70% suggests model degradation or threshold misalignment.
- **Review rate:** Fraction of cases routed to human review. Above the target range suggests model degradation or threshold too conservative; below suggests overconfidence.
- **Reviewer decision consistency:** Cohen’s κ between the same reviewer’s decisions on matched cases at different time points. Below 0.6 suggests reviewer fatigue or changing standards.
- **Outcome tracking:** Where outcomes are observable, track accuracy of decisions including human review, not just pre-review model predictions.

17 Almost-Autonomous: When AI Learns to Ask Better Questions

The most valuable capacity an AI system can develop is to locate its uncertainty precisely and ask about it efficiently.

— on asking better questions

17.1 The Student Who Got Better at Asking

Medical educators have a name for one of the hardest things to teach: **calibration**—the fit between how much a trainee knows and how much they believe they know. Track students through clinical-reasoning tasks and the two come apart easily; early students are frequently miscalibrated, confident where they should hesitate and hesitant where they are in fact right (Cleary et al., 2019).

The students who improve fastest are the ones who learn to close that gap—not by knowing more, exactly, but by locating their uncertainty precisely. Most learners arrive with questions they already know the answer to, designed to elicit the attending physician’s “yes, you’re right”—a way of seeking reassurance rather than knowledge. The more useful question sounds different: “I think it is likely A or B. I have ruled out C for this reason. What I do not know is whether this presentation is consistent with D—is that worth pursuing?” Not “I do not know”—that is too vague to be actionable. Not “I know”—that forecloses learning. But “I know A

and B and C, and the specific thing I do not know is X”—that is a question a teacher can answer.

This capacity—to locate uncertainty precisely and ask about it efficiently—is one of the most valuable things an AI system can develop. And a suite of techniques—some decades old, some very new—has emerged that begins to give AI systems something like this capacity.

17.2 Active Learning: The System That Learns to Use Your Time Well

Active learning starts from an observation: not all data points teach a model equally. Some examples confirm what the model already knows. Others teach it something new. Active learning is a framework for directing human labeling effort toward the examples that will improve the model most efficiently.

In a standard HITL setup without active learning, cases are routed to human review based on model uncertainty alone: uncertain cases go to humans, confident cases are handled automatically. This is better than random review. But it does not distinguish between two types of uncertain cases:

Near-boundary uncertain cases. The model is uncertain because the input is near the decision boundary—it could plausibly belong to either class. A human label on this case directly informs the model’s boundary.

Outlier uncertain cases. The model is uncertain because the input is far from any training examples—it is an extreme outlier, unlike anything the model has seen. A human label on this case is informative but may be less efficiently placed than a label on a near-boundary case.

Active learning adds a second layer of routing: among uncertain cases, which ones are worth prioritizing for human review?

Uncertainty sampling. Route the cases where the model is most uncertain. Simple and effective. Preferentially samples near the decision boundary.

Diversity sampling. Route cases that are dissimilar from existing training data. Expands coverage of the training distribution rather than just refining the boundary.

Query by committee. Maintain multiple models (an ensemble) and route cases where the models disagree most. Model disagreement is a proxy for the cases where the current training data is insufficient.

Expected model change. Route the cases that, if labeled, would change the model's parameters most significantly. Computationally intensive, and not always aligned with error reduction.

In practice, hybrid approaches combining uncertainty and diversity sampling tend to work well (Settles, 2009). The practical benefit: on tasks where it works well, active learning can reach the same model performance with substantially fewer labeled examples than random sampling—but the savings are task-dependent, ranging from dramatic to modest, and occasionally negative (Settles, 2009). For organizations where annotation is expensive, the gains, when they come, are worth pursuing.

17.3 The System That Improves Over Time

Active learning creates a virtuous dynamic: the model gets better at knowing what it does not know, which makes its queries to humans more

efficient, which makes the human labels more targeted, which makes the model improve faster.

This is the operational embodiment of Dimension 5 in the Five Dimensions framework: feedback integration. A system with well-designed active learning is not just using human labels to correct errors. It is systematically building competence in its weakest areas.

In practice, this means the character of human review changes over time in a well-designed HITL system:

Early phase Humans review many cases, including many that seem relatively straightforward. The model is young and uncertain about a lot.

Middle phase Humans review fewer cases, more concentrated near the current decision boundaries. The easy cases have been learned. The hard cases get more attention.

Late phase Humans review a small fraction of cases, concentrated on genuinely difficult edge cases and novel categories. The model handles everything else with high confidence. The human role is more expert and less routine.

This trajectory does not happen by accident. It requires attention to what cases are being routed and why. Systems that route randomly or by simple uncertainty thresholds without active learning integration do not exhibit this trajectory—they keep routing the same types of uncertain cases forever.

17.4 RLHF, DPO, and Constitutional AI: A General Reader's Guide

The techniques that have driven recent advances in large language models have a deep connection to human-in-the-loop design. Understanding them at a conceptual level is valuable for anyone thinking about the future of HITL systems.

Reinforcement Learning from Human Feedback (RLHF) is the technique behind the instruction-following capabilities of ChatGPT, Claude, and similar systems. The basic idea: train a language model to generate responses, show pairs of responses to human raters who indicate which is better, train a reward model on these preferences, and then use the reward model to fine-tune the generative model to produce responses humans prefer (Christiano et al., 2017; Ouyang et al., 2022).

RLHF is a form of HITL learning in the preference dimension: the model is being guided not by “correct labels” but by human preference signals. This makes the annotation task easier—judging which of two responses is better is often easier than generating the correct response from scratch—while capturing rich information about what humans find valuable, helpful, and safe.

The HITL dimension of RLHF is the human preference collection stage. The quality and diversity of the humans providing preferences shapes the values the model learns. If the preference raters are a homogeneous group (geographically, demographically, culturally), the model will learn to optimize for that group's preferences—which may not align with the preferences of the model's global user base.

Direct Preference Optimization (DPO) simplifies RLHF by eliminating the explicit reward model step. Instead of training a reward model and

then fine-tuning against it, DPO directly updates the model parameters to maximize the likelihood of preferred responses relative to rejected ones (Rafailov et al., 2023). The human feedback collection is the same; the technical training pipeline is simpler and more stable.

Constitutional AI (CAI) , developed by Anthropic, takes a different approach: instead of collecting human preferences for all fine-tuning, define a set of principles (the “constitution”) and train the model to evaluate its own outputs against those principles (Bai et al., 2022). Human oversight is still present—in the design of the constitution and in some feedback stages—but the scale of human input required is reduced because the model can apply the principles to self-critique its own outputs.

RLHF, DPO, and Constitutional AI all involve human judgment at a stage of the training process, not just the deployment process. The humans in the loop are shaping the model’s values and capabilities, not just catching its errors. This is a more fundamental form of human-AI collaboration than the case-by-case review of deployed outputs.

17.5 Proactive Assistance: The AI That Asks Before You Realize You Need Help

Most HITL systems are reactive: the AI processes a case, registers its uncertainty, and then asks for help. The human responds when asked.

A more sophisticated design is **proactive**: the AI anticipates when human input will be needed and asks before the situation becomes critical.

Consider a medical decision support system that monitors a patient’s vital signs continuously. Rather than only alerting when a threshold is breached (reactive), it can detect that a patient’s trend lines are heading

toward deterioration—that the current trajectory, if continued for four hours, will require intervention—and alert clinical staff while there is still time to act proactively. This is uncertainty about the future state, not just the current state. And acting on it early is often less costly and more effective than acting after a crisis.

For language models, proactive HITL means the system recognizes when a user's request is in a category where its outputs are less reliable, and flags this proactively: "I should mention that my information on this topic may be outdated—you may want to verify with a current source." This is not the system being asked to express uncertainty. It is the system recognizing its limitations (Kadavath et al., 2022) and surfacing them without being asked.

The engineering challenge of proactive HITL is knowing what the system does not know before it is asked. This requires the uncertainty quantification capabilities discussed in Chapters 2 and 3, applied not just to individual predictions but to the characteristics of the query type: "this type of question is where I tend to be less reliable."

17.6 The Frontier: AI That Knows What It Doesn't Know

The deepest challenge in HITL design—and the most active research frontier—is building AI systems that have accurate self-knowledge about their own limitations.

Current systems know something about their uncertainty on individual inputs: calibrated confidence scores, ensemble variance, Monte Carlo dropout estimates. These tell you how uncertain the model is about this specific prediction.

What current systems do not know well is the *structure* of their ignorance: where the training distribution is thin, what categories of questions

they systematically underperform on, what input characteristics predict poor calibration. This meta-level self-knowledge—knowing not just that you are uncertain here, but why you are uncertain and where else you will be uncertain for the same reason—would enable a fundamentally different quality of HITL.

A system with accurate meta-level self-knowledge could tell a human reviewer: “This case has the same characteristics as a category of cases where I have historically been wrong 30% of the time. My confidence score here does not fully capture this.” It would be asking for help not based on local uncertainty but based on category-level track record.

Research into AI “knowing what it does not know” is advancing along several fronts: epistemic uncertainty quantification in large language models, out-of-distribution detection, calibration on rare categories, and formal verification of model behavior. None of these is yet at the point where AI systems can reliably provide this kind of meta-level uncertainty communication. But the direction is clear, and the tools are improving.

17.7 What “Almost-Autonomous” Means and Why It’s the Right Goal

The title of this chapter claims that “almost-autonomous” is the right goal. This requires some justification.

The fully autonomous goal—AI that handles everything without human involvement—has appeal. It is simpler. It is faster. It requires no human infrastructure. But it fails in the places that matter most: novel situations, high-stakes edge cases, questions that require value judgment, contexts where accountability requires a human in the decision chain.

The purely human-dependent goal—AI that flags everything for human review—defeats the purpose of AI-assisted systems. It adds the cost of AI without capturing its benefits.

The “almost-autonomous” goal captures both: AI systems that handle the vast majority of cases efficiently and autonomously, that know their own limitations accurately enough to route appropriately, that ask for human help at the specific moments where it makes a material difference, and that improve over time through the feedback their asking generates.

This matters most as AI systems take on longer and more open-ended work. An agent that runs for hours—writing and revising code, gathering and synthesizing research, carrying a single task across many steps—faces a problem that a one-shot prediction does not: small errors compound. Context drifts, an early wrong assumption propagates, and by the end the output can be subtly off in ways that are hard to trace back to their source. The defense is not to watch every step, which would defeat the purpose, but to build in *checkpoints*: moments where a human re-orient, verifies what has accumulated, and corrects course before the divergence grows. For long-running systems, the human in the loop is less a reviewer of each action than a navigator who keeps the work pointed in the right direction—which is, not coincidentally, how this very book was written.

The Almost-Autonomous Ideal

This is the capacity those well-calibrated students had. They were not fully autonomous—they still had things to learn. They were not dependent on supervision for everything—they handled most cases on their own. They were almost-autonomous: capable and independent in their areas of competence, precise and honest about their areas of uncertainty.

That capacity—to be capable and honest at the same time—is what the best HITL systems aspire to. It is not the absence of human involvement. It is the intelligent management of human involvement, directed to where it creates the most value.

The next time you use an AI assistant and it gets something wrong, pay attention to whether the system expressed uncertainty before you noticed the error, or expressed confidence and was wrong. Keep track for a week. How often does the system's expressed confidence accurately predict when it is right vs. when it is wrong? This is a practical calibration audit—and it will give you a real sense of how close or far any particular system is from the “knows what it does not know” ideal.

Chapter Summary

Key Concepts

- Active learning directs human labeling effort toward the examples that improve the model most efficiently—on favorable tasks reaching equivalent performance with substantially fewer labels, though the savings vary
- RLHF, DPO, and Constitutional AI are HITL at training time—humans shaping model values, not just catching deployment errors
- Proactive HITL: systems that ask before the uncertainty becomes critical, rather than only when a threshold is breached
- The frontier: AI systems that know the structure of their ignorance, not just the magnitude of their uncertainty on individual cases

- “Almost-autonomous” captures the right goal: capable and independent in areas of competence, precisely honest about uncertainty

Key Examples

Active learning sample efficiency Substantial, task-dependent reductions in annotation needs through intelligent case selection

RLHF preference collection ChatGPT and Claude trained on human preference comparisons rather than gold labels—and the diversity implications for whose preferences are captured

Medical monitoring Detecting deterioration trends before crises—proactive rather than reactive HITL

Constitutional AI Scalable self-oversight through principle application, reducing but not eliminating human annotation requirements

Five Dimensions Check

- All five dimensions: this chapter describes the mature form of a system that has developed all five—detecting uncertainty, asking well, timing correctly, calibrating to stakes, and integrating feedback over time

In the next chapter, we turn from how the loop works to a question this book has so far left standing: who owns it—the data, the models, the labor, and the access.

Technical Appendix 17A: Active Learning Theory and Advanced Alignment Techniques

Chapter 17 describes the frontier of automation—where systems learn from increasingly sparse human feedback, align to human preferences rather than fixed labels, and detect their own limits. This appendix formalizes the theory behind each of those capabilities.

— Mathematical foundations for active learning, RLHF, and meta-uncertainty

17A.1 Active Learning: Formal Framework

17A.1.1 Query Complexity

Passive supervised learning in the agnostic (worst-case noise) setting requires $m = \Theta(d/\epsilon^2)$ labeled examples to achieve error ϵ for a hypothesis class of VC dimension d ; in the realizable case, $O(d/\epsilon)$ labels suffice. Active learning's gains are regime-dependent: under realizable or low-noise assumptions, label complexity can fall to $O(\theta d \log(1/\epsilon))$, where θ is the disagreement coefficient—an exponential improvement in $1/\epsilon$ —while in the fully agnostic worst case, no general improvement over d/ϵ^2 is possible.

Where it exists, this speedup arises because active learning concentrates queries near the decision boundary, where labels provide the most information about the target function f^* .

17A.1.2 Version Space Active Learning

The **version space** V_n is the set of hypotheses consistent with all n observed labels. Active learning strategies can be viewed as minimizing the expected version space volume after each query:

$$x^* = \arg \min_{x \in \mathcal{U}} \mathbb{E}_{y \sim p(y|x)} [\text{Vol}(V_{n+1}(x, y))]$$

Selecting the query that maximally shrinks the version space is equivalent to selecting the most informative query under an information-theoretic criterion.

17A.1.3 BADGE: Batch Active Learning by Diverse Gradient Embeddings

BADGE (Ash et al., 2020) combines uncertainty and diversity in a principled way:

1. For each unlabeled point x_i , compute the gradient of the loss, evaluated at the predicted label, with respect to the model parameters: $g_i = \nabla_{\theta} L(f_{\theta}(x_i), \hat{y}_i)$
2. Select a batch S by k -means++ initialization on the $\{g_i\}$ vectors

The gradient magnitude captures uncertainty (gradients are larger near the decision boundary); k -means++ selection enforces diversity across the batch.

```

1 import numpy as np
2 from sklearn.cluster import kmeans_plusplus
3
4 def badge_select(model, unlabeled_X, batch_size: int) ->
    np.ndarray:
5     """Select batch using BADGE: gradient embeddings +
        k-means++ init."""
6     probs      = model.predict_proba(unlabeled_X)
7     n_classes  = probs.shape[1]
8     n_samples  = unlabeled_X.shape[0]
9
10    pseudo_labels      = probs.argmax(axis=1)
11    gradient_embeddings = np.zeros((n_samples, n_classes))
12
13    for i in range(n_samples):
14        label      = pseudo_labels[i]
15        confidence = probs[i, label]
16        # softmax-level factor  $p - e_{\hat{y}}$  of the full
            last-layer
17        # gradient  $(p - e_{\hat{y}}) \times h(x_i)$ ; the
            hidden-activation
18        # factor  $h(x_i)$  is omitted in this simplified
            implementation
19        gradient_embeddings[i, label] = confidence - 1
20        for j in range(n_classes):
21            if j != label:
22                gradient_embeddings[i, j] = probs[i, j]
23
24    _, indices = kmeans_plusplus(gradient_embeddings,
        n_clusters=batch_size,
25                                random_state=42)
26    return indices

```

Listing 17.1: BADGE selection via gradient embeddings and k-means++

17A.2 RLHF: Mathematical Formulation

17A.2.1 The Reward Model

Given preference pairs $\mathcal{D} = \{(x, y_w, y_l)\}$, the reward model r_ϕ is trained by minimizing the loss:

$$\mathcal{L}_{\text{RM}}(\phi) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \left[\log \sigma(r_\phi(x, y_w) - r_\phi(x, y_l)) \right]$$

This is a Bradley-Terry model of preference:

$$P(y_w \succ y_l \mid x) = \sigma(r_\phi(x, y_w) - r_\phi(x, y_l))$$

17A.2.2 PPO Fine-Tuning

With a trained reward model, the policy π_θ is fine-tuned to maximize the objective:

$$\mathcal{L}_{\text{RLHF}}(\theta) = \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(y|x)} \left[r_\phi(x, y) - \beta \cdot \text{KL}(\pi_\theta(\cdot|x) \parallel \pi_{\text{ref}}(\cdot|x)) \right]$$

The KL penalty β prevents the fine-tuned model from drifting too far from π_{ref} , preserving general capabilities while aligning outputs with human preferences (Christiano et al., 2017; Ouyang et al., 2022).

17A.2.3 Direct Preference Optimization (DPO)

DPO (Rafailov et al., 2023) eliminates the explicit reward model and is trained by minimizing the loss:

$$\mathcal{L}_{\text{DPO}}(\theta) = -\mathbb{E}_{(x, y_w, y_l)} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right) \right]$$

DPO is mathematically equivalent to RLHF under certain assumptions but simpler to train: no reward model, no PPO loop.

17A.3 Meta-Uncertainty: Out-of-Distribution Detection

17A.3.1 Energy-Based OOD Detection

Liu et al. (2020) propose:

$$E(x; f) = -T \cdot \log \sum_c e^{f_c(x)/T}$$

where $f_c(x)$ is the logit for class c and T is a temperature parameter. In-distribution inputs have lower energy; OOD inputs have higher energy, enabling threshold-based routing to human review.

17A.3.2 Calibration Under Distribution Shift

Standard ECE is computed on a test set drawn from the same distribution as training. Under distribution shift, ECE typically degrades. A meta-calibration measure, defined over a specified family $\mathcal{P}$ of shifted distributions (e.g., a held-out set of deployment domains):

$$\text{ECE-shift} = \mathbb{E}_{P' \sim \mathcal{P}} [\text{ECE}(f, P')]$$

Models with lower ECE-shift are more reliably calibrated across deployment contexts—a critical property for HITL systems where deployment conditions evolve.

17A.4 Exercises

Exercise 17A.1 — Active Learning Comparison

On a binary classification dataset, implement uncertainty sampling, random sampling, and BADGE. Start with 100 labeled examples, query 50 per iteration for 10 iterations. Plot accuracy vs. annotation budget. What are the annotation savings at 85% accuracy?

Exercise 17A.2 — RLHF Simulation

Create a simple RLHF simulation: 1,000 synthetic text summaries with a known reward function (conciseness). Simulate human preferences with 10% noise. Train a reward model. Compare learned rankings to ground truth (Spearman ρ). At what noise level does the reward model fail to recover the ground-truth ranking?

Exercise 17A.3 — OOD Detection and Human Routing

Train an image classifier on CIFAR-10. Evaluate OOD detection on SVHN: compute softmax confidence and energy scores for both datasets. Which better separates in-distribution from OOD? At what threshold would you route to human review? Plot the precision/recall tradeoff for routing decisions as the threshold varies.

18 Who Owns the Loop?

The loop is a design object. It is also property.

— on who owns the loop

18.1 The People Who Taught the Machine Manners

In January 2023, an investigation by *Time* magazine introduced readers to some of the most consequential humans in the loop of modern AI—people almost none of those readers had ever heard of (Perrigo, 2023).

They worked in Nairobi, Kenya, employed through Sama, a San Francisco-based outsourcing firm, on a contract for OpenAI. Their job was to read and label textual descriptions of the worst things people do to each other—violence, abuse, hatred—so that a safety classifier could learn to recognize such content and keep it out of ChatGPT’s outputs. The work paid, by *Time*’s accounting, a take-home wage of roughly \$1.32 to \$2 per hour depending on seniority and performance. OpenAI paid Sama \$12.50 per worker-hour for the work, and Sama disputed *Time*’s figures, claiming effective pay of \$1.46 to \$3.74 per hour—a dispute that itself illustrates the opacity this chapter is about. Several workers described lasting psychological effects from months of concentrated exposure to the material. The contract ended early.

If the vocabulary of this book means anything, these workers were not peripheral. They were the loop. Their judgments are the human signal in the safety systems of one of the most widely used AI products ever deployed—Dimension 5, feedback integration, performed by hand. Every reader who has watched a chatbot politely decline to produce something vile has benefited from their labor, without seeing it, at a price that did not appear on any screen.

For seventeen chapters, this book has treated the human-in-the-loop system as a design object: something to be architected, calibrated, measured, and improved. That treatment is correct, and it is incomplete. A HITL system is also an economic object. The data has sources. The labels have authors. The model has a proprietor. The access has a gatekeeper. This chapter asks the question the design perspective quietly skips: not *does the loop work?* but *whose is it?*

18.2 The Ownership Gap

Begin with the raw material. A large language model is trained on text produced by hundreds of millions of people: forum posts, product reviews, news articles, encyclopedia entries, source code, poetry. An image model is trained on photographs that people took of their families, their cities, their weddings. Almost none of these people were asked. At the scale involved, they could not have been asked—and that impossibility has functioned, in practice, as permission. What was produced as expression was collected as resource.

Now consider the product. For the closed-weight systems that lead the field, the trained model is private property: its weights are generally protected as trade secrets, and even the composition of the training corpus is undisclosed—not the documents, not the sources, not the proportions. Open-weight models, whose parameters anyone may download, are a real

exception and an instructive one. Openness of weights has mostly relocated the gap rather than closed it: the corpora behind most open-weight models remain undocumented too. The handful of fully documented open corpora prove that disclosure is feasible wherever it is chosen. The people whose writing taught the model to write own neither the data assembled from their words nor the system trained on it. In most cases they cannot establish even the bare fact of their inclusion—whether their own sentences are inside the model at all is a question the proprietor is not obliged to answer, and one that is technically difficult to settle from the outside (the Technical Appendix discusses **membership inference**, the family of methods that attempts exactly this).

The gap between these two paragraphs is the ownership gap: contribution is nearly universal, ownership is extraordinarily narrow, and between them stands a wall of opacity.

The gap is no longer hypothetical or fringe. In December 2023, *The New York Times* sued OpenAI and Microsoft, alleging that millions of its articles were used to train models that now compete with it as a source of information (Grynbaum & Mac, 2023). Whatever the courts eventually settle, the suit put a structural question into the legal record: what, if anything, is owed to the people and institutions whose work trained the machine? Comparable questions have been raised by novelists, photographers, illustrators, and programmers. The disputes differ in their details; they share a shape. The value was aggregated from many; the asset is held by few; the terms were never negotiated because one side was never at the table.

States and publics have little more visibility than individuals. In most jurisdictions there is no general obligation to disclose what a commercial model was trained on, and the first legal obligations are only now arriving: the European Union’s AI Act requires providers of general-purpose models to publish a “sufficiently detailed summary” of their training con-

tent (European Parliament and Council of the European Union, 2024), and California’s training-data transparency law (AB 2013), whose first compliance deadline arrived in January 2026, requires developers to post documentation of the datasets behind their generative systems (California State Legislature, 2024). The first disclosures under these laws now exist—OpenAI, for one, has published the documentation California requires, though not, as of mid-2026, the Article 53(1)(d) training-content summary—which sharpens the point rather than settling it: disclosure is evidently feasible, what a summary must contain is already contested, and everywhere the obligations do not reach, the default remains opacity.

The reader may notice that this book has spent considerable energy celebrating Dimension 5—the system that learns from the responses it receives. The ownership gap is the shadow of that celebration. Feedback integration describes how the loop learns from everyone. Ownership determines who keeps what it learns.

The Loop Has an Owner

Every human-in-the-loop system in this book is also a property arrangement. Someone owns the model. Someone supplies the judgment. Someone lives with the output. These are rarely the same people, and when their interests diverge, the system’s design encodes whose interests govern. The design question and the ownership question are both answered in every deployed loop—whether or not anyone asks them out loud.

18.3 The Geography of the Loop

Chapter 14 told the story of the geodiversity weddings as a lesson in collection bias: classifiers trained on the field’s standard datasets—collected overwhelmingly from North America and Europe—stumbling on wedding

photographs from Ethiopia and Pakistan, a bride in a white Western dress labeled *bride* and *wedding* while a bride in Indian wedding clothing was labeled *costume* and *performance art* (Shankar et al., 2017; Zou & Schiebinger, 2018). Filed under collection bias, the lesson is about sampling. Stand on the other side of the same fact and it reads differently. The map of whose life counts as training data is not random. It tracks connectivity, language, wealth—and it correlates, uncomfortably, with the map of who owns the systems. Data is gathered globally; datasets are curated, models trained, and revenue booked in a small number of places. Scholars have given this pattern a name—**data colonialism**: the appropriation of human life as raw material, collected from everywhere, processed and owned somewhere quite particular (Couldry & Mejias, 2019). One need not adopt the full framework to notice the asymmetry it points at. Kate Crawford’s tracing of the AI supply chain extends the observation to the physical substrate: the minerals, the energy, the click-work, the territory—extraction all the way down, with the value pooling at the top (Crawford, 2021).

For HITL design, the geography has a specific consequence that earlier chapters let us name precisely. Uncertainty concentrates where training data is thin (Chapter 4). If the data is thin exactly where the datasets’ geography ends, then the people most poorly served by the model are also the people whose cases should be expected to flood the review queue—a conjecture no public audit of a deployed system has measured, which is its own indictment—to be judged, in most production systems, by reviewers hired where the owner operates, under guidelines the owner wrote. The loop is global at its input and narrow at its control points.

18.4 The Labor Behind the Labels

There is a name for the workforce this book has been describing since Chapter 7. Mary Gray and Siddharth Suri call it **ghost work**: the globally

distributed, largely invisible human labor—labeling, moderating, transcribing, verifying—that makes automated systems appear automatic (Gray & Suri, 2019). The work is skilled. The conditions are another matter: piece rates, no continuity, no benefits, no visibility, and in the case of content moderation, sustained exposure to traumatic material as part of the job description.

Here this book's own thesis turns and points at the industry. Its running claim has been that the human contribution to HITL systems is *skilled judgment*: that annotation quality determines model quality (Chapter 7), that annotators detect guideline failures before anyone else (Chapter 15), that the diversity of the rater pool sets the values a model learns (Chapter 17). The industry's own behavior now endorses that claim—unevenly. For scarce expert judgment in coding, medicine, and law, preference annotation has come to command professional rates. The laboratories concluded that the quality of the judgment sets the quality of the model. The conclusion has simply not traveled to the labor that is least visible: safety labeling and content moderation—the work of the Nairobi contract—remains priced as piecework and routed to wherever it is cheapest. The stratification tracks skill less than it tracks visibility, danger, and the bargaining power of the people doing the work—though Chapter 15's annotation-failure taxonomy applies to a moderation queue exactly as it does to a medical one. If the judgment is skilled, the terms follow. Wages that reflect skill rather than piecework arithmetic. Psychological support where the material is traumatic, as a cost of the system rather than an externality dumped on the worker. And a channel by which the people applying the guidelines can change them, because they are the first to know when the guidelines are wrong. The counterargument deserves its due: these jobs often pay above local alternatives, people queue for them, and a client's exit—as the end of the Nairobi contract showed—can hurt the very workers that scrutiny means to protect. The argument here is not abolition. It is terms.

The point is not sentimental. Chapter 17 made the technical version of the argument: RLHF models optimize for the preferences of whoever rates the outputs, and a homogeneous or hurried rater pool narrows the values the model encodes. The justice argument and the quality argument point in the same direction. A loop that treats its human judgment as a commodity input gets commodity judgment—and then ships it, at scale, as the values of the machine.

18.5 The Gate

There is a final asymmetry, newer than the others. Having been built from nearly everyone's data, the systems are not available to nearly everyone.

Some of the gating is commercial terms of service: frontier AI services are simply not offered in long lists of countries, for reasons ranging from sanctions law to risk appetite. Some of it is infrastructure: since October 2022, United States export controls have restricted which countries may buy the advanced processors that train and serve frontier models—aimed first at China, then widened (U.S. Department of Commerce, Bureau of Industry and Security, 2022). And in January 2025, the U.S. Department of Commerce went further, issuing a framework that sorted nearly every country on earth into tiers of eligibility for advanced chips and even for trained model weights (U.S. Department of Commerce, Bureau of Industry and Security, 2025). That particular framework was rescinded four months later by the new administration—announced by press release in May 2025, never formally replaced—which is, in its way, the most instructive fact about it. Access to the most capable systems of the era is a policy artifact, redrawable with each political season, and the people whose eligibility is being redrawn are not parties to the drawing.

Set the two flows side by side. Data was collected from everywhere; the collection recognized no tiers of eligibility. Capability is dispensed

selectively, on terms set by the owners and their governments. Whatever one thinks of any particular control—and there are serious arguments about security and misuse that this chapter does not dismiss—the asymmetry itself is the structural fact. At one end of the loop, everyone is a contributor. At the other, some are collaborators, some are customers, and some are data sources only.

18.6 Calibrated Collaboration, on Whose Terms?

The claim of this book fits in two sentences: the future is not full automation; it is calibrated collaboration. This chapter adds the rider that the word “collaboration” smuggles in. Collaboration is a relationship, and relationships have terms—and in the systems of the present, the terms are set almost entirely by one side.

The Five Dimensions, read one last time through the ownership lens:

Uncertainty detection Whose cases sit in the thin part of the distribution? The geodiversity failures are not randomly assigned; they fall on the people the data skipped.

Intervention design Who designs the question, and who only answers it? The owner designs the intervention; the annotator executes it; the affected person, most often, is its subject rather than its participant.

Timing Whose time is priced cheaply enough to interrupt? Review queues route to wherever judgment is least expensive, which is a statement about wages before it is a statement about design.

Stakes calibration Stakes for whom? The moderator’s trauma and the mislabeled bride’s erasure appear in no loss function. A system can be perfectly stakes-calibrated for its owner and blind to every other party’s stakes.

Feedback integration Contributions flow in; ownership flows back only to the few institutions positioned to litigate or negotiate. The dimension this book has praised most warmly is, economically, very nearly a one-way valve.

None of this argues that the systems should not exist, or that building them earns no return. Training a frontier model consumes real capital and carries real risk, and the builders are entitled to something for it. The argument is narrower and harder to wave away: opacity is a choice, not a necessity; labor conditions are terms, not weather; access tiers are policy, not physics. A field that has learned—painstakingly, over two decades—to calibrate uncertainty can learn to calibrate ownership: provenance that can be audited (Gebu et al., 2021), labor that is priced as the skilled judgment the field’s own results prove it to be, access rules made where the affected can contest them. The alternative is a loop that asks the world for help and keeps the answers.

Pick one AI system you use regularly and spend thirty minutes trying to answer four questions. What is publicly known about its training data? Who labeled or rated that data, where, and on what terms? In which countries is the system unavailable, and why? And if your own writing or photographs were in the training set, how would you find out? Score yourself on how many you could answer with confidence. The gaps you hit are not a failure of your research skills. They are the ownership gap, measured from the outside.

Chapter Summary

Key Concepts

- Every HITL system is a property arrangement as well as a design: someone owns the model, someone supplies the judgment, someone lives with the output—rarely the same people
- The ownership gap: training data is contributed by nearly everyone, models are owned by very few, and even the fact of one's inclusion is generally unknowable
- Dataset geography maps whose life counts as training data; the geodiversity failures fall on the populations the data skipped, who then face the model's uncertainty without a seat at its control points
- Annotation and preference labor is skilled judgment; expert domains now command professional rates while safety and moderation work stays priced as piecework—a stratification that tracks visibility and bargaining power, not skill
- Access to frontier AI is a policy artifact—export tiers, service geographies—revised season by season, while the data collection that built the systems was never tiered at all

Key Examples

Nairobi safety annotation The workers who trained ChatGPT's toxicity filters for under \$2 per hour—Dimension 5 as invisible, underpriced labor

The geodiversity brides Classifiers that recognized only Western weddings—collection bias seen from the ownership side

The New York Times lawsuit The ownership gap arriving in court: what is owed to those whose work trained the model?

The 2025 diffusion framework Every country tiered for access to chips and weights, then untiered by press release and never formally replaced—access as policy artifact

Five Dimensions Check

- All five dimensions have an ownership reading: whose uncertainty counts, who designs versus who answers, whose time is cheap, whose stakes are in the loss function, and who keeps what the feedback teaches

Final chapter: your role in this future—not despite AI capability, but because of it.

Technical Appendix 18A: Provenance, Attribution, and Access

Chapter 18 argues that every HITL system is a property arrangement: contributed data, supplied judgment, proprietary models, gated access. This appendix collects the technical machinery that makes those claims measurable—membership inference, corpus documentation, data valuation, machine unlearning, and architectures that relocate data ownership.

— Formal tools for the ownership questions of Chapter 18

18A.1 Membership Inference: Testing Inclusion from the Outside

Chapter 18 claims that an individual generally cannot establish whether their data was used to train a model. Membership inference is the family of techniques that attempts it, and its partial success is the precise measure of the opacity.

18A.1.1 The Attack Model

Given a trained model f_θ and a candidate record $z = (x, y)$, a **membership inference attack** is a function $\mathcal{A}(f_\theta, z) \in \{0, 1\}$ that guesses whether z was in the training set D . Its quality is measured by the *membership advantage*:

$$\text{Adv}(\mathcal{A}) = \Pr[\mathcal{A} = 1 \mid z \in D] - \Pr[\mathcal{A} = 1 \mid z \notin D]$$

An advantage of 0 means that attack learns nothing about membership (a model leaks nothing only if *no* attack achieves positive advantage); an advantage of 1 means membership is fully recoverable.

Shokri et al. (2017) introduced the shadow-model construction: train many “shadow” models on datasets the attacker controls, observe how those models behave on members versus non-members, and train an attack classifier on that behavioral difference. Yeom et al. (2018) showed that a far simpler signal often suffices: training members tend to have lower loss.

18A.1.2 The Loss-Threshold Attack

The simplest practical attack declares membership when the model’s loss on the record falls below a threshold τ (we write ℓ for the per-example loss, reserving $\mathcal{L}$ for aggregates):

$$\mathcal{A}_\tau(f_\theta, z) = 1[\ell(f_\theta(x), y) < \tau]$$

```

1 import numpy as np
2
3 def membership_advantage(model, members_X, members_y,
4                           nonmembers_X, nonmembers_y,
5                           threshold: float) -> float:
6     """Estimate membership advantage of a loss-threshold
       attack."""

```

```

7      # per_example_loss: array of the model's loss on each
      # example
8      loss_in = per_example_loss(model, members_X, members_y)
9      loss_out = per_example_loss(model, nonmembers_X,
      nonmembers_y)
10
11     tpr = np.mean(loss_in < threshold)  # members flagged
      # as members
12     fpr = np.mean(loss_out < threshold)  # non-members
      # wrongly flagged
13     return tpr - fpr
14
15     # Sweep thresholds; the max advantage measures leakage
16     def max_advantage(model, mX, my, nX, ny) -> float:
17         losses = np.concatenate([per_example_loss(model, mX,
18             my),
19                                     per_example_loss(model, nX,
20                                     ny)])
21     return max(membership_advantage(model, mX, my, nX, ny,
22         t)
23               for t in np.quantile(losses,
24               np.linspace(0.01, 0.99, 99)))

```

Listing 18.1: Loss-threshold membership inference audit

One methodological caution: choosing the threshold by maximizing advantage on the same samples used to report it overstates the leakage—the threshold should be selected on a disjoint calibration split (or on shadow data, per Shokri et al., 2017), with advantage reported on held-out records.

Two readings of the same measurement: to a privacy engineer, high advantage is a leak to be closed; to the excluded contributor of Chapter 18, high advantage is the only audit available. The tool that measures the harm is also the only instrument of transparency the outsider has.

18A.2 Corpus Documentation: What Is Actually in the Data

Geburu et al. (2021) proposed **datasheets for datasets**: standardized documentation covering motivation, composition, collection process, pre-processing, uses, distribution, and maintenance. The proposal treats a dataset the way engineering treats a component—nothing ships without a specification sheet.

Where datasheets do not exist, third-party audits partially reconstruct them. Dodge et al. (2021) audited C4, a widely used web-scraped training corpus, and documented its source-domain distribution, the side effects of its cleaning filters (which disproportionately removed text associated with minority dialects and identities), and the presence of machine-generated and benchmark-contaminating text. The audit demonstrates both that post-hoc documentation is possible and that it is no substitute for provenance recorded at collection time.

The EU AI Act’s requirement that general-purpose model providers publish a “sufficiently detailed summary” of training content (European Parliament and Council of the European Union, 2024) is, in effect, a legally mandated partial datasheet, as is California’s training-data transparency law (California State Legislature, 2024). The open technical question is what granularity makes such a summary auditable rather than ornamental: domain distributions, source lists, and license inventories can be checked; prose descriptions cannot.

18A.3 Data Valuation: Pricing the Contribution

If training data is a contributed input, what is any single contribution worth? Ghorbani and Zou (2019) adapt the Shapley value from cooperative game theory. Let $U(S)$ be the performance of a model trained on data

subset $S \subseteq D$. The **Data Shapley** value of point i is its average marginal contribution across all subsets:

$$\phi_i = \frac{1}{|D|} \sum_{S \subseteq D \setminus \{i\}} \binom{|D| - 1}{|S|}^{-1} [U(S \cup \{i\}) - U(S)]$$

The Shapley formulation has the fairness properties one would want from a pricing rule: contributions that never change performance are worth zero; identical contributions are worth the same; values sum to the total performance in excess of the empty-set baseline, $U(D) - U(\emptyset)$. Exact computation is exponential, but Monte Carlo estimates over sampled permutations converge well in practice.

The obstacle to using data valuation for actual compensation is not primarily technical. It is that valuation presupposes the very disclosure the ownership gap denies: one can only price contributions to a corpus one is allowed to see.

18A.4 Machine Unlearning: Withdrawing a Contribution

Ownership without the right of withdrawal is incomplete. **Machine unlearning** asks: given a trained model and a deletion request for record z , produce a model distributed as if z had never been in the training set—without paying for full retraining.

Exact unlearning demands distributional equivalence:

$$\mathcal{M}(D \setminus \{z\}) \stackrel{d}{=} \text{Unlearn}(\mathcal{M}(D), z)$$

Bourtoule et al. (2021) achieve this with the SISA architecture (Sharded, Isolated, Sliced, Aggregated): partition the training data into shards, train an isolated constituent model per shard, and aggregate predictions. A

deletion request retrains only the affected shard—and only from the slice checkpoint preceding the deleted record. The design trades some accuracy and storage for a deletion cost of $O(|D|/k)$ per request with k shards, rather than $O(|D|)$.

The architecture is instructive beyond its mechanics: unlearning is cheap only when the system was built, from the start, to remember where every contribution went. Provenance is not a compliance afterthought; it is the data structure that makes withdrawal computable.

18A.5 Relocating the Data: Federated Alternatives

The ownership gap is widest when data must be centralized to be useful. McMahan et al. (2017) showed it need not be. In **federated learning**, the data never leaves its owner; what travels is the model update. In federated averaging (FedAvg), each participating client k starts from the current global model θ_t and runs several epochs of local gradient descent on its own data to produce $\theta_{t+1}^{(k)}$; the server then aggregates, weighting by data volume:

$$\theta_{t+1} = \sum_{k=1}^K \frac{n_k}{n} \theta_{t+1}^{(k)}, \quad n = \sum_{k=1}^K n_k$$

The multiple local epochs are the point: a single local gradient step per round would reduce the scheme to distributed gradient descent (FedSGD) and forfeit the communication savings that make federation practical. Federated learning does not by itself resolve ownership—the aggregated model is still held by whoever runs the aggregation, and gradient updates can leak information about the underlying data. But it demonstrates that “the model needs everyone’s data in one warehouse” is an architectural choice, not a law of nature, which is the technical premise behind Chapter 18’s claim that opacity and centralization are choices.

Training compute, meanwhile, has grown at a pace—doubling roughly every six months in the modern deep learning era (Sevilla et al., 2022)—that concentrates frontier training in the small set of organizations able to fund it. Architectures can redistribute data custody; they cannot, by themselves, redistribute compute.

18A.6 Exercises

Exercise 18A.1 — Membership Inference Audit

Train a small classifier on half of a public tabular dataset, holding out the other half. Implement the loss-threshold attack and compute the maximum membership advantage across thresholds. Then repeat with strong L2 regularization and with early stopping. How does the advantage change, and what does that imply about the relationship between overfitting and membership leakage?

Exercise 18A.2 — Data Shapley on a Toy Corpus

On a dataset of 200 examples, estimate Data Shapley values by Monte Carlo permutation sampling (at least 500 permutations) with model accuracy as U . Corrupt the labels of 20 examples and re-estimate. Verify that the corrupted examples receive systematically lower (often negative) values. What fraction of permutations is needed before the ranking stabilizes?

Exercise 18A.3 — Documenting an Undocumented Corpus

Take a sample of at least 10,000 documents from a public web-scraped corpus. Produce a one-page datasheet in the style of Gebru et al. (2021): top source domains, language distribution, estimated fraction of boilerplate or machine-generated text, and any benchmark contamination you can detect. Note explicitly which datasheet questions you could not answer from the artifact alone—that residue is the case for provenance recorded at collection time.

19 Your Role in This Future

The human role is not diminishing. It is becoming more important, and more demanding.

— on the human’s enduring role

19.1 The Last Chapter Is the Beginning

There is a temptation, at the end of a book about human-AI collaboration, to write a rousing conclusion about how humans will always be irreplaceable—how our creativity, our empathy, our wisdom will forever distinguish us from the machines we build. It would be comforting. It would also be slightly dishonest.

The truth is more interesting and more demanding.

Machines are becoming capable in ways that are genuinely surprising. Large language models write text that, in many contexts, is indistinguishable from human writing. Image generators produce visual art. Music generators produce songs. Code generators produce software. The accelerating pace of AI capability development is real, and pretending otherwise changes nothing.

And yet: in domain after domain where AI capability has advanced, the demand for thoughtful human involvement has changed character more than it has fallen. Honesty requires acknowledging that headcount in the most exposed roles has fallen measurably—customer service, translation, and annotation work have all contracted since 2023 as AI systems absorbed

routine volume. But the judgment-bearing core of the human role has not followed the same curve. In the ways that matter most—in the decisions with real consequences for real people, in the moments where judgment requires grounding in a lived life, in the systems that need to earn trust—the human role is not diminishing. It is becoming more important, and more demanding.

This chapter is not a reassurance. It is an invitation.

19.2 What AI Cannot Replace (and Why That's Permanent)

Chapter 17 described a frontier: AI systems that know what they do not know, that ask for help at precisely the right moments, that improve over time through human collaboration. Chapter 18 asked who will own that frontier. Even in the ideal, there are things AI systems cannot do that humans can.

Accountability with consequence. When an AI system makes a decision that harms someone, there must be a human or institution that can be held accountable—not because we lack the technology to make AI more autonomous, but because accountability is a social relationship. A person who bears responsibility for a decision is motivated to exercise judgment carefully. An algorithm does not bear responsibility in the sense that creates this motivation. The human in the loop is not just a cognitive resource; they are the locus of accountability. This is why the Air Canada tribunal ruled that the company, not the chatbot, was responsible for the chatbot's error (Civil Resolution Tribunal of British Columbia, 2024). Accountability requires a human agent.

Value judgment in genuinely contested terrain. Many consequential decisions involve not uncertainty about facts but contestation about values. How should we balance privacy and security? Who should bear the cost of errors in criminal justice algorithms? What counts as harmful content in a culturally diverse global community? These are not questions that have correct answers waiting to be computed. They are questions on which reasonable people disagree, and where the resolution requires democratic deliberation, not optimization.

Representation of lived experience. A model trained on data about a community is not the same as a member of that community exercising judgment about it. A system trained on clinical records from a hospital is not the same as a patient navigating that hospital's care. The perspective of someone who lives with the consequences of a system's outputs is not reducible to a data signal. It is a form of knowledge that must be gathered through deliberate inclusion, not inferred from historical records.

Situated moral responsibility. When you review an AI's recommendation and confirm or override it, you are not just providing information to the system. You are taking responsibility for the outcome—not institutionally, but personally, in the first person. Institutional accountability assigns blame after the fact; it cannot be present in the moment of decision. What is present then is a particular person who will have to live with what the decision does, and who knows it while deciding. That kind of answerability belongs to moral life, not to system architecture.

19.3 New Jobs That HITL Creates

The emergence of AI as an economic force has produced a well-documented labor disruption narrative: AI will automate jobs, workers will need to

reskill, displacement is coming. This is true in some sectors. It is not the whole story. Nor is the rest of the story automatic: the new roles described below demand skills, credentials, and geographic access that displaced workers do not automatically hold, and no market mechanism guarantees that the people who lose the old jobs are the ones who get the new ones.

Human-in-the-loop AI systems create real and growing demand for people with specific skills.

AI trainers and annotation specialists label data, evaluate outputs, and provide the human signal that makes AI systems improve. Whether this is a permanent function in organizations that deploy learning AI systems, or a transitional labor market that shrinks as models pre-label and synthesize more of their own training data, is an open empirical question—and the answer matters to the people entering it. As of 2026, the Bureau of Labor Statistics does not yet have a job category for this work, but the employment is real and currently growing. Whether it is good work is a separate question—Chapter 18 documented the conditions under which much of it is currently done. The function is real today; the terms are not settled, and the people entering these roles will help settle them.

AI auditors and fairness analysts evaluate AI system outputs for bias, error patterns, and systematic failures. As AI deployment grows and regulatory requirements tighten, organizations increasingly need people who can rigorously evaluate AI system fairness and performance—a role that barely existed a decade ago and now exists at every major AI company. This is an interdisciplinary role combining statistics, domain knowledge, and ethics.

Human oversight coordinators manage the human review workflow in production HITL systems—staffing, training, quality control, escalation

protocols. These are operational roles that require both domain expertise and understanding of AI system behavior.

AI red teamers systematically attempt to find failure modes in AI systems—adversarial inputs, edge cases, ways the system can be manipulated or misused. Red teaming has become a standard practice for responsible AI deployment, and the work demands a combination of adversarial creativity and domain knowledge that is genuinely hard to find.

Conversational AI evaluators assess the outputs of language models on dimensions that automated metrics cannot capture: cultural sensitivity, factual accuracy in specialized domains, appropriate hedging, quality of uncertainty communication. These roles exist at every major AI company and are growing.

AI policy specialists work at the intersection of AI capability and governance—in government agencies, regulatory bodies, advocacy organizations, and within AI companies. The regulatory landscape for AI is developing rapidly, and people who understand both the technology and the policy terrain are rare—few training paths produce the combination.

19.4 Skills That Matter in a HITL World

The skills that make humans valuable in human-AI collaboration are not the same as the skills that made humans valuable in labor markets that did not include AI. Understanding this is practically important.

Domain expertise combined with AI literacy. The most valued position in almost every AI-adjacent field is a person with deep domain knowledge (medicine, law, finance, education, public health) who also

understands how AI systems work well enough to evaluate their outputs critically. Not a computer scientist—a doctor who understands what a medical AI is actually measuring. A lawyer who knows what a legal AI can and cannot reliably assess. Domain expertise alone is devalued by AI capability in the routine application domain. Domain expertise combined with AI literacy is dramatically more valuable.

Judgment about edge cases and novel situations. AI systems handle well-defined, well-populated problem categories efficiently. They struggle with novel situations, genuine ambiguity, and cases that lie outside their training distribution. Humans who can recognize when they are outside the AI's competence zone, and exercise sound independent judgment in those situations, are valuable.

Clear communication about uncertainty. In a world where AI systems produce confident-sounding outputs, humans who can clearly communicate the limits of their own and AI knowledge—who can say “I do not know, and here is what we would need to find out”—provide distinctive value.

Adversarial thinking. Understanding how AI systems fail, what their systematic blind spots are, how they can be gamed—this meta-knowledge about AI limitations is increasingly valuable in evaluation, red-teaming, and oversight roles.

Contextual judgment. The capacity to integrate information from multiple sources, hold contextual knowledge that does not fit in a prompt, recognize when stated facts do not fit lived reality—these are distinctly human capabilities that become more valuable, not less, as AI handles more of the routine information processing.

19.5 Civic Participation in AI Oversight

You are already in the loop. This is not a metaphor.

When you click “not spam” on a misclassified email, you are contributing to the training signal that shapes spam detection for millions of users. When you rate a product, you are contributing to recommendation algorithms. When you accept or reject an autocorrect suggestion, you are providing feedback that updates the model. When you report a phishing email, you are contributing to security AI training. The aggregate of billions of such micro-interactions is the human signal that AI systems rely on.

This happens mostly without your knowledge or consent, and with almost no transparency about how it shapes the AI systems you interact with. This is a governance problem. But it also points to something important: the relationship between humans and AI systems is not confined to explicit review workflows. It is embedded in everyday digital life.

Civic participation in AI governance extends this principle to the institutional level. Many decisions about how AI systems are designed, deployed, and regulated are made in processes that nominally invite public participation—regulatory comment processes, advisory committee hearings, corporate ethics board consultations, academic peer review of influential AI systems. These processes exist. They are underutilized by people who are not professionally engaged with AI governance.

The argument that AI governance is too technical for public participation is often wrong—not because the technical details are simple, but because the most important questions are not primarily technical. Should a recidivism prediction algorithm be used in sentencing? Should facial recognition be permitted in public spaces? Who should bear the cost when an AI system makes a systematic error? These questions require technical understanding at the margin, but they are fundamentally questions

about values, power, and accountability that are not resolved by technical expertise alone.

If you have read this book, you have enough technical literacy to participate meaningfully in many of these conversations. The choice of whether to do so is yours.

19.6 The Closing Invitation

Safety briefings usually break through only as comedy—the viral attendant, the seatbelt-demo riffs, a planeload of bored travelers caught laughing. But there is a quieter way to break the ritual that most passengers have learned to tune out. Picture the attendant who simply stops going through the motions: who makes eye contact and tells you, plainly, that they need you to pay attention—not because the rules require it, but because it might someday matter. Nothing about the information changes. What changes is that someone is genuinely present, asking for the same presence in return.

This is what good HITL design aspires to. Not the performance of oversight, but the reality of it. Not a human signature on an AI decision, but a human mind actually engaged with what the AI decided and why.

Adaptive Reserve Is Not Waste

In HITL systems, human review capacity is not merely a cost center to be minimized. It is adaptive reserve: the slack that allows a system to absorb uncertainty, distribution shift, edge cases, and failure recovery without breaking exactly when judgment matters most.

The machines are getting better. The questions they produce will get harder and more consequential, not easier and more routine. The cases

routed to human judgment will be the genuine edge cases, the ones that require the full weight of a situated, experienced, motivated human mind.

Your job is not to approve what the machine decided. Your job is to be present.

That is not a small thing. In a world of increasingly capable machines, genuine human presence—attention, judgment, accountability—is one of the most valuable things a person can offer. Not despite what the machines can do. Because of it.

Go back to the Try This from Chapter 1. Pay attention to how many times a piece of technology asks you something. Notice which ones feel genuinely like a request for your judgment, and which feel like theater. For the ones that feel like theater—the cookie banners, the empty confirmation dialogs, the permission requests you click through without reading—ask who benefits from the theater. For the ones that feel genuine—the fraud alerts, the medical AI flags, the uncertain autocorrect that waits for you to choose—ask how you can be as present and thoughtful as the system deserves. Then, for one of the theater ones: stop. Actually read it. Actually decide. You may surprise yourself.

Chapter Summary

Key Concepts

- The human's enduring role is not premised on AI limitations—it is premised on accountability, value judgment, lived experience, and moral responsibility

- HITL creates new roles: AI trainers, auditors, oversight coordinators, red teamers, evaluators, and policy specialists
- Most valuable HITL skills: domain expertise + AI literacy, edge-case judgment, uncertainty communication, adversarial thinking, contextual reasoning
- You are already in the loop: micro-interactions with AI systems aggregate into the human signal those systems rely on
- Civic participation in AI governance is accessible to anyone with basic AI literacy—and is underutilized
- Oversight capacity is adaptive reserve, not organizational waste; optimizing it away can make an AI system more fragile precisely when reality changes

Key Examples

Air Canada tribunal Accountability requires a human agent—the company, not the chatbot, bore responsibility for the chatbot’s confident error

Annotation specialists A real and growing workforce function—whose permanence, like its terms, is not yet settled

Public comment in AI regulatory processes A form of HITL at the governance level

The flight attendant who asked for genuine presence The model for what human-in-the-loop oversight should feel like

Five Dimensions Check

- All five dimensions completed: this chapter reflects the book's full argument—that well-designed human-AI collaboration, across all five dimensions, makes both the humans and the AI systems more effective.

Afterword

This book began with a Netflix prompt: “Are you still watching?”

The question seems trivial. But it encodes something profound: a system that was powerful enough to just keep playing, that chose instead to pause and ask. Not because it lacked confidence. Because it knew the limits of its confidence. And because it respected the human enough to check in.

That is the relationship this book has been about. Not between humans and machines that serve them. Between humans and systems smart enough to know when they need the human's help, and humble enough to ask for it.

The smartest machines still need us. Not as backup plans or temporary placeholders. As partners whose judgment they are designed to earn, respect, and return to.

Welcome to the loop.

Technical Appendix 19A:

Measuring Human Contribution

Chapter 19 asked what role humans should play in an AI-enabled future. This appendix provides the quantitative tools for measuring that role: the Human Contribution Index, complementarity analysis, civic AI literacy frameworks, and accountability architecture requirements.

— Formal foundations for Chapter 19

19A.1 Quantifying the Human Value-Add

19A.1.1 Human Contribution Index

In a HITL system, the human contribution can be formally measured as the difference between the system’s performance with and without human review:

$$\text{HCI} = \text{Acc}(f_{\text{HITL}}) - \text{Acc}(f_{\text{automated}})$$

Because the HITL system follows the model’s decision on unrouted cases and the human’s decision on routed ones, this difference factors exactly:

$$\text{HCI} = P(\text{routed}) \cdot (\text{Acc}_{\text{human, routed}} - \text{Acc}_{\text{model, routed}})$$

This is an identity, not a decomposition: HCI isolates the accuracy gap on routed cases, weighted by the routing rate. The two ingredients that determine it—*routing benefit* (were the routed cases genuinely harder for the model?) and *review quality* (how accurately does the human resolve them?)—remain entangled inside the single number. A useful report therefore lists the components separately: $P(\text{routed})$, $\text{Acc}_{\text{model, routed}}$ versus $\text{Acc}_{\text{model, unrouted}}$ (does routing select harder cases?), and $\text{Acc}_{\text{human, routed}}$ (review quality).

```

1 import numpy as np
2 from sklearn.metrics import accuracy_score
3
4 def compute_hci(y_true, model_preds, human_decisions,
5                 routed_mask):
6     """
7     y_true:          ground truth labels
8     model_preds:     model predictions for all cases
9     human_decisions: human labels for routed cases only
10    routed_mask:      boolean array, True where case was
11                      routed
12    """
13    hitl_decisions = model_preds.copy()
14    hitl_decisions[routed_mask] = human_decisions
15
16    hitl_acc = accuracy_score(y_true, hitl_decisions)
17    auto_acc = accuracy_score(y_true, model_preds)
18
19    human_acc_on_routed = (
20        accuracy_score(y_true[routed_mask], human_decisions)
21        if routed_mask.sum() > 0 else np.nan
22    )
23    model_acc_on_routed = (

```

```

22         accuracy_score(y_true[routed_mask],
23                        model_preds[routed_mask])
24     if routed_mask.sum() > 0 else np.nan
25 )
26
27 routing_rate = routed_mask.mean()
28 routing_benefit = (0.0 if routed_mask.sum() == 0 else
29                  routing_rate * (human_acc_on_routed
30                                -
31                                model_acc_on_routed))
32
33 return {
34     'hitl_accuracy': hitl_acc,
35     'auto_only_accuracy': auto_acc,
36     'human_accuracy_on_routed': human_acc_on_routed,
37     'model_accuracy_on_routed': model_acc_on_routed,
38     'routing_rate': routing_rate,
39     # NOTE: 'hci' and 'human_value_add' are identical by
40     # construction: hitl_acc - auto_acc equals
41     # routing_rate * (acc_human_routed -
42                     acc_model_routed)
43     'hci': routing_benefit,
44     'human_value_add': hitl_acc - auto_acc,
45 }

```

Listing 19.1: Computing the Human Contribution Index

19A.2 Human–AI Complementarity Analysis

The value of a HITL system depends on *complementarity* between human and AI capabilities:

$$C = P(\text{human correct} \mid \text{model wrong}) - P(\text{human correct} \mid \text{model right})$$

$C > 0$ means humans are more likely to be correct when the model is wrong — the system benefits from human involvement. $C \approx 0$ means human accuracy is independent of model correctness — routing review to model errors buys nothing over random review. $C < 0$ means humans and models tend to fail on the same cases — errors are correlated, which may reflect the same cognitive biases as the training data; as $C \rightarrow -1$, human review becomes perfectly redundant.

```

1  def complementarity_analysis(y_true, model_preds,
2      human_preds):
3      model_correct = (model_preds == y_true)
4      human_correct = (human_preds == y_true)
5
6      p_human_given_model_wrong = (
7          human_correct[~model_correct].mean()
8          if (~model_correct).sum() > 0 else np.nan
9      )
10     p_human_given_model_right = (
11         human_correct[model_correct].mean()
12         if model_correct.sum() > 0 else np.nan
13     )
14
15     return {
16         'complementarity':
17             p_human_given_model_wrong -
18             p_human_given_model_right,
19         'p_human_correct_given_model_wrong':
20             p_human_given_model_wrong,
21         'p_human_correct_given_model_right':
22             p_human_given_model_right,
23         'agreement_matrix': {
24             'both_right': (model_correct &
25                             human_correct).mean(),

```

```
21         'both_wrong':      (~model_correct &  
22                             ~human_correct).mean(),  
23         'human_only_right': (~model_correct &  
24                             human_correct).mean(),  
25         'model_only_right': (model_correct &  
                                ~human_correct).mean(),  
    }  
}
```

Listing 19.2: Complementarity analysis

If $C < 0$, humans and models fail on the same cases. This often occurs when the model was trained on human-generated data and has learned the same systematic biases; the more negative C , the more redundant the review. If instead $C \approx 0$, human accuracy is independent of model correctness: targeted review adds no more than random review. Diagnosis: compute complementarity by input subgroup to identify where human judgment does and does not add value.

19A.3 Civic AI Literacy Framework

Meaningful participation in AI governance does not require deep technical expertise. The following framework maps competency levels to participation forms:

Reading this book places you solidly at Level 2, with exposure to Level 3 concepts. Level 2 is sufficient for meaningful participation in most public governance processes.

Table 19.1: Civic AI literacy levels and participation forms

Level	Key concepts	Participation forms
Level 1 (Basic)	AI as statistical pattern matching; uncertainty is normal; models can be wrong	Noticing and reporting AI errors; informed consumer of AI products
Level 2 (Functional)	Training data determines behaviour; calibration vs. accuracy; fairness tradeoffs	Public comment on AI regulations; understanding AI explanations
Level 3 (Critical)	Feedback loops; impossibility theorems; accountability gaps	Auditing AI systems; advocacy for fairness standards; advising policymakers
Level 4 (Technical)	All of the above plus implementation	Red-teaming; designing audits; building oversight tools

19A.4 Accountability Architecture

A complete accountability architecture for a HITL system has four layers:

Decision Audit Trail (Layer 1)

Every automated and human-reviewed decision is logged with: timestamp, input features (or hash), model confidence, routing decision, human reviewer ID, human decision, and outcome when observable. Without this layer, no post-hoc analysis is possible.

Aggregate Monitoring (Layer 2)

Regular reporting of accuracy by demographic group, false positive and false negative rates, override rates, and reviewer agreement rates. This layer makes disparate impact visible before it becomes a compliance problem.

Independent Audit (Layer 3)

Periodic (quarterly or annual) review by parties independent of the operating team, with access to audit logs and aggregate statistics. Findings reported publicly for high-stakes systems. Independence is the critical property: internal teams cannot audit their own blind spots.

Contestability Pathway (Layer 4)

A documented, functional process by which individuals affected by a system's decision can: (a) receive a meaningful explanation; (b) challenge the decision with a human reviewer who has actual authority to override; (c) receive a response within defined time limits. A contestability page that leads to a contact form is not a contestability pathway.

19A.5 Exercises

Exercise 19A.1 — Human Contribution Audit

Using a classification dataset of your choice, simulate a HITL system:

1. Train a model; route cases with confidence < 0.7 to a simulated human.
2. Simulate human performance: correct 80% of the time on routed cases.
3. Compute HCI and complementarity scores.
4. At what confidence threshold is HCI maximised? Why?

Exercise 19A.2 — Accountability Architecture Audit

Choose an AI system you interact with regularly (spam filter, content recommendation, fraud detection):

1. Which layers of the accountability architecture does it have?
2. Is there a public description of how decisions are made?
3. Can you find a contestability pathway? How accessible is it?
4. Write a one-page accountability assessment.

Exercise 19A.3 — Civic Participation Simulation

Find a recent public comment period on AI regulation (FTC AI guidelines, NIST AI RMF updates, EU AI Act implementation):

1. Read the draft document's key provisions.
2. Write a public comment (250–500 words) using concepts from this book.
3. Identify: what fairness definition does the proposal implicitly adopt? What monitoring requirements does it include? What contestability provisions?

This appendix provides the quantitative and structural tools for measuring, communicating, and demanding human contribution in AI systems, as introduced in Chapter 19.

Bibliography

- Alvarez-Melis, D., & Jaakkola, T. S. (2018). On the robustness of interpretability methods.
- Angelopoulos, A. N., & Bates, S. (2023). Conformal prediction: A gentle introduction [Also arXiv:2107.07511]. *Foundations and Trends in Machine Learning*, 16(4), 494–591. <https://doi.org/10.1561/22000000101>
- Angwin, J., Larson, J., Mattu, S., & Kirchner, L. (2016). Machine bias: There's software used across the country to predict future criminals. and it's biased against blacks. <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>
- Ardila, D., Kiraly, A. P., Bharadwaj, S., Choi, B., Reicher, J. J., Peng, L., Tse, D., Etemadi, M., Ye, W., Corrado, G., Naidich, D. P., & Shetty, S. (2019). End-to-end lung cancer screening with three-dimensional deep learning on low-dose chest computed tomography. *Nature Medicine*, 25, 954–961. <https://doi.org/10.1038/s41591-019-0447-x>
- Ash, J. T., Zhang, C., Krishnamurthy, A., Langford, J., & Agarwal, A. (2020). Deep batch active learning by diverse, uncertain gradient lower bounds. *International Conference on Learning Representations (ICLR)*. <https://openreview.net/forum?id=ryghZJBKPS>
- Bai, Y., Kadavath, S., Kundu, S., Askill, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., et al. (2022). Constitutional ai: Harmlessness from ai feedback.
- Bainbridge, L. (1983). Ironies of automation. *Automatica*, 19(6), 775–779.
- Begoli, E., Bhattacharya, T., & Kusnezov, D. (2019). The need for uncertainty quantification in machine-assisted medical decision making. *Nature Machine Intelligence*, 1(1), 20–23. <https://doi.org/10.1038/s42256-018-0004-1>
- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? *Proceedings*

- of the 2021 ACM Conference on Fairness, Accountability, and Transparency, 610–623.
- Beyer, B., Jones, C., Petoff, J., & Murphy, N. R. (2016). *Site reliability engineering: How Google runs production systems*. O'Reilly Media.
- Bourtole, L., Chandrasekaran, V., Choquette-Choo, C. A., Jia, H., Travers, A., Zhang, B., Lie, D., & Papernot, N. (2021). Machine unlearning. *2021 IEEE Symposium on Security and Privacy (SP)*, 141–159. <https://doi.org/10.1109/SP40001.2021.00019>
- Brush, A. J. B., Lee, B., Mahajan, R., Agarwal, S., Saroiu, S., & Dixon, C. (2011). Home automation in the wild: Challenges and opportunities. *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, 2115–2124. <https://doi.org/10.1145/1978942.1979249>
- Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. *Proceedings of the 1st Conference on Fairness, Accountability and Transparency*, 81, 77–91.
- Bureau d'Enquêtes et d'Analyses pour la Sécurité de l'Aviation Civile. (2012, July). *Final report on the accident on 1st june 2009 to the airbus A330-203 operated by air france flight AF 447* (tech. rep.). BEA France.
- Bussone, A., Stumpf, S., & O'Sullivan, D. (2015). The role of explanations on trust and reliance in clinical decision support systems. *2015 International Conference on Healthcare Informatics*, 160–169. <https://doi.org/10.1109/ICHI.2015.26>
- California State Legislature. (2024). Generative artificial intelligence: Training data transparency (Assembly Bill 2013). https://leginfo.legislature.ca.gov/faces/billNavClient.xhtml?bill_id=202320240AB2013
- CBC News. (2016). Kitchener woman follows GPS right into Tobermory, Ont., harbour. <https://www.cbc.ca/news/canada/kitchener-waterloo/kitchener-woman-gps-drives-off-tobermory-boat-launch-1.3580930>
- CBS News. (2017). 6-year-old orders \$160 dollhouse, 4 pounds of cookies with Amazon's Echo Dot. <https://www.cbsnews.com/news/6-year-old-brooke-neitzel-orders-dollhouse-cookies-with-amazon-echo-dot-alexa/>
- Chouldechova, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2), 153–163. <https://doi.org/10.1089/big.2016.0047>

- Christiano, P. F., Leike, J., Brown, T. B., Martic, M., Legg, S., & Amodei, D. (2017). Deep reinforcement learning from human preferences. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 4299–4307. <https://arxiv.org/abs/1706.03741>
- Chung, Y., Kraska, T., Polyzotis, N., Tae, K. H., & Whang, S. E. (2019). Slice Finder: Automated data slicing for model validation. *IEEE 35th International Conference on Data Engineering (ICDE)*, 1550–1553. <https://doi.org/10.1109/ICDE.2019.00139>
- Civil Resolution Tribunal of British Columbia. (2024, February). *Moffatt v. air canada*. <https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html>
- Cleary, T. J., Konopasky, A., La Rochelle, J. S., Neubauer, B. E., Durning, S. J., & Artino, A. R. (2019). First-year medical students' calibration bias and accuracy across clinical reasoning activities. *Advances in Health Sciences Education*, 24(4), 767–781. <https://doi.org/10.1007/s10459-019-09897-2>
- Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46. <https://doi.org/10.1177/001316446002000104>
- Cohen, J. (1968). Weighted kappa: Nominal scale agreement provision for scaled disagreement or partial credit. *Psychological Bulletin*, 70(4), 213–220. <https://doi.org/10.1037/h0026256>
- Consumer Financial Protection Bureau. (2024). CFPB orders Apple and Goldman Sachs to pay over \$89 million for Apple Card failures. <https://www.consumerfinance.gov/about-us/newsroom/cfpb-orders-apple-and-goldman-sachs-to-pay-over-89-million-for-apple-card-failures/>
- Couldry, N., & Mejias, U. A. (2019). *The costs of connection: How data is colonizing human life and appropriating it for capitalism*. Stanford University Press.
- Court of Justice of the European Union. (2025). Judgment of 27 february 2025, CK v Magistrat der Stadt Wien (Dun & Bradstreet Austria), Case C-203/22. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:62022CJ0203>
- Crawford, K. (2021). *Atlas of AI: Power, politics, and the planetary costs of artificial intelligence*. Yale University Press.

- Cvach, M. (2012). Monitor alarm fatigue: An integrative review. *Biomedical Instrumentation & Technology*, 46(4), 268–277. <https://doi.org/10.2345/0899-8205-46.4.268>
- Dahl, M., Magesh, V., Suzgun, M., & Ho, D. E. (2024). Large legal fictions: Profiling legal hallucinations in large language models [Also arXiv:2401.01301]. *Journal of Legal Analysis*, 16(1), 64–93. <https://doi.org/10.1093/jla/laae003>
- Danziger, S., Levav, J., & Avnaim-Pesso, L. (2011). Extraneous factors in judicial decisions. *Proceedings of the National Academy of Sciences*, 108(17), 6889–6892. <https://doi.org/10.1073/pnas.1018033108>
- Dastin, J. (2018). Amazon scraps secret AI recruiting tool that showed bias against women. <https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G>
- Davani, A. M., Díaz, M., & Prabhakaran, V. (2022). Dealing with disagreements: Looking beyond the majority vote in subjective annotations. *Transactions of the Association for Computational Linguistics*, 10, 92–110. https://doi.org/10.1162/tacl_a_00449
- Davis, J., & Goadrich, M. (2006). The relationship between precision-recall and ROC curves. *Proceedings of the 23rd International Conference on Machine Learning (ICML)*, 233–240. <https://doi.org/10.1145/1143844.1143874>
- Dawid, A. P., & Skene, A. M. (1979). Maximum likelihood estimation of observer error-rates using the em algorithm. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 28(1), 20–28.
- Deng, A., Xu, Y., Kohavi, R., & Walker, T. (2013). Improving the sensitivity of online controlled experiments by utilizing pre-experiment data. *Proceedings of the 6th ACM International Conference on Web Search and Data Mining (WSDM)*, 123–132. <https://doi.org/10.1145/2433396.2433413>
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., & Li, F.-F. (2009). ImageNet: A large-scale hierarchical image database. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 248–255. <https://doi.org/10.1109/CVPR.2009.5206848>
- Dietvorst, B. J., Simmons, J. P., & Massey, C. (2015). Algorithm aversion: People erroneously avoid algorithms after seeing them err. *Journal of Experimental Psychology: General*, 144(1), 114–126. <https://doi.org/10.1037/xge0000033>

-
- Dodge, J., Sap, M., Marasović, A., Agnew, W., Ilharco, G., Groeneveld, D., Mitchell, M., & Gardner, M. (2021). Documenting large webtext corpora: A case study on the Colossal Clean Crawled Corpus. *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, 1286–1305. <https://doi.org/10.18653/v1/2021.emnlp-main.98>
- Dratsch, T., Chen, X., Rezazade Mehrizi, M., Kloeckner, R., Mähringer-Kunz, A., Püsken, M., Baßler, B., Sauer, S., Maintz, D., & Pinto dos Santos, D. (2023). Automation bias in mammography: The impact of artificial intelligence BI-RADS suggestions on reader performance. *Radiology*, 307(4), e222176. <https://doi.org/10.1148/radiol.222176>
- Dressel, J., & Farid, H. (2018). The accuracy, fairness, and limits of predicting recidivism. *Science Advances*, 4(1), eaao5580. <https://doi.org/10.1126/sciadv.aao5580>
- Drew, T., Võ, M. L.-H., & Wolfe, J. M. (2013). The invisible gorilla strikes again: Sustained inattention blindness in expert observers. *Psychological Science*, 24(9), 1848–1853. <https://doi.org/10.1177/0956797613479386>
- EEMUA. (2013). EEMUA publication 191: Alarm systems — a guide to design, management and procurement.
- Elkan, C. (2001). The foundations of cost-sensitive learning. *Proceedings of the Seventeenth International Joint Conference on Artificial Intelligence (IJCAI)*, 973–978.
- European Parliament and Council of the European Union. (2016). Regulation (EU) 2016/679 of 27 april 2016 on the protection of natural persons with regard to the processing of personal data (General Data Protection Regulation). <https://eur-lex.europa.eu/eli/reg/2016/679/oj>
- European Parliament and Council of the European Union. (2024). Regulation (EU) 2024/1689 of 13 june 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>
- Eykholt, K., Evtimov, I., Fernandes, E., Li, B., Rahmati, A., Xiao, C., Prakash, A., Kohno, T., & Song, D. (2018). Robust physical-world attacks on deep learning visual classification. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 1625–1634.

- Eyuboglu, S., Varma, M., Saab, K., Delbrouck, J.-B., Lee-Messer, C., Dunnmon, J., Zou, J., & Ré, C. (2022). Domino: Discovering systematic errors with cross-modal embeddings. *International Conference on Learning Representations (ICLR)*. <https://openreview.net/forum?id=FPCMqjI0jXN>
- Farid, H. (2021). An overview of perceptual hashing. *Journal of Online Trust and Safety*, 1(1). <https://doi.org/10.54501/jots.v1i1.24>
- Farina, D., Jiang, N., Rehbaum, H., Holobar, A., Graimann, B., Dietl, H., & Aszmann, O. C. (2014). The extraction of neural information from the surface EMG for the control of upper-limb prostheses: Emerging avenues and challenges. *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, 22(4), 797–809. <https://doi.org/10.1109/TNSRE.2014.2305111>
- Fawcett, T. (2006). An introduction to ROC analysis. *Pattern Recognition Letters*, 27(8), 861–874. <https://doi.org/10.1016/j.patrec.2005.10.010>
- Fleiss, J. L. (1971). Measuring nominal scale agreement among many raters. *Psychological Bulletin*, 76(5), 378–382. <https://doi.org/10.1037/h0031619>
- Flores, A. W., Bechtel, K., & Lowenkamp, C. T. (2016). False positives, false negatives, and false analyses: A rejoinder to “Machine Bias: There’s software used across the country to predict future criminals. And it’s biased against blacks.” *Federal Probation*, 80(2), 38–46.
- Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. *Proceedings of the 33rd International Conference on Machine Learning*, 1050–1059. <https://arxiv.org/abs/1506.02142>
- Gaube, S., Suresh, H., Raue, M., Merritt, A., Berkowitz, S. J., Lermer, E., Coughlin, J. F., Gutttag, J. V., Colak, E., & Ghassemi, M. (2021). Do as AI say: Susceptibility in deployment of clinical decision-aids. *npj Digital Medicine*, 4(31). <https://doi.org/10.1038/s41746-021-00385-9>
- Gebru, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H., Daumé III, H., & Crawford, K. (2021). Datasheets for datasets. *Communications of the ACM*, 64(12), 86–92. <https://doi.org/10.1145/3458723>
- Ghorbani, A., & Zou, J. (2019). Data Shapley: Equitable valuation of data for machine learning. *Proceedings of the 36th International Conference on Machine Learning*, 2242–2251. <https://proceedings.mlr.press/v97/ghorbani19c.html>

-
- Goddard, K., Roudsari, A., & Wyatt, J. C. (2012). Automation bias: A systematic review of frequency, effect mediators, and mitigators. *Journal of the American Medical Informatics Association*, 19(1), 121–127. <https://doi.org/10.1136/amiajnl-2011-000089>
- Goodfellow, I. J., Shlens, J., & Szegedy, C. (2015). Explaining and harnessing adversarial examples. *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1412.6572>
- Graham, P. (2002, August). *A plan for spam*. <https://www.paulgraham.com/spam.html>
- Gray, M. L., & Suri, S. (2019). *Ghost work: How to stop silicon valley from building a new global underclass*. Houghton Mifflin Harcourt.
- Gretton, A., Borgwardt, K. M., Rasch, M. J., Schölkopf, B., & Smola, A. (2012). A kernel two-sample test. *Journal of Machine Learning Research*, 13, 723–773.
- Grynbaum, M. M., & Mac, R. (2023). The Times sues OpenAI and Microsoft over A.I. use of copyrighted work. <https://www.nytimes.com/2023/12/27/business/media/new-york-times-open-ai-microsoft-lawsuit.html>
- Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural networks. *Proceedings of the 34th International Conference on Machine Learning*, 1321–1330. <https://arxiv.org/abs/1706.04599>
- Hara, K., Adams, A., Milland, K., Savage, S., Callison-Burch, C., & Bigham, J. P. (2018). A data-driven analysis of workers' earnings on Amazon Mechanical Turk. *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*, 1–14. <https://doi.org/10.1145/3173574.3174023>
- Hardt, M., Price, E., & Srebro, N. (2016). Equality of opportunity in supervised learning. *Advances in Neural Information Processing Systems (NeurIPS)*, 29, 3315–3323.
- Hauser, K. (2013). Recognition, prediction, and planning for assisted teleoperation of freeform tasks. *Autonomous Robots*, 35, 241–254. <https://doi.org/10.1007/s10514-013-9350-3>
- Hawaii News Now. (2023). Tourists following GPS directions mistakenly drive van into Hawaii harbor. <https://www.hawaiinewsnow.com/2023/05/02/tourists-mistakenly-drive-suv-into-hawaii-island-harbor/>

- He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770–778. <https://doi.org/10.1109/CVPR.2016.90>
- Holtzman, A., Buys, J., Du, L., Forbes, M., & Choi, Y. (2020). The curious case of neural text degeneration. *International Conference on Learning Representations (ICLR)*.
- Houlsby, N., Huszár, F., Ghahramani, Z., & Lengyel, M. (2011). Bayesian active learning for classification and preference learning. <https://arxiv.org/abs/1112.5745>
- Jain, S., Crețu, A.-M., & de Montjoye, Y.-A. (2022). Adversarial detection avoidance attacks: Evaluating the robustness of perceptual hashing-based client-side scanning. *31st USENIX Security Symposium (USENIX Security 22)*, 2317–2334. <https://www.usenix.org/conference/usenixsecurity22/presentation/jain>
- Javdani, S., Srinivasa, S. S., & Bagnell, J. A. (2015). Shared autonomy via hindsight optimization. *Robotics: Science and Systems (RSS)*. <https://doi.org/10.15607/RSS.2015.XI.032>
- Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., Tunyasuvunakool, K., Bates, R., et al. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*, 596, 583–589. <https://doi.org/10.1038/s41586-021-03819-2>
- Kadavath, S., Conerly, T., Askell, A., Henighan, T., Drain, D., Perez, E., Schiefer, N., et al. (2022). Language models (mostly) know what they know.
- Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., Gray, S., Radford, A., Wu, J., & Amodei, D. (2020). Scaling laws for neural language models.
- Kendall, A., & Gal, Y. (2017). What uncertainties do we need in Bayesian deep learning for computer vision? *Advances in Neural Information Processing Systems*, 30, 5574–5584. <https://arxiv.org/abs/1703.04977>
- Kirilenko, A., Kyle, A. S., Samadi, M., & Tuzun, T. (2017). The flash crash: High-frequency trading in an electronic market. *The Journal of Finance*, 72(3), 967–998. <https://doi.org/10.1111/jofi.12498>

-
- Koenecke, A., Nam, A., Lake, E., Nudell, J., Quartey, M., Mengesha, Z., Toups, C., Rickford, J. R., Jurafsky, D., & Goel, S. (2020). Racial disparities in automated speech recognition. *Proceedings of the National Academy of Sciences*, 117(14), 7684–7689. <https://doi.org/10.1073/pnas.1915768117>
- Koh, P. W., & Liang, P. (2017). Understanding black-box predictions via influence functions. *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 1885–1894.
- Kompa, B., Snoek, J., & Beam, A. L. (2021). Second opinion needed: Communicating uncertainty in medical machine learning [Article 4]. *npj Digital Medicine*, 4(1). <https://doi.org/10.1038/s41746-020-00367-3>
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. *Advances in Neural Information Processing Systems*, 25, 1097–1105.
- Kusner, M. J., Loftus, J., Russell, C., & Silva, R. (2017). Counterfactual fairness. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 4066–4076.
- Lakshminarayanan, B., Pritzel, A., & Blundell, C. (2017). Simple and scalable predictive uncertainty estimation using deep ensembles. *Advances in Neural Information Processing Systems*, 30, 6402–6413. <https://arxiv.org/abs/1612.01474>
- Landis, J. R., & Koch, G. G. (1977). The measurement of observer agreement for categorical data. *Biometrics*, 33(1), 159–174. <https://doi.org/10.2307/2529310>
- Lee, J. D., & See, K. A. (2004). Trust in automation: Designing for appropriate reliance. *Human Factors*, 46(1), 50–80. <https://doi.org/10.1518/hfes.46.1.50.30392>
- Lee, K., Lee, K., Lee, H., & Shin, J. (2018). A simple unified framework for detecting out-of-distribution samples and adversarial attacks. *Advances in Neural Information Processing Systems (NeurIPS)*, 31, 7167–7177.
- Li, F.-F. (2015, March). *How we're teaching computers to understand pictures*. TED2015, Vancouver. https://www.ted.com/talks/fei_fei_li_how_we_re_teaching_computers_to_understand_pictures
- Linder, J. A., Doctor, J. N., Friedberg, M. W., Reyes Nieva, H., Birks, C., Meeker, D., & Fox, C. R. (2014). Time of day and the decision to prescribe antibiotics.

- JAMA Internal Medicine*, 174(12), 2029–2031. <https://doi.org/10.1001/jamainternmed.2014.5225>
- Liu, W., Wang, X., Owens, J., & Li, Y. (2020). Energy-based out-of-distribution detection. *Advances in Neural Information Processing Systems (NeurIPS)*, 33, 21464–21475.
- Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765–4774.
- Magesh, V., Surani, F., Dahl, M., Suzgun, M., Manning, C. D., & Ho, D. E. (2025). Hallucination-free? assessing the reliability of leading AI legal research tools [Also arXiv:2405.20362]. *Journal of Empirical Legal Studies*, 22(2), 216–242. <https://doi.org/10.1111/jels.12413>
- Mata v. avianca, inc. [Opinion and order on sanctions (Hon. P. Kevin Castel)]. (2023).
- Mathur, N., Baldwin, T., & Cohn, T. (2017). Sequence effects in crowdsourced annotations. *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, 2860–2865. <https://doi.org/10.18653/v1/D17-1306>
- McAllester, D. A. (1999). PAC-bayesian model averaging. *Proceedings of the Twelfth Annual Conference on Computational Learning Theory (COLT)*, 164–170. <https://doi.org/10.1145/307400.307435>
- McMahan, B., Moore, E., Ramage, D., Hampson, S., & Agüera y Arcas, B. (2017). Communication-efficient learning of deep networks from decentralized data. *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, 1273–1282. <https://proceedings.mlr.press/v54/mcmahan17a.html>
- McNeil, B. J., Pauker, S. G., Sox, H. C., & Tversky, A. (1982). On the elicitation of preferences for alternative therapies. *New England Journal of Medicine*, 306(21), 1259–1262. <https://doi.org/10.1056/NEJM198205273062103>
- Michelmores, R., Wicker, M., Laurenti, L., Cardelli, L., Gal, Y., & Kwiatkowska, M. (2020). Uncertainty quantification with statistical guarantees in end-to-end autonomous driving control. *IEEE International Conference on Robotics and Automation (ICRA)*, 7344–7350. <https://doi.org/10.1109/ICRA40945.2020.9196844>

-
- Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient estimation of word representations in vector space. *International Conference on Learning Representations (ICLR) Workshop*.
- Mosier, K. L., Skitka, L. J., Heers, S., & Burdick, M. (1998). Automation bias: Decision making and performance in high-tech cockpits. *International Journal of Aviation Psychology*, 8(1), 47–63. https://doi.org/10.1207/s15327108ijap0801_3
- Mosqueira-Rey, E., Hernández-Pereira, E., Alonso-Ríos, D., Bobes-Bascarán, J., & Fernández-Leal, Á. (2023). Human-in-the-loop machine learning: A state of the art. *Artificial Intelligence Review*, 56, 3005–3054. <https://doi.org/10.1007/s10462-022-10246-w>
- Murphy, L. W., & Meta Platforms, Inc. (2020). Civil rights audit — final report. <https://about.fb.com/wp-content/uploads/2020/07/Civil-Rights-Audit-Final-Report.pdf>
- National Transportation Safety Board. (2019). *Collision between vehicle controlled by developmental automated driving system and pedestrian, tempe, arizona, march 18, 2018* (Highway Accident Report No. NTSB/HAR-19/03). National Transportation Safety Board. Washington, DC. <https://www.nts.gov/investigations/AccidentReports/Reports/HAR1903.pdf>
- New York Department of Financial Services. (2021). Report on Apple Card investigation. https://www.dfs.ny.gov/system/files/documents/2021/03/rpt_202103_apple_card_investigation.pdf
- Norman, D. A. (2013). *The design of everyday things: Revised and expanded edition*. Basic Books.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. (2022). Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems (NeurIPS)*, 35, 27730–27744.
- Paas, F. G. W. C., & Van Merriënboer, J. J. G. (1994). Variability of worked examples and transfer of geometrical problem-solving skills: A cognitive-load approach. *Journal of Educational Psychology*, 86(1), 122–133. <https://doi.org/10.1037/0022-0663.86.1.122>

- Parasuraman, R., & Manzey, D. (2010). Complacency and bias in human use of automation: An attentional integration. *Human Factors*, 52(3), 381–410. <https://doi.org/10.1177/0018720810376055>
- Parasuraman, R., & Riley, V. (1997). Humans and automation: Use, misuse, disuse, abuse. *Human Factors*, 39(2), 230–253. <https://doi.org/10.1518/001872097778543886>
- Perrigo, B. (2023). Exclusive: OpenAI used Kenyan workers on less than \$2 per hour to make ChatGPT less toxic. <https://time.com/6247678/openai-chatgpt-kenya-workers/>
- Pesapane, F., Latronico, A., Abbate, F., Penco, S., Rotili, A., Dominelli, V., Nicosia, L., Monzani, D., Grasso, R., Pravettoni, G., & Cassano, E. (2026). Evaluating cognitive biases in AI-assisted mammography interpretation: A simulation reader study of explainable AI across radiologist experience levels. *European Radiology*. <https://doi.org/10.1007/s00330-026-12666-6>
- Poly, T. N., Islam, M. M., Yang, H.-C., & Li, Y.-C. J. (2020). Appropriateness of overridden alerts in computerized physician order entry: Systematic review. *JMIR Medical Informatics*, 8(7), e15653. <https://doi.org/10.2196/15653>
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., & Finn, C. (2023). Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems (NeurIPS)*, 36, 53728–53741.
- Rajpurkar, P., Irvin, J., Zhu, K., Yang, B., Mehta, H., Duan, T., Ding, D., Bagul, A., Langlotz, C., et al. (2017). CheXNet: Radiologist-level pneumonia detection on chest X-rays with deep learning.
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). “Why Should I Trust You?”: explaining the predictions of any classifier. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1135–1144. <https://doi.org/10.1145/2939672.2939778>
- Ribeiro, M. T., Wu, T., Guestrin, C., & Singh, S. (2020). Beyond accuracy: Behavioral testing of NLP models with CheckList. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 4902–4912. <https://doi.org/10.18653/v1/2020.acl-main.442>

-
- Roberts, S. T. (2019). *Behind the screen: Content moderation in the shadows of social media*. Yale University Press.
- Roppolo, L. P., Westfall, A., Pepe, P. E., Nobel, L., Cowan, J., Kay, J. J., & Idris, A. H. (2009). Dispatcher assessments for agonal breathing improve detection of cardiac arrest. *Resuscitation*, 80(7), 769–772. <https://doi.org/10.1016/j.resuscitation.2009.04.013>
- Röttger, P., Vidgen, B., Hovy, D., & Pierrehumbert, J. (2022). Two contrasting data annotation paradigms for subjective NLP tasks. *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 175–190. <https://doi.org/10.18653/v1/2022.naacl-main.13>
- Sap, M., Card, D., Gabriel, S., Choi, Y., & Smith, N. A. (2019). The risk of racial bias in hate speech detection. *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 1668–1678. <https://doi.org/10.18653/v1/P19-1163>
- Sap, M., Swayamdipta, S., Vianna, L., Zhou, X., Choi, Y., & Smith, N. A. (2022). Annotators with attitudes: How annotator beliefs and identities bias toxic language detection. *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL)*, 5884–5906. <https://doi.org/10.18653/v1/2022.naacl-main.431>
- Searle, J. R. (1980). Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(3), 417–424. <https://doi.org/10.1017/S0140525X00005756>
- Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., & Batra, D. (2017). Grad-CAM: Visual explanations from deep networks via gradient-based localization. *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 618–626. <https://doi.org/10.1109/ICCV.2017.74>
- Sennrich, R., Haddow, B., & Birch, A. (2016). Neural machine translation of rare words with subword units. *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (ACL)*, 1715–1725. <https://doi.org/10.18653/v1/P16-1162>
- Settles, B. (2009). *Active learning literature survey* (tech. rep. No. 1648). University of Wisconsin-Madison, Department of Computer Sciences. <https://burrsettles.com/pub/settles.activelearning.pdf>

- Sevilla, J., Heim, L., Ho, A., Besiroglu, T., Hobbhahn, M., & Villalobos, P. (2022). Compute trends across three eras of machine learning. *2022 International Joint Conference on Neural Networks (IJCNN)*, 1–8. <https://doi.org/10.1109/IJCNN55064.2022.9891914>
- Shahriari, B., Swersky, K., Wang, Z., Adams, R. P., & de Freitas, N. (2016). Taking the human out of the loop: A review of Bayesian optimization. *Proceedings of the IEEE*, 104(1), 148–175. <https://doi.org/10.1109/JPROC.2015.2494218>
- Shankar, S., Halpern, Y., Breck, E., Atwood, J., Wilson, J., & Sculley, D. (2017). No classification without representation: Assessing geodiversity issues in open data sets for the developing world. *NIPS 2017 Workshop on Machine Learning for the Developing World*. <https://arxiv.org/abs/1711.08536>
- Shokri, R., Stronati, M., Song, C., & Shmatikov, V. (2017). Membership inference attacks against machine learning models. *2017 IEEE Symposium on Security and Privacy (SP)*, 3–18. <https://doi.org/10.1109/SP.2017.41>
- Skitka, L. J., Mosier, K. L., & Burdick, M. (1999). Does automation bias decision-making? *International Journal of Human-Computer Studies*, 51(5), 991–1006. <https://doi.org/10.1006/ijhc.1999.0252>
- Steiger, M., Bharucha, T. J., Venkatagiri, S., Riedl, M. J., & Lease, M. (2021). The psychological well-being of content moderators. *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, 1–14. <https://doi.org/10.1145/3411764.3445092>
- Strathern, M. (1997). 'improving ratings': Audit in the British university system. *European Review*, 5(3), 305–321. <https://doi.org/10.1017/S1062798700002660>
- Stureborg, R., Dhingra, B., & Yang, J. (2023). Interface design for crowdsourcing hierarchical multi-label text annotations. *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, 1–17. <https://doi.org/10.1145/3544548.3581431>
- Su, Z., Jiang, H., Yang, Y., Hou, X., Su, Y., & Yang, L. (2025). Acoustic features for identifying suicide risk in crisis hotline callers: Machine learning approach. *Journal of Medical Internet Research*, 27, e67772. <https://doi.org/10.2196/67772>
- Sweller, J. (1988). Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2), 257–285. https://doi.org/10.1207/s15516709cog1202_4

- The Boston Globe. (2021). Man following GPS navigation drives car into Charlton lake. <https://www.bostonglobe.com/2021/03/14/metro/man-following-gps-navigation-drives-car-into-charlton-lake/>
- Tversky, A., & Kahneman, D. (1974). Judgment under uncertainty: Heuristics and biases. *Science*, 185(4157), 1124–1131. <https://doi.org/10.1126/science.185.4157.1124>
- U.S. Department of Commerce, Bureau of Industry and Security. (2022). Implementation of additional export controls: Certain advanced computing and semiconductor manufacturing items. <https://www.federalregister.gov/documents/2022/10/13/2022-21658/implementation-of-additional-export-controls-certain-advanced-computing-and-semiconductor>
- U.S. Department of Commerce, Bureau of Industry and Security. (2025). Framework for artificial intelligence diffusion. <https://www.federalregister.gov/documents/2025/01/15/2025-00636/framework-for-artificial-intelligence-diffusion>
- U.S. Food and Drug Administration. (2021, January). Artificial intelligence/machine learning (AI/ML)-based software as a medical device (SaMD) action plan.
- Valiant, L. G. (1984). A theory of the learnable. *Communications of the ACM*, 27(11), 1134–1142. <https://doi.org/10.1145/1968.1972>
- Vapnik, V. N., & Chervonenkis, A. Y. (1971). On the uniform convergence of relative frequencies of events to their probabilities. *Theory of Probability and Its Applications*, 16(2), 264–280. <https://doi.org/10.1137/1116025>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 5998–6008. <https://arxiv.org/abs/1706.03762>
- Vigdor, N. (2019). Apple card investigated after gender discrimination complaints. <https://www.nytimes.com/2019/11/10/business/Apple-credit-card-investigation.html>
- Voice of America. (2017). Google super bowl commercial accidentally activates devices. <https://www.voanews.com/a/mht-google-super-bowl-ad-activates-google-home-devices/3708026.html>

- Warm, J. S., Parasuraman, R., & Matthews, G. (2008). Vigilance requires hard mental work and is stressful. *Human Factors*, 50(3), 433–441. <https://doi.org/10.1518/001872008X312152>
- Yang, J., Soltan, A. A. S., & Clifton, D. A. (2022). Machine learning generalizability across healthcare settings: Insights from multi-site COVID-19 screening. *npj Digital Medicine*, 5(1), 69. <https://doi.org/10.1038/s41746-022-00614-9>
- Yang, K., Qinami, K., Li, F.-F., Deng, J., & Russakovsky, O. (2020). Towards fairer datasets: Filtering and balancing the distribution of the people subtree in the ImageNet hierarchy. *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT*)*, 547–558. <https://doi.org/10.1145/3351095.3375709>
- Yang, R., & Newman, M. W. (2013). Learning from a learning thermostat: Lessons for intelligent systems for the home. *Proceedings of the 2013 ACM International Joint Conference on Pervasive and Ubiquitous Computing*, 93–102. <https://doi.org/10.1145/2493432.2493489>
- Yeom, S., Giacomelli, I., Fredrikson, M., & Jha, S. (2018). Privacy risk in machine learning: Analyzing the connection to overfitting. *2018 IEEE 31st Computer Security Foundations Symposium (CSF)*, 268–282. <https://doi.org/10.1109/CSF.2018.00027>
- Zaidan, O. F., Eisner, J., & Piatko, C. D. (2007). Using “annotator rationales” to improve machine learning for text categorization. *Proceedings of Human Language Technologies: The Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT)*, 260–267. <https://aclanthology.org/N07-1033>
- Zampieri, M., Malmasi, S., Nakov, P., Rosenthal, S., Farra, N., & Kumar, R. (2019). Predicting the type and target of offensive posts in social media. *Proceedings of NAACL-HLT 2019*, 1415–1420. <https://doi.org/10.18653/v1/N19-1144>
- Zech, J. R., Badgeley, M. A., Liu, M., Costa, A. B., Titano, J. J., & Oermann, E. K. (2018). Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: A cross-sectional study. *PLOS Medicine*, 15(11), e1002683. <https://doi.org/10.1371/journal.pmed.1002683>
- Zou, J., & Schiebinger, L. (2018). AI can be sexist and racist—it’s time to make it fair. *Nature*, 559(7714), 324–326. <https://doi.org/10.1038/d41586-018-05707-8>

Glossary

Active Learning

A machine learning approach where the algorithm can interactively query a human (or other information source) to label new data points, focusing on examples where it is most uncertain.

Aleatoric Uncertainty

Uncertainty arising from inherent randomness or noise in the data. Cannot be reduced by collecting more training data. Also called *data uncertainty* or *statistical uncertainty*.

Alarm Fatigue (Alert Fatigue)

A phenomenon where users become desensitized to frequent alerts or notifications, leading them to ignore or dismiss important warnings. A key failure mode when systems ask for help too often.

Alarm Flood

A condition in which an alarm system generates more alerts than operators can process—hundreds per operator per hour in poorly calibrated industrial plants—causing genuine alarms to be missed or ignored amid the noise. A documented contributor to operator error and industrial incidents.

Alert Fatigue

See *Alarm Fatigue*.

Bayesian Neural Network (BNN)

A neural network that represents uncertainty over its weights using probability distributions rather than point estimates, enabling principled uncertainty quantification.

Calibration

The degree to which a model's predicted probabilities match actual observed frequencies. A well-calibrated model that predicts 70% confidence should be correct approximately 70% of the time. The term is also used for human self-assessment: the fit between what a person knows and what they believe they know.

Cohen's Kappa (κ)

A statistic measuring inter-rater agreement for categorical labels, corrected for chance agreement. Values of $\kappa < 0$ indicate less agreement than chance; $\kappa = 1$ is perfect agreement; $\kappa \geq 0.61$ is generally considered substantial agreement.

Co-Adaptation

The mutual adaptation of a human and an AI system to each other over time: the user learns the system's behavior while the system personalizes to the user. Most visible in prosthetics and brain-computer interfaces, but present in every long-running HITL system.

Confidence Threshold

A predetermined probability level above which a system acts autonomously and below which it requests human intervention.

Data Annotation

The process of labeling raw data (text, images, audio) with ground-truth labels for use in supervised machine learning. Annotation quality directly determines the quality of any model trained on the data.

Data Colonialism

The appropriation of human life as raw material through data collection: gathered from everywhere, processed and owned somewhere quite particular. The term draws an analogy between historical resource extraction and the one-way flow of data from populations to the organizations that own the systems trained on it.

Dawid-Skene Model

A statistical model that jointly estimates true latent labels and individual annotator error rates from multiple noisy annotations using an EM algorithm. Accounts for annotators who are systematically biased toward particular labels.

Deep Ensemble

A collection of neural networks trained independently with different random initializations. Disagreement among ensemble members indicates epistemic uncertainty.

Epistemic Uncertainty

Uncertainty arising from lack of knowledge, which could theoretically be reduced with more training data or a better model. Also called *model uncertainty*.

Expected Calibration Error (ECE)

A metric for measuring how well-calibrated a model's probability predictions are, computed as the weighted average of the difference between accuracy and confidence across bins.

Framing Effect

A cognitive bias where the presentation of information (framing) influences judgment, independent of the underlying facts. In HITL systems, showing annotators a model's prediction before annotation

can anchor their judgment and reduce inter-annotator independence.

Feedback Integration

The fifth dimension of the HITL framework: how a system incorporates human responses to improve future performance.

Five Dimensions Framework

An analytical framework for evaluating HITL systems, comprising: (1) Uncertainty Detection, (2) Intervention Design, (3) Timing, (4) Stakes Calibration, and (5) Feedback Integration.

Graceful Degradation

A design principle where systems fail in ways that minimize harm, such as asking for help rather than acting on uncertain predictions.

Ghost Work

The globally distributed, largely invisible human labor—labeling, moderating, transcribing, verifying—that makes automated systems appear automatic. Coined by Mary Gray and Siddharth Suri.

Hallucination (AI)

When an AI system generates plausible-sounding but factually incorrect information with apparent confidence. A key failure mode in large language models.

Human-in-the-Loop (HITL)

A design paradigm where humans are incorporated into an automated system's decision-making process, particularly for handling uncertain or high-stakes situations.

Inter-Annotator Agreement (IAA)

A measure of how consistently multiple human annotators label the

same data. High IAA indicates the annotation task is well-defined; low IAA may signal ambiguous guidelines, genuinely contested cases, or annotator quality issues. Commonly measured with Cohen's κ or Fleiss' κ .

Intervention Design

The second dimension of the HITL framework: how a system presents its request for human input, including the clarity of the question and available response options.

Label Noise

Errors or inconsistencies in training labels, typically introduced by annotator mistakes, ambiguous guidelines, or inherently subjective tasks. Label noise can systematically degrade model performance, especially when it is correlated with input features.

Machine Unlearning

Techniques for removing the influence of a specific training record from a trained model without full retraining, so the result behaves as if the record had never been included. The technical prerequisite for a meaningful right of data withdrawal.

Membership Inference

A family of techniques for determining, from a model's behavior alone, whether a specific record was part of its training data. To a privacy engineer, a leak to be closed; to an excluded contributor, the only available audit of inclusion.

Monte Carlo Dropout

A technique for estimating uncertainty by running multiple forward passes with dropout enabled during inference and measuring the variance in predictions.

Ownership Gap

The structural distance between who contributes to an AI system (training data from nearly everyone, judgment from annotators and raters) and who owns the result (the model as private property), with opacity—undisclosed corpora, unverifiable inclusion—standing between the two.

Overconfidence

A failure mode where systems act autonomously despite high uncertainty, often leading to errors that could have been prevented by human intervention.

Proactive Assistance

A HITL design pattern in which the AI system anticipates when human input will be needed and asks before the situation becomes critical, rather than only after a threshold is breached.

Reliability Diagram

A visualization tool for assessing model calibration by plotting predicted confidence against observed accuracy across confidence bins.

Shared Autonomy

A control arrangement in which a human and an AI system each contribute to executing a task—the AI handling low-level control, the human high-level intent—with the boundary shifting based on task difficulty, system confidence, and communication latency.

Stakes Calibration

The fourth dimension of the HITL framework: how well a system adjusts its behavior based on the potential consequences of being wrong.

Timing

The third dimension of the HITL framework: when a system chooses to request human input, including synchronous (immediate) vs. asynchronous (batched) interventions.

Uncertainty Detection

The first dimension of the HITL framework: a system's ability to recognize when it lacks confidence in its predictions or decisions.

Uncertainty Quantification

The process of estimating and communicating how confident a model is in its predictions, enabling appropriate decisions about when to act autonomously vs. request human input.

Index

A

Active learning 440
adversarial examples 278
agonal breathing 121
alarm fatigue 515
alarm flood 421, 515
aleatoric uncertainty 53, 515
alert fatigue .. 15, 124, *see* alarm
 fatigue
attention 251
attentional allocation 332
Automation bias 164, 338
automation bias ... 154, 280, 331
Automation complacency ... 338
automation complacency ... 330

B

Bayesian neural network ... 516
blameless post-mortem 398

C

CAI 258
calibrated 272
calibrated accuracy 58

calibrated trust 332
calibrating 19
Calibration 303
calibration 60, 439, 516
calibration session 187
co-adaptation 426, 516
Cohen's kappa 182, 189, 516
confidence signals 203
confidence threshold 516
confused 57
consistency trap 333
Constitutional AI 258
convolutional neural network
 276, 277

D

data annotation 516
data colonialism 461, 517
Data Shapley 473
datasheets for datasets 472
Dawid-Skene model 517
decision boundary 102
deep ensemble 517

Diarization Error Rate 432
 Direct Preference Optimization
 258
 DPO 258

E

empirical risk minimization . . 89
 epistemic uncertainty . . . 53, 517
 expectation confirmation . . . 333
 Expected Calibration Error . 404
 expected calibration error . . 517
 Explainable AI 157

F

F1 score 303
 fairness impossibility theorem
 360
 federated learning 474
 Feedback Integration 132
 feedback integration . . . 204, 518
 filters 291
 five dimensions framework . 518
 five whys 398
 framing effect 518

G

Gender Shades 356
 generalization gap 79, 90
 genuine ambiguity 185
 ghost work 461, 518
 Goodhart's Law 302

graceful degradation 518
 Grad-CAM 158

H

hallucination 253, 518
 hallucinations 14
 HITL Loop 211
 HITL Pipeline 211
 human-in-the-loop 9, 518

I

influence functions 409
 information gain 137
 instruction problem 181
 inter-annotator agreement . . 519
 Intersection over Union 293
 intervention design 519

L

label noise 519
 labels are not neutral 180
 large language models 252
 LIME 158

M

Machine unlearning 473
 machine unlearning 519
 mean Average Precision 293
 membership inference . 459, 519
 membership inference attack 470
 Monte Carlo dropout 519

N

noise disagreement 185

O

optimal query 138

overconfidence 520

overfitting 80

overtrust 332

ownership gap 520

P

perceptual hashing 282

perfectly calibrated 67

pipeline 202

precision 302

proactive 444

proactive assistance 520

R

recall 302

Reinforcement Learning from
Human Feedback . . 205,
257

reliability diagram 520

residual connections 292

reward hacking 258

reward model 257

RLHF 257

S

say-do gap 154

scaled dot-product attention 267

SHAP 158

shared autonomy 424, 520

Slice finder 405

slice-based evaluation 405

spatial hierarchy 292

spot failure 392

stakes calibration 520

statistical auditing 363

systematic failure 393

T

the divergence trap 306

threshold 100

timing 521

tokens 265

transformer 251

trust-attention-response
triangle 124

Two-factor authentication . . 126

U

uncertainty detection 521

uncertainty quantification 13, 521

underfitting 81

undertrust 332

V

vectors 250

version space 452

W

Word Error Rate 431

Word2Vec 250

wrong 57

About the Author

Lazaros Toumanidis conceived, directed, and edited *Human in the Loop: Misunderstood*—but he did not write it alone, or, in the traditional sense, write it at all. The text was produced in collaboration with a range of AI tools: drafted, expanded, and refined through human-in-the-loop iteration, with human curation, fact-checking, and correction throughout.

The book is, in that sense, its own example—made the way it argues good systems should work: not a machine acting alone, not a human doing everything by hand, but a person and their tools each doing what they do best, with a human in the loop who stays accountable for the result.

This project brings together the Winter textbook, the Summer reading edition, the Autumn pamphlet, the Spring build edition, and the companion materials that accompany them.

If you want to follow updates, corrections, or future iterations of the project, use the contact details below.

Connect with the author:

- Websites: laztoum.com | what-if.io | limen-os.io
- Email: laztoum@protonmail.com